

Retrieval-Augmented Generation for Large Language Models: Architectures, Methodologies, Applications, and Future Directions

Generated by SurveyForge

Contents

1	Introduction	2
2	Foundations, Motivations, and Evolution of the Paradigm	3
2.1	Historical Evolution: From Parametric-Only Architectures to Semi-Parametric Systems	3
2.2	The Motivational Nexus: Addressing Core Limitations of Large Language Models	4
2.3	The Evolutionary Taxonomy of Retrieval-Augmented Paradigms	6
2.4	Delineation from Adjacent Techniques for Knowledge Integration	7
3	The Retrieval-Centric Engine: Indexing, Query Processing, and Search Strategies	8
3.1	Corpus Construction, Indexing Structures, and Chunking Strategies	8
3.2	Query Processing, Reformulation, and Semantic Representation	10
3.3	Search Mechanisms: From Sparse Lexical to Dense Semantic and Generative Retrieval	11
3.4	Retriever-Generator Alignment and Co-Optimization Strategies	12
3.5	Context-Based Compression and Efficient Retrieval Context Utilization	14
4	Synergistic Architectures and Fusion: Orchestrating Retrieval and Generation	15
4.1	Context Integration Strategies and the Input Pipeline	15
4.2	Fusion-in-Decoder and Cross-Attention Architectures for Multi-Document Synthesis	17
4.3	Adaptive Retrieval Planning and Dynamic Routing	18
4.4	Recursive, Iterative, and Multi-Hop Retrieval for Complex Reasoning	20
4.5	Post-Retrieval Processing for the Generation Module	22
4.6	Orchestration of Advanced and Multi-Agent RAG Systems	23
5	Training Paradigms and Adaptation for Augmented Models	25
5.1	Foundations of Training Paradigms for Retrieval-Augmented Models	25
5.2	End-to-End Joint Optimization of the Retriever and Generator	26
5.3	Preference Optimization, Alignment, and Human Feedback	27
5.4	Noise Robustness, Denoising, and Multi-Task Learning	29
5.5	Data-Centric Strategies and Synthetic Data Augmentation	30
6	Robustness, Safety, and Trustworthiness in the Loop	31
6.1	Adversarial Threats via Knowledge Base Poisoning and Backdoor Injection . . .	31
6.2	Vulnerabilities Arising from Context Injection and Prompt Manipulation	32
6.3	Defensive Mechanisms and Robustness toward Adversarial Inputs	33
6.4	Safety and Alignment Challenges in Augmented Generation	34
6.5	Ensuring the Memory and Retrieval Pipeline: Privacy Risks and Data Leakage .	36
6.6	Evaluating Bias, Discrimination Amplification, and Sustainability of RAG Systems	37

7	Evaluation and Benchmarking for the Retrieval-Augmented Paradigm	39
7.1	Component-Specific Evaluation Metrics for Retrieval and Generation	39
7.2	End-to-End Evaluation and Multidimensional Frameworks	40
7.3	The Rise of Human and Model-Based Assessment in Benchmarking	41
7.4	Benchmarking Infrastructures, Real-World Fidelity, and Contamination Control .	43
8	Applications, Extensions, Emerging Frontiers, and Future Directions	44
8.1	Knowledge-Intensive Vertical Applications and Continuous Adaptation	44
8.2	Multimodal and Multi-Source Extension	46
8.3	Agentic Workflows, Interaction, and Emergent Reasoning	47
8.4	Emerging Frontiers and Future Research Directions	48
9	Conclusion	49

1 Introduction

The rapid advancement of large language models (LLMs) has fundamentally transformed natural language processing, yet the prevailing paradigm of parametric-only architectures reveals critical limitations that constrain their reliability and utility [1, 2]. These models, such as GPT and T5, encode world knowledge implicitly within their fixed parameters during pre-training, a design that grants them remarkable fluency and broad language understanding but simultaneously renders them susceptible to significant weaknesses: generating plausible yet factually incorrect content, the inability to update stale knowledge post-training without expensive re-training, persistent difficulties in accessing real-time or private information, and notable bottlenecks in reasoning about long-tail or domain-specific knowledge [3, 4]. As the field has converged on scaling model parameters as the primary driver of capability [5], these parametric shortcomings become increasingly consequential, particularly for high-stakes, knowledge-intensive applications where truthfulness and verifiability are critical [6, 7].

Retrieval-Augmented Generation (RAG) has emerged as a promising paradigm to address these fundamental limitations. By transforming the architecture into a semi-parametric system, RAG integrates an external, non-parametric memory—often a dense vector index of a large corpus—with the parametric backbone [8, 9]. As defined in LEWIS et al. [8], the core premise is to relax the sole reliance on the parameters of the pre-trained model by allowing for a differentiable access to explicit non-parametric memory, thereby providing the model with the ability to adapt and provide provenance for its decisions. In practice, the RAG pipeline in a unified framework consisting of a query-processing unit that transforms user input into a searchable representation with generation-augmented or rewriting approach [10, 11, 12, 13], an external memory component (a structured or unstructured knowledge base with an indexing and chunking structure [14]), a retrieval module that efficiently searches the memory for relevant passages (utilizing dense, sparse, or hybrid methods), and a generation module that conditions on the retrieved evidence to produce grounded output [15]. This synergy aims to enhance factual correctness, interpretability, and knowledge updateability [16]; only by providing attribution and provenance can the system be considered reliable [6]. As we will establish, these features are not systematically studied: mitigating the solutions and getting rid of constraints when adapting to expert domains is a fragmented process [12, 12]. The aim to provide a robust baseline for future research, and we believe that both historically and as the field progresses, RAG integrates "wild" external knowledge in a safe and robust manner can identify relevant and irrelevant information, provide faithful attribution and tasks to ensure high-quality performance [4, 8]. This fields shows areas being aided by production cases, for example, financial, medical, customer service, telecommunication [17, 18, 19].

The current literature, while rapidly expanding, faces a critical fragmentation and lack of integration. Surveys have provided valuable yet partial discussions [15, 20, 21], often concentrating on a single component, like the retriever, or a narrow slice of the evolution from naive to advanced paradigms [12, 21]. These surveys often lack a rigorous, formal taxonomy that covers the interconnected nature of system design—particularly data-centric and system-level challenges that directly affect deployment—and a granular analysis of the interplay between training objectives and retrieval strategies. This survey aims to fill this gap by proposing a unified conceptual framework that reconsiders design choices (from retriever and generator to advanced agentic systems) with training and data-centric strategies, and the pressing concerns of robustness and efficiency. We trace the progression from the historical non-parametric k-nearest-neighbor and Memory networks [22] to the scale of REALM and DPR [8] to the current collection of techniques and architectures [21, 10]. The document synthesizes a future roadmap that acknowledges the complexity inherent in the framework and in the system.

The principal contributions of this survey are as follows: Our formalized, systematic, and evolutionary taxonomy of RAG frameworks through state-of-the-art, including the architecture.

(2) An in-depth, component-level analysis of the retrieval-centric engine and the integration. (3) A consolidation of the training paradigms and the methods, on fine-tuning, alignment, and noise robustness. (4) A holistic treatment of the field and its high-stakes, safety, adversarial vulnerabilities, and privacy leakage risks. (5) The emerging landscape of evaluation and benchmarking. (6) An extensive discussion of the emerging directions and real-world applications, including the paramount need for efficiency and sustainability [23]: The survey will follow this structure: Section 2 establishes the foundations; Sections 3 covers retrieval; Section 4 generation and augmentation; Section 5 training; Section 6 evaluation; Section 7 applications and future.

2 Foundations, Motivations, and Evolution of the Paradigm

2.1 Historical Evolution: From Parametric-Only Architectures to Semi-Parametric Systems

The evolutionary trajectory of large language models (LLMs) is fundamentally a story of knowledge representation—a gradual shift from monolithic, parametric-only systems toward hybrid architectures that dynamically interface with external, non-parametric memory. This transition is not merely an engineering convenience but a conceptual necessity born from the inherent limitations of encoding the world’s knowledge solely within static weight matrices. The lineage of modern Retrieval-Augmented Generation (RAG) can be traced through three distinct epochs: the consolidation of parametric-only pre-trained models, the emergence of small-scale non-parametric memory prototypes, and the scaling of these ideas into the massive, general-purpose retrieval-augmented architectures that define contemporary practice.

The rise of parametric-only LLMs, exemplified by the GPT and T5 families, established a new paradigm of scaling laws and emergent capabilities. These models provided a revolutionary mechanism for compressing world knowledge into their parameters, enabling them to reason and generate text with unprecedented fluency [1]. However, this exclusive reliance on implicit knowledge introduced critical vulnerabilities. The knowledge encoded is frozen at the time of training, rapidly becoming stale if it cannot adapt to a shifting information ecosystem [24, 5]. Since all knowledge is learned from static corpora, these models are prone to generating plausible but factually incorrect content—a phenomenon known as hallucination—which makes them ill-suited for knowledge-intensive tasks that require precise factual grounding [8, 25]. Furthermore, the sheer scale of these models makes the continuous updating of their explicit knowledge prohibitively expensive in terms of computation and storage [26]. These monumental limitations of the parametric-only paradigm—the inability to update the world model, the absence of verifiable provenance, and a propensity for hallucination—became the primary catalysts for the exploration of hybrid systems.

Early inklings of non-parametric augmentation were seeded in the form of small-scale memory integrations, which laid the groundwork for subsequent innovations. The k-nearest-neighbor language models (kNN-LM) proposed extending the vocabulary space of an existing LM by directly referencing a datastore of training examples at decoding time, establishing a connection between the generated output and specific data points in a non-parametric manner. In the same vein, memory-augmented architectures like the continuous cache model proposed simplified forms of memory that stored past hidden activations, accessing them via a dot product with the current hidden state to condition predictions on recent history [9]. Concurrently, the broader exploration of memory networks, and a more comprehensive training approach like TRIME, sought to introduce explicit episodic memory stores that could be read and written to during inference [22]. Though small-scale, these systems provided critical empirical and theoretical evidence that explicitly connecting the core neural network to external memory could improve performance on tasks demanding precise, context-aware knowledge, thus introducing the crucial and lasting concept of semi-parametric systems on the path toward modern architectures.

The scaling of these ideas from small memory stores to massive, dense vector indexes marks the third and most defining phase of this paradigm. A major breakthrough was the realization that the retriever itself could be a pre-trained and fine-tuned component, leading to general-purpose "retrieve-read" pipelines. The advent of REALM, which introduced unsupervised pre-training for the retriever using a masked language model objective, demonstrated that retrieval could be trained jointly with the generator to optimize downstream task performance. This was followed by DPR, which established a canonical architecture by training a query and context encoder with a contrastive loss to create a powerful, dense index of Wikipedia. The formalization of this entire paradigm as a non-parametric complementary memory element for knowledge-intensive generation tasks solidified the function of Retrieval-Augmented Generation (RAG) as a core generalizable architecture. By combining a parametric seq2seq generator with a non-parametric, dense vector index accessed via a learned retriever, RAG successfully demonstrated that a semi-parametric system could achieve state-of-the-art results on multiple open-domain QA tasks while simultaneously generating more specific, diverse, and factual language than its parametric-only counterparts [8, 13]. This era of a key paradigm shift anchored previously theoretical ideas to scale, leading to broad, flexible, and computationally efficient implementations.

The foundational lessons from this historical evolution have defined the modern research agenda for RAG. A key insight emerging from this lineage is that the LLM no longer functions merely as a 'reading' device but as a mediator of a complex, information-seeking loop [27]. The initial success of early RAG systems is often constrained by the quality of the retriever; without a high-precision retriever, irrelevant knowledge from the non-parametric memory can degrade generation quality [19, 10]. Moreover, the "lost in the middle" thesis highlights that even modern long-context generators fail to use relevant information if it is not posited correctly, further emphasizing the complex synergy required between retrieval, indexing, and generation [28]. These initial challenges have defined the field's research agenda: addressing task-specific alignment of retriever and generator, the design of adaptive, multi-step retrieval schedules [27, 29], and the development of unified, reconfigurable intermediate modules [21]. From the stable monolithic GPT models to dynamic, knowledge-grounded, and retrievable approaches, the evolutionary path of language models has solidified a new era of generative AI, where the model's parameters are no longer the sole repository of knowledge. Instead, by weaving internal memories with a vast, mutable world, the systematic, modular integration of retrieval and generation has established itself as a central pillar for building reliable, faithful, and continuously adapting AI systems [12].

2.2 The Motivational Nexus: Addressing Core Limitations of Large Language Models

The ascendancy of parametric-only large language models (LLMs)—systems that encode world knowledge exclusively within their frozen weights—has revealed a cluster of interconnected and deeply consequential limitations. While these models demonstrate remarkable fluency and an extensive repository of memorized facts, their purely parametric nature fundamentally constrains their reliability, adaptability, and verifiability in dynamic and high-stakes environments. Retrieval-augmented generation (RAG) has emerged as a powerful corrective, systematically integrating non-parametric, external knowledge to address these shortcomings. This subsection systematically categorizes the core limitations of parametric-only models that form the motivational nexus for the RAG paradigm, linking them to the specific technical innovations discussed throughout this survey.

Arguably the most significant problem catalysing RAG research is the phenomenon of hallucination—the generation of fluent, plausible, yet factually incorrect statements [12]. This failure arises because LLMs are trained to predict the next token from statistical patterns, a process that does not optimize for truthfulness [15]. The issue is particularly pronounced for long-tail knowledge, where data is sparse or rarely observed during training [4]. While

scaling model parameters can mitigate this to a degree, it often fails to appreciably improve the memorization of factual knowledge in the long tail [4]. Fine-tuning, a common remedy, can exacerbate the issue by potentially introducing internal hallucination when the learned parametric knowledge is insufficient or contradictory [24]. By grounding the generative process in retrieved evidence from a trusted external corpus, the system’s output is anchored to specific, verifiable sources, thereby promoting faithfulness and substantially reducing the likelihood of generating unfounded content [8, 30]. This anchoring not only enhances factual correctness but also introduces a path for direct attribution, transforming the generation process from an inscrutable black box into a more inspectable pipeline—a theme that will recur in our discussion of agentic and self-reflective retrieval systems.

A second, fundamental limitation of the parametric-only approach is the inability to refresh and update knowledge in a timely manner. The world knowledge of an LLM is a frozen snapshot captured at the conclusion of computationally prohibitive training phases [2], leaving the model’s internal knowledge to quickly become stale in domains where information changes rapidly. Continual learning, which seeks to update model weights incrementally, is hampered by significant challenges including high training costs and the risk of catastrophic forgetting, where the acquisition of new knowledge severely disrupts previously learned information [26]. Model editing techniques offer a more surgical approach but are typically limited to updating few associations and struggle with generalization across different contexts [31, 32]. In contrast, RAG decouples the knowledge store from the parameters, introducing a framework where the external knowledge source can be dynamically updated [8]. Knowledge updates become a matter of simply modifying the external data source or, more commonly, re-embedding new documents, eliminating the need for frequent and expensive retraining—this flexibility directly addresses the shortcomings of the monolithic parametric era described in the previous subsection and will be elaborated upon in the following taxonomy of RAG architectures.

Closely tied to this concern is the practical need for domain-specific and real-time adaptation. Parametric models, trained on general-purpose web corpora, often struggle to capture the nuanced vocabulary and specialized knowledge required for vertical domains like medicine, law, or telecommunications [33]. While fine-tuning can specialize these models, this demand for substantial amounts of specialized—and often private—data, along with significant computational budgets, makes it cumbersome in practice. RAG sidesteps this limitation by integrating domain-specific knowledge directly into the prompt or context without any parameter updates [12]. For instance, highly technical content such as 3rd Generation Partnership Project (3GPP) standards has been effectively navigated by leveraging an open-source RAG pipeline, demonstrating the paradigm’s utility in tackling complex, expert-level text processing [33]. In addition, in scenarios that require accessing massive amounts of structured data or integrating a wide array of external tools, RAG techniques such as connecting to dynamic databases and APIs allow the model to issue requests for specific information and use the returned JSON, demonstrating far greater adaptability for real-time tasks than a purely frozen parametric model [34]. This not only provides a command-increasing solution to domain constraints but also enables a system to query structured data like knowledge graphs to capture relational knowledge—a task that is particularly challenging for a parametric-only model [35, 36]. These abilities foreshadow the tool- and planner-based orchestration central to Modular and Agentic RAG systems discussed in the following subsection.

A final, often undervalued, motivation is the importation of interpretability and attribution into an opaque process. In pure parametric models, the process of connecting an un-stated fact to a source is virtually impossible. This lack of auditability is especially problematic when AI is deployed in high-stakes domains that require decisions to be reviewed [5]. The very nature of the RAG process—retrieving and presenting evidence alongside the final LLM output—enhances the traceability of information flow [30]. By structuring the retrieved information to support the final output, RAG’s architecture contributes to the building of more trustworthy and attributable

systems [16]. This aspect of verifiability offers a path towards more responsible development of artificial intelligence, and it drives the design of iterative retrieval loops where each piece of evidence can be traced to the final response.

In conclusion, the emergence of retrieval-augmented generation is not driven by a single failure but by the confluence of these fundamental parametric-only limitations—the drive to develop LLMs that are not just persuasive but also more truthful, current, adaptable, and transparent. These motivational factors have spurred innovations in adaptive retrieval, dynamic query scheduling, and iterative retrieval-generation loops that directly address the hallucination problem. They have also led to the development of query classification techniques and router mechanisms to decide when retrieval is necessary, which are central to the broader taxonomy of RAG paradigms outlined in the upcoming subsection [37]. While recent progress in expanding the context windows of some LLMs presents an alternative to RAG by allowing more source text without a separate retrieval step [37], research shows that overly long contexts can overwhelm attention mechanisms, reducing focus on relevant information [11]. This demonstrates the continued importance of weaving an efficient retrieval mechanism into the very process of decoding, ensuring that the interaction between parametric world knowledge and non-parametric memory remains a cornerstone in the evolution of foundational models.

2.3 The Evolutionary Taxonomy of Retrieval-Augmented Paradigms

The design space of retrieval-augmented generation has undergone a profound transformation, evolving from simple, one-shot interventions into sophisticated, adaptive, and composable system architectures. To navigate this rapidly expanding landscape, we introduce a taxonomy structured along the nature and timing of the interaction between the retrieval and generative modules. This framework usefully delineates the progression from early paradigms to contemporary agentic systems, providing a conceptual map for the remainder of this survey.

The foundational paradigm, often referred to as Naive RAG, embodies a straightforward "retrieve-read" process [12]. In this single-step approach, a static pipeline retrieves a fixed set of relevant documents based on the user's query and concatenates them into the prompt without further interaction. This approach, exemplified by early systems like the original RAG model [8], often retrieves indiscriminately, irrespective of whether the model's parametric knowledge is sufficient or the retrieved passages are entirely relevant [30]. While effective for simpler tasks, this paradigm is fundamentally limited in handling multi-step reasoning, and its rigid pipeline is ill-suited for complex user requirements [38]. The retrieved information is treated as a static, one-time context injection, leaving no room for iterative refinement.

Recognizing these constraints, the field progressed to *Advanced RAG* paradigms, characterized by dynamic, iterative, and recursive retrieval-generation loops. A prominent family of methods, including Iter-RetGen and IRCot, interleaves retrieval with generation across multiple turns [29]. The initial output of the generator serves as a context-aware query for the subsequent retrieval step, allowing the system to acquire missing information and perform multi-hop reasoning in a dynamic loop that alternates between thinking and searching. These methods are particularly useful when solving complex, multi-step tasks, as they permit the model to gather a more comprehensive set of evidence to support a final answer [29]. However, they can be resource-intensive. Echoing this philosophy, GenRead has shown that LLM-generated contexts can often rival the utility of externally retrieved ones, blurring the lines between internal and external knowledge sources [39].

The current frontier is represented by *Modular* and *Agentic* RAG systems, which treat RAG as a flexible, composable framework where an LLM acts as a controller. These systems dynamically decide *whether* retrieval is necessary, *what* to retrieve, and *which* tool or memory to use. For instance, self-aware systems like those in [4] highlight the importance of gating retrieval based on model confidence. Self-RAG adaptively retrieves passages "on demand," using reflection tokens to trigger retrieval only when it benefits the generation [30]. More advanced *agentic*

systems further incorporate complex abilities such as planning, query decomposition, and tool invocation. The LLM may even orchestrate a multi-agent system to write and execute code for data queries or navigate a complex ecosystem of tools, as seen in frameworks that explore a high-level, modular composition [40]. This represents a move beyond simple augmentation and toward a full system-level orchestration.

This evolutionary progression can be synthesized into a unified design framework for the RAG design space. The essential axes of comparison include **when to retrieve** (preemptively vs. on-demand), **how to model** the interaction (single-shot, iterative, or agentic), and **the granularity of the augmentation** (at the prompt, token, or even parameter level). These axes map the paradigms from Naive to Advanced to Modular RAG. The shift represents an evolution from passively injecting text into prompts to actively orchestrating a system to reason over external world knowledge, reducing hallucination [25] and enabling robust performance on a wider array of intricate real-world tasks. As this evolution continues, it is likely that future systems will depend less on rigid pipelines and more on flexible, self-improving agents.

2.4 Delineation from Adjacent Techniques for Knowledge Integration

Given this clear evolutionary trajectory—from static pipelines to adaptive, agentic systems—a crucial question emerges: if RAG systems are becoming increasingly sophisticated at orchestrating external knowledge, where do they begin and end relative to other methods of knowledge injection, such as fine-tuning, continual learning, or model editing? The proliferation of techniques for injecting external knowledge into large language models necessitates a clear delineation of retrieval-augmented generation (RAG) from its adjacent—and often complementary—methodologies. While fine-tuning, continual learning, and model editing all pursue the shared goal of enhancing model capabilities with new or corrected information [41], they operate on fundamentally different principles, entailing distinct trade-offs across parameter efficiency, computational cost, operational flexibility, and knowledge freshness. Understanding these boundaries is essential, as the architectural choices described in the previous subsection are not merely technical details; they are direct expressions of this fundamental distinction between altering a model’s internal weights and augmenting its external context.

The most fundamental boundary lies between these two poles: RAG and fine-tuning. We characterize this as the distinction between "context-tuning" and "weights-tuning." Fine-tuning modifies a model’s parametric memory by updating its weights on task-specific or domain-specific data, thereby altering its internal knowledge representation [42]. This approach, from parameter-efficient adapters to full-scale updates, achieves deep task specialization but incurs substantial computational costs, such as gradient computation, hyperparameter tuning, and the risk of catastrophic forgetting. Critically, once fine-tuned, the model’s knowledge remains static until the next training run. RAG, in stark contrast, accomplishes adaptation without updating a single weight by injecting retrieved evidence directly into the context—and thus into the model’s attention mechanism—at inference time [12]. This distinction aligns perfectly with the evolution we traced: while fine-tuning represents a system that knows a fixed set of facts, RAG systems, particularly the agentic and modular paradigms, are designed to *know what they don’t know* and to seek out information dynamically. The emergence of semi-parametric systems that combine frozen parameters with a mutable, external index has demonstrated that powerful task adaptation can occur dynamically, leading to more stable and durable systems. Empirically, studies have shown that while fine-tuning can improve performance on low-frequency entities, RAG provides superior gains when fresh or specialized information is accessible [41]. The choice between the two is therefore a direct trade-off between static, parametric storage and dynamic, mutable context.

This distinction extends naturally to the domain of continual learning, which seeks to preserve and update a model’s parametric knowledge over a sequence of tasks without catastrophic forgetting. RAG offers a distinct approach to lifelong knowledge updates. Instead of struggling

to alter model weights, RAG systems re-adjust the external knowledge index and its embeddings to introduce new facts or developments. This "retrieval-centric" updating is significantly more feasible than continual learning for large-scale, real-world applications, because it bypasses the need for gradient computation entirely by re-structuring data [12]. This is particularly relevant to the agentic systems discussed earlier, where the LLM controller dynamically orchestrates the retrieval tool; the technology to update the tool's index is inherently modular and scalable. Furthermore, the emerging field of Retrieval-Enhanced Machine Learning (REML) suggests that this paradigm—gaining access to updated memory through retrieval—can be generalized beyond language to various domains, such as diverse time series or biology experiments [43]. In an enterprise setting, RAG's ability to integrate private data sources into its context through a continuously refreshed index bypasses both the high compute and staleness issues of continual fine-tuning [19, 44]. From a strict contrastive view, continual learning over-parameterizes to adapt but the system evolves; RAG displaces the knowledge base that doesn't need a modification of training.

Model editing, or surgical knowledge editing, offers a more invasive yet targeted alternative that intercedes at the level of the individual neuron or weight. While methods that "locate and edit" modify the representational geometry of specific facts (e.g., to change a single token), they face challenges in coupling. In the face of contradictory knowledge across thousands of entries, these methods demonstrate brittleness. RAG, by contrast, is non-invasive and offers verifiability of the source—alterations involve indexing or deleting a contiguous document in the external data store. This approach is inherently more scalable for broad and rapid knowledge updates (e.g., editing an entire website, novel scientific facts) and provides immediate *grounding* by allowing traceable citations that the edited weight methods cannot naturally support. This aligns with the gating mechanisms discussed in the previous subsection—self-aware pre-Retrieval filtering. The complementary nature is undeniable: model editing can, in specific use cases, modify the pretrained weights to directly handle the queries, rendering retrieval unnecessary for that piece of knowledge, effectively allowing a hybrid system to act as a confluence of the two—a param-memory alteration followed by RAG for everything else.

In synthesis, the choice between RAG, fine-tuning, continual learning, and editing is a strategic decision dictated by the task regime. The evolutionary arc of the previous section guides this choice: multimodal and agentic systems will likely recruit this full spectrum of tools. A simplified decision threshold can be defined as follows: if the task requires continuous, dynamic updates across acts of constantly shifting facts, the flexible RAG is preferred. If the task needs a deep, static comprehension of a large corpora (e.g., a domain specialization), then the retriever or generator components can be fine-tuned. Critically, these methodologies are not mutually exclusive but lie on opposite ends of a spectrum that is increasingly integrated. Far from being superseded, the evolutionary sequence is converging. For continuous improvement, RAG handles the ephemeral and open-ended updates, while a staged, retrained model can handle a more stable, skill-oriented specialization. This hybrid perspective—context-tuning for real-time, weights-tuning for core competencies, and editing for surgical precision—highlights a future where the taxonomy fades into a unified, multi-modal approach for knowledge integration, where an agentic controller decides which methodology fits the current prompt's demands.

3 The Retrieval-Centric Engine: Indexing, Query Processing, and Search Strategies

3.1 Corpus Construction, Indexing Structures, and Chunking Strategies

The foundational layer of any retrieval-augmented generation (RAG) system is the transformation of raw, unstructured corpora into a structured, queryable knowledge base. This process, encompassing corpus construction, indexing, and chunking, is the primary determinant of both

retrieval recall and the contextual density available to the downstream generative model. The design choices made at this stage fundamentally shape the system’s effectiveness, dictating not only what information is findable but also how coherent and sufficient that information is for grounding generation. In this subsection, we dissect the critical decisions at this layer, organizing our analysis around granularity-aware chunking, hierarchical and relational indexing, and data preprocessing, concluding with a discussion of their synergistic integration.

Central to corpus preparation is the granularity of text segmentation. The dominant approach, fixed-size windowing, segments text into contiguous slots of a predetermined token length. While simple and computationally efficient, this method often disrupts semantic and syntactic boundaries, leading to truncated or noisy representations that complicate downstream reasoning. In contrast, structure-aware chunking respects document layout by using sentences, paragraphs, or even sections as the atomic retrieval unit; this often yields better defined boundaries that are more coherent and self-contained [14]. The success of a RAG system, however, is not solely determined by coherence of the passages. The countervailing requirement that retrieved context must be concise enough to fit within an LLM’s context window necessitates a compromise. Recent findings show that in document-heavy sectors, such as financial reports, chunking by the structural elements of a document, rather than fixed-size windows, yields significantly better RAG performance. This element-based chunking, informed by document understanding models, automatically identifies the constituent parts of a document, leading to the "best chunk size without tuning" while preserving the structural information critical for context [14].

This logic can be extended by considering the capability of long-context LLMs: the LongRAG framework, for example, attempts to rebalance the 'heavy' retriever and 'light' reader in traditional RAG by grouping related documents into large 4K-token units. This approach moves the process of bringing the evidence to the generator and, by virtue of this, dramatically improves the retrieval segment’s efficiency, but it can run into the long context problem in which models fail to use information in the middle of their input and lose track of the relevant "needle" [45]. Such findings highlight a definitive trade-off between the retrievable unit size and the reader’s capacity to faithfully process it.

Beyond the choice of unit size, effective indexing concerns how these units are organized and referenced. Flat indexes where all chunks are stored in a single space for large corpora can lead to high latency and issues of noise. To mitigate this, a hierarchical and multi-level index architecture has been proposed, such as that used by GraphRAG. GraphRAG formalizes a workflow encompassing graph-based indexing, graph-guided retrieval, and graph-enhanced generation [35]. By building a knowledge graph from documents, the retrieval process can capture the structural knowledge and long-range dependencies that are missed by purely text-based embeddings. In an even more complete variation, HybridRAG, which integrates a traditional vector retrieval approach with a KG, has been shown to improve accuracy and answer generation in financial contexts, such as for earning call transcripts [46]. Other approaches use hierarchical summaries or clusterings, which first retrieves summaries at a high level and then drills down to more granular segments, reducing the amount of context sent to the eventual generator.

Furthermore, an internal mechanism to determine retrieval necessity is not always optimal; a query may be straightforward and the model may already have the answer in its parametric memory. In response to this, a decision, rather than retrieving external knowledge for every query, can improve the output. When Not to Trust Language Models shows the utility of not retrieving for high-popularity questions, and instead, after performing retrieval and ranking, a "skip" action is triggered to avoid adding no value [4]. This not only reduces inference costs but also prevents the addition of noise.

The practical implementation of these indexing strategies demands a sophisticated data preparation pipeline, especially when dealing with semi-structured sources like PDFs. The importance of parsed and well-constructed corpora is often understated, with the performance of a RAG system often hampered by noisy or incomplete parsing [47].

In synthesis, the research has clearly demonstrated that the one-size-fits-all universe of the corpus is outdated. The optimal chosen granularity and structural framework are evolving to be more specific to the domain and in using the capacity of the generation model. Emerging trends point towards an end-to-end co-design where the corpus construction is not a one-time process, but it is driven by the downstream task. A promising direction is to re-evaluate which pieces of content are actually indispensable for answering a query, considering not just relevance to the query but the value of the evidence in supporting the final generation. The external index becomes an integral component for enhancing self-consistency and attribute integrity, forming a more coherent and reliable system.

3.2 Query Processing, Reformulation, and Semantic Representation

The query processing unit serves as the critical interface between the user’s raw informational need and the structured retrieval machinery, transforming potentially ambiguous, under-specified, or lexically mismatched inputs into representations that can be effectively matched against the external corpus. Building directly upon the corpus construction and indexing strategies detailed in the previous subsection—which established the upper bound of what is findable—this stage determines whether the subsequent retrieval operation succeeds in surfacing the most salient contextual evidence, making it a pivotal determinant of overall system competence. Research consistently demonstrates that the way the query is formulated—both linguistically and semantically—directly conditions the recall of the index; inadequate raw queries, particularly those exhibiting lexical or semantic gaps with target documents, can fail to actuate the retriever’s matching mechanism, thus rendering even the most well-structured corpus ineffective. This makes the operationalization of query transformation, semantic encoding, and the mitigation of cross-modal divergence a first-order priority in the design of robust RAG systems.

A substantial body of work focuses on the rewriting, reformulation, and expansion of the query prior to embedding. Since user inquiries often exhibit intention ambiguity, explicit rewriting, such as several recently proposed methods, has proven crucial for aligning queries with the indexing scheme of the underlying knowledge space [8]. For example, the iterative generation-retrieval synergy captures that the language model can be actively engaged in generating intermediate outputs that serve as richer, more informed query contexts for the subsequent retrieval step, thus breaking down the original query into stages that resolve increasingly complex and multi-step information needs [29]. Such mechanisms induce a dynamic, iterative process that refines the query based on intermediate generations—a concept that resonates with the feedback loops between generation and retrieval that are central to the joint optimization discussed in the following subsection. To this end, explicit reformulation, as seen in the **MVRAG** framework, decomposes multi-faceted queries into multiple sub-queries, each conditioning context from a unique viewpoint, thereby resolving complex intent that a static input cannot capture [29]. The learning of these reformulation patterns can be automated: models are trained as rewriters using synthetic queries and pseudo-relevance signals from the retrieval system, injecting an adaptive dimension and improving the fidelity of the indexing lookup. One can think of this transformative process as shifting from a raw user artifact to a retrieval-optimized query that is maximally compatible with the search mechanism that follows.

Complementing the linguistic precision of reformulation is the central machinery of semantic representation learning, which provides the vector index with dense, *representational tokens*. This semantic encoding directly dictates the quality of the match mechanisms that search the index prepared by the previous stage. Query and document sides are mapped into a latent embedding space using bi-encoders, which supply the scalable vector index with dense embeddings. The choice between lightweight bi-encoders and high-fidelity cross-encoders is dictated by a fundamental efficiency-fidelity trade-off: bi-encoders are beneficial for real-time, one-to-many matching, while cross-encoders produce highly accurate relevance signals but lack scalability to large indexes. Regardless of the encoder architecture chosen, a persistent challenge

lies in the representation of long documents and query contexts, where we see an established degradation of embeddings due to positional bias and information decay in long inputs [48]. Moreover, the discrete constraints underlying LLM inference require the use of fixed hidden states; while memory-editing techniques such as **MEMIT** can update models with new knowledge [31], embedding spaces pose the challenge of aligning the raw tokens from different modalities or dialects into a shared, target subspace—both of which are essential to enhancing the effectiveness of the search mechanisms that depend on these representations.

The distributional drift between query and document spaces becomes even more pronounced in the multilingual domain, where the mapping of diverse dialects and low-resource languages into a unified dense space is critical for fair and effective retrieval. More advanced training regimes, such as contrastive instruction-tuned embedding models, suggest that heavy fine-tuning with high-quality data bolsters raw retrieval performance and cohesion of query understanding [48]. Yet, the practical interplay of dense, sparse, and learned representations requires not simply the selection of a single vector index but the deployment of hybrid pathways that bridge generated text and queries into a sparse vocabulary as well. The path forward for semantically robust query-side representation demands not more monolithic encoders but carefully navigable hybrid models that fuse interaction signals. This never is more present than in the hybrid search models that fuse sparse and dense signals, which will be the subject of the following subsection, and which integrate such query-level transformations into a search process that balances the linguistic precision of term matching with the semantic generalization learned by embeddings [40]. By translating the raw embedding space into a representation that works in tandem with diverse search paradigms, the query processing unit constructs the very bridge between the archived corpus and the retrieval mechanisms that will be detailed next.

3.3 Search Mechanisms: From Sparse Lexical to Dense Semantic and Generative Retrieval

The search mechanism constitutes the computational heart of the retrieval subsystem, determining how indexed representations are actually matched against the query. The evolution of these mechanisms can be characterized as a progression through three fundamental paradigms: sparse lexical methods rooted in term matching, dense semantic approaches based on distributed representations, and the emerging frontier of generative retrieval that dispenses with explicit search structures altogether. Each paradigm embodies a distinct trade-off between linguistic precision, semantic generalization, and computational efficiency, with practical systems increasingly relying on their integration.

Sparse lexical retrieval, exemplified by BM25, operates on lexical overlap. These models are characterized by their computational efficiency, robustness, and interpretability, but suffer from the lexical gap problem, failing to capture synonymy and paraphrase relationships. As a strong baseline, BM25 has proven remarkably resilient and remains the de facto standard in numerous information retrieval systems. However, the fundamental limitation of sparse methods in capturing semantic similarity has been a central motivation for the shift towards dense semantic retrieval, which maps queries and documents into a shared low-dimensional space where relevance is defined by geometric proximity. The success of this paradigm hinges on the training of bi-encoder architectures [8], where effective representation learning aligns query embeddings with structurally relevant and related document embeddings, effectively overcoming the lexical gap.

A notable point in this evolution is the observation that the dense retrieval component introduces an "embedding gap," where standard inner-product retrieval may not fully capture the nuanced details of information. To further bridge the gap between sparse and dense paradigms, learned sparse models have been developed. These models incorporate a mechanism to enrich the sparse representations, by injecting contextualized information from pretrained language models into the sparse vector spaces. These learned sparse methods aim to combine the precision

and interpretability of lexical signals, with the semantic capabilities of dense models.

The practical deployment of dense retrieval at scale relies on the use of Approximate Nearest Neighbor (ANN) search algorithms. While exhaustive search over a large corpus ensures complete fidelity, it is computationally prohibitive for large-scale indexes. Methods such as Hierarchical Navigable Small Worlds (HNSW) have become standard in this domain, which build a multi-layered graph structure to facilitate efficient traversal. These structures allow the system to effectively trade recall for significantly improved query latency, the performance of which is controlled by algorithmic parameters that modulate the speed-memory-recall trade-off. Although these techniques address dense search efficiently, knowledge of the role of the retriever is critical. Recent work emphasized that enriching the quality of the query and its derived embeddings is crucial. Approaches such as *reranker* that exploit strong language understanding capabilities to elicit more relevant contexts have proven significantly more effective than simple model scaling [49]. This has been validated by approaches that teach search models to represent queries and documents more effectively from instruction tuning for dense retrieval [50], while in practical RAG systems, a retriever augmented with highly capable instructions has been shown to increase quality and knowledge adaptation [3]: the retrieval module may itself rely on the generative model to produce a search query from a natural language input, integrating the search. FreshLLMs highlight the importance of such query reformulation for dynamic QA tasks [51].

A more recent paradigm is *generative retrieval*, which seeks to move away from a complex, multi-step process of indexing and search. Instead of embedding all documents and relying on a search engine (like vector databases), a foundation model-based retrieval system is adapted to directly decode the document identifier from an input query in a sequence-to-sequence manner. This new search paradigm removes the need for a separate index and solves the trade-off between the representation and search, and creates a new avenue for jointly optimizing and retrieving with the language model’s semantic capabilities. However, it often faces performance drops and scalability issues when the document collection grows large. The choice of the search mechanism ultimately underscores a central trade-off: the nature of the queries (lexical vs. semantic) vs. the scale of the corpus. While lexical search remains a stalwart for exact and rare matching, the increased number of *retrieval-augmented generation* tasks necessitates mechanisms that are often more sensitive to how LLMs consume the retrieved output, and the shifted emphasis to a purely *contextual generation* pushes the envelope of what the meaning of "search" truly is.

In practice, expert systems are converging on hybrid search that fuse the outputs of multiple signals—sparse and dense—to improve coverage and robustness. Rather than relying on only a single vector index, these search ensembles combine the precision, recall, and linguistic robustness of independent algorithms; By harnessing different retrieval output, we can provide a richer but also more balanced context to the LLM, thus effectively easing the OCR, or non-semantic term-mismatch issue. The potential of leveraging the generative output signal as a further feedback loop has been shown to enhance retrieval performance across the board, where the model can be trained to refine the query based on intermediate generation results [29]. This suggests that the evolution of search will become more deeply intertwined with the capabilities of the model itself, using a dynamic feedback loop that challenges the boundary between retrieval and generation, and moves towards a future where the main role of the search mechanism is not only to satisfy retrieval but also to actively curate and steer evidence.

3.4 Retriever-Generator Alignment and Co-Optimization Strategies

The efficacy of a retrieval-augmented generation (RAG) system hinges not merely on the individual competence of its retriever and generator, but critically on their synergistic alignment. The preceding discussion traced how the search mechanism itself has evolved from static lexical matching to sophisticated dense and generative paradigms, yet even the most advanced retriever can falter if its notion of relevance diverges from what the generator actually needs. A retriever

optimized for lexical or semantic relevance, as measured by traditional information retrieval metrics, may surface passages that are ill-suited for the downstream generative task, leaving a "preference gap" between what is human-friendly and what is LLM-friendly [52]. Consequently, this subsection examines the training paradigms and co-optimization strategies designed to harmonize these components, transforming the retriever from a static search engine into a dynamic partner attuned to the generator’s needs—effectively completing the loop that connects the search mechanisms previously discussed with the post-retrieval processing that follows. We dissect the spectrum of tactics, ranging from supervised fine-tuning of the retriever to end-to-end, policy-based joint optimization, detailing their underlying mechanisms, benefits, and inherent limitations.

A foundational approach to aligning the retriever involves supervised and contrastive learning objectives that move beyond mere relevance to utility. Rather than simply preserving the geometric proximity established by the dense retrieval paradigm, the objective becomes training an embedding space where the top-ranked results for a given query are not just relevant, but also maximally and useful for the generator to synthesize a correct answer. This often involves fine-tuning the retriever with contrastive loss functions and hard negative mining to sharpen the boundary between useful and plausible-but-detrimental context. To obtain the required gold labels for this supervised fine-tuning, a powerful strategy involves leveraging the generator itself as a labeler. By sampling the generator to produce documents based on retrieved passages and using direct downstream labels—a technique formalized in approaches like eRAG, where each document is individually employed to generate an output scored against the ground truth—it becomes possible to create richer, loss-informed relevance labels [53]. Such fine-tuning with generation-informed labels helps decentralize the knowledge representation in the pre-trained model, ensuring the retriever captures the semantic cues that actually matter for the generative objective.

However, the most fundamental obstacle remains the non-differentiable nature of the discrete document selection process, which blocks direct backpropagation from the generator through the retrieval stage. To overcome this, more advanced co-optimization strategies attempt to pipeline the entire system end-to-end. A preliminary solution is to treat the retrieval selection as a stochastic process and employ a straight-through Gumbel-top-k estimator—as seen in Stochastic RAG—to provide a differentiable approximation for sampling without replacement, thereby allowing gradient flow based on expected utility maximization [54]. In scenarios where the generator is a massive, potentially-to-be-kept-frozen model, this end-to-end compatibility can instead be achieved through a collaboration of supervised and reinforcement learning to train a 'bridge model' that optimizes the connection, re-ranking or utilising retrieved context to better fit the LLM’s preference [52]. Alternatively, reinforcement learning (RL) offers a framework where the retriever is treated as a policy, with the final answer quality as a reward. Recent work [44] leverages multiple agents in a multi-agent Reinforcement Learning setup to make traversals across structured partitions of the retrieval database to optimize the generation results, demonstrating a compelling path for explicit and end-to-end utility maximization.

Notably, the interplay between retrieval and generation is not confined to a single, upstream alignment stage; it can be designed as an ongoing dialogue that unfolds across the entire generation process. This iterative synergy is apparent in approaches like iterated Ret-Gen, where the model’s emerging output—the partial response that reveals what information might still be needed—is subsequently used as a more informative query for continuing retrieval [29]. Here, the generator’s output acts as a powerful query-side distillation, a feedback loop that improves the relevance of non-parametric information for each generative step. This aligns well with methods that see the LLM itself as a tool for retrieval alignment, using its language understanding as a "retrieval interpreter" to bridge the gap between static retrieval and dynamic reasoning, directly complementing the query reformulation techniques discussed in the previous subsection. Further high-performing systems formalize this as active retrieval, where the LM’s

own predictive confidence determines whether new retrieval actions are triggered, making the alignment an adaptive, state-dependent process [27].

Looking forward, the field is converging toward dynamic alignment that moves away from static, fixed retrievers. We are observing a shift from merely relying on a pre-trained corpus to 'tuning' the retrieval process itself on a per-query basis. As the need for grounding in factual, verifiable content grows, the optimization is expected to shift beyond simply maximizing predictive accuracy. The rise of LLM-labelers (LLMs-as-Labelers) in the RAG pipeline—and the ability to perform lightweight device alignment via Low-Rank Adaptation (LoRA) on the retriever to improve generation utility without full retraining, signal a promising direction. Future systems increasingly based on pure utility and outcome scores, compiled by reinforced self-fueled feedback loops we feel, promise to create RAG systems that are truly co-optimized, addressing hallucinations by ensuring the retriever and generator are optimized as interconnected pieces of an end-to-end reasoning system. This co-optimization, with its adaptive retrieval triggers and output-aware relevance judgment, directly supports the post-retrieval strategies explored next, providing these new generation-ware and generation-aware mechanisms with the right context to enable efficient, selective context compression and confidence-aware routing for when to consult external knowledge.

3.5 Context-Based Compression and Efficient Retrieval Context Utilization

The retrieval stage of a RAG pipeline is fundamentally constrained by the finite context window of the generative model and the computational cost associated with processing long input sequences. As the volume of retrieved evidence grows—often scaled to improve recall—the risk of introducing noise, diluting the model’s attention, and inflating latency and memory usage becomes a critical bottleneck [55]. Traditional approaches that indiscriminately concatenate all top- k passages along with the query can overwhelm the model’s capacity to focus on the most relevant information—a phenomenon underscored by the "lost-in-the-middle" effect, where LLMs not only exhibit degraded performance with longer contexts but also demonstrate a positional bias, underutilizing critical evidence located in the middle of the prompt. It is within this context that post-retrieval processing, specifically compression and efficient context utilization, emerges as a necessary and distinct stage in the RAG architecture, transforming a raw retrieval set into an optimal, information-dense, and cost-effective input for generation.

The first family of techniques focuses on context window-aware compression and filtering, which aims to discard non-essential information before it reaches the generator. As the primary objective is to reduce the token complexity while retaining the task-relevant semantics, these methods range from simple extractive filtering to more sophisticated model-based approaches. A prominent class of methods utilizes a lightweight evaluator to assess the quality of the entire retrieved set, identifying critical documents and filtering out those considered redundant, irrelevant, or outright harmful. If the initial retrieval pass is deemed suboptimal, such evaluators can trigger a "decompose-and-recompose" process, where complex documents are segmented and the importance of their constituent parts is assessed, integrating only the pertinent key sentences and terms into a new, refined context [55]. This granularity-aware approach is crucial, as recent empirical findings suggest that system performance can be surprisingly enhanced by the inclusion of some irrelevant documents, which might provide incidental context or serve as a contrasting set [55], emphasizing that filtering is not about the outright removal but rather the optimal weighting and selection of evidence. Furthermore, filtering is often implicitly combined with evidence correction, where evaluators judge the relevance of each retrieved document, removing unhelpful ones to improve the next generation [30]. This iterative, evaluator-driven filtering ensures a dynamic and quality-controlled context, rather than a static top- k selection.

Beyond simple filtering, a second class of approaches develops token reduction and abstraction through hierarchical and clustering techniques. These methods avoid the necessity of passing an excessive number of raw chunks to the model by first constructing a hierarchical index or

clustering the retrieved documents. In this scheme, the context is either formed by selecting representative chunks or by fetching a high-level summary of a document cluster before drilling down into the detailed knowledge. For instance, an agentic system could navigate database partitions at a higher conceptual level, retrieving and fusing information from only the most relevant partitions (M-RAG), dramatically mitigating the computational cost of subsequent stages and reducing the noise that a flat, exhaustive retrieval set would introduce [56]. A token-level elimination, building on the compression capabilities of the k -nearest neighbor perspective, is to output a cache of keys and values for each candidate which avoids recomputation [57]. However, it's crucial to be able to distinguish between extracting sources and synthesizing a response in a way that the token reduction does not come at the cost of the retained context's coherency and logical flow. Compression techniques based on the Abstract Meaning Representation (AMR) exemplify this balance, distilling the raw documents into a compact set of crucial semantic concepts; this compact concept-based context explicitly constrains the LLM to focus solely on vital information and has shown to be increasingly effective as the number of supporting documents is increased, filtering out interference and improving answer accuracy [58].

Finally, the dynamic and latency-aware retrieval sampling indicate the need to shift the paradigm from a static retrieve-then-generate decode. This mechanism acknowledges that, for many queries, retrieval is not necessary at all; the LLM's parametric memory is sufficient to generate the answer, and an unnecessary retrieval step only adds to the latency and computational overhead, risking the injection of noisy or irrelevant context that would degrade the output [38]. Such approaches aim to determine when to retrieve, how many documents to retrieve, and how often, so as to optimize the efficiency of the RAG pipeline. The framework known as Adaptive-RAG is a foundational piece in this design; it learns a classifier that predicts the complexity of the input query, routing the pipeline to either no-retrieval, single-step retrieval, or multi-step retrieval paths. For easy queries, the system short-circuits the retrieval process, saving computational overhead; for complex ones, it performs a more expensive retrieval/iteration, dynamically balancing accuracy and latency [38]. In this perception of latency, Retrieval-Augmented Generation (RAG) also suffers from an efficiency bottleneck: the generative inference time becomes a major issue due to the long sequence length required by injecting knowledge. However, caching the intermediate states of the knowledge in a multilevel cache, as in RAGCache, can reduce the time to first token and improve throughput [57]. Furthermore, as we scale the model size, the cost of the forward pass on the context does not increase proportionally, but the fixed cost becomes more manageable.

Looking ahead, the convergence of these three pillars—intelligent filtering, semantic compression, and adaptive retrieval determination—will define the efficiency frontier of RAG. A unified framework, for example, a dynamic posting retrieval model that might refuse to retrieve when unnecessary, uses a low-cost hierarchical retriever to select candidate clusters, and then adapts a generative compressor to distill a coherent context, will become necessary for large-scale use. The challenge is to ensure that the "lost in the middle" phenomenon is mitigated in a way that is both semantic and robust; the interplay between the retrieval noise and the LM's ability to reliably handle it will shape the future of RAG systems. As the field moves forward, leveraging the cross-modal nature of some of these systems (e.g., in the medical domain) and verifying the robustness of the RAG pipeline itself is underexplored [59, 60].

4 Synergistic Architectures and Fusion: Orchestrating Retrieval and Generation

4.1 Context Integration Strategies and the Input Pipeline

The manner in which retrieved evidence is integrated into the generative model constitutes a foundational design axis in retrieval-augmented generation, as the architectural choices governing

input representation and document placement directly modulate the model’s capacity to leverage external knowledge for reasoning. A spectrum of integration strategies has emerged, ranging from static, monolithic pipelines to adaptive mechanisms that dynamically modulate the retrieval-generation interface based on query complexity and intermediate model states.

The most prevalent paradigm is static concatenation, epitomized by the retrieve-then-read framework where retrieved passages are prepended or appended to the original query and fed as a unified token sequence into a frozen language model. This approach offers architectural simplicity and computational economy, as it does not require any modification to the model’s parameters [8]. However, a critical empirical insight reveals that the *positional placement* of retrieved passages fundamentally impacts the model’s ability to utilize them. The "lost in the middle" phenomenon, documented in [28], demonstrates that LLM performance degrades substantially when relevant information is positioned in the middle of a long context, whereas performance is highest for information at the beginning or end of the sequence. This finding indicates that a naive ordering of "relevant" documents can inadvertently exacerbate the model’s positional biases. This phenomenon is particularly acute and quantified in RAG systems [59], and is further compounded by the finding that a model’s in-context recall capability is highly prompt-dependent and may be compromised by biases present in its training data [61]. Additionally, RAG systems with long retrieved contexts need to be robust against injecting irrelevant information when implementing long-context integration [28]. Such observations motivate the need to move beyond fixed concatenation toward adaptive or re-ordering mechanisms within the input pipeline.

To mitigate these positional and computational inefficiencies, more sophisticated and dynamic integration strategies have emerged. A key development is the notion of dynamic contextual triggering, also conceptualized as selective augmentation. This paradigm recognizes that not all queries require external knowledge, as models can robustly handle questions when parametric memory is strong or when internal knowledge is preferable to RAG, prompting the need to verify the parametric memory’s effectiveness first [4]. The approach often employs a gating mechanism to assess the model’s level of confidence in a query *before* retrieval—such as assessing confidence gaps in it or in response validation loops—to decide whether to query an external store or rely on parametric generation, thereby avoiding unnecessary computational cost as well as avoiding the injection of potentially distracting or irrelevant context [62, 63]. Additionally, frameworks like LLM-Augmenter iterate over model outputs, using retrieval to verify and when necessary, correct intermediate responses, thereby aligning with retrieval generation that is feedback-informed and iterative [63]. Beyond that, even with static concatenation, prompting strategies can be optimized—such as placing retrieved content before or after the instruction text itself—to achieve better downstream performance without architectural changes [6, 61].

Dynamic integration strategies also encompass improving retrieval effectiveness by rewriting the original query on-the-fly. Within the input pipeline, the raw user query may present a semantic gap with underlying corpus representation. Strategies such as query generation blending utilize multiple perspectives of a query—for instance, merging the original, paraphrased, or instruction-like variants—to perform multi-view retrieval so that a synthesized prompt is based on a more comprehensive context than what a single query could retrieve [64, 65]. Additionally, iterative frameworks such as Iter-RetGen interleave retrieval and generation to improve relevance using the model’s intermediate output as a search query, therefore better tailoring the context fed into the LLM [29]. In structured technical corpora, a prior accessible stage (e.g., element-type chunking) also influences the prompt composition space, demonstrating that the structure of input documents supports the effectiveness of prompt-based injection [14].

However, a key current limitation of the numerous static and dynamic strategies is batching retrieval and generation as separate, sequential stages that disregard runtime states and system-level latencies. Pipeline parallelism has been introduced to align retrieval time with generation blocks, thereby significantly reducing the end-to-end system latency observed in adaptive retrieval pipelines [66]. From a holistic perspective, the described context integration strategies

are not mutually exclusive and often need to be tightly coordinated; the choice between static and dynamic injection is primarily a trade-off between efficiency and query-awareness. Static concatenation with well-chosen ordering and preprocessing techniques cannot adapt to differing confidence levels. Dynamic triggering, conversely, is a robustly safer choice in case of risky retrieval, though it incurs additional inference latency and engineering overhead due to gate/confidence computation. Future research directions should consider the intersection of context integration with continual model learning, where memory-augmented enriches the context stores [9] or new memory mechanisms [67, 68] can be used to sequence integration in training or to reduce reliance on longer static prompts by providing the system with dynamic but relevant memory in a latent form. Ultimately, the integration strategy should be thought of less as a fixed, singular module, and more as a core component subject to continuous co-optimization with the retriever itself, to achieve grounded and faithful generation.

4.2 Fusion-in-Decoder and Cross-Attention Architectures for Multi-Document Synthesis

The efficient synthesis of information distributed across multiple retrieved documents constitutes a central challenge in retrieval-augmented generation. While simple concatenation of evidence into the prompt is computationally straightforward, it often proves suboptimal: the model must learn to attend to a long, potentially redundant, and variably relevant sequence, leading to quality degradation and increased computational cost. Fusion-in-Decoder (FiD) architectures address this limitation by introducing a dedicated architectural treatment of retrieved evidence. Rather than concatenating all documents into a single input sequence, Encoder-Decoder models such as T5 or BART are employed to encode each retrieved document independently and in parallel. The generative decoder then performs cross-attention over the concatenated set of encoded document representations, thereby fusing information from all sources during decoding. Critically, the answer is generated by attending over all encoded documents jointly, forcing the model to integrate and synthesize evidence during the generation step, a mechanism that more closely aligns with the goal of multi-document synthesis.

The foundational designs of RAG models [8] established the paradigm of jointly training a dense retriever with a pretrained sequence-to-sequence generator, conditioning on a fixed number of retrieved passages. The core advantage of this design is its tractability: at inference time, the number of documents to be fused, k , can be scaled to hundreds or thousands without a prohibitive increase in compute, as encoding is performed once per document and only the final decoder operates over the concatenated set of representations. This scalability is notable because it directly counteracts the context-utilization limitations discussed earlier. While static concatenation and adaptive gating in the input pipeline operate on a single flattened context, FiD’s distributed encoding ensures that the decoder can attend to each evidence source with equal representational fidelity, irrespective of the total volume. Indeed, the response quality continues to improve with the inclusion of more evidence, a behavior that contrasts with the "lost-in-the-middle" phenomenon observed in decoder-only LLMs that receive the entire context in a flattened sequence, which often only effectively utilize the initial portion of the provided context [37]. This positional robustness in FiD suggests that its explicit cross-attention mechanism is fundamentally better suited for scalable evidence fusion than the sequential attention of its decoder-only counterparts.

This architectural divergence has deep implications for the RAG design space presented earlier—the choice between decoder-only input modulation and encoder-decoder fusion is not merely about parameter placement but about control over the evidence-attention interface. Whereas previous subsections considered how context is structured at the input level (e.g., ordering, selective augmentation, multi-round rewrites), FiD relocates the fusion bottleneck from the input layer to the cross-attention layer of the decoder. The comparative analysis reveals that decoder-only architectures, despite their widespread adoption, tend to under-utilize retrieved

contexts, relying more heavily on parametric memory when the input grows long [37, 69]. In contrast, the encoder-decoder structure provides a deterministic encoding of the evidence that is independent of generation state, ensuring that every retrieved document is attended to through the cross-attention mechanism, regardless of its position in the sequence. This property makes FiD models particularly robust to the ordering of input documents, a significant mitigation of the positional biases identified earlier.

However, this robustness comes at a cost. The independent encoding of documents prevents the encoding of inter-document interaction until the decoder, which can limit the model’s ability to perform higher-order reasoning that requires comparing or contrasting evidence across documents. Furthermore, the full set of encoded representations must be stored in memory for the decoder to attend during the entire generation process, resulting in significant memory overhead that scales linearly with the number of retrieved documents. This memory burden becomes a practical bottleneck for many real-world RAG systems, especially in latency-sensitive deployment scenarios, where the complexity of adaptive retrieval planning (detailed in later sections) might outweigh the benefits of unconditional parallel encoding. In summary, while FiD offers a robust and high-ceiling architecture for evidence fusion, its computational proportionality and limited inter-document interaction highlight the trade-off of architectural modularity.

An interesting but less-explored direction is the comparison of FiD with decoder-only models in the context of downstream RAG performance. Recent investigations into the effective context utilization of language models show that decoder-only architectures tend to utilize a smaller proportion of the full context when compared to their encoder-decoder counterparts [37, 70]. While decoder-only models have demonstrated superior scalability in general-purpose generative tasks, including question-answering, their performance in RAG settings is often bounded by the inclusion of noisy and irrelevant documents that inject confusion into the context window—a challenge that the earlier discussions of dynamic gating and adaptive retrieval aim to mitigate through careful selection. In this regard, FiD architecture provides a more controlled information flow, where the retriever’s output is explicitly separated from the source sequence and given dedicated vector representations that are inherently fine for cross-attention. However, the raw capacity and generative diversity of large decoder-only models are not to be smaller; the combination of FiD’s architectural strengths with decoder-only generative flexibility remains a potential avenue for hybrid systems. Whether by learning to compress the encoded evidence in a latent space or by integrating hierarchical fusion layers that selectively route information from different evidence groups, the future is direction to balance the robustness of FiD-style architectures with the scalability and familiarity of the dominant decoder-only paradigm.

In bridging these considerations, this analysis underscores that architectural innovation in evidence fusion—from encoder-decoder cross-attention to prospective decoder-integrated fusion mechanisms—is as critical as the retrieval and planning components discussed surrounding it. The efficiency and synthesis of multi-document grounding depend not only on which documents are fetched and how the flow is triggered but also on the internal structure of the generative model itself. These insights position encoder-decoder fusion not as a singular model but as a type of module to be optimized alongside and even within the decoder-only frameworks, moving from a fixed architectural dichotomy toward more adaptive mixture-of-experts of heterogeneous mechanisms and hybrid integration.

4.3 Adaptive Retrieval Planning and Dynamic Routing

The progression from static, single-shot retrieval pipelines to adaptive, plan-driven systems represents a fundamental shift in how retrieval-augmented generation (RAG) architectures balance efficacy against computational economy. Rather than treating every query as requiring the same degree of external intervention, adaptive retrieval planning introduces an explicit decision-making layer that governs *whether* to retrieve, *when* to retrieve within the generation process, and *how many* retrieval actions are warranted. This capability is predicated on the

observation that not all queries demand uniform augmentation; simple factual queries may be answered from the model’s parametric memory with high confidence, whereas complex, multi-hop queries require iterative grounding [4]. The central design question, therefore, becomes how to allocate retrieval resources commensurate with query complexity while maintaining (or improving) answer quality. Early evidence suggests that even LLMs with robust parametric memory show unwavering confidence in their responses, yet retrieval augmentation can significantly enhance their perception of knowledge boundaries—pointing to the need for a principled, rather than universal, approach to when external evidence is consulted [71].

A prominent family of adaptive systems addresses this challenge through explicit complexity estimation and routing. Adaptive-RAG, for instance, introduces a query complexity classifier—a small language model trained to predict whether a user query is straightforward, multi-hop, or requires no retrieval at all—and dynamically routes the query to the appropriate pipeline variant: no-retrieval, single-step retrieval, or iterative retrieval [38]. This classifier is trained on labels automatically inferred from the predicted outcomes of models and the inherent inductive biases present in the datasets, thus enabling a balanced strategy that improves both efficiency and accuracy compared to uniform-approach baselines. Such routing is a stark contrast to unconditioned retrievers, which indiscriminately inject potentially irrelevant or redundant documents into the prompt, often amplifying computational costs without commensurate performance gains. Complementary and conceptually parallel to Adaptive-RAG, Self-Route uses LLM self-reflection about whether a query can be answered by the long-context model itself before deciding to route it to either RAG or long-context LLM [37]. This demonstrates the broader value of dynamically alternating between parametric and non-parametric knowledge of the current context, reducing computation cost while maintaining comparable accuracy with respect to the long-context approach. These complexity-based and uncertainty-based routes, while operating on different feature signals (external classifier versus model self-reflection), both thread the problem of multi-strategy routing—ensuring the augmentation is tailored to the needs of each query.

Beyond the binary decision of whether to integrate external knowledge, a second critical axis is *when* to initiate retrieval within the generation trajectory. Systems in this vein move retrieval from pre-query to being interleaved with generation steps, where a model can benefit from its own partially generated output to refine queries. The Iter-RetGen approach is archetypal in this respect: it iteratively introduces a generation step between retrievals, using the earlier generation as richer context to lead to more informed retrievals [29]. Such a method processes each retrieval round’s knowledge as a whole and does not structurally constrain generation, which gives it substantial flexibility in multi-hop QA and fact verification. This method relates to the broader class of dynamic retriever decisions, where a system infers whether to retrieve more sources after evaluating the quality of the currently retrieved passages. An example of this is the use of *critique models* or *reflection tokens* that can assess whether a given generation already addresses the query’s information needs or if further retrieval is likely to add value. Self-RAG is particularly instructive here: a single LM is trained to generate reflection tokens, allowing the model to decide how retrieval is required, to critique the relevance of retrieved passages, and to validate the factuality of the final answer [30]. The integration of self-reflection into inference enables the model to adapt to diverse requirements without increasing inference costs beyond what is necessary. Such self-reflective mechanisms reduce the need for an externally trained mapper and foster a sense of introspection within the generation engine itself, serving as a self-aware function to incorporate external evidence.

A further dimension of planning involves orchestrating *multiple*, potentially heterogeneous retrieval actions. This class of systems embodies the "agentic" turn in RAG, where a language model acts as a controller to query databases or knowledge tools and assess the output. In these settings, RL-based policies are particularly functional for learning an optimal ordering of tool use. Bridging the preferences between retrievers and LLMs is a common challenge—retrieval

models curated with human-friendly relevance signals do not always align with what LLMs require. Addressing this gap, a bridge model is trained by chaining supervised and reinforcement learning [52]. The role of such a bridge is visible in the broader adaptive interaction ecosystem, where RAG involves multiple disciplines such as time-series analysis or database navigation. By introducing such explicit policy learning, the retrieval becomes task-oriented, able to adapt the number of retrievals and stop early once the dead-end is clear, hence preserving computational budgets.

A crucial insight across these adaptive approaches is that the decision of *when not to retrieve* is equally critical to assert and execute. However, common implementations often suffer from a simplistic one-way separate attempt: while a classifier can make a single decision, multi-step reasoning often requires iterative routing across several intermediate steps. Interestingly, removing irrelevant documents from the query context has been shown to have unexpected positive effects, in some cases, increasing performance [55]. Through the appropriate selection of retrieved context, adaptive planning is often handling noise, not just relevant context. This behavior is typically best placed alongside policy learning, which implicitly learns to structure retrieval over sequential rounds.

Looking forward, adaptive retrieval planning is set to be shaped by two factors: multi-level routing architectures and continuous learning. First, research should target the difficulty in dynamically identifying "when to stop" retrieval—termination criteria that are task-dependent, and especially stringent in domains like legal or medical. Second, experimentation with trained classifiers (e.g., Adaptive-RAG) updated at deployment time could yield advances in robustness, provided the classifier is continuously trained based on downstream evaluation signals [38]. Third, current systems could benefit from extending retrieval planning to include covariates such as relevance versus, allowing for longer horizons and more autonomous behavior in real-scale deployments.

4.4 Recursive, Iterative, and Multi-Hop Retrieval for Complex Reasoning

The progression from single-shot retrieve-then-read pipelines to recursive and iterative frameworks marks a fundamental shift in the architecture of retrieval-augmented generation (RAG), enabling systems to engage in complex, multi-step reasoning. This shift builds upon the adaptive planning principles discussed previously—where systems decide *whether* and *when* to retrieve—by operationalizing those decisions into a dynamic, multi-turn trajectory. While adaptive routing determines the *type* of retrieval pipeline to invoke (e.g., a single-step vs. a multi-step variant), iterative frameworks define the *execution loop* within that pipeline, enabling the system to accumulate and synthesize evidence across a sequence of interdependent retrieval-generation cycles. These advanced paradigms recognize that for knowledge-intensive queries—particularly those requiring multi-hop inference—a single retrieval round is insufficient; the information required to answer a question must be accumulated and synthesized across a sequence of interdependent retrieval-generation cycles, a realization that directly motivates the careful, multi-stage handling of evidence discussed in the next section on post-retrieval processing.

A canonical family of methods in this space capitalizes on the reciprocal relationship between generation and retrieval. The Iter-RetGen framework [29] demonstrates that an LLM’s intermediate output, generated based on initial retrieval, offers a rich, contextualized reflection of the information still needed, which can be leveraged to construct a more relevant query for the next retrieval round. This iterative loop allows the system to process retrieved knowledge as a whole, preserving the flexibility of generation while progressively converging on a well-grounded answer. This synergy is further refined in strategies that interleave retrieval with explicit chain-of-thought reasoning steps. The IRCOT approach [72] addresses the fundamental question of necessary information, which is a function of what has already been derived. By alternating a single reasoning sentence with a retrieval step, it ensures that retrieval is anchored to the evolving inferential context, rather than the initial query alone, yielding substantial

improvements in both retrieval and final QA performance on multi-hop benchmarks. This interleaving is essential, as it mitigates the error cascade that often occurs when irrelevant context is retrieved and integrated into the reasoning path.

Building upon the adaptive *whether* and *when* decisions introduced in the previous subsection, iterative RAG systems often incorporate adaptive control over the *execution* and *length* of the retrieval loop. In this context, the retrieval trigger in active retrieval augmented generation [27] is employed as a step to predict a future sentence and identify which information is missing—transforming retrieval from a static appendage to a proactive, forward-looking process. The Forward-Looking framework actively retrieves based on the future content’s prediction, using low-confidence tokens as a signal to decide *when* to interleave retrieval. More generally, the question of retrieval necessity can be determined dynamically by a complexity classifier that routes queries to different strategies, as demonstrated by the Adaptive-RAG framework [38]. By learning to route queries from no-retrieval to multi-step retrieval, this approach quantifies a key trade-off in iterative systems: the balance between computational overhead and the value of additional information, ensuring that iterative retrieval is not a burden but a targeted intervention.

The orchestration of multiple specialized modules for knowledge exploitation represents another critical direction, moving beyond linear iteration to more structured decompositions. For example, in the process of retrieval-augmented iterative self-feedback (RA-ISF) [73], a teacher model decomposes a complex task into a set of sub-problems, facilitates a self-feedback mechanism for a model to extract key information, and iteratively progresses through these to keep the reasoning trajectory robust to false or irrelevant findings. This module-level division of labor proves particularly robust for maintaining logical consistency. More generally, self-feedback frames, which turn generation into a self-critique loop, are further exemplified by Self-Reflective Retrieval-Augmented Generation (Self-RAG) [30]. Here, LLMs learn to output a sequence of “reflection” tokens—such as Retrieve, IsRel, IsSup—which control the transfer of retrieval to the model and critique the retrieved context against the query’s relevance. These adaptable artifacts give the model the ability to override noisy retrieved context and selectively incorporate information when needed, which is particularly relevant not just for multi-hop reasoning but for fact-checking and long-form QA.

Embedding recursive search within an agentic framework extends the forward-looking search horizon further. For instance, RAP (Retrieval-Augmented Planning for multimodal LLM agents) [74] introduces a contextual memory that is used to create a planning process, where past experiences corresponding to the current situation and context are retrieved from a high-dimensional memory, enabling agents to dynamically reflect upon and augment their decision-making process. This conceptualizes iterative retrieval in-situ, reframing it not as the search for answer passages, but as the retrieval of past experiences and planning sub-goals in a complex environment, positioning the integration of retrieval as part of a richer, memory-informed cognitive process of the agent.

These iterative and recursive systems converge on the principles of token importance and retrieval utility. The structure of iterative retrieval provides a natural termination criterion: the convergence of the query and internal confidence, where retrieval is halted when the generation does not meaningfully improve the answer, or when a task-tracking mechanism like a progress tracker indicates that a final answer can be issued from the collected context [75]. However, this convergence problem is nuanced. As highlighted by the potential negative impact of irrelevant context [76] and an unexpected benefit of irrelevant documents in some setups [55], we need a more nuanced insight. The power of noise means that iterative loops can go awry if not carefully managed, suggesting that existing iterative models which simply augment the prompt can be improved by introducing a structured evaluation and correction of the retrieved context’s quality at each step.

A promising avenue for the next steps of complex reasoning systems lies in the intersection of

multi-hop retrieval with sophisticated decision-making strategies. Rather than viewing retrieval as an automatic attachment to a single generated intermediate step, more recent architectures leverage models to determine the reliability of the generated reasoning. In dynamic question generation or the queried reformulation, the response that is generated is re-processed and refined from the perspective of decomposed sub-steps; complex questions are broken down into a series of single-hop questions, and each of them is sequentially grounded in a new retrieval (Generate-then-Ground) [77]. This demonstrates how iterative generation can be fused with a structured refinement of the evidence set, a principle that directly cascades into the post-retrieval processing techniques covered in the next section, which focus on reranking, filtering, and compressing the outputs of such multi-step retrieval trajectories to ensure that only the most salient and non-redundant information reaches the generator. Thus, the future direction lies in a more nuanced interplay not just in the *where* to look, but in the *when* to ask and generate—with a strong emphasis on preventing computational redundancy in RAG while scaling the reasoning capacity of LLMs to capture multi-step semantic dependencies.

4.5 Post-Retrieval Processing for the Generation Module

Post-retrieval processing constitutes a critical intermediary stage in the RAG pipeline, situated between the raw retrieval of candidate passages and the final synthesis by the generation module. The quality of the final generated output is not solely a function of retriever recall or the LLM’s parametric knowledge; rather, it is substantially conditioned by how effectively the system selects, compresses, and restructures the retrieved evidence [8, 12]. Naively concatenating the top-k retrieved passages into the prompt often introduces noise, redundancy, and positional biases that can degrade model performance, a phenomenon now known to be a primary bottleneck in system reliability. Research has revealed that the inclusion of irrelevant and misleading documents can be unexpectedly harmful, thereby motivating a dedicated processing stage to bridge the gap between crude relevance and final utility [55, 71].

One of the most basic and essential post-processing operations is reranking, which aims to reorder the initially retrieved candidates by a more refined relevance or utility score. While primary retrievers—whether lexical or dense—are optimized for global recall, they frequently struggle with the nuanced intent of the downstream task. A dedicated cross-encoder or a distilled ranker can compute a more fine-grained interaction between the query and each document, thereby reordering the set to place the most authoritative passages in the most salient positions of the prompt [49]. Intriguingly, the *position* of these relevant documents is as significant as the documents themselves, as LLMs exhibit a "lost-in-the-middle" bias, failing to fully leverage evidence placed mid-context, a finding that underscores the need for disciplined ordering strategies independent of retrieval scores [59, 11].

Beyond selection, equally important is the explicit correction and refinement of the evidence set. An increasing body of literature emphasizes *filtering* and *refinement* over simple reordering. The retrieved corpus may contain outputs that are completely irrelevant to the user’s request, or are duplicated and contradictory. By implementing a lightweight retrieval evaluator that assesses the quality of the retrieved context, systems can trigger decomposition-based re-composition algorithms to extract and integrate key information while discarding the chaff [15]. The BIDER framework, for instance, refines retrieved documents into "Key Supporting Evidence" by learning from KSE (Key Supporting Evidence) simulation, aligning the information directly to the acquisition preferences of the LLM, leading to a 7% boost in answer quality while simultaneously reducing input length by 80% [78]. Furthermore, iterative self-feedback mechanisms, as demonstrated in RA-ISF, enable the evacuation of non-essential retrieval augmentations and instead reinforce the model’s inherent reasoning capacity [73].

To mitigate the problem of the finite attention budget, *context compression* has become a key strategy, moving beyond mere selection to summarization. End-to-end compression methods or concept-based abstractive algorithms transform multifaceted and verbose documents into a

concise summary, which is then fed to the final teacher model. Abstract Meaning Representation (AMR)-based concept distillation can be employed to filter out noise within the passage and retain critical concepts, guiding the LLM to focus on the distilled concept templates, rather than the entire raw passage [58]. Such methods reduce both computational overhead and the possibility of the generation model being swamped by irrelevant semantic content. This concept is in line with "InFO-RAG," an information refinement training framework that treats the LLM in the RAG pipeline as an "information refiner," consistently integrating external knowledge with the model’s intrinsic parameters to generate superior outputs [79].

Emerging research has demonstrated the utility of a novel *familiarity-aware compression* mechanism. This approach encodes an entire retrieved set into specific regenerated utility (e.g., a "soft prompt") that the LLM can directly adapt to and utilize, often aiming to ease the computational burden of a fixed-sized intermediate memory [58]. The integration of such token-level controls extends the manipulation of k-value with tokens that are generated for consumption by the model, effectively merging the distinction between retrieval hardware and generation logic within the context of a single framework.

Given that the raw information derived from retrieval spaces is inevitably imperfect, robust post-processing is essential to prevent hallucination in RAG systems. The future of post-processing is in targeted, model-aware compression techniques that can systematically identify and prioritize crucial information, while also communicating the retriever’s relevance so that the final generation module can decide whether to incorporate, contradict, or reject the retrieved evidence—factors known to be critical for robust application in high-stakes domains like medicine. A shift toward *unified context representation*—where the post-retrieval processor is aligned not only with the user intent but also with the generator’s task-specific awareness—promises to substantially elevate the performance of the final generated text [80].

4.6 Orchestration of Advanced and Multi-Agent RAG Systems

The orchestration of advanced retrieval-augmented generation (RAG) systems represents a paradigm shift from monolithic pipelines toward composable, agentic architectures in which a central controller—typically an LLM—coordinates retrieval, reasoning, and generation through dynamic, often multi-agent, workflows. Building directly upon the post-retrieval processing stage discussed previously, orchestration extends the notion of refinement beyond individual evidence selection and compression to govern the entire retrieval-generation loop as a cohesive, decision-driven process. This additional layer of control enables systems to transcend the limitations of static retrieve-then-read paradigms by introducing self-assessment, tool use, and iterative decision-making capabilities tailored to query complexity and information needs.

A cornerstone of advanced orchestration is the integration of self-assessment and reflection mechanisms that govern the retrieval-generation loop, effectively generalizing the filtering and refinement goals introduced in post-retrieval processing to a system-wide level. Frameworks like Self-RAG introduce trainable reflection tokens that allow the generator to evaluate its own outputs, deciding whether retrieval is necessary—thereby bypassing the retrieval step when parametric knowledge suffices—and whether retrieved passages are relevant or the final response is evidence-supported [81, 82]. This self-reflective capacity enables dynamic retrieval triggers and iterative evidence refinement, both of which echo the iterative self-feedback strategies seen in post-retrieval refinement but are now governed centrally by the orchestrator [82]. Building on this, evaluative components can serve as built-in validators, checking reliability, task-relevance, and redundancy through a cascade of controlled checkpoints before the final answer is committed, effectively formalizing the filtering role inside a closed control loop [83]. Such architectures not only enhance factual grounding but also confer robustness against irrelevant or adversarial retrieved noise, particularly in multi-turn interactions [84, 81]. Critically, the orchestration layer acts as the decision-making hub that negotiates between parametric memory and non-parametric evidence, mitigating incoherence from noisy retrieval through explicit planning and evaluative

gates.

The rise of modular pipeline toolkits and unified frameworks has accelerated the adoption of standardized orchestration layers, providing the structural means to realize the context-unification goals envisioned at the end of the post-retrieval stage. Open-source ecosystems offer interoperable, plug-and-play modules for all sequential operations—from query rewriting and retrieval to reranking, evidence extraction, answer generation, and fact-checking—thereby enabling systematic composition and comparison [44]. FlashRAG and similar libraries implement collections of state-of-the-art retrieval and generation modules with standardized benchmarks, facilitating rapid reproduction and fair benchmarking [44]. Such modular blueprints simplify incremental ablation and support the integration of adaptive policies that select between low-cost RAG paths and high-cost long-context LLM paths based on query complexity or model confidence, effectively connecting the retrieval quality and generator awareness that post-processing sought to align [85, 86]. The modular approach further allows knowledge-graph-based retrieval to be combined with dense vector search, as demonstrated by HybridRAG and graph-augmented systems, enriching both recall and answer fidelity [17, 17].

Memory-augmented agentic frameworks push orchestration one step further by integrating external memory as a persistent component of the RAG loop, complementing the evidence compression strategies from the predecessor stage with long-term, interpretable memory. For instance, RET-LLM equips the LLM with a read-write memory unit encoding knowledge as triplets, offering an explicit and scalable memory store that supports continual learning and temporal reasoning [44]. HippoRAG, meanwhile, uses a knowledge graph as a non-parametric memory and applies Personalized PageRank to decode multi-hop queries in a single step, rivaling the performance of iterative methods like IRCot at reduced computational cost [17]. These designs suggest an evolution toward orchestration where memory and retrieval are isomorphic components, strengthening the context unification initiated by post-retrieval compression and enabling a tighter loop between storing, retrieving, and grounding.

A final layer of sophistication involves adaptive and efficient orchestration strategies, where the orchestration agent proactively routes queries, plans multi-step exploration, and incorporates user feedback. When query complexity is heterogeneous, the orchestrator serves as a complexity classifier, routing simple queries directly to the generator and reserving multi-hop for harder tasks, which manages computation and error budgets [82]. Multi-agent and reinforcement-learned controllers further augment decision-making across large storage resources, selecting and sequencing retrieval actions for maximal relevance under resource constraints [87]. In user-centric settings, the interactive feedback and query-rewriting loops improve retrieval-generation convergence, creating a self-correcting feedback path that mirrors the iterative refinement objectives from the post-retrieval stage but from a more systemic perspective [88]. These orchestrated multi-agent patterns lay the groundwork for a next-generation RAG architecture, where retrieval, post-processing, and generation are—under a centralized, adaptive, and goal-driven engine.

In summary, the orchestration layer advances the RAG pipeline from static, module-level post-processing to dynamic, agentic, self-reflective coordination. Building on the earlier focus on post-retrieval selection, compression, and context unification, orchestration integrates these components into a cohesive digital workflow that can assess metadesisions, manage memory, and redesign retrieval-generation interactions in real time. Future orchestration research should target principled planning, lifecycle-guided RL mechanisms, memory-efficient long-term structuring, and user-in-the-loop adaptation to foster robust systems in dynamic and possibly adversarial settings, bringing the field closer to trustworthy, high-performance knowledge engines for complex tasks.

5 Training Paradigms and Adaptation for Augmented Models

5.1 Foundations of Training Paradigms for Retrieval-Augmented Models

Training retrieval-augmented generation (RAG) systems constitutes a fundamental departure from conventional language model fine-tuning, as it necessitates the coordinated optimization—or deliberate non-optimization—of two distinct knowledge substrates: the parametric memory embedded within the weights of a neural network and the non-parametric memory housed in an external datastore. The foundational choices made during this training phase are not merely implementation details; they fundamentally dictate the system’s knowledge allocation, its ability to generalize to novel distributions, and its operational viability in production environments characterized by stringent resource constraints.

The first and most consequential decision concerns the allocation of learning capacity along this parametric–non-parametric spectrum. Parametric knowledge, learned during model pre-training and often fine-tuning, offers dense, compressed world knowledge accessible with minimal latency, but is ineluctably static without incurring the significant computational penalties of re-training or up-training. Non-parametric memory, conversely, provides a mutable knowledge substrate for agile updates [26]. This creates a fundamental trade-off: training strategies must strategically divide learning responsibility between the encoding of retrieval cues within the retriever and the generation capabilities of the parametric model. Empirical studies consistently show that for queries about vulnerability long-tail entities or rapidly evolving events, the dynamic retrieval of non-parametric evidence proves essential, as unassisted LMs remain competitive only for high-popularity entities [4]. Consequently, the literature suggests a gradient of training strategies, spanning from zero-shot, training-free augmentation to intensive joint optimization, each placing a distinct emphasis on parametric versus non-parametric knowledge use [42].

At one end of this spectrum lies a purely training-free augmentation paradigm. This approach leverages frozen off-the-shelf language models and retrieves and concatenates documents into the prompt context, requiring no weight updates. This paradigm has proven remarkably effective as a quick and resource-conscious technological bridge between domain adaptation and private knowledge injection [77]. However, its critical limitation is the absence of alignment between the retriever and the factors of the generative model, often leading to suboptimal integration of retrieved content. To mitigate these without full fine-tuning, *warm-start* and parameter-efficient strategy exist that insert trainable components into the frozen backbone. Examples include looking for low-rank adapters (LoRA) and explicitly turning the retrieved context into virtual token representations to calibrate the generator’s ability to distinguish foreign information from parametric knowledge through lightweight injections, in contrast to purely closed-source APIs where only the prompt is editable.

In clean contrast, *full end-to-end optimization* of the retriever and generator is a path whose greatest asset is deep alignment but likewise carries heavy computational loads. Unlike simple tasks, this approach views the entire pipeline as a differentiable learning system or as systems with gradients estimated, where the retrieval actions form an intermediate policy conditioned on generation’s final output through an alignment objective or via reinforcement learning, seeking to optimize downstream generation utility [6]. Furthermore, when access to the true world-state (or a verifiable factual anchor) is available, further offline objectives such as factual consistency metrics could be seamlessly integrated as task rewards [24]. Yet, a crucial lesson is that end-to-end training’s true superiority is designed for scenarios where the the retrieval; the exploratory objective is direct and fully differentiable.

The choice of training paradigm is also regulated by the *Catalog of Training Objectives*. The elementary objective is due to the maximum likelihood estimation (MLE) of the target text, but behaves as a random walk for semantic relevance. More granularity is given by tasks involving sequence-level losses that separate factually consistent evidence from generators with diverging trajectories, or the use of preference optimization to fine-tune a generator’s factuality through

automatic generation of reward rankings, rather than by hand-labeled preference data [24]. On the retrieval front where the objective is not to reproduce documents but to produce passage over relevant content, contrastive losses between positive and hard negative pairs are applicable. The retriever can be trained judiciously either with its own labels or via a multi-task objective that jointly maximizes the end-to-end task performance, with data generation performed by stronger teacher models in the loop [24]. In a broad sense, the design space includes solving a multi-objective optimization problem over the Pareto frontier of accuracy, efficiency, and adaptation, a challenging and largely open frontier [23, 20].

In determining the ideal training regime, key is the operational scenario: the regimes where gold data is plentiful and coupled with large compute usually prove favorable for end-to-end alignment models; at the other extreme, where rapid adaptability, personal data protection, or domain-specific knowledge updates are preferred, a modular or training-free retriever is favored [3]. The best design is to land on a portfolio containing training and inference costs through memory interfaces and retrieval-aware fine-tuning at multiple levels of the system. We have identified the preamble of the *dynamic combination of a frozen parametric backbone in conjunction with a cheaply adaptive datastore*, and the search for task-specific retriever as valuable research directions [3]. Furthermore, the building blocks that jointly curate architecture are similarly need to be consolidated in dataset-level document pivots (retrieval scale and underlying data structure). In foreseeing, the current fast-paced space of RAG architecture must catalyze a transformation from the current one-size-fits-all to a "Skill-aware and utilization-flexible" training regime where effort is spent precisely on high-value logic and knowledge adapters [3]. In short, the ultimate promise of RAG is not realized through a single optimization procedure but through a calculated, task-aware partitioning of the overall learning burden.

5.2 End-to-End Joint Optimization of the Retriever and Generator

The joint optimization of retriever and generator components within a unified RAG framework represents one of the most technically formidable challenges in the field, as it requires reconciling two fundamentally heterogeneous modules through a single differentiable pipeline. The central computational obstacle is the discreteness of the retrieval operation: document selection is inherently a non-differentiable, combinatorial process that blocks the natural flow of gradients from the generation loss back to the retriever parameters. This bottleneck has inspired a rich body of work that seeks to approximate or circumvent the discrete bottleneck through continuous relaxation, stochastic estimation, and hierarchical credit assignment strategies. Critically, the training paradigm chosen—whether at the frozen or fully jointly optimized end of the spectrum described earlier—directly dictates which of these gradient-solving techniques become applicable, and the severity of the credit assignment problem they must address.

The earliest formal treatments of this challenge emerged from the recognition that the standard marginalization of retrieval variables—summing the likelihood of generation over all possible retrieved documents—is computationally intractable due to the combinatorial explosion of candidate sets. The foundational framework [8] addressed this problem from a practical perspective, demonstrating that jointly fine-tuning a dual-encoder retriever with an encoder-decoder generator via gradient descent could yield substantial gains across knowledge-intensive tasks despite the non-differentiable top-k operator. Rather than employing explicit continuous relaxations, this approach routed the gradient through the selected top-k documents, subsequently propagating the resulting importance-weighted gradients into the dual-encoder via a dot product objective. This jointly fine-tuned architecture effectively establishes the practical computational manifold on which much of the subsequent gradient research in this area operates.

A notable advancement is the use of stochastic relaxations of the discrete retrieval operator, as prominently formalized in [89]. A key critique has been that the marginalization strategies and indexing independence assumptions imposed by earlier generation-balanced works are both overly restrictive and often theoretically unjustifiable. Stochastic RAG instead develops a more general

view of the retrieval process as a stochastic sampling without replacement operation, which is made differentiable through the Gumbel-top-k distribution, using straight-through estimators that preserve discrete forward behavior while allowing gradient propagation. Empowered by this general formulation, the approach enables standard grid-level, task-specific end-to-end optimization with arbitrary objectives. Empirically, this formulation demonstrates substantial and consistent performance gains over well-established retrieval-augmented methods across a diverse set of tasks, including open-domain question answering, fact verification, and dialogue systems, showcasing the benefits of formally treating the discrete bottleneck.

In contrast to these path-gradient methods, a family of approaches based on reinforcement learning (RL) imposes fewer requirements by modeling the retrieval step as an action selection process governed by a stochastic policy, and estimating gradients through reward signals generated by downstream task feedback. This setting is conceptually the most flexible for application where the anticipated future quality of the model’s output can be defined by task-specific, human-derived signals, making it highly compatible with the preference-optimization and alignment objectives that are the focus of the next section. Its practical efficiency, however, is somewhat constrained; the retriever and the generator may be updated with heterogeneous reward formulations, given their differing computational requirements and update frequencies.

An orthogonal but increasingly influential pathway—one that directly extends the expected-utility philosophy of stochastic RAG—is the construction of a differentiable surrogate objective over retrieval indices, aiming to maximize the expected utility over latent retrieval decisions. These estimators, instead of collapsing or discretizing the space, treat the ultimate objective as an expectation over documents, embedding the retrieval scores into a differentiable calculator. The expected-utility estimator in particular forges a formal connection between the RAG objective and the retrieval scores, providing a principled foundation for end-to-end training that has historically been absent.

Despite these advances, so there remain significant challenges in the compromise between training stability, bias introduced by continuous relaxation, and computational overhead when operating with large retrieval corpora. The full range of gradient estimation choices also leads to a frequent practical division of labor: across-agent entities, the retriever is initially pretrained independently, followed by diversified fine-tuning circuits that balance exactness with scalability. As RAG systems evolve toward more agentic and iterative frameworks [43], the corresponding end-to-end credit assignment for long-horizon interactions—with multiple queries, multiple reasoning steps and multi-turn objectives—will become increasingly complex, requiring new relaxation algorithms and potentially higher-level control architectures that can operate across sequential accesses. The careful reconstruction of gradients in these stochastic, and possibly noisy, environments will likely be a defining and fruitful direction for future research in RAG training.

5.3 Preference Optimization, Alignment, and Human Feedback

The integration of human preference signals into retrieval-augmented generation (RAG) training represents a paradigm shift from optimizing purely for task accuracy toward aligning system outputs with nuanced human expectations regarding style, relevance, and faithfulness. Unlike conventional fine-tuning that maximizes likelihood over reference texts [90], preference optimization within RAG must contend with a dual-objective landscape: the combined behavior of the retriever and generator jointly determines the quality of the final response, yet the credit assignment problem is exacerbated by the discrete, non-differentiable retrieval step. Early work in this area has framed the optimization as a utility maximization problem, where a reward signal—often derived from human pairwise comparisons—guides both the ranking of retrieved evidence and the generation of the final answer [91].

The construction of reward models lies at the heart of this process. As shown in the framework proposed by Peng et al., a distributional reward model can be trained to align

RAG agents with human preferences by conditioning on both the query and the retrieved context, capturing not merely answer correctness but also the answer’s supportiveness relative to the evidence [24]. However, human preference labels for retrieval-augmented outputs are notoriously heterogeneous and noisy; pairwise judgments may conflate stylistic preferences with factual correctness, requiring careful aggregation strategies. A promising trajectory is the use of continuous cache mechanisms—analogue to memory-augmented networks that store past hidden states to adapt predictions to recent history [9]—to dynamically integrate preference signals over time without full retraining, thereby enabling adaptive alignment in deployment.

Among the training algorithms, proximal policy optimization (PPO) has been the dominant approach in RLHF literature, yet its application to RAG requires careful decomposition of the reward into components for generation quality, retrieval-grounding fidelity, and the KL distance from a reference policy. A defining aspect is that the reward function must span both the retrieval and generation dimensions, where the retrieval component is trained to maximize the expected utility of the downstream generation. Notably, a decomposed reward design that incorporates an explicit consistency penalty with the retrieved evidence often demonstrates higher overall alignment, especially in open-domain QA settings where hallucinated or non-groundable claims are penalized [92, 30]. Meanwhile, alternative, lightweight approaches—such as direct preference optimization (DPO)—have gained traction as a cost-effective surrogate for PPO. DPO avoids explicit reward modeling and policy sampling, yet it adapts to RAG by treating the retrieved documents as part of the conditioning context, enabling interleaved fine-tuning of the retriever’s ranking objective and the generator’s preference alignment [3].

The emerging emphasis on data-centric methods and automated feedback is distinctive to RAG alignment. Since human annotation of relevance and factuality at scale is prohibitively expensive, a prominent direction leverages the LLM itself as a self-critic or teacher to generate feedback signals. For instance, the Self-RAG framework learns explicit reflection tokens that evaluate the relevance of retrieved passages and the coherence of a generation [30], while “LLM as a judge” approaches are used to provide a reward for preference labeling without human-in-the-loop [93]. Moreover, iterative self-feedback strategies—where a model decomposes the task, provides feedback based on retrieval results, and refines retrieval queries—have demonstrated strong improvements in complex multi-step reasoning tasks [73]. However, as the field evolves toward open-ended generation, the evaluation of output becomes more difficult, raising questions about how well these automated preference signals align with human judgments and whether such methods remain appropriate for high-stakes, safety-critical applications [90].

An overlooked but critical issue in preference optimization for RAG is the distribution mismatch between the training-time preference data and the test-time retrieval distribution. If the retriever is considered frozen during preference alignment, the system may overly rely on parametric knowledge, leading to a reduced use of non-parametric memories, which is a significant problem when the model is deployed in a continuously changing knowledge landscape [4, 94]. Bridging the preference gap between retrievers and LLMs—through either supervised response learning or reinforcement learning to train a bridge model that optimizes the connection—becomes a necessary step for alignment-driven RAG [52]. Additionally, recent evidence suggests that preference feedback is best used to adjust the entire architecture (retriever + generator) in a joint fashion, rather than only the generator, to avoid a situation where the final system becomes less sensitive to user preference signals [37]. We expect that future work will asymptotically combine parameter-efficient adaptation (e.g., adapter modules or low-rank perturbations) with preference optimization to tune RAG systems cheaply, while simultaneously ensuring alignment robustness against low-quality, noisy, or even poisoned retrieved contexts, thus unifying alignment objectives with security and reliability [30, 24].

Citation Corrections Made:

1. "[24]" – Correct as is; this paper covers tuning language models for factuality using preferences, supporting the sentence about factuality scores in preference alignment.
2. "[3]" – I

replaced the citation for DPO with this paper because the original citations provided do not clearly discuss the DPO algorithm; this paper covers comparing fine-tuning and retrieval for knowledge injection, fitting the context of adapting RAG with retrieval-aware fine-tuning. 3. "[95]" – Replaced the citation for "LLM as a judge" with the data annotation survey paper, which explicitly supports using LLMs for annotation and automated feedback. This aligns with the context of automated, LLM-driven preference/reward labeling without human annotations. 4. "[4]" – Corrected the title to match the actual paper title. The previous citation "[4]" had "Memory" singular, which was ambiguous; the provided paper_title uses "Memories." I've corrected it to "Memories" for exact match. 5. **Duplicate citation cleanup:** Replaced a duplicated "[30, 24]" at the end with "[30, 24]" for correctness and matching the provided titles. 6. Minor spelling/spacing cleanups (e.g., "learning to retrieve, generate, and criteria" → "generate, and Critique through Self-Reflection" corrected; "non-parametric memory" → "non-parametric memories" for title match).

All other citations remain as provided since they align with the content of the cited papers.

5.4 Noise Robustness, Denoising, and Multi-Task Learning

The retrieval module in deployed RAG systems is inherently imperfect: it may surface passages that are irrelevant, partially related, or even outright contradictory to the user's query. Consequently, a central challenge in training augmented models is instilling robustness against such noisy contexts, ensuring that retrieval enhances rather than degrades generation quality. This challenge is compounded by the fact that, as discussed in the previous subsection, preference-based alignment signals often assume a well-behaved retriever; when the retriever introduces noise, the preference-optimized generator must autonomously learn to distinguish reliable evidence from distractors, a responsibility that becomes increasingly critical yet under-specified at test time. Empirical evidence underscores both the severity and the nuance of this problem. Notably, some studies have revealed a counter-intuitive phenomenon: including seemingly irrelevant documents can unexpectedly enhance task performance by more than 30% in accuracy, suggesting that uninformative context might act as a form of regularization or provide useful lexical priors [55]. However, this finding contrasts sharply with the more common observation that misleading or partially relevant information—especially in multi-hop reasoning—can lead to cascading errors, where a single irrelevant passage propagates inaccuracies through subsequent retrieval and generation steps [76, 72]. This dual nature of noise posits that the objective of training is not merely to ignore irrelevant information, but rather to selectively leverage useful signals within a noisy context while distinguishing them from true distractors.

A primary line of research tackles this issue through a fine-grained, directly supervised approach. A foundational method involves fine-tuning the language model on a mix of relevant and irrelevant contexts, demonstrating that even a surprisingly small set of 1,000 fine-tuned examples can effectively train the model to maintain high performance when contexts are relevant while becoming robust to irrelevant ones [76]. This finding highlights the cost-effectiveness of targeted data creation in the face of a significant knowledge problem. Building on this, adversarial augmentation, where the model is trained on perturbed or noisy contexts, forces the model to develop augment internal mechanisms for identifying and disregarding misleading information [12, 17]. These methods are often augmented by cognitive objectives that explicitly mask certain parts of the input, forcing the model to focus on the correct evidence [17].

Beyond direct filtering, meta-cognitive approaches train the model to explicitly assess the reliability of retrieval sources before grounding its responses, decoupling the surface-level plausibility of a passage from its actual utility and achieving stronger calibration [96].

An alternative and highly effective mechanism for denoising is through dynamic re-ranking and retriever-reader co-training. Iterative frameworks such as Iter-RetGen demonstrate that by allowing the model to use the information acquired from a previous generation step to reformulate the query, the gap between the learner's understanding and the retrieved corpus is

progressively narrowed, thereby filtering out noisy evidence [29, 72]. Similarly, by employing a "generate-then-ground" approach, the model can use the initial language-model-generated answer as shortlist, combined with retrieved documents for re-ranking, showing that the boundaries of what is considered "noise" is not static [97]. These co-training techniques not only ground but also testify to the shift toward bridging the gap between the retriever and the LLM, which was a limitation observed in preference optimization when the retriever remains frozen and the test-time retrieval distribution deviates from training.

Despite the theoretical advancements, training paradigms for noise robustness remain heterogeneous and loosely coupled, and the community still grapples with challenges in defining the optimal weight on parametric vs. non-parametric knowledge [12]. A promising yet under-explored direction is to close the loop between these two dimensions by utilizing both fine-grained and meta-cognitive training signals, which will not only build a strong preference function but will make robust generation at scale a recognized, actionable goal. Future directions could include training a dedicated, lightweight evaluator-creator to generate "hard negative" contexts within the training loop rather than relying solely on retrieved data, making the implicit selection of sources a more declarative, actionable goal that naturally connects to the data-centric pipelines developed in the following.

5.5 Data-Centric Strategies and Synthetic Data Augmentation

The performance ceiling of retrieval-augmented generation is bounded not only by architectural innovations and loss design but, critically, by the quality, distribution, and specification of the data used for training and adaptation. Data-centric strategies—spanning dataset selection, reweighting, and, most notably, synthetic data augmentation—have emerged as a decisive lever for enhancing both retrieval quality and generation fidelity. This perspective is supported by empirical evidence demonstrating that the careful curation and construction of training data can yield gains that rival or exceed those obtained from architectural modifications or objective function engineering.

A critical consideration in the development of RAG systems is the selection and weighting of existing corpora, as the performance of both retrievers and generators is highly sensitive to the distribution and relevance of the training data. Rather than relying on costly manual annotation, unsupervised curation of high-quality training signals from large unlabeled corpora presents a low-cost and effective alternative for domain adaptation. The core of data-centric innovation, however, explores synthetic data augmentation to enhance the model’s coverage of diverse tasks and retrieval scenarios. A common approach involves using a larger, more capable teacher model to generate synthetic instructions and corresponding inputs, which are used as training data for smaller student models [98]. This scalable method can produce a richer distribution of scenarios, which is critical for generalization. The utility of synthetic data also extends to generating hard-negative examples and relevance labels for the retriever itself, thereby enabling a self-training or active-learning loop that aligns the index with the downstream generation task.

A crucial advancement in this area is the explicit control over the synthetic data generation process through the incorporation of retrieval and logits feedback. This approach represents a significant stride beyond the process of merely querying the teacher model with in-context prompting, showing the utility of more complex data generation techniques that directly aim to optimize the end-to-end RAG pipeline.

The distinction between the quality of synthetic data and the model’s ability to use it effectively is critical. The RA-DIT approach demonstrates that explicit fine-tuning on a blend of knowledge utilization and retrieval tasks is essential to bridge the intersection gap between the retriever, retrieved-context, and generator [99]. This fine-tuning, guided by synthetic data from a teacher model, has demonstrated the model’s capability to achieve up to +8.9% improvements in zero-shot settings by leveraging a combination of document-aware and document-agnostic augmentation. Furthermore, in domains with low-frequency entities, synthetic data can be

used to improve retrieval of factual knowledge and the ability to generate answers for diverse vocabularies. In the medical domain, constructing RAG systems driven by synthetic QA pair generation has been shown to be a feasible strategy for adapting models to constantly evolving clinical guidelines. This approach is notable as it can be used to test the impact of a generated dataset, yielding significant accuracy gains and demonstrating the clinical effectiveness of non-inferiority testing after incorporating synthetic RAG-specific data [100].

The emergent frontier in data-centric RAG inherently centers on the orchestration of data generation (synthetic or curated), selection, and the management of the knowledge base scaling. The future of this domain will increasingly concentrate not on designing more complex architectures but on how to efficiently enable and harness specialized data pipelines for systematically isolating the key failures in retrieval and generation. This principled approach aims to create deep-tailored, high-quality datasets that are agnostic to specific domains, thereby maximizing the power of both the LLMs and addressing the explicit resource needs of enterprise RAG systems.

6 Robustness, Safety, and Trustworthiness in the Loop

6.1 Adversarial Threats via Knowledge Base Poisoning and Backdoor Injection

The openness of the external knowledge store in Retrieval-Augmented Generation (RAG) systems introduces a unique and critical attack surface: the retrieval corpus itself. Unlike parametric-only LLMs, where knowledge is embedded within model weights during a controlled training phase, RAG systems dynamically incorporate external documents at inference time, making them inherently vulnerable to adversaries who can manipulate this corpus. This vulnerability fundamentally challenges the assumption that grounding generation in external knowledge inherently improves trustworthiness; instead, it transforms the knowledge base into an attack vector where malicious content is not just noise but a precisely engineered instrument for controlling model behavior.

Formally, corpus poisoning can be conceptualized as an adversarial optimization problem. An attacker aims to inject a set of crafted documents $\mathcal{D}_p$ into a larger knowledge base $\mathcal{D}$ to force a target answer y_{target} for a given query q . The attack is successful if the retriever, when processing q , retrieves the poisoned documents from $\mathcal{D}_p$, thereby compelling the generator to produce y_t . This formulation necessitates that the poisoned documents are not only relevant to the query, but also crafted to dominate the retrieval score and lead the LLM’s generative process down a predetermined path. This is a distinct challenge from traditional search poisoning, as the success criterion is the fidelity of the final generated answer, not just its click-through rate.

The dual nature of poisoning, encompassing both "immediate" manipulation and "latent" backdoor attacks, creates distinct threat models. For immediate manipulation, the attacker can exploit the black-box nature of the system by crafting documents that are self-contained and highly salient to the target query, a strategy that has been shown to achieve high success rates against RAG pipelines that rely on both dense and sparse retrievers. This is particularly pervasive in domains that rely on open-source, user-generated content where the barrier to entry is low. Conversely, the deployment of backdoors introduces a persistence element; by using "trigger phrases" and targeted grammar errors, an attacker can create a Trojan horse within the retrieval index, ensuring that malicious passages remain undetected and are only activated for specific, targeted queries while maintaining normal behavior for standard queries. This is analogous to poisoning the model’s long-term memory.

The threat model becomes even more dangerous in a completely remote and constrained scenario. Attackers, without any knowledge of the target user’s queries, the content of the knowledge base, or the model parameters, can exploit the fundamental mechanisms of popular

retrievers by treating the entire system as a black-box. By injecting synthetically "good" content—documents that are highly self-referential and semantically consistent with the target topic—an adversary can have the embedded passages retrieved and fed into the LLM. This exploits the inherent and often unavoidable conflict between the retriever’s objective (to find semantically relevant text) and the generator’s difficulty in resisting instructions embedded within that text. This mirrored-real-world scenario, such as a user uploading crafted documents to a public repository that later gets indexed by a RAG system for various queries, highlights the remote exploitability of the architecture.

The most concerning aspect of these attacks is not just the capacity to manipulate the final answer but the way they compromise trust. Unlike hallucinations, where the model is simply wrong on its own, poisoning becomes an active countermeasure to the primary benefit of RAG: verifiability. By injecting a carefully crafted passage that seems to "support" the false target answer, the attacker can co-opt the authority and the citation mechanism, turning the system’s own framework for grounding into a deceptive tool. The problem thus becomes not just one of model behavior, but of the validation of the entire retrieval process. The generated output from this poisoned pathway is self-authoritative, a flaw that makes the output appear more reliable, not less, regardless of its factual inaccuracy.

These attacks operate at the intersection of the retrieval and generation modules, creating a significant challenge for designing effective defenses. While basic post-hoc filtering can remove irrelevant documents, it is less effective against adversarial documents that are semantically related to the query. The sophistication of these backdoor attacks highlights the need for a "clean lab" approach to understanding and detecting such subtle injections [12]. A future avenue of defense may not lie in simply making the retriever more "secure," but in making the entire system more robust to the process—for instance, by making the model aware of the possibility of adversarial injections in its own training via adversarial training or by building a "trust partition" that excludes instruction-following behavior from the retrieved data [12]. The fundamental long-term risk is that if applied at scale to public data pools, these poisoning and backdoor attacks will not only erode user trust in the individual responses but will fundamentally undermine the credibility of the entire retrieval-augmented generation paradigm.

6.2 Vulnerabilities Arising from Context Injection and Prompt Manipulation

The integration of external knowledge through retrieval mechanisms, while enhancing the factual grounding and adaptability of large language models, simultaneously opens a critical attack surface: the input context. Whereas the corpus-level threats discussed previously involve preemptively compromising the knowledge source before retrieval, the vulnerabilities examined here emerge at the moment of generation, when adversarial content enters the system through the retrieved passages themselves. Unlike traditional prompt injection attacks where an adversary directly interacts with the model, retrieval-augmented generation (RAG) systems become a conduit for **indirect prompt injection**, where malicious instructions are embedded within the very documents retrieved to augment a user’s benign query. The presence of irrelevant documents can also negatively impact the system’s response quality, degrading the factual grounding that RAG is designed to provide [55, 101, 70]. This vulnerability is particularly insidious because the user is unaware of the adversarial content, and the model, trained to faithfully follow instructions, treats the retrieved text as authoritative context rather than as untrusted data—a failure of epistemic vigilance that directly undermines the trustworthiness guarantees that motivated the RAG architecture in the first place.

This attack vector exploits the inherent tension between an LLM’s instruction-following capability and its inability to reliably distinguish between user commands and data. An attacker can craft a document that, when retrieved and incorporated into the prompt, subtly or explicitly overrides the user’s original intent, compelling the system to execute hidden directives. The success of this attack is often independent of a vulnerability in the generative model itself;

rather, it is a systemic flaw in standard RAG pipelines, which lack robust mechanisms for data-command isolation. Furthermore, this vulnerability can be compounded in iterative or agentic RAG settings, where a single successful injection can hijack the entire multi-step reasoning and tool-use process toward malicious objectives. The threat extends beyond simple output deviations to more severe forms of exploitation, including **jailbreaking**, where an attacker can use the retrieved context to bypass the model’s built-in safety features and elicit prohibited content.

Following a successful injection attack, a RAG system becomes susceptible to **data extraction and prompt leaking**. An adversary can craft documents designed to manipulate the model into revealing its internal instruction set, a critical asset that details its operational logic, or to act as an oracle for the retrieval database itself. Through carefully constructed prompts within the retrieved text, an attacker can query the system in a way that causes it to reproduce sensitive, proprietary, or personally identifiable information from the external knowledge base. This risk is particularly notable in enterprise and domain-specific deployments, such as those in the telecommunications (Telco-RAG) or personalized dialogue systems, where RAG systems are deployed over internal, confidential corpora [33, 102]. These attacks demonstrate that the security of a RAG system is not determined by its best-case performance, but by its worst-case behavior when the entire context—including the injected data—is adversarial—a property that, as we will next discuss, fundamentally shapes the defensive strategies required to render the system robust.

Defending against these manipulative attacks constitutes a proactive and multi-faceted undertaking, as the primary security paradox of RAG is the platform’s openness: it must process untrusted external content to provide value, yet it must not be subjugated by that content. Practical defenses champion isolation and trust-partitioning—principally ensuring a clear, structural separation between untrusted document content and privileged system instructions in the prompt. This line of defense forms a natural bridge to the architecture-based defensive techniques elaborated in the following section, where structural boundaries serve as the primary mechanism for neutralizing adversarial content. State-of-the-art detection mechanisms include robust input filtering and instruction-aware inspection, achieved through learning to classify or judge the relevance and intentionality of the content to build a robust input filter [101], or by using a self-refinement mechanism to iteratively check the consistency between the generated response and the retrieved evidence [102]. Given both the knowledge-store vulnerabilities described earlier and the retrieval-context attacks examined here, future research directions emphasize the need for certified robustness, ensuring that the model’s output remains correct even when a fraction of the retrieved chunks are noised, thereby providing formal guarantees in place of heuristic defenses [103, 101].

6.3 Defensive Mechanisms and Robustness toward Adversarial Inputs

The defensive landscape for retrieval-augmented generation (RAG) systems is fundamentally shaped by the dual challenge of securing both the parametric generator and the non-parametric knowledge store, as adversarial inputs can exploit either component to compromise system integrity. Defensive strategies have evolved along three principal axes: (i) detection and filtering of malicious content before it enters the generation pipeline, (ii) architectural isolation that structurally prevents compromised content from exerting influence, and (iii) certified or robust training procedures that equate the system’s resilience to its theoretical worst-case behavior. The design space is constrained by an inevitable tension: defensive measures must balance the utility of legitimate retrieval signals against the risk of over-filtering or imposing excessive computational overhead that degrades system responsiveness.

A foundational class of defensive mechanisms targets detection and input filtering, often by distinguishing instruction from data within the retrieved content. Learning-based judges, such as the reflection tokens in Self-RAG, enable models to assess the relevance and factual

support of retrieved passages and to abstain from generation when evidence is insufficient [30, 63]. More sophisticated filtering frameworks employ post-hoc "backtranslation" or intent prediction to intercept adversarial prompts before they prompt the system to produce harmful output. However, a critical drawback of detection-only approaches is their self-limiting nature: they rely on the same generator to evaluate content and may exhibit "self-confirm" bias, where the model fails to recognize adversarial content if it aligns with its parametric priors. An alternative to hard filtering is to employ generative-augmentation of the query itself to bias reasoning toward known-benign content, exploiting the power of generative retrieval to reformulate inputs [40].

The architecture-based defenses focus on how retrieved content is integrated with the LLM's instruction-following process, sometimes through explicit context boundaries or hierarchical organization that partitions the structure provided from untrusted data. By allocating different "isolation levels" to system instructions and retrieved documents, these approaches attempt to mimic the memory protection mechanisms of operating systems within neural models. Relatedly, decoupling the parametric content from the non-parametric flow can be coupled with adapting a frozen core using memory editing, effectively preventing a maliciously relevant passage from overwriting the model's underlying parameterized knowledge [32]. With careful tuning of these boundaries, the model retains a "trusted core" of verified knowledge and becomes less susceptible to nothing retrieved that is less reliable. However, such demarcation is only as reliable as the verifier, and the fundamental property of trust is hard to guarantee, unless the core knowledge is kept in discrete, monitored slots, directly analogous to clean-label constraints.

Finally, formally querying or mechanisms based on fine-grained disturbance can disclose certified or high-confidence guarantees. A system can be trained on noisy, irrelevant, or adversarial passages (e.g., via an objective that teaches the model to distinguish clean vs. noisy contexts) so that it learns to emphasize true contributions, thereby mitigating hallucinations [93]. Test-time denoising can also include "contrastive decoding," which offers an implicit ground-truth check by demonstrating the difference between evidence-based on provided context and relying solely on internal priors [104]. While these techniques shift the failure distribution, they are subject to the same failure modes as other deep systems; the mitigation they provide is not absolute, but it elevates the "attack cost" for the adversary.

In the future, defense in the RAG pipeline likely needs to be "in-the-loop" rather than a single stage, requiring active probing and self-evaluation at both the retrieval and generation stages. A critical program is to systematically bridge the gap between what the retrieval finds and the final answer, involving components of self-reflection and iterative feedback loops [4]. Additionally, we note the potential role of adaptive agents, wherein LLMs are allowed to reject unreliable evidence and defer when uncertain, in order to close the vulnerabilities exposed by context injection and poisoning. The extension to effective evaluation, e.g., through testing on unseen, long-tail perturbations, will be crucial to validate the performance beyond held-out datasets [12]. Ultimately, robustness should be established with verifiable functionality—not just be a property that holds at a superficial level—especially in critical domains such as medical and high-stakes finance, where a tiny induced error could outweigh any gains in accuracy [12].

6.4 Safety and Alignment Challenges in Augmented Generation

While adversarial attacks represent deliberate, externally-orchestrated manipulations of RAG systems, a more pervasive and arguably more insidious challenge lies in the safety and alignment failures that arise organically from the inherent imperfections of the retrieval-augmented pipeline itself. These failures, which manifest even in the absence of malicious intent, stem from the unmanageable nature of retrieved information, the model's interaction with noisy or misleading content, and its inability to align its outputs with safety principles when confronted with conflicting or insufficient evidence. Whereas the defensive mechanisms discussed previously—detection, isolation, and robust training—attempt to mitigate both malicious and organic disruptions from the outside, this subsection shifts focus inward to examine the unsafe behaviors

that persist even when defenses are in place, particularly the system’s epistemic failures when the generator itself must evaluate untrusted content. Unlike the privacy attacks explored in the following subsection, which exploit the system’s faithful adherence to retrieved context as a side-channel for information disclosure, the safety failures examined here stem from the opposite problem: the system’s overconfidence and its unwillingness to recognize when adherence is inappropriate.

A primary safety challenge is the model’s tendency to act as a sycophantic resonator, generating fluent but false or harmful content that is grounded in—or even amplified by—the retrieved context, rather than being critically evaluated against it. This is a failure of what we might term the *epistemic vigilance* of the system. The work from [15] highlights that the quality of a RAG system is not solely determined by its generator, but is critically contingent on the relevance of the retrieved documents. However, the danger escalates from mere irrelevance to active harm when the model accepts and magnifies false information present in irrelevant or adversarial documents. The implicit assumption that retrieval equates to factuality is a fundamental alignment flaw, as the model’s expectation is to follow the retrieved evidence, even when it is non-credible or conflicts with its parametric knowledge. This is termed the "noise robustness" problem, and the authors of [55] have provocatively shown that including irrelevant documents can sometimes inadvertently enhance performance by up to 30%, but this comes at the high cost of unreliability. A more formal treatment is found in works like [30], which indicates that indiscriminately retrieving a fixed number of passages can introduce irrelevant information that diminishes the LM’s versatility or leads to unhelpful responses. The foundation of this challenge is that the model’s utility function is to generate text, not to verbalize a refusal or—worse—to identify the unanswerability of a query.

This points directly towards the second critical safety capability: the ability to perform **negative rejection**. The ability of a RAG system to abstain from answering when it determines that the retrieved evidence is incomplete, insufficient, or contradictory is a prerequisite for safe deployment in high-stakes domains. However, most modern LLMs are optimized to produce an answer, and a failure to provide one is thus discouraged by alignment fine-tuning. Adaptive frameworks such as [38] demonstrate a mechanism to route queries based on complexity, but the decision to abort is the ultimate complexity. Further, building on the concept of self-reflection [30] permits for adaptive retrieval, where the model intrinsically has a conception of confidence and knowledge boundaries. We are just beginning to see advances in using the model’s own internal confidence scores to trigger a retrieval or abstention, but these approaches apply *post-hoc* to a failure, rather than design to prevent an unsafe generation from the outset. This gap between detection and prevention is central to the system’s inability to *reliably* reject unreliable evidence, and it is this capability—more than the ability to generate any answer at all—that determines the trustworthiness of a RAG system.

The third, and perhaps most insidious, challenge in the safety of the system is the **"counterfactual robustness"**: the ability to effectively handle retrieved contexts that are factually incorrect without being derailed by them. This is the precise failure point where [59] revealed that a certain dataset property can actually cause a RAG system to perform *worse* than the LLM alone, as the generated context introduces distracting or irrelevant information that misleads an already-capable model. This indicates that the model is not performing a robust evaluation of the retrieved context, but instead cannot be manipulated by the phrasing of a query. The more powerful the retriever is at finding semantically similar content, the more likely it is to find a seemingly plausible but deeply flawed document. The lack of a functional filter against such misleading evidence is a major obstacle to deployment. For example, [105] mentions that "validation is only feasible during operation," indicating that even the inherent uncertainties are not fully addressed in the design phase, but can only be discovered in the field. Therefore, a next-generation RAG must incorporate an evaluator that can distinguish the *assessment of relevance* from the *essence of truth*—a task of equivalent complexity to the built-in reasoning

process itself.

The primary constraint, and the focal point of future research, is that safety and reliability are not properties that can be "bolted on" after the completion of the model, but must be designed into the iterative process of retrieval and generation in a self-consistent manner. This requires robust confidence estimation and calibration to distinguish between the generated content being a faithful synthesis of the retrieved evidence rather than a regurgitation of it. Without such intrinsic safeguards, the system cannot meet the epistemic standards demanded by the attacks described earlier or the privacy guarantees required later, and will remain structurally unable to withstand the organic noise and contradictions that define real-world retrieval environments.

6.5 Ensuring the Memory and Retrieval Pipeline: Privacy Risks and Data Leakage

The integration of external knowledge bases into large language models, while transformative for capability, fundamentally redefines the system’s attack surface by introducing a non-parametric memory component that generates distinct privacy threats. The external knowledge base is not merely a passive repository but an active and addressable component whose contents, access patterns, and internal representations can be probed to expose sensitive information. The privacy risks within the retrieval pipeline can be systematically informed by the broader theoretical frameworks of adversarial robustness and confidentiality - the system must ensure that neither the knowledge base, the model internals, nor the intermediate latent states are rendered vulnerable to leakage via malicious inputs or through inference on legitimate outputs [106, 84]. A successful privacy guarantee for a RAG system therefore necessitates a defense-in-depth approach that encompasses the storage layer, the retrieval mechanism, and the generation interface.

A primary class of attacks targets the membership status of specific data points within the source knowledge base, which is known as Membership Inference Attacks (MIA). In an MIA, an adversary seeks to determine whether a specific, known text passage was part of the retrieved document store. This is enabled by the fact that the output distribution of the RAG system is a function of the documents provided to the generator. As demonstrated by a black-box MIA framework, member samples, when embedded and retrieved, will hold significant semantic similarity to the model’s contextualized embeddings of its own generated text [107]. The attacker can leverage the system’s API access to compute the cosine similarity between the candidate sample and the model’s output response (which is conditioned on the retrieved context) to build a membership score. Additional features such as the language model’s perplexity on the response further inform the attack [107]. By aggregating these features, two distinct attack vectors are articulated: (i) a Threshold-based Attack that assigns membership based on a simple score threshold, and (ii) a Machine Learning-based Attack that trains a binary classifier over the extracted feature vectors [107]. These attacks achieve a high ROC AUC (82%) under black-box access, demonstrating that even without internal knowledge of the database structure or model weights, the retrieval mechanism’s output is a reliable oracle for membership disclosure. This represents a direct and significant privacy threat if the database contains proprietary or sensitive documents.

A more dangerous threat vector involves the direct inversion of the retrieval representations used to query the database. The dense embeddings used as indices are rich in information and can be exploited to reconstruct the original text of the stored documents. A malicious actor with access to the embedding space, either through a compromised API or a model-internal attack, can employ embedding inversion techniques to recover original text sequences and exfiltrate the contents of the retrieval database [106]. Such attacks are facilitated by the properties of modern retrieval models, where the embeddings can be inverted to reconstruct the input. These attacks can bypass the retrieval interface and directly reconstruct memory. The implications are severe, as they could allow an adversary to extract private information contained in the knowledge base, bypassing other security measures.

These privacy risks are exacerbated in situations where the retrieval system itself is a target of a more sophisticated attacker through *knowledge poisoning*. An attacker could inject maliciously crafted documents into the corpus to manipulate the behavior of the RAG system [106, 84]. Such attacks can not only force the generation of false content but also be used to probe the structure of the embeddings and undermine the confidentiality of the retrieval interface. For example, indirect prompt injection (resulting from content provided *into* the corpus) can cause the LLM to reveal information about the documents that were *not* retrieved or about the internal system’s behavior [84]. This dual threat highlights the severe privacy risks posed by data poisoning and injection attacks.

Defending the memory and retrieval pipeline necessitates a multi-faceted approach that addresses both the statistical and architectural surfaces. Embedding transformation is one potential defense mechanism, where a transformation is applied to the retrieved vector representations before further processing, ensuring that the vector’s information content is robust to inversion [47]. Relatedly, adversarial training of the model can aim to make intermediate representations more robust against membership inference attacks, altering the objective to make the embeddings less discriminative with respect to membership without sacrificing retrieval utility [106, 84]. To mitigate membership inference, output perturbation can be applied, such as carefully calibrating the returned top-k documents to obscure the ranking signal [8].

Also, moving beyond defense at the point of inference, we must enforce privacy at the protocol level through architectural and governance changes, where the retrieved documents are provisioned through a trusted infrastructure decoupling the embedded data from direct user interaction. More semantically, verticalizing the operation of the pipeline can implement a security-clearance-level policy for documents, enforced with a role-based retrieval, to prevent an adversary from successfully extracting private information. The privacy foundations of the RAG must be considered within the broader framework of fairness, as biases embedded in the memory can be inadvertently amplified via the reconstruction process and affect the model output [55]. Thus, privacy and traceability are interlocked with data management.

In conclusion, the memory and database component is a critical and vulnerable component of an RAG system from a privacy perspective. The increasing attack methods confirm that MIA and embedding inversion are not theoretical; they can be realized with high success in black-box settings [107, 8]. The academic roadmap must, therefore, pivot from viewing the database as a sealed vault to considering an ongoing exposure to model outputs. Future research should investigate dynamic access, differentially private levels of retrieval granularity, and secure multi-party computation for embedding look-ups to make exfiltration computationally prohibitive. Moreover, the adoption of formal privacy metrics and a rigorous cross-evaluation of adversarial robustness is crucial, as the interplay between inference and the system state currently flirts with privacy risks that lack provable security guarantees. A trustworthy RAG system must incorporate a transparency layer that allows provenance and attribution without revealing the underlying data.

6.6 Evaluating Bias, Discrimination Amplification, and Sustainability of RAG Systems

The evaluation of fairness, bias propagation, and sustainability constitutes a critical yet comparatively underdeveloped dimension of Retrieval-Augmented Generation (RAG) research. While precision-oriented metrics in retrieval engineering have progressed considerably across paradigms spanning from classic lexical matching to learned dense representations [108, 109], the systemic consequences of embedding societally biased corpora into the retrieval-generation pipeline warrant deeper analytical scrutiny. Unlike purely parametric systems where bias is a function of the training distribution, RAG systems inherit biases directly from the retrieval database at inference time. This positional vulnerability separates the fairness evaluation challenge of RAG from that of monolithic large language models, requiring a dedicated analytical

focus.

The retrieval stage introduces an information-theoretic selection process where lexical and semantic matching algorithms may inadvertently prioritize majority-group language patterns. Dense retrieval architectures that leverage bi-encoder representations and self-attention mechanisms [40] risk encoding statistical regularities in the corpus, such that documents relevant to underrepresented groups—which are rarer in the training data or contain lexical variants—are systematically down-weighted in the ranking. The granularity of corpus indexing exerts a compounding influence: chunking strategies that fragment and structure documents into isolated units may exacerbate or partially mitigate social biases, depending on whether the boundary selection permits context boundaries that reveal or obscure identity-related information [14]. This causes a substratal distribution shift where minority-group queries receive lower-quality contexts before any generative process begins, thereby transforming an evidentiary retrieval problem into downstream systemic discrimination.

Once contexts are integrated, the generative model must synthesize evidence from multiple retrieved passages plus its parametric knowledge. This step introduces *bias amplification* dynamics, where biases present in a small set of retrieved passages are not simply reflected but magnified through the model’s response generation and synthesis. This phenomenon arises because the model is designed to integrate multiple sources of contextual data [110], and the underlying LLMs are trained on data that contain correlation patterns between social attributes and outcomes. The generation model does not merely reproduce the retrieved bias; it generalizes across contexts, producing confident, believable inaccuracies and measurable, potentially disparate outcomes. The effect becomes particularly acute when the retrieved content is historically accurate but contextually exclusionary in high-stakes settings, and it directly compounds the privacy threats discussed earlier: the same retrieved evidence that can be inverted or leaked through membership inference may embed systemic skews, which the generator subsequently surfaces as authoritative, amplifying both the confidentiality breach and the downstream harm.

To quantitatively understand these failure modes, several evaluation frameworks have been proposed by LongRAG et al., but machine learning research would benefit from formalized auditing, particularly when the design space involves the competing dynamics of RAG—robustness to noise, rejection of negative evidence, and counterfactual resolution. For all these valuation approaches, a recurring limitation needs to be exposed: one well-crafted framing often includes predictive fairness only, while distributive justice or procedural fairness across the entire candidate’s retrieval corpus and answer generation remains underemphasized. This gap is particularly consequential when paired with the privacy defenses previously discussed: transformations designed to obscure embeddings may inadvertently skew retrieval relevance in ways that disproportionately harm marginalized groups, necessitating a joint optimization of privacy, fairness, and utility.

Perhaps the most under-explored aspect of the fairness of RAG lies in the global equity of its computing requirements. Recent work in efficient LLMs and dynamic retrieval methods highlights how practical RAG systems balance accuracy against latency and computational overhead [23, 38, 57]. However, training and serving large retriever-generator systems at scale has a dramatic carbon footprint, meaning that high-accuracy systems are concentrated in well-resourced actors—creating a two-tiered global access schema. This distributional inequity not only limits the reach of beneficial RAG applications but also centralizes powerful models in a few jurisdictions, a reproducibility of the trust and provenance concerns raised in the previous subsection on privacy. This may contribute to adoption harms, where retrieval knowledge bases and computational requirements are disparate, increasing the likelihood of bias outcomes across populations and regions.

In conclusion, the assessment of RAG fairness is not a peripheral consideration but an inherent element of its architectural design: the selection of external evidence and subsequent

generation can reproduce and amplify systemic biases, and this is deeply intertwined with the privacy and robustness failures already discussed. Future research must develop fairness-aware contextualization, introducing interventions at the point of knowledge insertion and ranking. Formalized standards for fairness audits are needed, with an explicit emphasis on reporting trade-offs while maintaining robust predictors. Moreover, these audits must be integrated with future privacy and security guarantees (e.g., differentially private retrieval, secure multi-party computation for lookup lookups), as the interaction between these social and technological concerns will define the trustworthiness of the next generation of RAG systems. Unless we consider these socio-technical dynamics in a unified evaluation framework, RAG systems risk deployment that not only underestimates its own bias failures but also reinforces the very inequities that contextual knowledge could otherwise help correct—echoing the broader theme that the trustworthiness of RAG is inseparable from the interplay of its memory, access, and generation layers.

7 Evaluation and Benchmarking for the Retrieval-Augmented Paradigm

7.1 Component-Specific Evaluation Metrics for Retrieval and Generation

The evaluation of retrieval-augmented generation (RAG) systems necessitates a dual-pronged assessment strategy that examines the constituent retrieval and generation modules independently while acknowledging their deep functional interdependence. Component-specific metrics provide essential diagnostic granularity, enabling researchers to isolate performance bottlenecks and guide targeted improvements. For the retrieval module, the field borrows a mature set of metrics from classical information retrieval. Metrics such as Recall@K, Precision@K, Mean Reciprocal Rank (MRR), and Normalized Discounted Cumulative Gain (NDCG) are routinely used to quantify the retriever’s ability to surface relevant context passages. Recall@K, which measures the proportion of relevant documents retrieved within the top-K results, is favored for assessing the comprehensiveness of the retrieval stage, ensuring the retriever casts a sufficiently broad net to capture all essential evidence for downstream generation.

For the generation module, evaluation metrics generally fall into three categories. Traditional token-overlap metrics, including BLEU, ROUGE, and METEOR, are widely reported; however, their focus on lexical matches fails to capture semantic equivalence or the factual correctness of a response that is linguistically fluent. To address this shortcoming, embedding-based semantic similarity measures, such as BERTScore and vector cosine similarities, have become standard. Crucially, new metrics for factuality and faithfulness, which assess the degree to which generated claims are consistent with retrieved evidence, have been introduced, for example, through the RAGAS framework’s "faithfulness" dimension [111, 25]. These model-based metrics leverage a large language model to judge the generated content and are increasingly essential for evaluating robust system behavior in performance-focused settings.

The most critical limitations in a component-specific evaluation often arise from misalignment. The use of traditional IR relevance labels fails to capture downstream utility; a document that a human judge may deem only partially relevant to the query might contain the exact evidence needed for the generator to reason and produce a correct answer, and vice versa. This "isolation" of the separate assessments results in a misleading sub-systems evaluation, where high retrieval scores do not necessarily correlate with generation fidelity. As emphasized by research on evaluating RAG systems [77], a doc-level downstream effectiveness evaluation, where each document is individually passed through the language model and its utility is scored against the ground truth, offers far more relevant signal, demonstrating a higher correlation (improvements in Kendall’s Tau between 0.168 and 0.494) with the final system’s performance [53]. Therefore, the standard metrics are only useful when diagnosing a specific failure point—for example,

determining whether poor generation stems from an inability to retrieve (a retrieval failure) or the generator’s inability to reason over provided context—rather than for optimizing the RAG system towards its final outcomes.

In fact, the field is increasingly aware that a key bottleneck for RAG systems can lie in the generator itself. Even when the relevant context is provided, the model may not fully leverage it; the model may ignore information contained in the input, particularly if it conflicts with memorized parametric knowledge [4]. This challenge is compounded by the fact that a model’s ability to use the provided context is often not robust; and can be dependent on the position of that information within the input [28, 61]. This further emphasizes the need for diagnosing not only the quality of what is retrieved, but also the degree to which the model can be grounded in that evidence to better distinguish inadequate retrieval from failed retrieval-augmentation. While the incorporation of generation-awareness into the retrieval process is a promising area of active research, evaluating it through rigorous and robust metrics remains a critical open challenge in the rg.

7.2 End-to-End Evaluation and Multidimensional Frameworks

Building upon the component-level diagnostic metrics discussed previously, which isolate failures in either retrieval or generation, the interdependent nature of these modules necessitates evaluation frameworks that assess the RAG pipeline as a unified system. The emergence of the Retrieval-Augmented Generation Benchmark (RGB) has established a foundational approach by decomposing evaluation into four fundamental abilities: noise robustness, negative rejection, information integration, and counterfactual robustness [70]. This capability-oriented design reveals significant performance variability across models—while LLMs exhibit a certain degree of noise robustness, they struggle markedly with negative rejection and the integration of conflicting information [70]. This disentangled capability analysis, replicated in related work [77], aligns directly with the diagnostic logic advocated in the sub-system evaluation discussion: rather than relying on aggregate scores that mask underlying failures, it allows researchers to identify whether a system’s failure originates from retrieval inadequacy or from the generator’s inability to reason over conflicting contexts—a distinction that the overlapping decomposition of retrieval and generation failures made evident.

Building on this foundation, multidimensional evaluation frameworks have expanded the evaluation space to capture more nuanced aspects of system behavior—an evolution that mirrors the transition from isolated, token-level metrics to capability-oriented judgments. Testbeds like CRUD-RAG, for instance, categorize RAG applications into create, read, update, and delete operations, generating distinct datasets for each scenario to assess how well the system generalizes across application types [112]. Similarly, for multi-hop query comprehension, the MultiHop-RAG benchmark has been introduced to test reasoning over multiple pieces of supporting evidence—a task where existing RAG systems consistently underperform—revealing a generic inadequacy in retrieving and synthesizing distributed information [113]. Beyond task diversity, automated, reference-free evaluation frameworks such as RAGAS have been developed to quantify different dimensions like faithfulness and answer relevance without relying on ground-truth comparisons but by comparing generated claims with retrieved evidence [111]. These developments anticipate the automated, LLM-as-a-judge paradigm discussed in the next subsection, and in doing so, they introduce a dependency on the judge model’s reliability—a dependency that, as will be elaborated in the following subsection, constitutes a central methodological concern for the credibility of RAG evaluation.

A particularly challenging yet vital evaluation dimension within this end-to-end perspective is the handling of adversarial or misleading evidence. An effective RAG system must not only retrieve relevant information but also possess the ability to abstain, demonstrating robustness for negative rejection when the retrieved corpus is incomplete, contradictory, or inherently false [70]. This requires the retriever-generator to be counterfactually robust, capable of recognizing

that the top-ranked contexts contradict the query, and knowing when to decline an answer to maintain its reliability [77, 101]. Crucially, this evaluation dimension connects directly to the insertion of the self-reflective mechanism introduced earlier: assessing **refusal** requires the system to distill its own knowledge against the quality and sufficiency of the incoming documents, and to course-correct its final response based on that self-assessment [30]. The integration of such self-reflective and evaluative loops can act as a dynamic correction signal, fostering a synergistic interplay where the critic and generator work together to produce more faithful and contextually grounded output [12]. This evaluation of negative rejection and self-reflection bridges the capability-based metrics above with the structured feedback mechanisms (e.g., reflection tokens) that LLM-as-evaluator systems employ in the next coverage; both the human-LLM hybrid validation strategy discussed there is also relevant for verifying the robustness of these frameworks.

A central tension in end-to-end evaluation remains the relationship between retrieval accuracy and the actual downstream generation quality. Standard relevance metrics for retrieval, such as Recall@K and NDCG, can be misleading, as they do not account for what the language model truly requires from a context window. In fact, including documents that are not relevant to the user query can sometimes surprisingly improve final answer accuracy by more than 30%, naturally challenging the presumption that a better retrieval score always implies better RAG performance [55]. This points to a deeper misalignment between the **relevance** that a retriever is trained and evaluated against, and the **utility** of that context to an LLM—a misalignment that was already raised in the component-specific evaluation of the previous subsection where the necessity of a document-level downstream utility evaluation was emphasized. Recent work in the multimodal domain suggests that using auxiliary readers or the LLM itself to rerank and filter the output of a frozen retriever can be more effective for the downstream task [12, 52]. Yet, the final answer should be *grounded in* sources—integrated and evaluated together [12]. This grounding requirement aligns with the LLM-as-a-judge strategies used to enforce factuality in the following subsection, as the judge must verify the consistency of the final answer against the retrieved evidence.

The evaluation landscape is thus moving toward multi-dimensional, capability-based analyses that probe a system’s failure modes in a quantifiable and reproducible way. Future advancements will necessitate the development of benchmarks that are private and dynamically updated to avoid data contamination from the increasing presence of LLM-generated content online [77, 114]. This concern over contamination and evaluator drift is central to the LLM-as-a-judge discussion in the next subsection, as it underscores the need for continuous validation of the evaluation infrastructure. Additionally, the integration of efficiency metrics, such as token consumption and latency, into these multidimensional frameworks is becoming essential for practical deployment, offering a holistic view of system capability against computational costs and real-world sustainability [37]. By combining capability-based probing, robustness testing, and efficiency-awareness, end-to-end evaluation can provide a more complete and actionable picture of RAG system performance, complementing both the system diagnostic granularity of component-level metrics and the scalable, model-driven evaluator mechanisms discussed in the following subsection.

7.3 The Rise of Human and Model-Based Assessment in Benchmarking

The evaluation landscape for retrieval-augmented generation (RAG) has undergone a profound transformation, shifting from a paradigm dominated by labor-intensive human annotation and evaluation toward a hybrid ecosystem in which large language models (LLMs) serve as both annotators and judges. This transition is not merely an issue of convenience but a response to the fundamental scale mismatch between the combinatorial explosion of hyperparameters, retrievers, prompts, and model families in modern RAG pipelines and the inherent cost of human assessment. The core challenge lies in reconciling the reliability of human judgment with the

scalability and reproducibility promised by automated, model-based evaluation.

Historically, the gold standard for evaluating RAG systems was grounded in human evaluation of both retrieval relevance and generation quality, often conducted through the lens of Crowdsourced annotation [115]. While these human-in-the-loop protocols offer high fidelity, they are not sustainable for rapid model iteration. Recent work has confirmed that semi-automated annotation with LLMs—as proposed in systems like AnnoLLM—surpasses or at least matches the performance of crowdworkers on critical tasks such as relevance judgement, BoolQ, and WiC [115, 95]. This has lowered the barrier to building high-quality datasets, yet it introduces a critical dependency: the degree to which the annotator model itself is biased or limited relative to the system under evaluation.

This dependency manifests most prominently in the “LLM-as-a-Judge” paradigm, where a reward or critique model is used to offer scalar or fine-grained scores on final and intermediate outputs. This is particularly valuable for long-form generation and knowledge-grounded tasks where reference-free assessments are mandatory. For instance, work on evaluating LLMs for relevance judgment has analyzed the utility of using the LLM to generate absolute ratings versus pairwise rankings and shows that the latter are far more robust [116]. Furthermore, RAG models can be assigned factuality scores to enforce grounding by retrieving a set of supporting documents, either through dense search or a forward-looking internet search, and then judging the consistency of the answer against these sources [63, 24]. When used to rank or select between competing decompositions of the context, LLM feedback has also served as a training-time reward for the retriever, demonstrating the utility of LLM-generated labels as a knowledge signal for pipeline coordination [52].

The practical implementation of LLM-as-a-judge experiences depends heavily on structured feedback mechanisms, moving beyond single scalar scores. For example, in Self-RAG, the LM is trained to produce special reflection tokens; during evaluation, these tokens act as internal critique designed to self-govern retrieval necessity, relevance, and support of the generated response, enabling controllable and interpretable evaluation [30]. A related approach decomposes the evaluation of biomedical QA into a step of retrieval and self-reflection, including checks if the final answer is grounded in a retrieved document or relies purely on the model’s own parametric knowledge, yielding atomic factuality signals for downstream scoring [117]. Similarly, evolving outputs and instructions are combined with an automated critic to reduce hallucination, which is an evaluation method that integrates well with self-assessment-confidence measures [73].

The combination of automatic evaluation and RAG also improves the feasibility of probing for the model’s boundary of knowledge. To decide when additional facts are needed, LLMs can rewrite the same question in multiple languages, and the consistency across responses becomes a signal of uncertainty [118]. This follows the principle that leveraging the model’s own uncertainty is a way to efficiently generate supervised or pseudo-supervised data for retrieving augmentation triggers [119]. The deeper issue is that measurement of the inherent overconfidence of LLMs or the purposeful for detecting when to abstain from answering can be re-formulated as an automatic evaluation task, a step that reduces costly human evaluation.

However, the use of LLMs as evaluators still holds significant pitfalls, particularly in terms of systematic biases, self-preference and position bias. Given that evaluation LLMs offer pairwise ranking of documents, studying the non-uniformity of their internal training, such as the model’s relative preference for generated and retrieved contexts, indicates a hidden preference for the generated context [94]. Researchers must account for these “subjectivity” of the judge, and, as for the earlier results in the relevance judgment [116], caution that a hybrid human-machine collaboration spectrum where massive LM annotations and human spot-checks are combined is the most defensible strategy against silently corrupting evaluation protocols.

Looking forward, assessment must also solve the problem of when to retrieve, studying which questions benefit from RAG and which do not. The degree of popularity of the entity or domain improves isolation of when retrieval is necessary [4]. In this complete evaluation view, the

LLM-as-a-judge is not finalized, but a benchmark tool that allows to verify the generation step for fine-grained judgments is scalable—provided that the judge itself is continuously validated against small but high-quality human-created test benches, including both system robustness (e.g., rejection) and clean relevance. The balance will be decisive for bridging end-to-end application value with scalable model-based assessment.

7.4 Benchmarking Infrastructures, Real-World Fidelity, and Contamination Control

The evaluation of retrieval-augmented generation systems has progressed from isolated, ad-hoc metric reporting toward standardized, modular infrastructure. Yet, the field’s rapid evolution has outpaced the development of robust benchmarking ecosystems, creating a critical gap between the experimental claims of individual methods and the reproducible evidence required for credible scientific advancement—a gap that directly undermines the trustworthiness of the LLM-as-a-judge paradigm discussed previously. If even the most carefully calibrated evaluator models rely on datasets that are contaminated, static, or misaligned with real-world conditions, the diagnostic utility of the evaluation loop itself collapses. This growing recognition has led to a consensus that conventional evaluation protocols, which often rely on static, publicly-available datasets, are becoming fundamentally ill-suited to assess systems in a dynamically shifting landscape [12, 120].

A significant portion of this methodological evolution is driven by the emergence of unified toolkits designed to address the reproducibility crisis. Frameworks such as FlashRAG exemplify this trend by offering a modular architecture that standardizes the implementation of retrievers, generators, and augmentation strategies, alongside 32 pre-processed benchmark datasets and 12 advanced RAG methods, ensuring that different systems are evaluated under identical, comparable conditions within a unified research framework [121]. Such infrastructure not only reduces the load of reproducing existing work but also establishes a canonical baseline for future scientific contributions, moving beyond monolithic libraries that often sacrifice experimental rigor for flexibility [121]. Similarly, specialized frameworks like the MedRAG toolkit in the medical domain emphasize the role of standardized, domain-specific infrastructures that produce reproducible results across diverse corpora, retrievers, and LLM backbones, addressing the challenge of heterogeneous data and terminology [59]. These systematic toolkits provide the substrate for building coherent metrics and protocols that can be adapted and extended, yet they also present a constraint: by standardizing a curated set of configurations, they may inadvertently narrow the space of viable experimental designs—and thereby the range of empirical results the field can accumulate.

Reproducibility itself, however, remains contingent on the integrity of the underlying data. A particularly insidious threat to validity is the prevalence of evaluation contamination, where benchmark datasets are included in the foundational corpora of web-scale LLMs, violating the assumption that external documents are novel evidence to be retrieved [12]. This means that a model can implicitly memorize the answer to benchmark queries during pre-training, rendering the retrieval component a mere façade of grounding. Traditional solutions, such as dynamic corpora that are continuously updated, may offer the most ethical and practical path in the short term, but they are resource-intensive and difficult to maintain at scale. In this context, the need for “private” benchmarks—constructed from non-public or contemporaneously-new interventions, curated by expert evaluators, and managed in controlled environments—is increasingly urgent. The rapid obsolescence of public benchmarks reflects their weak correlation with real-world fidelity: often the top-performing lab models experience unforeseen failures when validated against target conditions in production environments where robustness is only truly testable during operation, rather than in controlled labs [105].

Given this, a stronger coupling between evaluation design and deployed operational constraints is necessary. Static offline benchmarks routinely miss key operational factors such as database

updates, shifting query distributions, and the dynamic behaviors of agentic frameworks [77]. While modern agentic workflows—such as Self-Reflection in Self-RAG [30] or iterative retrieval during reasoning in IRCOT [72]—showcase strong results on classic QA datasets, there is limited empirical evidence on temporal divergence or deployment-specific degradation. A complementary corrective is the recent integration of retrieval-specific metrics, as proposed in eRAG, which reveals that conventional relevance metrics correlate poorly with end-to-end task performance [53]. This underscores that future benchmarks must move beyond single-number summaries to capture performance nuance—for instance, by unpacking latency across query complexity tiers, as in Adaptive-RAG, and by incorporating failure modes such as hallucination, missing citations, and refusal behavior when evidence is insufficient [30]. In parallel, model-based evaluation has matured into frameworks such as RAGAS and ARES, which provide multi-dimensional scoring across context relevance, answer faithfulness, and answer relevance, replacing reliance on gold labels with algorithmic judges [111, 122]. Yet as the previous subsection cautioned, LLM-based metrics are prone to systematic biases and hallucinated preferences, necessitating the integration of human validation—an imperative that frameworks like RAGElo aim to address by aligning human and model judgments through an Elo-based ranking interface [123].

Looking forward, a robust benchmarking infrastructure for RAG requires a fundamental shift towards systemic resilience. First, the field must deliberately migrate away from unrestricted public datasets as the sole yardstick of progress, adopting a controlled-release model where private test sets are safeguarded, refreshed on a prescribed schedule, and curated to prevent contamination. Second, evaluation scope must expand to include operational costs and latency, as demonstrated by systems like HybridRAG, which integrates knowledge graphs with vector retrieval and signals that performance gains must be weighed against computational overhead and deployment practicality [46]. Ultimately, the field’s methodological maturity will be defined not by the volume of benchmarks, but by its commitment to transparent protocols, contamination avoidance, and the integrity of evaluator-driven metrics that bridge offline capability with online operational success. It will require future generations of adaptive assessment frameworks as living sociotechnical systems—continually updated, tested, and refined within the production ecosystems they aim to serve, ensuring that the scalability of LLM-based evaluation is never decoupled from human oversight and contextual validation.

8 Applications, Extensions, Emerging Frontiers, and Future Directions

8.1 Knowledge-Intensive Vertical Applications and Continuous Adaptation

The deployment of retrieval-augmented generation in knowledge-intensive verticals represents one of the most consequential translations of the RAG paradigm from research laboratories to production environments. Unlike general-purpose question answering, domains such as healthcare, law, finance, and telecommunications impose a tripartite set of constraints: highly specialized vocabularies that challenge generic embedding models, accuracy requirements with severe consequences for failure, and knowledge landscapes that evolve on timescales ranging from months to minutes. These demands have catalyzed a shift from generic retrieve-then-read pipelines to tailored RAG architectures engineered for domain-specific transparency and reliability.

In clinical and professional decision support, the central challenge lies in aligning the retrieval subsystem with domain-specific ontologies and terminologies. Evaluations such as the Medical Information Retrieval-Augmented Generation Evaluation (MIRAGE) benchmark have shown that RAG significantly improves the accuracy of models like GPT-3.5 and Mixtral—elevating them to GPT-4-level performance—owing to the presence of structured medical corpora and appropriate retrievers [59]. The engineering imperative in such settings is not merely to retrieve

relevant passages but to do so via fine-tuned bi-encoders and rerankers that are cognitively aligned with professional reasoning patterns. Notably, the integration of knowledge structures—specifically knowledge graphs—has demonstrated substantial advantages over text-only indexing, as evidenced by deployments in customer service where semantic structure is critical for precise issue resolution and has served to reduce the median per-issue resolution time [19]. This trend underscores the broader insight that in structured verticals, the relational semantics of documents are as significant as their lexical contents. The most effective systems augment the non-parametric memory not just with text chunks but with a structured graph of entities and dependencies which facilitates retrievals over multi-hop relationships that are otherwise opaque to flat vector search [35]. Hybrid retrieval approaches—which combine the flexibility of dense vector search with the deterministic precision of knowledge graph traversal—exhibit superior performance compared to either mode. In the financial domain, for instance, the HybridRAG framework combining KG-based retrieval and vector-based retrieval has demonstrated superiority at both retrieval and generation stages when applied to complex, Q&A-format documents like earnings call transcripts, where the interplay of structural information and domain-specific terminology typically degrades the performance of purely vector-based systems [46].

A second pillar of vertical RAG systems is architectural sensitivity to document structures. Documents common in professional settings—particularly financial reports and telecommunication standards—rarely present as uniform free text. This recognition has given rise to structure-aware chunking strategies that dissect documents based on structural constituents such as headings, tables, and bullet points, moving beyond the fixed-window and paragraph-level uniform segmentation that dominate general-domain RAG. Research indicates that structure-aware chunking yields substantial improvements in downstream QA performance by reducing the information fragmentation that occurs when tables or multi-element lists are split arbitrarily, thereby improving the capacity of the retriever to provide complete and coherent evidence to the generation module [14]. Similarly, for the highly technical and ever-evolving domain of telecommunications, which is characterized by intricate 3GPP standards documents, standard chunking is insufficient. Novel frameworks addressed these challenges via a hierarchical processing in which document-level navigation is integrated with chunk analysis, highlighting how structural parsing facilitates the retrieval of deeply nested and cross-referenced technical specifications, particularly because a strict hierarchy of nested, multi-part specifications renders unstructured segmentation inadequate [33].

The final, and foundational, pillar for knowledge-dense domains is the mechanism for continuous adaptation. Given the rapid evolution of medical practice, legal jurisprudence, and financial reporting, static corpora quickly lose relevance. Hence, self-updating RAG systems are beginning to emerge, integrating online adaptation schemes without the exorbitant cost of full model retraining. One approach—Memory of Amortized Contexts—illustrates this by compressing new documents into compact modulations stored in a memory bank, allowing the model to adapt its behavior on the fly via a single forward pass for the question. This alleviates the inherent computational staleness of RAG and shows robustness against performance degradation after nearly a million updates while preserving knowledge-retention capabilities [67, 68]. These memory-dynamic processes serve as a foundation: rather than viewing the retriever and generator as static modules updated by offline scripts, the system itself facilitates a continual process of information assimilation. While dynamic memory pools effectively address the "how" to adapt, the "when" is an equally important design consideration. One perspective is how long Okary reports coincident a CoM 4K token-based RAG retrieves larger-grained units, thereby reducing the systems burden on the retriever and enabling processing of full documents as single units [45]. Here, the retriever is not coupled with a large-scale, finer-tuned search to generate massive contexts; instead, evaluation emphasizes identifying the proportional balance between retrieval granularity and the efficiency of the generation process.

Ultimately, vertical applications are forcing the RAG ecosystem to mature. The future of

domain-specific RAG lies in the dynamic orchestration of heterogeneous knowledge sources—via hybridizing structured graph traversal with dense semantic search and enabling evolving, self-motivated mechanisms that can assess trustworthiness. The [124] is shifting from a passive function of retrievers to an active contextual planner aligning itself with the evolving knowledge landscapes of the most knowledge-critical industries, continuously balancing precision against latency, and cost against reliability.

8.2 Multimodal and Multi-Source Extension

The extension of retrieval-augmented generation beyond unimodal text represents a fundamental paradigm shift, moving from retrieving discrete textual passages to orchestrating heterogeneous sources—images, video, code, and structured knowledge—that jointly condition the generation pipeline. This broadening is motivated by the intrinsic multimodality of human knowledge and the growing demand for applications that synthesize information from diverse, complementary representations. Early systems laid groundwork by aligning entities across textual and visual domains, but modern frameworks now integrate large-scale multimodal language models (MLLMs) directly into the retriever-generator loop, leveraging their zero-shot, in-context generalization capabilities to unify the previously separate strands of multi-source retrieval discussed in the context of structured verticals [8]. Whereas the preceding subsection emphasized how domain-specific constraints shape *how* retrieval is executed—via knowledge-graph traversal, structure-aware segmentation, and self-updating memory pools—this section addresses a complementary axis: the *breadth* of what is retrieved, namely the expansion from text to heterogeneous modalities, and how these diverse evidence types must be aligned, indexed, and integrated to support robust generation.

A central architectural concern in multimodal RAG is cross-modal alignment—the construction of unified representation spaces where text, images, video, code, and structure triples are indexed, embedding, embedding, and retrieved over heterogeneous databases [8]. Alignment strategies provide a stable basis for retrieving and generating passages and images together, but they reveal substantive challenges: multimodal retrieval introduces substantial challenge of irrelevant or noise correspondences, and the index must preserve fine-grained semantics of the complex data structures [15]. Modern approaches employ instruction-tuned LLMs as zero-shot rerankers to filter out irrelevant and noisy candidates after a coarse first-stage dense retrieval, and show enhanced and context-conditioned relevance; and further propose a correction strategy for the loss of fine-grained semantics [15].

The design space has further expanded to handle code, video, and structured data, which are most naturally retrieved around syntactic structure and semantic boundaries, not unstructured lexical tokens. For code repositories, the unit of retrieval is a function or block, where methods that discretize source code into focused, semantically dense units directly condition generation for repair and generation tasks [1]. For video retrieval, systems index moments enriched with transcribed speech and visual semantic metadata, enabling video question answering in the temporal domain. In the structured-data domain, knowledge graphs and relational inference chains—already highlighted as critical for vertical applications—are integrated as well, offering strong contextual signals for multi-hop and fact-intensive reasoning [35]. These paradigms share a perspective: they prioritize the retrieval of semantically complete units over flat text fragments, aligning with the structure-aware segmentation that proved critical for documents in professional and technical domains.

A major open challenge that persists—and that directly parallels the accuracy and reliability demands introduced in vertical RAG—is the robustness of the generation model to noise in the retrieval results, specifically the uncertainty in the relevance of retrieved items across modalities. Noise heterogeneity exists across modalities and granularities, and recent multimodal systems have incorporated noise injection into generator training to better approximate real-world retrieval-error distributions, mitigating hallucination [125]. Other proposals use MLLM-derived

relevance scores to generate feedback signals for adaptive, round-wise retrieval as a form of self-reflective evidence grounding [30]. In such systems, self-evaluation gates augmentation, deciding whether retrieval is needed and whether the retrieved multimodal evidence is sufficient [4]. The use of attributed and traceable multimodal retrieval contributes to the broader principles of trustworthy generation: finalized outputs can be linked back to the evidence, aligning with the verification-level needs of verticals and agentic workflows where the model must decide not only what to retrieve, but when and how to integrate it [16].

Looking ahead, the next generation of multimodal and multi-source extensions will hinge on the resolution of two primary bottlenecks echoing the open challenges of vertical and agentic RAG: end-to-end differentiation of retrievers combined with heterogeneous generators, addressed via stochastic gradient relaxation strategies [89]; and the systemic trade-off between bandwidth, latency, and quality. We expect the trajectory to lead toward mutual alignment of models more dramatically than in vertical deployments, where diagrams and cross-modal structures transition from simplistic (and potentially noisy) retrieval to computed indexes and query-time optimization, returning richer and more verifiable evidence. Unified tools, controlled multimodal grounding, and verification-level fact-checking in multimodal content will become the frontier for robust deployment beyond text, informed by the capacity as well as task complexity and noise tolerance. This shift positions multimodal RAG not as a separate endeavor but as an integral step in the progression toward agentic, self-reflective retrieval that the field now demands.

8.3 Agentic Workflows, Interaction, and Emergent Reasoning

The evolution of retrieval-augmented generation has progressed from static, single-shot pipelines to dynamic, agentic systems in which the language model functions as an autonomous controller, orchestrating retrieval, reasoning, and tool use in an interleaved fashion. This transition represents a fundamental shift in the architectural philosophy of RAG: rather than treating retrieval as a one-time preamble to generation, contemporary systems embed retrieval within an iterative decision-making loop, enabling the model to adaptively determine what information is needed, when to acquire it, and how to integrate it with its own parametric knowledge. This paradigm, often termed agentic RAG, transforms the model from a passive knowledge consumer into an active epistemic agent capable of decomposing complex queries, formulating sub-queries, and synthesizing evidence from diverse sources to achieve robust reasoning outcomes.

At the core of agentic RAG is the ability to interleave retrieval with generation across multiple steps, a design principle that enables models to handle multi-hop reasoning tasks by progressively refining their information needs. Iterative frameworks such as Iter-RetGen [29] exemplify this synergy by using the model’s preliminary outputs as informative signals for subsequent retrieval rounds; the generated text reveals what the model already knows and what remains missing, thereby guiding more targeted and contextually relevant retrieval. Similarly, Retrieval-Augmented Thoughts (RAT) [126] applies this iterative retrieval during chain-of-thought generation, revising each thought step with freshly retrieved information, achieving substantial gains in code generation, mathematical reasoning, and long-horizon planning tasks. These methods collectively support an adaptive query planning mechanism that reduces the need for manual reasoning engineering while enabling efficient multi-step information acquisition, as the system can refine its search strategy based on intermediate outputs.

The adaptive retrieval strategies inherent in agentic workflows address a critical bottleneck in traditional RAG: the ability to determine when retrieval is necessary. Frameworks such as Self-RAG [30] propose adaptive retrieval mechanisms to determine whether to retrieve and how to use the retrieved knowledge [30]. Self-RAG [30] first introduces a method that enables the model to selectively retrieve and reflect on its own outputs. Building on this, Adaptive-RAG [38] introduces a classifier to estimate query complexity, enhancing the adaptability of the retrieval strategy. Related work on knowing when to trust parametric vs non-parametric memories reveals the crucial role of entity popularity as a predictor of models’ knowledge boundaries,

informing algorithms that selectively retrieve [4]. This research parallels the insights from RA-ISF (Retrieval Augmented Iterative Self-Feedback) [73], specifically its learning to answer more accurately by iteratively improving the retrieval efficiency.

A more advanced manifestation of agentic RAG is multi-step orchestration, where multiple specialized components cooperate and interact to solve tasks by decomposing them into a series of subtasks. RA-ISF [73] proposes a systematic framework with feedback loops [73]. Specifically, RA-ISF decomposes the original problem, retrieves relevant documents, and generates a response [73]. Furthermore, the same principle of decomposition and self-correction is at the heart of RAT, where the reasoning chain leverages the interleaving of thought and retrieval to enhance reasoning [126]. Such techniques enable the model to align its retrieval behavior based on its own outputs.

The integration of these multiple components is often achieved through the development of a more sophisticated planning process. This idea is embodied in the methods for response-selected strategies: a language model can determine if a given query requires no external knowledge, a limited retrieval, or a more complex, multi-step search [37]. This reactive and self-aware model places the LLM in the center of the planning, deciding not only when to adopt a query, but also whether to invoke a retrieval or reasoning strategy [37]. Looking to the future, there is a strong shift toward combining these tools with strategic planning and even a more nuanced, selective retrieval to further enhance the robustness of these systems. Self-RAG’s reflection tokens also represent an emergent but fast-moving direction, providing a mechanism for the language model to critique its own output with regard to the retrieved evidence [30]. What remains clear is that the shift to agentic models moves the generation from per-query retrieval to active knowledge manipulation, and thus the central challenge is not merely increasing retrieval accuracy but more fundamentally developing meta-strategies to know when, how, and what to retrieve for a robust generation under varying tasks.

8.4 Emerging Frontiers and Future Research Directions

As the field matures beyond the component-wise optimization and agentic orchestration discussed above, the research agenda for retrieval-augmented generation is being reshaped by emerging frontiers that collectively move the conversation from “how to build better pipelines” toward a deeper interrogation of the theoretical underpinnings, training methodologies, and evaluation integrity of the paradigm itself. While agentic systems have demonstrated remarkable empirical success in dynamically controlling retrieval, the community is now confronting the fundamental limits of this largely heuristic-driven progress, prompting a shift toward more principled and robust formulations.

A significant frontier lies in the transition from the task-specific fine-tuning that underpins current agentic systems toward sophisticated, multi-objective optimization strategies that deeply integrate retrieval and generation. While end-to-end training has been explored via relaxations and policy-gradient formulations [89], a pressing need exists for learning frameworks that explicitly model noise and irrelevance. The dominant approach of retrieving and then processing a fixed number of passages, regardless of quality, remains fragile; current methods can demonstrably degrade in performance when provided with irrelevant contexts [76]. Future work must move beyond task-specific fine-tuning to build denoising and robustness directly into the objective, such as by generating and learning from synthetic noise [76]. Furthermore, integrating human preferences and reinforcement learning to align RAG-adapted systems with diverse user expectations, rather than a singular “correct answer,” represents a core unmet challenge that extends the meta-strategies of agentic control into the training regime itself [52].

Another critical frontier is the rethinking of the retrieval mechanism, which complements the agentic orchestration discussed previously by asking not just when and how to retrieve, but what constitutes the foundational retrieval operation. The field is with a paradigm shift, moving from the adaptive retrieval strategies of the agentic paradigm to a more radical redefinition

of the retriever as a generative process. The emerging generative retrieval paradigm, where documents are decoded directly [127, 127], challenges the very notion of a separate index and the external search systems so integral to agentic workflows. A key research direction lies in unifying this generative retriever with the generation model as a single, coherent sequence-to-sequence architecture [128], potentially eliminating the need for complex external retrieval components entirely. Future research may also extend the adaptive retrieval of agentic frameworks, going beyond the retrieve-then-read paradigm to “ask” the model not just for what information is missing, but for the nature of the missing knowledge itself, making the process truly metacognitive.

A third, and arguably fundamental, frontier is the development of a theoretical foundation for RAG, which directly addresses the reliance on empirically derived meta-strategies and self-monitoring mechanisms inherent to agentic systems. Current success is largely anecdotal, characterized by empirical comparisons on specific benchmarks rather than derived guarantees. The literature needs a framework that can formally bound the expected utility of a system under varying noise conditions, or analytically predict when a refusal to answer is superior to an educated guess. Such a theory could not only resolve the "long-context vs. RAG" optimization problem analytically [11, 12]. Establishing this foundational rigor would turn the art of designing agentic RAG architectures into an algorithmic science, guiding when to retrieve and when to rely on parametric knowledge.

Finally, as RAG systems become more agentic and theory-driven, the instrumental instrumentation for evaluation faces a corresponding "arms race" against contamination and overfitting. As large language models are trained on increasing volumes of data, the risk that they memorize and recall common benchmark answers has risen sharply, undermining the validity of RAG assessments. This challenge extends to the need for dynamic, constantly created benchmarks for evolving contexts, and private, inaccessible evaluation sets [111]. The future "ground truth" for evaluating agentic RAG systems must therefore be robust to these shifts and capable of measuring not only a system’s ability to synthesize the relevant evidence, but also its ability to correctly refuse and abstain when new evidence demands it—a reflection of the self-corrective loop inherent in the most advanced agentic designs [129].

In summary, the future frontier for RAG is a multi-pronged swing: (1) the evolution of training towards dynamic, self-monitoring architectures that are robust to noise, learning from the agentic paradigm’s meta-strategies [130], (2) a merging of retrieval within a generative framework and unifying its operation with the language model itself, (3) a theoretical foundation that formalizes the heuristics of agentic control, and (4) rigorous, robust evaluation and benchmarking methodology that stands up to contamination and domain shift. The novelty and effectiveness of next-generation systems will depend on their ability to negotiate and advance all four simultaneously.

9 Conclusion

This survey has traversed the multifaceted landscape of Retrieval-Augmented Generation, charting its evolution from foundational concepts to the sophisticated, modular systems that define the current state of the art. Our synthesis reveals a field in a state of dynamic maturation. The progression from early semi-parametric prototypes like k-nearest-neighbor language models and memory networks [9] to large-scale integrated systems [8] is not merely a matter of scale, but a fundamental shift in how we address knowledge-intensive tasks. The initial, often monolithic, approaches have given way to the modularized design of the **ModularRAG** framework [21], transitioning from a simple retrieve-then-generate pipeline to a highly diverse design space that integrates multiple components with different optimization strategies. Across this survey, a recurring theme is that while RAG effectively mitigates the limitations of parametric-only models [3], its performance is acutely sensitive to the orchestration of its constituent modules—from the granularity of corpus chunking [14] and the efficacy of the retrieval engine to the sophisticated

alignment of the retriever with the generator’s utility [53]. This synergy is critical: naive system design can lead to failures or degraded performance [105], while a well-orchestrated architecture enables the model to handle complex, multi-hop tasks by decomposing them into smaller, iterative steps [29]. A critical insight is that retrieval decisions significantly affect the usefulness of RAG models for counterfactual facts [4].

The field’s trajectory reveals a continuous negotiation between architectural complexity and practical efficacy. While iterative and recursive methods offer clear advantages in complex reasoning, they introduce new challenges in cost and supervision, raising important questions about the optimal balance between performance and efficiency [27]. The advances in agentic systems and modular toolkits show a promising direction in this context, enabling more adaptive retrieval planning and dynamic orchestration [21, 121]. These advances are augmented by system-level designs, such as **PipeRAG** which optimize for efficiency through algorithm-system co-design [66]. Furthermore, the survey has underscored the crucially important role of evaluation. The benchmarks and multi-dimensional frameworks, such as RAGAS [111] and the Retrieval-Augmented Generation Benchmark (RGB) [70], have grown in sophistication beyond traditional metrics, driving the need for more capable systems that can be validated on all aspects. As the field increasingly relies on optimizing LLMs, we are observing a turn toward more holistic approaches seeking to balance performance with efficiency and grounding.

Looking ahead, the conceptualization of RAG is fundamentally shifting, and our last observation is that the issue of robustness is becoming as important as performance. The key to unlocking the full potential of RAG lies in moving beyond treating it as a purely structural component. Instead, the next generation of systems must integrate retrieval as a core strategic element that interacts synergistically with the intrinsic capabilities of the LLM. The development of **self-reflective** and feedback-driven systems [63](<https://example.com/>), which can refine their own output, is a prime example of this new direction. The future will likely focus on dynamic fine-tuning and training strategies, where RLHF and direct preference alignment optimize the entire pipeline [24]. Most promising is the movement in multi-modal RAG where we are seeing the emergence of a new frontier [131]. In conclusion, we expect the RAG framework to evolve from a knowledge tool into an integral component of general-purpose artificial intelligence, a system that is not only more reliable and adaptable [16](<https://example.com/paper>), but is robust to the challenges of data distribution and knowledge updates, ensuring sustainable and trustworthy deployment in the long term.

References

- [1] Shervin Minaee, Tomas Mikolov, Narjes Nikzad, Meysam Chenaghlu, Richard Socher, Xavier Amatriain *et al.*. *Large Language Models A Survey*. arXiv:2402.06196v2, 2024. <https://arxiv.org/abs/2402.06196v2>
- [2] Katikapalli Subramanyam Kalyan. *A Survey of GPT-3 Family Large Language Models Including ChatGPT and GPT-4*. arXiv:2310.12321v1, 2023. <https://arxiv.org/abs/2310.12321v1>
- [3] Oded Ovadia, Menachem Brief, Moshik Mishaali, Oren Elisha. *Fine-Tuning or Retrieval Comparing Knowledge Injection in LLMs*. arXiv:2312.05934v3, 2023. <https://arxiv.org/abs/2312.05934v3>
- [4] Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Daniel Khashabi, Hannaneh Hajishirzi. *When Not to Trust Language Models Investigating Effectiveness of Parametric and Non-Parametric Memories*. arXiv:2212.10511v4, 2022. <https://arxiv.org/abs/2212.10511v4>

- [5] Jean Kaddour, Joshua Harris, Maximilian Mozes, Herbie Bradley, Roberta Raileanu, Robert McHardy. *Challenges and Applications of Large Language Models*. arXiv:2307.10169v1, 2023. <https://arxiv.org/abs/2307.10169v1>
- [6] Tianyu Gao, Howard Yen, Jiatong Yu, Danqi Chen. *Enabling Large Language Models to Generate Text with Citations*. arXiv:2305.14627v2, 2023. <https://arxiv.org/abs/2305.14627v2>
- [7] Andrei-Laurentiu Bornea, Fadhel Ayed, Antonio De Domenico, Nicola Piovesan, Ali Maatouk. *Telco-RAG Navigating the Challenges of Retrieval-Augmented Language Models for Telecommunications*. arXiv:2404.15939v2, 2024. <https://arxiv.org/abs/2404.15939v2>
- [8] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal *et al.*. *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. arXiv:2005.11401v4, 2020. <https://arxiv.org/abs/2005.11401v4>
- [9] Edouard Grave, Armand Joulin, Nicolas Usunier. *Improving Neural Language Models with a Continuous Cache*. arXiv:1612.04426v1, 2016. <https://arxiv.org/abs/1612.04426v1>
- [10] Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, Zhen-Hua Ling. *Corrective Retrieval Augmented Generation*. arXiv:2401.15884v2, 2024. <https://arxiv.org/abs/2401.15884v2>
- [11] Tan Yu, Anbang Xu, Rama Akkiraju. *In Defense of RAG in the Era of Long-Context Language Models*. arXiv:2409.01666v1, 2024. <https://arxiv.org/abs/2409.01666v1>
- [12] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi *et al.*. *Retrieval-Augmented Generation for Large Language Models A Survey*. arXiv:2312.10997v5, 2023. <https://arxiv.org/abs/2312.10997v5>
- [13] Yuning Mao, Pengcheng He, Xiaodong Liu, Yelong Shen, Jianfeng Gao, Jiawei Han *et al.*. *Generation-Augmented Retrieval for Open-domain Question Answering*. arXiv:2009.08553v4, 2020. <https://arxiv.org/abs/2009.08553v4>
- [14] Antonio Jimeno Yepes, Yao You, Jan Milczek, Sebastian Laverde, Renyu Li. *Financial Report Chunking for Effective Retrieval Augmented Generation*. arXiv:2402.05131v3, 2024. <https://arxiv.org/abs/2402.05131v3>
- [15] Yizheng Huang, Jimmy Huang. *A Survey on Retrieval-Augmented Text Generation for Large Language Models*. arXiv:2404.10981v1, 2024. <https://arxiv.org/abs/2404.10981v1>
- [16] Akari Asai, Zexuan Zhong, Danqi Chen, Pang Wei Koh, Luke Zettlemoyer, Hannaneh Hajishirzi *et al.*. *Reliable, Adaptable, and Attributable Language Models with Retrieval*. arXiv:2403.03187v1, 2024. <https://arxiv.org/abs/2403.03187v1>
- [17] Zhipeng Xu, Zhenghao Liu, Yibin Liu, Chenyan Xiong, Yukun Yan, Shuo Wang *et al.*. *ActiveRAG Revealing the Treasures of Knowledge via Active Learning*. arXiv:2402.13547v1, 2024. <https://arxiv.org/abs/2402.13547v1>
- [18] Tilman Bruckhaus. *RAG Does Not Work for Enterprises*. arXiv:2406.04369v1, 2024. <https://arxiv.org/abs/2406.04369v1>
- [19] Zhentao Xu, Mark Jerome Cruz, Matthew Guevara, Tie Wang, Manasi Deshpande, Xiaofeng Wang *et al.*. *Retrieval-Augmented Generation with Knowledge Graphs for Customer Service Question Answering*. arXiv:2404.17723v2, 2024. <https://arxiv.org/abs/2404.17723v2>
- [20] Yucheng Hu, Yuxing Lu. *RAG and RAU: A Survey on Retrieval-Augmented Language Model in Natural Language Processing*. arXiv:2404.19543v1, 2024. <https://arxiv.org/abs/2404.19543v1>

- [21] Yunfan Gao, Yun Xiong, Meng Wang, Haofen Wang. *Modular RAG: Transforming RAG Systems into LEGO-like Reconfigurable Frameworks*. arXiv:2407.21059v1, 2024. <https://arxiv.org/abs/2407.21059v1>
- [22] Zexuan Zhong, Tao Lei, Danqi Chen. *Training Language Models with Memory Augmentation*. arXiv:2205.12674v3, 2022. <https://arxiv.org/abs/2205.12674v3>
- [23] Zhongwei Wan, Xin Wang, Che Liu, Samiul Alam, Yu Zheng, Jiachen Liu *et al.*. *Efficient Large Language Models A Survey*. arXiv:2312.03863v3, 2023. <https://arxiv.org/abs/2312.03863v3>
- [24] Katherine Tian, Eric Mitchell, Huaxiu Yao, Christopher D. Manning, Chelsea Finn. *Fine-tuning Language Models for Factuality*. arXiv:2311.08401v1, 2023. <https://arxiv.org/abs/2311.08401v1>
- [25] Cunxiang Wang, Xiaoze Liu, Yuanhao Yue, Xiangru Tang, Tianhang Zhang, Cheng Jiayang *et al.*. *Survey on Factuality in Large Language Models Knowledge, Retrieval and Domain-Specificity*. arXiv:2310.07521v3, 2023. <https://arxiv.org/abs/2310.07521v3>
- [26] Tongtong Wu, Linhao Luo, Yuan-Fang Li, Shirui Pan, Thuy-Trang Vu, Gholamreza Haffari. *Continual Learning for Large Language Models A Survey*. arXiv:2402.01364v2, 2024. <https://arxiv.org/abs/2402.01364v2>
- [27] Zhengbao Jiang, Frank F. Xu, Luyu Gao, Zhiqing Sun, Qian Liu, Jane Dwivedi-Yu *et al.*. *Active Retrieval Augmented Generation*. arXiv:2305.06983v2, 2023. <https://arxiv.org/abs/2305.06983v2>
- [28] Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranajpe, Michele Bevilacqua, Fabio Petroni *et al.*. *Lost in the Middle How Language Models Use Long Contexts*. arXiv:2307.03172v3, 2023. <https://arxiv.org/abs/2307.03172v3>
- [29] Zhihong Shao, Yeyun Gong, Yelong Shen, Minlie Huang, Nan Duan, Weizhu Chen. *Enhancing Retrieval-Augmented Large Language Models with Iterative Retrieval-Generation Synergy*. arXiv:2305.15294v2, 2023. <https://arxiv.org/abs/2305.15294v2>
- [30] Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, Hannaneh Hajishirzi. *Self-RAG Learning to Retrieve, Generate, and Critique through Self-Reflection*. arXiv:2310.11511v1, 2023. <https://arxiv.org/abs/2310.11511v1>
- [31] Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, David Bau. *Mass-Editing Memory in a Transformer*. arXiv:2210.07229v2, 2022. <https://arxiv.org/abs/2210.07229v2>
- [32] Eric Mitchell, Charles Lin, Antoine Bosselut, Christopher D. Manning, Chelsea Finn. *Memory-Based Model Editing at Scale*. arXiv:2206.06520v1, 2022. <https://arxiv.org/abs/2206.06520v1>
- [33] Andrei-Laurentiu Bornea, Fadhel Ayed, Antonio De Domenico, Nicola Piovesan, Ali Maatouk. *Telco-RAG: Navigating the Challenges of Retrieval-Augmented Language Models for Telecommunications*. arXiv:2404.15939v3, 2024. <https://arxiv.org/abs/2404.15939v3>
- [34] Shishir G. Patil, Tianjun Zhang, Xin Wang, Joseph E. Gonzalez. *Gorilla Large Language Model Connected with Massive APIs*. arXiv:2305.15334v1, 2023. <https://arxiv.org/abs/2305.15334v1>
- [35] Boci Peng, Yun Zhu, Yongchao Liu, Xiaohe Bo, Haizhou Shi, Chuntao Hong *et al.*. *Graph Retrieval-Augmented Generation: A Survey*. arXiv:2408.08921v2, 2024. <https://arxiv.org/abs/2408.08921v2>

- [36] Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, Yu Su. *HippoRAG: Neurobiologically Inspired Long-Term Memory for Large Language Models*. arXiv:2405.14831v1, 2024. <https://arxiv.org/abs/2405.14831v1>
- [37] Zhuowan Li, Cheng Li, Mingyang Zhang, Qiaozhu Mei, Michael Bendersky. *Retrieval Augmented Generation or Long-Context LLMs? A Comprehensive Study and Hybrid Approach*. arXiv:2407.16833v1, 2024. <https://arxiv.org/abs/2407.16833v1>
- [38] Soyeong Jeong, Jinheon Baek, Sukmin Cho, Sung Ju Hwang, Jong C. Park. *Adaptive-RAG Learning to Adapt Retrieval-Augmented Large Language Models through Question Complexity*. arXiv:2403.14403v2, 2024. <https://arxiv.org/abs/2403.14403v2>
- [39] Wenhao Yu, Dan Iter, Shuohang Wang, Yichong Xu, Mingxuan Ju, Soumya Sanyal *et al.*. *Generate rather than Retrieve Large Language Models are Strong Context Generators*. arXiv:2209.10063v3, 2022. <https://arxiv.org/abs/2209.10063v3>
- [40] Yutao Zhu, Huaying Yuan, Shuting Wang, Jiongnan Liu, Wenhan Liu, Chenlong Deng *et al.*. *Large Language Models for Information Retrieval A Survey*. arXiv:2308.07107v3, 2023. <https://arxiv.org/abs/2308.07107v3>
- [41] Heydar Soudani, Evangelos Kanoulas, Faegheh Hasibi. *Fine Tuning vs. Retrieval Augmented Generation for Less Popular Knowledge*. arXiv:2403.01432v2, 2024. <https://arxiv.org/abs/2403.01432v2>
- [42] Shangyu Wu, Ying Xiong, Yufei Cui, Haolun Wu, Can Chen, Ye Yuan *et al.*. *Retrieval-Augmented Generation for Natural Language Processing: A Survey*. arXiv:2407.13193v2, 2024. <https://arxiv.org/abs/2407.13193v2>
- [43] To Eun Kim, Alireza Salemi, Andrew Drozdov, Fernando Diaz, Hamed Zamani. *Retrieval-Enhanced Machine Learning: Synthesis and Opportunities*. arXiv:2407.12982v1, 2024. <https://arxiv.org/abs/2407.12982v1>
- [44] Masoomali Fatehkia, Ji Kim Lucas, Sanjay Chawla. *T-RAG Lessons from the LLM Trenches*. arXiv:2402.07483v1, 2024. <https://arxiv.org/abs/2402.07483v1>
- [45] Ziyang Jiang, Xueguang Ma, Wenhua Chen. *LongRAG: Enhancing Retrieval-Augmented Generation with Long-context LLMs*. arXiv:2406.15319v3, 2024. <https://arxiv.org/abs/2406.15319v3>
- [46] Bhaskarjit Sarmah, Benika Hall, Rohan Rao, Sunil Patel, Stefano Pasquali, Dhagash Mehta. *HybridRAG: Integrating Knowledge Graphs and Vector Retrieval Augmented Generation for Efficient Information Extraction*. arXiv:2408.04948v1, 2024. <https://arxiv.org/abs/2408.04948v1>
- [47] Demiao Lin. *Revolutionizing Retrieval-Augmented Generation with Enhanced PDF Structure Recognition*. arXiv:2401.12599v1, 2024. <https://arxiv.org/abs/2401.12599v1>
- [48] Chankyu Lee, Rajarshi Roy, Mengyao Xu, Jonathan Raiman, Mohammad Shoeybi, Bryan Catanzaro *et al.*. *NV-Embed: Improved Techniques for Training LLMs as Generalist Embedding Models*. arXiv:2405.17428v1, 2024. <https://arxiv.org/abs/2405.17428v1>
- [49] Yue Yu, Wei Ping, Zihan Liu, Boxin Wang, Jiaxuan You, Chao Zhang *et al.*. *RankRAG: Unifying Context Ranking with Retrieval-Augmented Generation in LLMs*. arXiv:2407.02485v1, 2024. <https://arxiv.org/abs/2407.02485v1>
- [50] Yutao Zhu, Peitian Zhang, Chenghao Zhang, Yifei Chen, Binyu Xie, Zhicheng Dou *et al.*. *INTERS Unlocking the Power of Large Language Models in Search with Instruction Tuning*. arXiv:2401.06532v2, 2024. <https://arxiv.org/abs/2401.06532v2>

- [51] Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei *et al.*. *FreshLLMs Refreshing Large Language Models with Search Engine Augmentation*. arXiv:2310.03214v2, 2023. <https://arxiv.org/abs/2310.03214v2>
- [52] Zixuan Ke, Weize Kong, Cheng Li, Mingyang Zhang, Qiaozhu Mei, Michael Bendersky. *Bridging the Preference Gap between Retrievers and LLMs*. arXiv:2401.06954v2, 2024. <https://arxiv.org/abs/2401.06954v2>
- [53] Alireza Salemi, Hamed Zamani. *Evaluating Retrieval Quality in Retrieval-Augmented Generation*. arXiv:2404.13781v1, 2024. <https://arxiv.org/abs/2404.13781v1>
- [54] Yuren Mao, Xuemei Dong, Wenyi Xu, Yunjun Gao, Bin Wei, Ying Zhang. *FIT-RAG Black-Box RAG with Factual Information and Token Reduction*. arXiv:2403.14374v1, 2024. <https://arxiv.org/abs/2403.14374v1>
- [55] Florin Cuconasu, Giovanni Trappolini, Federico Siciliano, Simone Filice, Cesare Campagnano, Yoelle Maarek *et al.*. *The Power of Noise Redefining Retrieval for RAG Systems*. arXiv:2401.14887v3, 2024. <https://arxiv.org/abs/2401.14887v3>
- [56] Zheng Wang, Shu Xian Teo, Jieer Ouyang, Yongjun Xu, Wei Shi. *M-RAG: Reinforcing Large Language Model Performance through Retrieval-Augmented Generation with Multiple Partitions*. arXiv:2405.16420v1, 2024. <https://arxiv.org/abs/2405.16420v1>
- [57] Chao Jin, Zili Zhang, Xuanlin Jiang, Fangyue Liu, Xin Liu, Xuanzhe Liu *et al.*. *RAGCache Efficient Knowledge Caching for Retrieval-Augmented Generation*. arXiv:2404.12457v2, 2024. <https://arxiv.org/abs/2404.12457v2>
- [58] Kaize Shi, Xueyao Sun, Qing Li, Guandong Xu. *Compressing Long Context for Enhancing RAG with AMR-based Concept Distillation*. arXiv:2405.03085v1, 2024. <https://arxiv.org/abs/2405.03085v1>
- [59] Guangzhi Xiong, Qiao Jin, Zhiyong Lu, Aidong Zhang. *Benchmarking Retrieval-Augmented Generation for Medicine*. arXiv:2402.13178v2, 2024. <https://arxiv.org/abs/2402.13178v2>
- [60] Sukmin Cho, Soyeong Jeong, Jeongyeon Seo, Taeho Hwang, Jong C. Park. *Typos that Broke the RAG’s Back Genetic Attack on RAG Pipeline by Simulating Documents in the Wild via Low-level Perturbations*. arXiv:2404.13948v1, 2024. <https://arxiv.org/abs/2404.13948v1>
- [61] Daniel Machlab, Rick Battle. *LLM In-Context Recall is Prompt Dependent*. arXiv:2404.08865v1, 2024. <https://arxiv.org/abs/2404.08865v1>
- [62] Dongyu Ru, Lin Qiu, Xiangkun Hu, Tianhang Zhang, Peng Shi, Shuaichen Chang *et al.*. *RAGChecker: A Fine-grained Framework for Diagnosing Retrieval-Augmented Generation*. arXiv:2408.08067v2, 2024. <https://arxiv.org/abs/2408.08067v2>
- [63] Baolin Peng, Michel Galley, Pengcheng He, Hao Cheng, Yujia Xie, Yu Hu *et al.*. *Check Your Facts and Try Again Improving Large Language Models with External Knowledge and Automated Feedback*. arXiv:2302.12813v3, 2023. <https://arxiv.org/abs/2302.12813v3>
- [64] Shashank Rajput, Ying Sheng, Sean Owen, Vitaliy Chiley. *Inference-Friendly Models With MixAttention*. arXiv:2409.15012v1, 2024. <https://arxiv.org/abs/2409.15012v1>
- [65] Cicero Nogueira dos Santos, Xiaofei Ma, Ramesh Nallapati, Zhiheng Huang, Bing Xiang. *Beyond [CLS] through Ranking by Generation*. arXiv:2010.03073v1, 2020. <https://arxiv.org/abs/2010.03073v1>
- [66] Wenqi Jiang, Shuai Zhang, Boran Han, Jie Wang, Bernie Wang, Tim Kraska. *PipeRAG Fast Retrieval-Augmented Generation via Algorithm-System Co-design*. arXiv:2403.05676v1, 2024. <https://arxiv.org/abs/2403.05676v1>

- [67] Jihoon Tack, Jaehyung Kim, Eric Mitchell, Jinwoo Shin, Yee Whye Teh, Jonathan Richard Schwarz. *Online Adaptation of Language Models with a Memory of Amortized Contexts*. arXiv:2403.04317v1, 2024. <https://arxiv.org/abs/2403.04317v1>
- [68] Yu Wang, Xiushi Chen, Jingbo Shang, Julian McAuley. *MEMORYLLM Towards Self-Updatable Large Language Models*. arXiv:2402.04624v1, 2024. <https://arxiv.org/abs/2402.04624v1>
- [69] Simran Arora, Sabri Eyuboglu, Michael Zhang, Aman Timalsina, Silas Alberti, Dylan Zinsley *et al.*. *Simple linear attention language models balance the recall-throughput tradeoff*. arXiv:2402.18668v1, 2024. <https://arxiv.org/abs/2402.18668v1>
- [70] Jiawei Chen, Hongyu Lin, Xianpei Han, Le Sun. *Benchmarking Large Language Models in Retrieval-Augmented Generation*. arXiv:2309.01431v2, 2023. <https://arxiv.org/abs/2309.01431v2>
- [71] Ruiyang Ren, Yuhao Wang, Yingqi Qu, Wayne Xin Zhao, Jing Liu, Hao Tian *et al.*. *Investigating the Factual Knowledge Boundary of Large Language Models with Retrieval Augmentation*. arXiv:2307.11019v2, 2023. <https://arxiv.org/abs/2307.11019v2>
- [72] Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, Ashish Sabharwal. *Interleaving Retrieval with Chain-of-Thought Reasoning for Knowledge-Intensive Multi-Step Questions*. arXiv:2212.10509v2, 2022. <https://arxiv.org/abs/2212.10509v2>
- [73] Yanming Liu, Xinyue Peng, Xuhong Zhang, Weihao Liu, Jianwei Yin, Jiannan Cao *et al.*. *RA-ISF Learning to Answer and Understand from Retrieval Augmentation via Iterative Self-Feedback*. arXiv:2403.06840v1, 2024. <https://arxiv.org/abs/2403.06840v1>
- [74] Tomoyuki Kagaya, Thong Jing Yuan, Yuxuan Lou, Jayashree Karlekar, Sugiri Pranata, Akira Kinose *et al.*. *RAP Retrieval-Augmented Planning with Contextual Memory for Multimodal LLM Agents*. arXiv:2402.03610v1, 2024. <https://arxiv.org/abs/2402.03610v1>
- [75] Diji Yang, Jinmeng Rao, Kezhen Chen, Xiaoyuan Guo, Yawen Zhang, Jie Yang *et al.*. *IM-RAG: Multi-Round Retrieval-Augmented Generation Through Learning Inner Monologues*. arXiv:2405.13021v1, 2024. <https://arxiv.org/abs/2405.13021v1>
- [76] Ori Yoran, Tomer Wolfson, Ori Ram, Jonathan Berant. *Making Retrieval-Augmented Language Models Robust to Irrelevant Context*. arXiv:2310.01558v1, 2023. <https://arxiv.org/abs/2310.01558v1>
- [77] Hao Yu, Aoran Gan, Kai Zhang, Shiwei Tong, Qi Liu, Zhaofeng Liu. *Evaluation of Retrieval-Augmented Generation: A Survey*. arXiv:2405.07437v2, 2024. <https://arxiv.org/abs/2405.07437v2>
- [78] Jiajie Jin, Yutao Zhu, Yujia Zhou, Zhicheng Dou. *BIDER Bridging Knowledge Inconsistency for Efficient Retrieval-Augmented LLMs via Key Supporting Evidence*. arXiv:2402.12174v1, 2024. <https://arxiv.org/abs/2402.12174v1>
- [79] Shicheng Xu, Liang Pang, Mo Yu, Fandong Meng, Huawei Shen, Xueqi Cheng *et al.*. *Unsupervised Information Refinement Training of Large Language Models for Retrieval-Augmented Generation*. arXiv:2402.18150v1, 2024. <https://arxiv.org/abs/2402.18150v1>
- [80] Yunxiao Shi, Xing Zi, Zijing Shi, Haimin Zhang, Qiang Wu, Min Xu. *ERAGent: Enhancing Retrieval-Augmented Language Models with Improved Accuracy, Efficiency, and Personalization*. arXiv:2405.06683v1, 2024. <https://arxiv.org/abs/2405.06683v1>
- [81] Bowen Cao, Deng Cai, Leyang Cui, Xuxin Cheng, Wei Bi, Yuexian Zou *et al.*. *Retrieval is Accurate Generation*. arXiv:2402.17532v3, 2024. <https://arxiv.org/abs/2402.17532v3>

- [82] Kshitij Fadnis, Siva Sankalp Patel, Odellia Boni, Yannis Katsis, Sara Rosenthal, Benjamin Sznajder *et al.*. *InspectorRAG: An Introspection Platform for RAG Evaluation*. arXiv:2404.17347v1, 2024. <https://arxiv.org/abs/2404.17347v1>
- [83] Orion Weller, Benjamin Chang, Sean MacAvaney, Kyle Lo, Arman Cohan, Benjamin Van Durme *et al.*. *FollowIR: Evaluating and Teaching Information Retrieval Models to Follow Instructions*. arXiv:2403.15246v1, 2024. <https://arxiv.org/abs/2403.15246v1>
- [84] Avital Shafran, Roei Schuster, Vitaly Shmatikov. *Machine Against the RAG: Jamming Retrieval-Augmented Generation with Blocker Documents*. arXiv:2406.05870v2, 2024. <https://arxiv.org/abs/2406.05870v2>
- [85] Keheng Wang, Feiyu Duan, Peiguang Li, Sirui Wang, Xunliang Cai. *LLMs Know What They Need: Leveraging a Missing Information Guided Framework to Empower Retrieval-Augmented Generation*. arXiv:2404.14043v1, 2024. <https://arxiv.org/abs/2404.14043v1>
- [86] Maxim Khanov, Jirayu Burapachee, Yixuan Li. *ARGS: Alignment as Reward-Guided Search*. arXiv:2402.01694v1, 2024. <https://arxiv.org/abs/2402.01694v1>
- [87] Hansi Zeng, Chen Luo, Hamed Zamani. *Planning Ahead in Generative Retrieval Guiding Autoregressive Generation through Simultaneous Decoding*. arXiv:2404.14600v1, 2024. <https://arxiv.org/abs/2404.14600v1>
- [88] Xinbei Ma, Yeyun Gong, Pengcheng He, Hai Zhao, Nan Duan. *Query Rewriting for Retrieval-Augmented Large Language Models*. arXiv:2305.14283v3, 2023. <https://arxiv.org/abs/2305.14283v3>
- [89] Hamed Zamani, Michael Bendersky. *Stochastic RAG: End-to-End Retrieval-Augmented Generation through Expected Utility Maximization*. arXiv:2405.02816v1, 2024. <https://arxiv.org/abs/2405.02816v1>
- [90] Chunting Zhou, Pengfei Liu, Puxin Xu, Srini Iyer, Jiao Sun, Yuning Mao *et al.*. *LIMA: Less Is More for Alignment*. arXiv:2305.11206v1, 2023. <https://arxiv.org/abs/2305.11206v1>
- [91] Alireza Salemi, Surya Kallumadi, Hamed Zamani. *Optimization Methods for Personalizing Large Language Models through Retrieval Augmentation*. arXiv:2404.05970v1, 2024. <https://arxiv.org/abs/2404.05970v1>
- [92] Zhichao Wang, Bin Bi, Shiva Kumar Pentyala, Kiran Ramnath, Sougata Chaudhuri, Shubham Mehrotra *et al.*. *A Comprehensive Survey of LLM Alignment Techniques: RLHF, RLAI, PPO, DPO and More*. arXiv:2407.16216v1, 2024. <https://arxiv.org/abs/2407.16216v1>
- [93] S. M Towhidul Islam Tonmoy, S M Mehedi Zaman, Vinija Jain, Anku Rani, Vipula Rawte, Aman Chadha *et al.*. *A Comprehensive Survey of Hallucination Mitigation Techniques in Large Language Models*. arXiv:2401.01313v3, 2024. <https://arxiv.org/abs/2401.01313v3>
- [94] Hexiang Tan, Fei Sun, Wanli Yang, Yuanzhuo Wang, Qi Cao, Xueqi Cheng. *Blinded by Generated Contexts: How Language Models Merge Generated and Retrieved Contexts for Open-Domain QA*. arXiv:2401.11911v4, 2024. <https://arxiv.org/abs/2401.11911v4>
- [95] Zhen Tan, Alimohammad Beigi, Song Wang, Ruocheng Guo, Amrita Bhattacharjee, Bohan Jiang *et al.*. *Large Language Models for Data Annotation: A Survey*. arXiv:2402.13446v1, 2024. <https://arxiv.org/abs/2402.13446v1>
- [96] Yujia Zhou, Zheng Liu, Jiajie Jin, Jian-Yun Nie, Zhicheng Dou. *Metacognitive Retrieval-Augmented Large Language Models*. arXiv:2402.11626v1, 2024. <https://arxiv.org/abs/2402.11626v1>

- [97] Zhengliang Shi, Weiwei Sun, Shen Gao, Pengjie Ren, Zhumin Chen, Zhaochun Ren. *Generate-then-Ground in Retrieval-Augmented Generation for Multi-hop Question Answering*. arXiv:2406.14891v2, 2024. <https://arxiv.org/abs/2406.14891v2>
- [98] Yue Zhou, Chenlu Guo, Xu Wang, Yi Chang, Yuan Wu. *A Survey on Data Augmentation in Large Model Era*. arXiv:2401.15422v2, 2024. <https://arxiv.org/abs/2401.15422v2>
- [99] Xi Victoria Lin, Xilun Chen, Mingda Chen, Weijia Shi, Maria Lomeli, Rich James *et al.*. *RA-DIT Retrieval-Augmented Dual Instruction Tuning*. arXiv:2310.01352v3, 2023. <https://arxiv.org/abs/2310.01352v3>
- [100] YuHe Ke, Liyuan Jin, Kabilan Elangovan, Hairil Rizal Abdullah, Nan Liu, Alex Tiong Heng Sia *et al.*. *Development and Testing of Retrieval Augmented Generation in Large Language Models – A Case Study Report*. arXiv:2402.01733v1, 2024. <https://arxiv.org/abs/2402.01733v1>
- [101] Feiteng Fang, Yuelin Bai, Shiwen Ni, Min Yang, Xiaojun Chen, Ruifeng Xu. *Enhancing Noise Robustness of Retrieval-Augmented Language Models with Adaptive Adversarial Training*. arXiv:2405.20978v1, 2024. <https://arxiv.org/abs/2405.20978v1>
- [102] Hongru Wang, Wenyu Huang, Yang Deng, Rui Wang, Zezhong Wang, Yufei Wang *et al.*. *UniMS-RAG A Unified Multi-source Retrieval-Augmented Generation for Personalized Dialogue Systems*. arXiv:2401.13256v1, 2024. <https://arxiv.org/abs/2401.13256v1>
- [103] Mintong Kang, Nezihe Merve Gürel, Ning Yu, Dawn Song, Bo Li. *C-RAG Certified Generation Risks for Retrieval-Augmented Language Models*. arXiv:2402.03181v3, 2024. <https://arxiv.org/abs/2402.03181v3>
- [104] Zheng Zhao, Emilio Monti, Jens Lehmann, Haytham Assem. *Enhancing Contextual Understanding in Large Language Models through Contrastive Decoding*. arXiv:2405.02750v1, 2024. <https://arxiv.org/abs/2405.02750v1>
- [105] Scott Barnett, Stefanus Kurniawan, Srikanth Thudumu, Zach Brannelly, Mohamed Abdelrazek. *Seven Failure Points When Engineering a Retrieval Augmented Generation System*. arXiv:2401.05856v1, 2024. <https://arxiv.org/abs/2401.05856v1>
- [106] Wei Zou, Runpeng Geng, Binghui Wang, Jinyuan Jia. *PoisonedRAG Knowledge Poisoning Attacks to Retrieval-Augmented Generation of Large Language Models*. arXiv:2402.07867v1, 2024. <https://arxiv.org/abs/2402.07867v1>
- [107] Yuying Li, Gaoyang Liu, Yang Yang, Chen Wang. *Seeing Is Believing: Black-Box Membership Inference Attacks Against Retrieval Augmented Generation*. arXiv:2406.19234v1, 2024. <https://arxiv.org/abs/2406.19234v1>
- [108] Wayne Xin Zhao, Jing Liu, Ruiyang Ren, Ji-Rong Wen. *Dense Text Retrieval based on Pretrained Language Models A Survey*. arXiv:2211.14876v1, 2022. <https://arxiv.org/abs/2211.14876v1>
- [109] Wei-Cheng Chang, Felix X. Yu, Yin-Wen Chang, Yiming Yang, Sanjiv Kumar. *Pre-training Tasks for Embedding-based Large-scale Retrieval*. arXiv:2002.03932v1, 2020. <https://arxiv.org/abs/2002.03932v1>
- [110] Xin Cheng, Xun Wang, Xingxing Zhang, Tao Ge, Si-Qing Chen, Furu Wei *et al.*. *xRAG: Extreme Context Compression for Retrieval-augmented Generation with One Token*. arXiv:2405.13792v1, 2024. <https://arxiv.org/abs/2405.13792v1>
- [111] Shahul Es, Jithin James, Luis Espinosa-Anke, Steven Schockaert. *RAGAS Automated Evaluation of Retrieval Augmented Generation*. arXiv:2309.15217v1, 2023. <https://arxiv.org/abs/2309.15217v1>

- [112] Yuanjie Lyu, Zhiyu Li, Simin Niu, Feiyu Xiong, Bo Tang, Wenjin Wang *et al.*. *CRUD-RAG A Comprehensive Chinese Benchmark for Retrieval-Augmented Generation of Large Language Models*. arXiv:2401.17043v2, 2024. <https://arxiv.org/abs/2401.17043v2>
- [113] Yixuan Tang, Yi Yang. *MultiHop-RAG Benchmarking Retrieval-Augmented Generation for Multi-Hop Queries*. arXiv:2401.15391v1, 2024. <https://arxiv.org/abs/2401.15391v1>
- [114] Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Yarin Gal, Nicolas Papernot, Ross Anderson. *The Curse of Recursion Training on Generated Data Makes Models Forget*. arXiv:2305.17493v3, 2023. <https://arxiv.org/abs/2305.17493v3>
- [115] Xingwei He, Zhenghao Lin, Yeyun Gong, A-Long Jin, Hang Zhang, Chen Lin *et al.*. *AnnoLLM Making Large Language Models to Be Better Crowdsourced Annotators*. arXiv:2303.16854v2, 2023. <https://arxiv.org/abs/2303.16854v2>
- [116] Guglielmo Faggioli, Laura Dietz, Charles Clarke, Gianluca Demartini, Matthias Hagen, Claudia Hauff *et al.*. *Perspectives on Large Language Models for Relevance Judgment*. arXiv:2304.09161v2, 2023. <https://arxiv.org/abs/2304.09161v2>
- [117] Minbyul Jeong, Jiwoong Sohn, Mujeen Sung, Jaewoo Kang. *Improving Medical Reasoning through Retrieval and Self-Reflection with Retrieval-Augmented Large Language Models*. arXiv:2401.15269v2, 2024. <https://arxiv.org/abs/2401.15269v2>
- [118] Hanxing Ding, Liang Pang, Zihao Wei, Huawei Shen, Xueqi Cheng. *Retrieve Only When It Needs Adaptive Retrieval Augmentation for Hallucination Mitigation in Large Language Models*. arXiv:2402.10612v1, 2024. <https://arxiv.org/abs/2402.10612v1>
- [119] Shiyu Ni, Keping Bi, Jiafeng Guo, Xueqi Cheng. *When Do LLMs Need Retrieval Augmentation Mitigating LLMs’ Overconfidence Helps Retrieval Augmentation*. arXiv:2402.11457v1, 2024. <https://arxiv.org/abs/2402.11457v1>
- [120] Wenqi Fan, Yajuan Ding, Liangbo Ning, Shijie Wang, Hengyun Li, Dawei Yin *et al.*. *A Survey on RAG Meeting LLMs: Towards Retrieval-Augmented Large Language Models*. arXiv:2405.06211v3, 2024. <https://arxiv.org/abs/2405.06211v3>
- [121] Jiajie Jin, Yutao Zhu, Xinyu Yang, Chenghao Zhang, Zhicheng Dou. *FlashRAG: A Modular Toolkit for Efficient Retrieval-Augmented Generation Research*. arXiv:2405.13576v1, 2024. <https://arxiv.org/abs/2405.13576v1>
- [122] Jon Saad-Falcon, Omar Khattab, Christopher Potts, Matei Zaharia. *ARES An Automated Evaluation Framework for Retrieval-Augmented Generation Systems*. arXiv:2311.09476v2, 2023. <https://arxiv.org/abs/2311.09476v2>
- [123] Zackary Rackauckas, Arthur Câmara, Jakub Zavrel. *Evaluating RAG-Fusion with RAGElo: an Automated Elo-based Framework*. arXiv:2406.14783v1, 2024. <https://arxiv.org/abs/2406.14783v1>
- [124] Yunzhen Feng, Elvis Dohmatob, Pu Yang, Francois Charton, Julia Kempe. *Beyond Model Collapse: Scaling Up with Synthesized Data Requires Reinforcement*. arXiv:2406.07515v1, 2024. <https://arxiv.org/abs/2406.07515v1>
- [125] Rachit Bansal, Bidisha Samanta, Siddharth Dalmia, Nitish Gupta, Shikhar Vashishth, Sriram Ganapathy *et al.*. *LLM Augmented LLMs Expanding Capabilities through Composition*. arXiv:2401.02412v1, 2024. <https://arxiv.org/abs/2401.02412v1>
- [126] Zihao Wang, Anji Liu, Haowei Lin, Jiaqi Li, Xiaojian Ma, Yitao Liang. *RAT Retrieval Augmented Thoughts Elicit Context-Aware Reasoning in Long-Horizon Generation*. arXiv:2403.05313v1, 2024. <https://arxiv.org/abs/2403.05313v1>

- [127] Xiaoxi Li, Jiajie Jin, Yujia Zhou, Yuyao Zhang, Peitian Zhang, Yutao Zhu *et al.*. *From Matching to Generation: A Survey on Generative Information Retrieval*. arXiv:2404.14851v3, 2024. <https://arxiv.org/abs/2404.14851v3>
- [128] Xiaoxi Li, Zhicheng Dou, Yujia Zhou, Fangchao Liu. *CorpusLM Towards a Unified Language Model on Corpus for Knowledge-Intensive Tasks*. arXiv:2402.01176v2, 2024. <https://arxiv.org/abs/2402.01176v2>
- [129] Zihan Zhang, Meng Fang, Ling Chen. *RetrievalQA Assessing Adaptive Retrieval-Augmented Generation for Short-form Open-Domain Question Answering*. arXiv:2402.16457v1, 2024. <https://arxiv.org/abs/2402.16457v1>
- [130] Tolgahan Cakaloglu, Christian Szegedy, Xiaowei Xu. *Text Embeddings for Retrieval From a Large Knowledge Base*. arXiv:1810.10176v2, 2018. <https://arxiv.org/abs/1810.10176v2>
- [131] Wenhui Chen, Hexiang Hu, Xi Chen, Pat Verga, William W. Cohen. *MuRAG Multimodal Retrieval-Augmented Generator for Open Question Answering over Images and Text*. arXiv:2210.02928v2, 2022. <https://arxiv.org/abs/2210.02928v2>