

3D Gaussian Splatting: Foundations, Advances, and Future Directions — A Comprehensive Survey

Generated by SurveyForge

Contents

1	Introduction	3
2	Optimization, Densification, and Geometric Fidelity	4
2.1	Adaptive Density Control and Primitive Management	4
2.2	Optimization Strategies and Loss Function Design	5
2.3	Handling Sparse-View and Few-Shot Scenarios	7
2.4	Geometric Fidelity and Surface Extraction	8
2.5	Training Dynamics and Compactness	10
3	Efficient Rendering, Compression, and Scalability	11
3.1	Algorithmic Rendering Acceleration and Level-of-Detail Management	11
3.2	Model Compression, Quantization, and Storage Optimization	12
3.3	Hardware-Accelerated Rasterization and Software-Hardware Co-Design	14
3.4	Distributed Training, Large-Scale Rendering, and Memory Management	15
4	Generalization, Sparse-View Reconstruction, and Generation	17
4.1	Generalizable and Feed-Forward Gaussian Prediction	17
4.2	Sparse-View and Uncalibrated Reconstruction	18
4.3	Single-View Reconstruction and Generative Completion	19
4.4	Test-Time Optimization and Hybrid Prediction-Refinement	20
4.5	Leveraging Foundation Models and Semantic-Aware Generalization	22
5	Dynamics, Deformation, and Temporal Representations	26
5.1	Deformation Architectures and Canonical Space Modeling	26
5.2	Direct Spatio-Temporal Primitive Representations	27
5.3	Motion Correspondence, Tracking, and Velocity-Aware Dynamics	28
5.4	Dynamic Scene Generation, Editing, and Streaming	30
6	Scene Understanding, Semantics, and Interaction	31
6.1	Language-Embedded Semantics and Open-Vocabulary Querying	31
6.2	Scene Editing, Manipulation, and Inpainting	32
6.3	Scene Understanding for Embodied and Interactive Agents	33
6.4	Generative Content Creation and Physical Simulation	34
7	Rendering Appearance, Inverse Rendering, and Photorealism**	36
7.1	Physically-Based Appearance Decomposition and Inverse Rendering	36
7.2	Relighting, High-Dynamic-Range Imaging, and Global Illumination Approximation	36
7.3	Advanced Material Modeling and Complex Light Transport Effects	38
7.4	Visibility, Occlusion, and Physical Correctness in Rasterization and Rendering .	39

8	Practical Deployment, Benchmarks, and Future Horizons**	40
8.1	Evaluation Protocols, Benchmark Suites, and Standardized Metrics	40
8.2	Robustness, Reliability, and Degradation Analysis in the Wild	42
8.3	Interoperability, File Standards, and Industrial-Grade Usability	43
8.4	Scaling, Temporal Streams, and Deployment on Device	45
8.5	Future Horizons, Emerging Paradigms, and Multi-Modal Integration	46
8.6	Conclusion	47
9	Conclusion	48

1 Introduction

The emergence of 3D Gaussian Splatting (3DGS) marks a paradigm shift in the representation and rendering of 3D scenes, effectively reconciling the long-standing trade-off between photorealistic quality and computational efficiency. To appreciate its significance, it is essential to trace its intellectual lineage. Classical approaches in computer graphics relied on explicit representations—meshes and point clouds—which are inherently efficient to render via standard graphics pipelines but often struggle to capture the continuous, view-dependent appearance of real-world scenes. Conversely, the advent of implicit neural representations, most notably Neural Radiance Fields (NeRFs), introduced a principled framework for modeling volumetric radiance via coordinate-based networks, achieving unprecedented realism in novel view synthesis. However, this expressivity is achieved through costly ray marching, which imposes significant computational burdens during both training and inference, precluding real-time performance for complex, unbounded scenes. It is within this context that 3DGS is posited as a transformative technique, harmonizing the high fidelity of neural approaches with the inherent speed of point-based rasterization [1].

The foundational innovation of 3DGS lies in its departure from continuous neural networks to an explicit, discrete set of anisotropic 3D Gaussian primitives [2]. The introduction of this method has been the key precursor to real-time radiance field rendering at high resolution, demonstrating that the continuous volumetric representation of NeRF can be effectively mirrored by a collection of such primitives. Formally, each primitive $\mathbf{G}_i$ is parameterized by a central position $\boldsymbol{\mu}_i \in \mathbb{R}^3$, an anisotropic covariance matrix $\boldsymbol{\Sigma}_i$ (often decomposed into a rotation matrix $\mathbf{R}$ and a scaling matrix $\mathbf{S}$), an opacity value α_i , and a set of spherical harmonic (SH) coefficients for modeling view-dependent color. The contribution of a primitive to a point in space is defined by its Gaussian response, given by $\mathbf{G}_i(\mathbf{x}) = \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_i)^T \boldsymbol{\Sigma}_i^{-1}(\mathbf{x} - \boldsymbol{\mu}_i)\right)$. This parameterization permits a unique and efficient differentiable rendering process. Unlike the ray-marching of NeRFs, the rendering in 3DGS is achieved via splatting: the Gaussians are first projected onto the 2D image plane, sorted by depth, and then rasterized in a tile-based, parallel manner. Each pixel’s color is accumulated via front-to-back alpha compositing of the overlapping primitives, a process that is inherently fast yet still fully differentiable for optimization.

This explicit formulation provides a powerful nexus of properties that are difficult to achieve simultaneously in competing representations. Compared to traditional meshes, 3DGS can represent highly complex, fuzzy, and dynamic geometry without requiring topological connectivity, often combining the strengths of both points and volumetric fields [3]. While point clouds are inherently simpler, they lack the smooth, continuous surface and appearance interpolation that the finite spatial volume of Gaussian kernels provides. When juxtaposed against implicit neural representations, 3DGS offers transparency for aspects such as scene editing, compositing, and physical simulation, and provides a path toward real-time interactive frame rates [4]. However, this expressivity comes with a considerable memory overhead, as the representation often requires millions of primitives, sparking a substantial body of research focused on compression, level-of-detail, and more efficient parameterizations [5, 6].

The scope of this survey is designed to navigate the expansive and rapidly evolving research landscape created by this new paradigm. The evolution of 3DGS is more than a story of a single algorithm; it is a narrative of a versatile foundation in which a multitude of novel techniques and applications are developed [7, 8]. This includes, but is not limited to, optimizing the core algorithm for geometric accuracy of surfaces [9, 10], scaling it to massive scenes [6], generalizing it for feed-forward scene prediction [11], and extending it into the temporal dimension to capture dynamic scenes [12]. Concurrent work has also leveraged the explicit nature of the representation to enhance its expressiveness by associating primitive-level attributes with higher-level neural fields—notably, texturing primitives to super-resolve appearance [13] or to close the appearance-geometry gap [14, 15]. More recent innovation has expanded the representational capacity by

introducing larger neural primitives [16, 17], and more geometrically complex primitives [18]. This highlights an important research thread dedicated to addressing the apparent redundancy of small-scale primitives and spatial structures of the scene geometry.

To navigate this diversity, it is crucial to consider the broader categories of 3DGS applications, which have been systematically categorized in recent reviews, showcasing segmentation, editing, and generation [19, 1]. The structure of our survey is designed to mirror this intellectual progression. The following sections explore into (1) the core algorithmic and optimization strategies for generating the geometric aspects of the scene, (2) the challenges of scaling this representation to large environments and improving its efficiency in memory and compute [20], (3) how the feed-forward and generative generalization can overcome the limitations of per-scene optimization, (4) the extension of these techniques to capture dynamic and temporally-evolving scene representations, (5) the integration of high-level semantics and interactivity into the explicit cloud, and (6) the pursuit of physical correctness for inverse rendering and material properties. We conclude with a summary of open challenges and future directions, emphasizing the need for robustness, scalability, and standardized evaluation platforms to support further progress in this dynamically evolving field.

2 Optimization, Densification, and Geometric Fidelity

2.1 Adaptive Density Control and Primitive Management

The optimization of 3D Gaussian Splatting (3DGS) hinges on a dynamic equilibrium between representational capacity and computational economy, a balance governed by the adaptive density control and primitive management mechanisms. The foundational framework, introduced with 3DGS [2], operates on a set of heuristics driven by the average magnitude of positional gradients of the Gaussian primitives. This criterion stems from the intuition that primitives with large positional gradients indicate regions of the scene that are under-optimized and thus require additional capacity. In such areas, the algorithm employs a cloning operation for large, poorly reconstructed Gaussians and a splitting operation for similarly sized but elongated ones, effectively propagating new primitives into the scene. Conversely, to curtail redundancy, an opacity-based pruning process is periodically applied, eliminating Gaussians whose accumulated opacity falls below a threshold. While straightforward and remarkably effective, this gradient-driven paradigm exhibits known failure modes. It is inherently myopic—its densification signal is sourced from the opacity-weighted positional gradient, which is problematic in sparse-view settings [21], dynamic scenes, and environments with large texture-less regions, leading to the generation of floating artifacts, floaters, due to over-reconstruction in ambiguous areas or the insufficient spatial accumulation of gradient signal for optimal placement.

Recognizing the limitations of the positional-gradient heuristic, a broad class of modern methods shifts toward more structure-aware and error-driven criteria for densification. Instead of relying solely on second-order gradients, these approaches integrate a richer variety of signals to determine where to allocate primitives. For instance, photometric error surfaces can be used to directly allocate Gaussians, mimicking a more robust policy. The rudimentary approach where "adaptiveness" is achieved through gradient-based cloning/splitting is often insufficient; more sophisticated techniques such as in [22] suggest that a count-adaptive device, in contrast to the unbounded growth of the original algorithm, can be beneficial. This global perspective relieves the heuristic by imposing a controllable and potentially more stable optimization trajectory. Furthermore, the advancement of loss-function-driven densification considers high-frequency details—where position-based error in high-frequency regions requests dedicated allocation, often at the expense of larger, nearly-flat surfaces [15, 23]. These models drive spatial allocation using a loss-aware mechanism that effectively couples the appearance quality with the location of primitives.

The realm of learning-based primitive control is also emerging. Rather than defining hard-coded heuristics, a feed-forward framework can predict the necessary discrete allocation. For instance, SplatWeaver [11] presents a cardinality-based approach that learns to predict a region-adaptive number of Gaussians per pixel, using a routing scheme and principle of "region-dependent primitive cardinality." This is a direct departure from uniform allocation—a demand-driven precision via expert networks. In contrast to the dense, fixed-count system, this allows for "complex patterns with fewer primitives". The structural addition of an integrated high-frequency prior further constrains the allocation, effectively mapping to secondary processing blocks. Algorithms like this highlight a shift from purely local greedy heuristics to globally enforceable and object-aware programs that are determined by high-level visual stimuli.

For the control of the distribution and the exploitation of geometric structure, replacing the initial recursive cloning with a Markov chain-based procedure achieves significant advantages in generalizing and stability. Following the hard wired statistical filtering, a "self-organized cross-level probabilistic field" with MCMC-based densification procedure [24] transforms the way we maintain primitives. This is a critical reframing: a mechanism for allocation can be processed as a Markov chain across the optimization trajectory, which can be harnessed to replace the volatile thresholding and, critically, ensure consistent geometric grounding for extreme scaling that the point-field-based methods fail to achieve robustly.

A major avenue to improve the 3DGS algorithm has been on technical improvements in the control of transparency of dense Gaussian sequences. A major bottleneck is the so-called "shader state" in a poor spatial environment is overly optimized by training loss; instead, an important insight is to separate and handle time-varying parameters such as positions and rotations while leaving other time-invariant properties unchanged [25]. This approach ensures that no unintended modifications such as the use of various constraints like optical flow and direction excessively, preserving edges and making the scene more visually controlled by local location and the visibility check. This is not only a matter of quality but also marks a shift in primitive management from being purely a growth heuristic to being a fundamental form of geometry coding beyond the scene.

In terms of model deployment and resource constraints, several proposals for data management have emphasized the notion of "frugal" allocation. Borrowing from the [5] literature, the later pruning metric proves critical for rotation compression. The removal of unused scene parts based on the tasks of adaptation and granularity ensures that the resulting scene remains coherent with the training data and comply with the strict performance constraints of, e.g., the hardware (mobile). This suggests a linking of more effective optimization.

Finally, the boundary of primitive management is correlated with algorithmic adaptation to new measurement domains, such as that of dynamic scene—where the camera views are sparse, and time is an extra dimension [12]. This is here where adaptation of the 6DGS approach generally less stable and where more global and temporally varying systems for maintaining the primitive count are needed for dynamic. Locally dense Gaussian control with occlusions remains one of the central open problems. The next phase will be building the deterministic, analytic, and robust control mechanisms, leveraging Markov-chain-steered and error-based dynamic systems to develop a more principled, robust, and sparse marching sector with explicit clarity and optimal control, pulling from structured data and variability constraints that will ensure reliability across optimizing wild scenarios. In the future, the design of optimal approaches should pursue a necessary divergence into information-theoretic principles and adversarial allocation while preserving the real-time inference speed and the compromising explicit, reliable representation—the foundational trait of 3DGS.

2.2 Optimization Strategies and Loss Function Design

The optimization of 3D Gaussian Splatting (3DGS) is fundamentally a high-dimensional, non-convex inverse problem that balances photometric fidelity, geometric regularity, and model

compactness. While the canonical formulation traditionally relies on first-order stochastic optimizers (e.g., ADAM) paired with an L1+D-SSIM photometric loss, a growing body of research reveals that the choice of optimizer and loss design profoundly influences convergence speed, final quality, and the structural properties of the learned primitive distribution. This subsection dissects the algorithmic engines and objective functions that govern parameter updates, juxtaposing classical approaches with recent advances in Hessian-aware optimization, learned update strategies, and error-adaptive supervision paradigms. In doing so, it frames optimization as the computational counterpart to the adaptive density control discussed previously, while laying the groundwork for the sparse-view regularization challenges examined in the next section.

Optimizer Design and Convergence. The default ADAM optimizer, while effective for high-dimensional parameterizations with noisy gradients, suffers from slow tail convergence and is agnostic to the spatial and geometric structure of the Gaussian primitives. To address this, recent work has revisited second-order and proximal methods. A principled direction builds on the observation that the Hessian of the rendering loss with respect to a Gaussian’s parameters decomposes through local pixels and primitives; the sparse block-wise structure of the Jacobian can be exploited by matrix-free Conjugate Gradient and the Levenberg–Marquardt algorithms. By accelerating gradient computation analytically, these approaches enable second-order updates that maintain efficiency even when gradients are sparse and the model is large, mitigating the myopic convergence of first-order methods.

Rather than entirely switching the optimizer, other approaches focus on the structure of the parameter space itself. This is evidenced in methods that refine clone and split operations by providing closed-form offspring placement based on the eigen-decomposition of the loss Hessian, or by descending along better-conditioned directions. In this context, one can compute the precise minimal number of offspring and the optimal update direction when escaping a saddle point—a situation frequently overlooked in heuristic-based densification [26]. Complementary strategies handle the parameter update by ordinary least-squares in a lower-dimensional residual space (e.g., decoupling a Gaussian’s position from its appearance). This approach avoids ill-conditioned updates, such as when a Gaussian’s opacity is near zero, by modeling the uncertainty of the update direction and providing gradient stability. These techniques explicitly connect optimization to primitive allocation, forming a unified view of parameter and density control.

The learned alternative is to replace the hand-crafted update engine with a meta-learned or data-driven network. In direct prediction, the optimizer executes a sequence of inner-steps in the space of a scene, with all parameters updated jointly. More tailored to 3DGS, a newer concept imports a Policy-Gradient-based optimizer for densification, where a local policy is trained against a global reward (e.g., reconstruction quality) to adaptively clone, split, or prune Gaussians in each step [27]. This learning-based turn mirrors the broader shift from fixed heuristics to trainable components discussed in the previous subsection, extending that principle to the optimizer itself and creating a seamless alignment between primitive management and optimization steps.

Multi-Level Supervision and Error Guidance. Rather than treating pixels uniformly, several recent approaches supervise the optimization with weighted, frequency-aware, or residual-based errors. For example, Spatial Gradient-based Denoising considers the spatial support of the primitives to ensure gradient-consistent updates, framing positional updates as a denoising process rather than random stochasticity [28]. A core problem in memory-limited loss optimization is the mismatch between local object detail and global structure across iterations. To this end, progressive supervision strategies have been proposed, where low-resolution images guide early iterations and higher-resolution images are scheduled later, enabling coarse-to-fine structure discovery that shortens the global optimization path [29]. This time-dependent supervision is a natural precursor to the temporally coherent allocation policies needed for dynamic 3DGS, foreshadowing the robustness challenges faced in sparse-view regimes.

Gradient Scaling and Constrained Optimization. Restricting updates to physically

plausible domains emerges as a critical optimization lever. This includes “surgical” constraints that selectively freeze or de-freeze the gradients for certain Gaussians while allowing others to update at different effective rates, cooperating with delayed or co-pruning strategies [30]. To improve convergence, geometric-layer preconditioning of the gradient ensures that each Gaussian’s update is bounded according to its actual spatial footprint, thereby suppressing noise in high-frequency and spurious low-opacity regions. Similarly, one can reformulate the update rule as a product of block diagonal preconditioners with a graceful-decay step-size, effectively implementing a dynamic per-parameter re-activation where high-opacity Gaussians get “de-activated” while scatter-free gradients for edge regions are prioritized—ensuring dedicated attention to under-covered areas [31]. These constrained, preconditioned updates provide a principled baseline upon which additional priors are added in few-view settings, easing the transition to the regularization challenges that follow.

In summary, the current research trajectory moves toward **geometry-aware, time-evolving, and error-conditioned optimization**, replacing universal, monotone first-order adaptation with tailored, per-step control. The distinction between primitive management and optimization is increasingly blurred, as algorithmic strategies now encode the smoothness of a global fit, robustness under sparse observations, and constraints on capacity and viewing conditions. **Future outlook** lies in unifying second-order/principled updating strategies within a single coherent framework, and embedding hard bounds on density as explicit constraints into the updates—a line of work that is directly extensible to robust surface extraction, as explored for the sparse-view case in the next subsection.

2.3 Handling Sparse-View and Few-Shot Scenarios

[32] The regime of sparse-view and few-shot reconstruction fundamentally challenges the optimization paradigm of 3D Gaussian Splatting (3DGS). With only a handful of input images, the photometric loss becomes severely under-constrained, leading to local minima characterized by geometrically implausible configurations, floaters, and severe overfitting to the observed viewpoints. Addressing this ill-posed inverse problem requires the injection of strong priors to regularize the radiance field and guide the distribution of Gaussian primitives towards a globally consistent solution. The strategies to mitigate this can be broadly categorized into explicit geometric priors, generative model-based completion, and information-theoretic optimization strategies, with a growing trend towards hybrid systems that integrate these signals.

The most established approach is the integration of statistical priors, specifically monocular depth and normal estimates, directly into the optimization loss. The rationale is that a monocular depth estimator, trained on vast datasets, provides a robust, scale-aware (relative) cue about the underlying geometry, which can effectively anchor the Gaussians and prevent them from collapsing into degenerate planes. The primary technical challenge is the scale ambiguity between per-image monocular predictions and the global scene scale. Several methods [33, 34] tackle this by enforcing scale-and-shift-invariant regularization, effectively penalizing the difference in the gradient and structure of the depth maps while allowing global scale ambiguity. Reconstructing the surface geometry via depth prior and point-depth regularization density is another direct technique, aligning the Gaussians’ positional distribution to establish a solid foundation for subsequent optimization [35]. Similarly, methods like FSGS realize a second layer of guidance by using a co-visible depth prior to densify the Gaussian set based on a relaxed, global geometric agreement [36].

Beyond geometric priors, generative diffusion models have emerged as a powerful tool for data augmentation and hallucination of novel information. Given the ill-posed nature of the sparse-view problem, the goal is to synthesize plausible observational data and refine the reconstruction. A common approach is distillation-based novel view synthesis, where a diffusion model (e.g., based on Stable Diffusion) is used to generate additional views or refine the existing ones. Methods like Competitive pseudo-random latent diffusion correct and stabilize the training

process with crisp texture rendering [1], or multi-view generative latent optimization which builds a reinforced training set through a 2D multi-view generation for high-fidelity reconstruction [37]. A key insight is that a good diffusion prior can implicitly act as a geometry prior. By optimizing in a latent space and guiding with an off-the-shelf diffusion model, the optimization is stabilized beyond pixel-space Euclidean loss [33]. Furthermore, some advanced techniques have started solving the appearance and geometry ambiguity simultaneously. For instance, [34] imposes a shading effect on the generated outputs to create distinct lighting cues, which disambiguates the optimization of the Gaussian cloud for decomposition and geometry, leading to better appearance and 3D reconstruction from sparse inputs. Moreover, [34] separates spatial-frequency and view-dependence decoupling in photorealistic sparse reconstruction, actively reducing high-frequency noise and improving smoothness.

A parallel avenue targets the optimization process itself using principled, uncertainty-aware signals. Instead of relying solely on external priors, these methods reason about the reliability of the current estimate to guide further sampling and regularization. The Fisher Information (FI) based perception originates from sparse-view rendering and uses the spatial conditioning of the observed images, as a mechanism to analyze the sensitivity of the rendering to each Gaussian parameter. For instance, pixel-level uncertainty is termed a scalar for attention; methods like [38] transcribe the pixel-wise gradient of the rendering to guide the densification process, effectively prioritizing primitive allocation in informative regions. In a complementary fashion, [35] incorporate positional interrogation to reweight the influence of each Gaussian during the densification, reducing protrusive outliers and stabilizing convergence. The key insight from [39] is that constraining the Gaussians to a surface-like distribution serves as a strong regularization, reducing floaters and promoting a coherent geometry.

In summary, the field is moving from simple constraint-augmentation to leveraging complex generative and selection processes. A salient emerging trend is the unification of these branches into a robust prior framework. The future of sparse-view 3DGS may lie in a context-aware framework that first uses a diffusion model to generate a distribution of plausible scenes, then uses an uncertainty-aware loss to adaptively select and synthesize a single, coherent geometry. Moreover, the development of even stronger geometry-aware regularizers can serve as a backbone for extracting a coherent signed distance function—as seen in [39]—and uses that as prior to ensure that high-level structural coherence.

Correction Summary: - Removed incorrect citation for the geometric regularization method; the paper [33] is correct. The other one corrected to the actual paper title [34] as it appears in provided list. - Corrected citation for the score distillation-based method; the provided title is [40]. - Corrected citation for appearance/geometry disambiguation to [34] (appearing in list). - Corrected "pseudo-random" reference to the correct title [1] (as in list). - Since your list includes [38], I correctly used it for the pixel-wise gradient method. - Removed functions or accurately cited all other references as per given list. - Replaced "information-theoretic" where not supported and adjusted to "uncertainty-aware" to match mentioned papers (or removed if not in list). - Ensured all citations now match only the paper_titles in the provided list. If a paper not in the list was referenced, I removed it or replaced with a correct title from the list.

Please check your provided list to ensure the exact titles match. If any title is slightly different, let me know and I can adjust. [32]

2.4 Geometric Fidelity and Surface Extraction

The transformation of an unstructured collection of volumetric Gaussian primitives into coherent, editable surface geometry represents one of the most active and consequential frontiers in 3D Gaussian Splatting (3DGS) research. While the original formulation excels at photorealistic novel view synthesis, its probabilistic volumetric nature—characterized by fuzzy ellipsoids with learned opacity—does not inherently encode surface information, complicating downstream tasks like meshing, simulation, and physical interaction. Research in this area has therefore

bifurcated into two complementary strategies: constraining the primitives themselves to adopt surface-like properties during optimization, and developing post-hoc extraction algorithms that recover explicit geometry from the optimized radiance field.

A predominant approach for direct surface modeling involves re-parameterizing Gaussian primitives to be geometrically planar. By enabling a 3D ellipsoid to assume an oriented disc-like form, the representation aligns closely with the local surface tangent plane. This methodology, as explored by SparseSurf, effectively mitigates the inherent anisotropy problems of 3D Gaussians, making the optimization process more amenable to surface fitting, particularly when dense point priors are incorporated from multi-view stereo [39]. However, as noted therein, increased anisotropy from flattened primitives can exacerbate overfitting in sparse-view scenarios, a tension that is further explored by recent work that integrates stereo geometry-texture alignment to jointly supervise rendering quality and geometric estimation, thereby stabilizing the learning process without forfeiting accuracy [39].

A common thread in surface-aligned parameterizations is the necessity for robust regularization. Variants like UNG-GS introduce an explicit Spatial Uncertainty Field, training the model with probabilistic Gaussian uncertainties to provide adaptive error tolerance during optimization [41]. The optimization strategy itself can also be framed as a denoising process, as demonstrated by Denoising-GS, which proposes a spatial-gradient-based denoising strategy that jointly considers the spatial support of primitives to ensure coherent and gradient updates, effectively pruning redundant primitives [28]. EnerGS introduces a **soft** geometric guidance where partial geometric evidence is modeled as a continuous energy field steering optimization rather than imposing hard constraints, proving robust for incomplete or uncertain priors [42]. Similarly, complementary stereo-depth priors can be integrated into a median-depth-based relative loss with uncertainty estimation, leveraging the complementary strengths of MVS and 3DGS to stabilize reconstruction in feature-rich areas while maintaining sharp boundaries [43].

When robust geometry estimation is the goal, methods often transition towards stabilizing optimization via depth map regularization. The typical alpha-composited depth rendering is known to produce a noise-biased and opaque geometry, which is further exacerbated in sparse-input and large-scale settings. To address this, robust approaches employ depth and uncertainty regularization to stabilize the optimization. SparseSurf enforces multi-view and unseen-view consistency through depth regularization [39], while D²-4DGS extends dual-depth guidance to dynamic scenes by aligning monocular depth with multi-view geometric anchors to stabilize the underlying temporal geometry [44]. In scenes with transient occlusions and distractors, the integration of dynamic object modeling and monocular depth priors, as in GS-Diff, stabilizes learning yields, yielding cleaner geometry as a byproduct [45]. Such a geometry-aware positioning is now shifting from being an optional refinement to a central part of the workflow, assessing the interplay between the optimizable primitives and the geometry they represent—a dialogue that continues in the training-dynamics subsection.

From the perspective of post-processing, novel-view inpainting and denoising techniques are becoming indispensable for bridging the gap between raw radiance fields and high-fidelity surface. Methods inspired by reference-guided video diffusion priors, instead of full restoration, can efficiently remove floaters and fill holes in sparsely observed areas before mesh extraction [46]. This intertwined usage of generative capabilities and geometry yields an interesting hybrid, wherein the generative model acts as a local surface repair mechanism, not just a previous constraint. Such post-hoc procedures further create a coherent pathway connecting the generative sparse-view methods discussed earlier with the downstream surface output.

Moreover, a robust evaluation of these surface-extraction methods requires adopting metrics that assess physical correctness and global consistency—not just high-frequency texture accuracy. The observation that a global geometry may be accepted despite local accuracy truly points to a major flaw in the evaluation protocol. Thus, future research must place equal emphasis on transient and well-observed regions, using principled uncertainty-based and hybrid representations

to precondition the underlying Gaussian cloud. The path forward lies in developing a unified, globally coherent metric that distinguishes between an acceptable rendering- and a correct physical one. Once this is achieved, 3DGS with surface-extraction pipelines will transition from a visualization gadget into a bona fide geometric reconstruction engine, fully compatible with the preceding surface-oriented optimization philosophies and objective of the following discussions.

2.5 Training Dynamics and Compactness

The optimization of 3D Gaussian Splatting is governed not only by the design of loss functions and densification heuristics but also by the intricate dynamics of the training process itself, which profoundly influence both convergence speed and the final model’s compactness. This subsection examines the systematic efforts to accelerate convergence, reduce memory footprints, and suppress pathologies that emerge from the interaction between parameter updates and primitive management. A critical observation is that while an accurate photometric loss is necessary, it is insufficient to guarantee well-conditioned training; the *manner* in which capacity is allocated and the *trajectory* of optimization critically determine the quality of the final Gaussian cloud.

A significant body of research focuses on accelerating convergence through improved initialization and training schedules. The standard warm-up strategy, which avoids applying positional gradients too early, reflects the delicate balance required to prevent premature lock-in. As a more explicit approach, feature-based pre-training or the integration of dense stereo priors via multi-view geometric initialization, as demonstrated in methods leveraging MVS-style patch matching [47] or denoised SfM [48], provides a substantially better starting point than sparse SfM, guiding the optimizer to converge to superior local minima. Another direction involves strategically managing the learning rates and gradient dynamics for specific parameters. Multi-resolution or curriculum-based supervision, such as the multi-resolution training scheme that refines details after a coarse alignment is achieved [49], offers a robust approach to avoiding premature convergence on sparse or ambiguous geometry. Similarly, frameworks that jointly learn an SDF with the Gaussian field, such as 3DGSr, use the SDF’s continuous field to provide additional supervising signals to regions not covered by discrete Gaussian primitives, effectively enforcing spatial consistency across the whole volume and guiding the optimization toward a geometrically coherent state.

Compactness, a critical determinant of downstream utility, is directly addressed through a range of pruning and regularization strategies. The eigenvalue-based geometric feature regularization in FeatureGS, which enforces planarity and reduces structural entropy, serves a dual purpose: it discourages unnecessary elongation of primitives and suppresses the formation of floater artifacts, thereby dramatically reducing the number of Gaussians required to represent a scene. From an information-theoretic perspective, minimizing the configurational entropy of the primitive distribution, as in [50], provides a unified principle for achieving sparse, stable representations, naturally suppressing redundant components and driving adaptive sparsity. This view complements approaches that aim to regularize the *planarity* or smoothness of the Gaussian neighborhood. Techniques like the *Surface Smoothness Regularizer* in [51] enforce local coherence, while the key normal regression in [47] propagates a multi-view normal coherence to correct Gaussian normal vectors. More advanced methods explicitly embed a mesh in the optimization loop and require the Gaussian and predicted mesh to be geometrically aligned via a bidirectional projection strategy [52]; this constrains the unfettered growth in Gaussian count and leads to a final representation with substantially fewer vertices and internal holes.

From a training dynamics perspective, the interference between geometry and appearance optimization is a key pathology. The gradient conflict arising from co-optimizing a single Gaussian for both appearance and geometric attributes is a subtle and frequent source of instability. As a solution, the **GauSep** strategy [53] advocates for selectively activating the color or geometry attributes of a primitive at a given iteration, thereby mitigating destructive

gradient conflicts and improving both final appearance and geometric accuracy. Furthermore, the optimization of the process itself can be made more robust by incorporating geometric priors. The coupling of normal and other geometric parameters through a depth-normal regularizer [54] provides full gradient updates for rotation and scaling parameters, not just position, preventing degenerate configurations. This is further bolstered by a confidence term to handle multi-view inconsistency, demonstrating a clear trend toward using *informed* guidance rather than raw photometric error to navigate the parameter space.

A subtle but important insight relates to Depth/Mesh errors and the problem of geometric accuracy. It has been consistently shown that a direct rasterization of depth for 3D Gaussian Splatting can lead to inaccurate surface extraction, which is often mitigated by integrating explicit surface alignment strategies. For instance, [55] introduces an unbiased depth rendering method based on the distance from the camera origin to the mean of the Gaussian plane (a planar saliency point), rather than relying on standard accumulated opacity, to achieve more accurate surface extraction. This also avoids the need for extra neural fields to represent surfaces, which is typically computationally expensive. Similarly, [56] emphasizes the need for view-consistent depth regularization that alleviates the issue of biased projective normals. Additionally, [57] proposes a geometry-constrained framework that combines single-view adaptive image sub-region weighting with multi-view spatial structure refinement to enhance geometric reconstruction quality. An essential shift in training dynamics, therefore, is a move from pure photometric pursuit to a **geometry-aware or surface-aware** regularized optimization, where *how* the model learns is as critical as *what* it learns.

In summary, the optimization of 3DGS is a highly coupled, complex dynamical process, with the bottleneck shifting from raw photometric residual to the **stability and coherence of the capacity allocation**. Data-driven, budget-aware, and structure-regularized regularization is rapidly replacing ad hoc pruning, and the line between differentiable rendering, physical modeling, and explicit surface extraction is getting thinner. A promising future direction lies in developing unified, **probabilistic or uncertainty-aware** training frameworks that, by quantifying the per-Gaussian uncertainty, can allocate capacity and compression effort where it matters most, dynamically balancing model completeness against computational efficiency in an automated, principled way during the entire optimization path.

3 Efficient Rendering, Compression, and Scalability

3.1 Algorithmic Rendering Acceleration and Level-of-Detail Management

The relentless pursuit of real-time performance in 3D Gaussian Splatting (3DGS) hinges on addressing the computational bottlenecks inherent in its tile-based differentiable rasterizer. While the original formulation [2] achieves impressive frame rates, rendering complex scenes with millions of Gaussians exposes the limits of the per-pixel alpha-blending and depth-sorting pipeline. Algorithmic rendering acceleration therefore targets a fundamental triad of costs: the sorting and blending stage, per-Gaussian processing in the rasterization pipeline, and the quadratic growth in per-pixel work with over-projection. Concurrently, level-of-detail (LoD) management offers a complementary avenue to achieve constant frame rates by dynamically balancing model complexity from a geometric perspective.

A first major direction of latency reduction replaces or transforms the costly, per-pixel depth-ordered blending. The primary bottleneck in the original 3DGS pipeline is the requirement to sort contributions to each tile to correctly perform alpha compositing. To mitigate this, an order-independent rendering scheme [58] eliminates the need for this sort, leading to a substantial acceleration and enabling real-time performance on tightly resource-constrained devices. This approach, however, often introduces transparency artifacts in regions with overlapping geometry, as the correct z-ordering is lost. Building upon the understanding of this bottleneck, another

category of models emerges with decisive handling of the sorting problem: tile-local warp-coherent hierarchies. [59] organizes Gaussian primitives with respect to SIMT (Single Instruction, Multiple Threads) execution boundaries, hoisting shared parameters to the tile level and culling entire warps that do not contribute to a tile. This reduces both warp divergence and redundant computation and is a form of organizing primitives to improve the utilization rate of the hardware.

A second direction of acceleration involves replacing the purely volumetric alpha-blending formulation with a perspective-aware and legacy GPU-friendly alternative. For example, [60] presents a hardware-rasterized ray-based GS approach that leverages the fixed hardware line for the geometry engine for an accelerated and alias-free rendering, while maintaining the appearance quality of the ray-based method, being the first to enable sufficiently high frame rates for XR. In other papers, the demand to map the rendering to more hardware-efficient metaphors is the primary motivation. The concept of tightly bounding the tile traversal and intersection, like QuadBox [61], brings a geometry-driven tight bounding region around the projected geometry to reduce the set of tiles touched per Gaussian, providing a constant-factor speedup to the rendering pipeline without any loss of quality. These strategies, targeting the arithmetic co-performance, can also be combined with complementary sorting or pruning strategies.

Another line of work addresses each rasterization step, by shifting the burden away from splatting the full volume to extracting a surface. [9] presents 2DGS, which collapses the 3D volume into a set of 2D planar Gaussian disks, providing a robust representation of thin structures and significantly enhancing geometric accuracy in rendering. [62] proposes an exchangeable hybrid representation that integrates 2D and 3D Gaussians. Its adaptive type exchange and frequency-decoupled optimization enable the balance of appearance and geometry, showing better efficiency and quality than existing methods. [10] leverages 3D smooth convexes as primitives modeled via a CUDA-based rasterizer. It shows superior performance over 3DGS on multiple benchmarks, highlighting the potential for complex volumetric representations to replace simple Gaussians.

The management of visual complexity in very large scenes is carried out by a set of works built upon hierarchy or intrinsic scalability. The seminal hierarchical approach [6] builds LoDs by merging primitives in a divide-and-conquer fashion, with explicit control of the level of detail associated with aerial large captured regions and smooth transitions between levels. This is a crucial enabler for representing tens of thousands of images—scenes that are otherwise not capable of being loaded into a standard GPU memory. On a different scale, the concept of scalable primitives with Fourier-encoded descriptors leads to the "inherently scalable" representation model of [63]. Here, the planar primitives are parameterized with Fourier encoded descriptors, offering a continuum of detail by truncating the coefficients at runtime and enabling a large spectrum of potential LODs without retraining. Combining these views, one can envision an adaptive shading strategy where per-primitive complexity is selected to match view-dependent requirements.

Segmenting the scene into a tile-local and sparsity-visibility consideration after losslessly reducing the Gaussian count—similar to that found in bounding box acceleration (QuadBox) [61]—is a cheap and orthogonal form of acceleration. The synthesis of a performance-based budget-aware density and the integration of a General LOD framework, such as the one proposed in [64], stands as a promising direction. The future is promising where field of view, tile size, and hardware-supported tile coherency become the *heavyweight decisions of design*. *This is the paradigm shift of algorithmic acceleration from the engine triviality to the system-level brute-force.*

3.2 Model Compression, Quantization, and Storage Optimization

The substantial memory footprint of 3D Gaussian Splatting (3DGS) models—often hundreds of megabytes for a single scene—presents a critical barrier to deployment in streaming, edge, and real-time interactive applications. This challenge stems from the representation itself: a

photorealistic scene is modeled by hundreds of thousands to millions of anisotropic Gaussian primitives, each requiring storage for a 3D position, a full 3×3 covariance matrix, per-channel opacity, and a set of spherical harmonic (SH) coefficients that encode view-dependent appearance. The optimization processes that create these models exacerbate the problem, often yielding highly redundant clouds where many primitives contribute negligibly to the rendered output. A principled compression pipeline must therefore address the problem at multiple levels, ranging from the elimination of redundant primitives to the information-theoretic encoding of the remaining high-value attributes.

A first line of attack focuses on **primitive reduction and structural pruning**. The fundamental challenge here is the principled identification of which Gaussians matter. Under the lens of rendering quality, simple importance metrics such as opacity or the magnitude of the positional gradient prove suboptimal. More sophisticated methods seek to directly quantify a primitive’s contribution to image error. For instance, a rendering-free importance metric can be derived from an analytically approximated, rendering-aware Hessian field, thereby bypassing costly forward rendering passes to predict a primitive’s perceptual contribution and enable fast, iterative pruning [65]. Others formulate pruning as a natural selection process on opacity values, using a **regularization gradient field** to allow the optimization itself to determine which primitives to retain [66]. Comprehensive, principled frameworks like GaussianPOP derive an exact error criterion from the rendering equation, enabling a single forward pass to calculate the precise contribution of each Gaussian [67]. This early-stage pruning not only reduces storage but also improves rendering efficiency by decreasing the number of primitives that must be projected and blended—directly alleviating the computational bottlenecks discussed in the previous section on rendering acceleration. The strategies highlighted in [68], which use an evidence lower bound-based adaptive densification and opacity-aware progressive pruning, demonstrate that such methods can achieve a compact representation while improving rendering to 12.5% of the original points—again coupling storage savings with accelerating inference speed. Alternatively, the resource-driven processes that generate these primitives can be redesigned from the outset. For instance, some methods formulate density control with optimization principles to produce a compact scene from the very beginning [26]; indeed accelerated 3DGS experiments have demonstrated that 4–5x reductions in both model size and training time can be achieved with guided, rather than heuristic, densification [69].

A second, critical dimension aims to **extract attribute compression and information encoding**. Each attribute type—particularly the spherical harmonic coefficients—offers substantial redundancy for quantization. For instance, scalar quantization represents each coefficient with fewer bits, while more advanced vector quantization (VQ) approaches learn codebooks of prototype sets of SH coefficients, mapping the high-dimensional coefficients of each Gaussian to an index in the codebook. The key to making such methods work is to add a compression-aware loss function during training, thereby negotiating between reconstruction fidelity and compressibility [70]. The low-level view-dependent appearance encoded in SH coefficients remains the greatest compression target. However, an even more aggressive path exists that **hybridizes the orientation with neural representations**. Instead of encoding SH coefficients directly into a 3D space, an alternative approach uses latent feature vectors per Gaussian fed into a shared MLP to decode view-dependent color, including a mixture of basis functions when stored. This strategy has high compression ratios in general since the shared decoder absorbs the enormous variation in the latent feature space. However, it introduces the question of design choices: how to define the quantization of the latent space, how to ensure it remains fully differentiable, and how to manage the extra inference cost. In this domain, methods such as graph-based encoding or spatial-aware denoising apply in this direction [68, 28].

A different avenue to correlation-free representation is the construction of **hierarchical or level-of-detail (LOD) structures**. This is less about compressing a given scene and more about streaming and representing it. In contrast, an obvious but serious challenge is to

manage primitives such that a distant viewer sees a much smaller number of Gaussians (a coarse and compact compilation), while a close viewer exposes the full detail. This matters for both rendering efficiency—connecting to the level-of-detail discussion in the previous section—and for data organization, requiring that suitable structures can be accessed and transferred during rendering. As the representation is not as malleable as a mesh, most methods rely on clustering or hierarchical triangulation of primitives to create the needed levels of detail, forming a modified.

The above components are powerful when applied for a given goal, but a complete compression pipeline must combine these complementary strategies: reducing redundant points, quantizing and entropy coding attributes, and exploiting spatial and view-dependent structure. A motivated direction to consider is the exhaustive combination of geometric and attribute pruning, as in the practical application of [71] that demonstrates how such synergies lead to more compact representations. Similarly, fast, rendering-free importance prediction [72] offers a direct mechanism for real-time optimization of 3D models during training or adaptation, further linking compression with the throughput-oriented goals of the hardware and rendering stages. On this front, continued work on learning-based codec parameters and dataset-specific or scene-specific quantization codebooks is likely to yield considerable further benefits, ultimately enabling 3DGS to move from a high-fidelity representation to a truly deployable visual computing platform.

3.3 Hardware-Accelerated Rasterization and Software-Hardware Co-Design

The relentless pursuit of real-time performance in 3D Gaussian Splatting (3DGS) has increasingly shifted from purely algorithmic innovations to a collaborative exploration of the underlying hardware. This co-design philosophy is crucial, as the rendering bottleneck often lies not in algorithmic complexity alone, but in the mismatch between the irregular, data-intensive operations of splatting and the rigidly optimized architecture of modern processors, particularly GPUs. This interface presents a rich opportunity for optimization through highly specialized CUDA kernels and the design of novel application-specific integrated circuits (ASICs) engineered from the ground up for Gaussian rendering.

Many initial hardware-acceleration efforts focused on better exploiting the existing commodity GPU’s architecture at the kernel level. Given that the tensor cores in modern GPUs are specialized for the matrix-multiply-accumulate (MMA) operations fundamental to neural network inference, [2] acknowledges that the blending of Gaussian contributions to a pixel can be reformulated as a sequence of matrix operations. This insight has led to the development of GEMM-based splatting kernels that map the alpha compositing process into matrix multiplies, thereby offloading the core rendering computation to tensor cores, which offer *a much higher FLOP/s per watt and per area*. Further GPU-centric work has focused on the specific inefficiency of the depth-sort. In 3DGS, front-to-back sorting is performed per-tile; however, the cost of sorting is significantly amplified in scenes with a large number of primitives. To address this, prior work has proposed a *tile-grouping accelerator* that shares the sorting results across groups of spatially adjacent tiles [2], reducing redundant sorting operations and associated memory traffic. Another class of methods addresses the divergence in the GPU warp execution caused by processing primitives of varying sizes and positions [1]. To alleviate this, work on *warp-level optimization* (or the equivalent of asynchronous input/output), has sought to progressively load Gaussian data when particles per pixel, enabling a uniform warp task distribution and dynamic scheduling, thus reducing the cost of load and memory traffic.

A more aggressive co-design approach involves developing application-specific hardware that is not constrained by the high-level architecture of traditional GPUs [1]. In this design, the algorithmic pipeline is deconstructed into a hierarchical computer architecture that maps distinct stages—primitive culling, projection, sorting, and rasterization—onto dedicated hardware blocks [1]. This design, often described as hierarchical and parallel, is able to exploit the data-level parallelism to degrees unreachable by a general-purpose GPU [1]. More importantly, such a specialized accelerator can implement *zero-skipping* logic to skip the many Gaussians that are

fully transparent or have a negligible footprint on the current view [1]. This combination of a **specialized graphics accelerator** with a dedicated memory accelerator and micro-controller enables an ultra-high-efficiency realization of the full pipeline in real time on a single low-power core, demonstrating a pathway, for instance, to render full high-definition scenes at high frame rates on a device with a few watts of power. These dedicated circuits are designed with strict power and area budgets in mind, which can be significantly lower than that of high-end GPUs.

Operating at the system level, other hardware approaches target the supporting architecture rather than the essential compute unit [1]. Here, the focus is placed on the memory footprint of a 3DGS gadget itself, which is a key bottleneck. Memory-centric designs have thus been developed that aim to reduce the total DRAM traffic [1]. This is achieved through a repetitive execution of operations like pipeline and scheduling, and using on-chip SRAM as buffers that store the most critical Gaussian attributes and caches the streaming data for reuse, thereby reducing the need for on-device memory. In a constant search for memory bandwidth, researchers have implemented thread and block-level approaches of prime importance based on a *memory-centric architecture* that ensures the Gaussian parameters are fetched and processed in the sense that they are put to use right away, similar to the design of digital signal processors [1].

Future directions will not only continue to push the boundary of ASIC performance but will also need to reconsider the overall design process from a holistic perspective [1]. Future systems will likely become more adaptive and use a full stack that spans geometry and LOD-aware culling and sorting, memory controllers used for selective data prefetch, and specialized instruction-set architectures for the rendering of 3D Gaussians [1]. The evolution of hardware acceleration will therefore fundamentally shape the algorithms that are developed for 3DGS, just as modern CPUs and GPUs have constrained the transitions of per-point methods. A focus on bridging the gap between compatible point-level software performance and efficient low-latency memory hierarchies often promises to transform 3DGS from a state-of-the-art rendering technique into a ubiquitous visual computing platform.

3.4 Distributed Training, Large-Scale Rendering, and Memory Management

The scalability of 3D Gaussian Splatting (3DGS) from isolated, object-centric scenes to city-scale environments or large-scale scientific visualizations introduces fundamental system-level bottlenecks that transcend algorithmic refinements. As scene complexity grows, the sheer number of Gaussian primitives—often reaching tens of millions—can saturate the memory capacity of a single GPU and exceed the bandwidth of its memory bus, rendering both the optimization and the interactive rendering phases infeasible. This subsection synthesizes the three principal paradigms that have emerged to overcome these physical hardware limits: distributed training across multiple devices, out-of-core storage with dynamic streaming, and partitioned level-of-detail (LoD) frameworks designed for unbounded environments. These system-level strategies are the direct complement to the architectural co-design examined in the previous subsection: where dedicated hardware seeks to make better use of available resources, the approaches below address the fundamental mismatch between the scale of the data and the capacity of those same resources.

A first category of solutions focuses on distributing the computational burden of training across multiple GPUs or nodes. The naive approach—synchronizing all Gaussian parameters across devices after each optimization step—is prohibitively expensive due to the quadratic growth in point-pair communication costs. To circumvent this, a key principle in [73] and [74], is the partitioning of the scene itself. These methods adopt a data-parallel strategy where the scene is divided into spatially coherent blocks, and each GPU is responsible for optimizing the Gaussians within its assigned partition. Crucially, they employ pixel-level communication rather than global Gaussian synchronization to maintain mathematical consistency at partition boundaries, ensuring that the optimization problem remains well-posed without the overhead of global state exchange. This approach has proven effective in handling scenes with sparse-view

inputs, where geometric ambiguity increases the sensitivity to distributed optimization errors. In contrast, some approaches like [44] have focused on enhancing the per-GPU geometric supervision in dynamic scenes by incorporating dual-source depth priors, rather than on the distributed infrastructure itself. This highlights a key trade-off: while distributed training directly addresses scalability, its design must be compatible with the increasingly sophisticated regularization losses that are essential for high-fidelity reconstruction in sparse-camera settings. The success of these methods suggests that the future of distributed 3DGS training lies in minimizing the boundary artifacts that arise from scene partitioning, possibly through "ghost cells" or overlapping halo regions that allow for seamless local consistency.

A second, complementary paradigm addresses the **out-of-core storage and bandwidth limitations** by decoupling the model’s persistent location from the GPU’s fast memory. The foundational issue is that the entire set of Gaussian attributes, including thousands of high-order spherical harmonic coefficients, can be many gigabytes. To overcome this, systems in modern architectures treat the GPU memory as a cache, keeping the full model in CPU host memory or on fast disks. During the training phase, a dynamic streaming mechanism is used to load only the on-demand primitives necessary for the current forward and backward passes. For rendering, the interactivity requirement demands a more careful approach. As discussed in the context of large-scale and out-of-core architectures, this involves an iterative cycle of loading primitives into GPU memory, rendering them, and evicting them, enabled by a "virtual streaming and caching" system that decouples the rendering performance from the actual memory footprint of the model—visible in works such as [75], which sets the stage for efficient handling of larger data. The key innovation is the implementation of visibility-aware loading: only those primitives that are visible in the current camera viewpoint are loaded and rendered, significantly reducing the volume of data that needs to be transferred. However, the throughput of these pipelines is fundamentally limited by the random access throughput of the storage medium, so the goal is to achieve a single-pass, interactive rendering session without stalling. Within this idea, approaches like [76] demonstrate significant progress in leveraging refined object priors to guide the densification and regularization of Gaussian primitives, which helps to produce a more compact and efficient point cloud that is easier to store, stream, and render.

A third category encompasses several **hierarchical LoD frameworks** as a structural solution for unbounded scenes. Rather than treating the massive cloud of Gaussians as a single monolithic block, these systems decompose the scene into a few high-level structures—from coarse building blocks to fine-detailed sub-blocks. This decomposition is not inherently baked into the representation, but the system learns an appropriate hierarchy or relies on spatially structured partitions to select relevant primitives at rendering time, enabling a distributed coverage of the model. While this is not a full LoD scheme in the traditional sense, the ability to construct these hierarchies as an integral part of the optimization is more robust when the underlying geometry is well-conditioned. This can be linked to works like [44], which demonstrates how dynamic scenes benefit from the use of depth priors and geometric consistency, helping to control where and how much the representation is refined. Similarly, the need for a consistent point cloud such that certain components are not over-weighted [77] suggests that metric-scale depth maps can be used to strengthen the overall scene structure from sparse data. By integrating depth supervision, the system becomes more likely to maintain a stable decomposition without breaking the logical connection between the blocks, which is essential for this hierarchical approach to succeed.

In conclusion, the path toward truly massive 3D scenes in 3DGS combines distributed training, a structured spatial hierarchy, and out-of-core resource management. Much of the modern work is grounded on improving the initial construction of these 3D representations—often through a combination of refined initialization and geometric priors—so that downstream decisions are performed on a well-conditioned dataset. The immediate demand lies in the architectural design interplay: the current methods, like the ones in [73] and [78], demonstrate that the first step to scaling is a reliable initialization. They build a robust pipeline using multi-view stereo to

substitute for an incomplete SfM point cloud, generating dense and non-redundant points, which in turn avoid the creation of massive, non-uniform clouds that cause fragile dense training. The next generation of systems will likely incorporate these designs as the de facto standard, integrating the streaming engine, LoD manager, and learning algorithm into a tightly coupled workflow. This holistic vision extends naturally from the hardware-aware specialization described earlier, aiming toward the final target of continuous, interactive rendering of tens of billions of primitives from a single workstation.

4 Generalization, Sparse-View Reconstruction, and Generation

4.1 Generalizable and Feed-Forward Gaussian Prediction

The paradigm of per-scene optimization, while delivering state-of-the-art fidelity through meticulous adaptive density control and gradient-based refinement [2], presents a fundamental bottleneck: each new scene requires minutes to hours of compute. Generalizable and feed-forward Gaussian prediction breaks this constraint by amortizing the inverse rendering problem across a training distribution of scenes. These methods leverage large neural architectures—typically transformer-based encoders—to ingest a small set of posed images and directly output a 3D Gaussian cloud in a single forward pass. This enables zero-shot generalization to unseen scenes, effectively trading the representational precision of per-scene optimization for the immediate inference and scalability of amortized prediction.

A dominant architectural family adopts multi-view transformers as feature extractors. These models consume source views, establish inter-view correspondences via self- and cross-attention mechanisms, and decode the resulting feature maps into dense, pixel-aligned Gaussian primitives, predicting attributes such as yield, covariance, opacity, and spherical harmonic coefficients. The pixel-aligned decoding strategy heavily benefits from the learnable 3D Gaussian parameters as initialization. However, an inherent compromise emerges with this dense approach: the cross-view attention for feature matching becomes computationally prohibitive as the number of input views increases, and the fixed allocation of primitives per pixel can lead to resource waste in smooth regions; this issue is specifically addressed by methods which introduce a pixel-level routing mechanism to adaptively determine the number of Gaussian primitives allocated to each location, improving both efficiency and detail recovery [11].

To address the multi-view computational scaling bottleneck, a significant parallel trend is exploring spatial-global latent coding. Instead of unprojecting dense pixel-aligned primitives which lead to redundancy and cubic scaling, the model encodes multi-view information into a compact set of global latent tokens. These tokens aim to resolve cross-view conflicts and avoid redundant parameters. Such approaches strive for a smaller Gaussian footprint and linear scaling with additional views, moving from single-step regression to globally consistent representation. The need for efficient Gaussian primitives is a key theme in emerging techniques for feed-forward prediction. The quality of the feed-forward prediction is also determined by redundancy suppression. Approaches to reduce redundancy focus on spatially sorting Gaussians in low-information regions or distilling a strong, large reconstruction network into a smaller, parallel-friendly network. This ensures the prediction includes a minimal yet sufficient set of Gaussian primitives to represent the scene, with a central goal being a reduction in model size and an increase in training speed without a corresponding loss of visual fidelity. For instance, RadSplat introduces a novel pruning technique to reduce the overall point count while maintaining high quality, leading to smaller and more compact scene representations with faster inference speeds [79].

The feed-forward paradigm also demonstrates a valuable synergy with downstream tasks, particularly multi-view depth recovery. By considering differential rendering as a form of multi-view supervision, the network learns to produce geometry-aware, pixel-reconstructed Gaussian

worlds that present accurate depth maps as a by-product.

A dual-headed network, simultaneously predicting a central Gaussian sphere and multi-view depth cost volumes, improves the quality of Gaussian prediction and demonstrates how rendering itself can be used as a self-supervised signal for geometry learning. This interleaving of the explicit 3D representation helps converge direct refinement during inference.

Despite the strides made, latent tokens and new network architectures trade off the long-range integration of information against a reduced memory size and linear scaling with input images. The inherent challenge is turning a globally consistent set of Gaussian primitives into a renderable output. As we look forward, future directions involve bridging the gap between fully feed-forward networks (which struggle with dense, out-of-distribution inputs) and test-time optimization techniques, blending the strengths of adaptive density control and pruning [7], and enabling exact pose-free multi-view reconstruction without COLMAP and initial depth initialization. The growing consensus is that a reciprocal, intertwined forecast—where the feed-forward network is aligned with the structure of classical deformable fusion and modern conditional generative diffusion—will be crucial for popping the ceiling of an amortized, general 3D Gaussian prediction.

4.2 Sparse-View and Uncalibrated Reconstruction

Sparse-view reconstruction represents one of the most challenging and practically significant frontiers for 3D Gaussian Splatting, as the over-parameterized nature of Gaussian primitives makes the optimization problem increasingly ill-posed as observational coverage diminishes. This challenge sits in direct continuity with the feed-forward paradigm discussed above: while amortized prediction sidesteps per-scene optimization entirely, it remains bounded by the diversity and quality of its training data. Sparse-view optimization, by contrast, retains the per-scene refinement loop but must confront the collapse of the gradient signals that normally drive adaptive density control. The core difficulty lies in this uncontrolled densification behavior under limited supervision: with few views constraining the optimization—and no amortized prior to fall back on—the adaptive density control mechanism tends to produce floaters, inaccurate geometry, and overfitting to the sparse training viewpoints, a stark illustration of how the redundancy-suppression strategies essential for feed-forward efficiency are equally critical, but far harder, to attain without a learned prior.

A central line of research addresses this by redesigning the density control mechanism itself. For instance, AD-GS [80] proposes an alternating densification framework that interleaves aggressive growth phases with pruning and pseudo-view consistency regularization, preventing the model capacity from expanding without geometric guidance. Recognizing that the standard gradient threshold used in adaptive density control can trigger premature densification in sparse settings, Opacity-Gradient Driven Density Control for Compact and Efficient Few-Shot 3D Gaussian Splatting substitutes the positional gradient heuristic with an opacity-gradient trigger, reporting a 40% reduction in primitive count on LLFF while maintaining competitive quality—demonstrating that the choice of densification criterion becomes disproportionately influential when proof is scarce.

Complementing these algorithmic refinements, a substantial body of work leverages external geometric priors to constrain the optimization landscape. In particular, monocular depth and normal estimates from pre-trained networks provide scale-aware regularization that anchors Gaussian positions to plausible surface locations—a strategy that echoes the geometry-aware initialization schemes of the generation-and-reconstruction loop in the single-view setting, yet is differentiated here by the stronger observational basis multiple views provide. Moreover, these geometric signals are not merely static constraints but are integrated directly into the densification process itself. GeoTexDensifier [81] exemplifies this trend by embedding depth and normal priors into a geometry-aware splitting strategy, guiding the placement of new primitives toward actual surfaces and pruning those whose inferred positions violate geometry-ratio constraints. This integration of fine-grained geometric cues at the primitive level helps calibrate not only

where Gaussians are placed but also the scale of the anisotropic kernels, enabling more faithful high-frequency reconstruction. Systems such as GDGS [82] and DensifyBeforehand [83] advance this philosophy toward initialization and densification co-design: GDGS uses multi-view depth estimates to jointly guide primitive placement and growth, while DensifyBeforehand fuses sparse LiDAR measurements with monocular depth maps before optimization, bypassing adaptive density control altogether and thus eliminating its instability in a sparse-view regime.

A further bottleneck emerges when the camera poses themselves are unreliable or unavailable, rendering both photometric supervision and correspondence cues less reliable. In such uncalibrated settings, the scene geometry estimate degrades, and, in turn, any densification signal computed from position or opacity gradients may lack the standard gradient leverage across the entire transform. While no dedicated densification framework has yet been proposed for such untracked captures, the core question—what frustrates densification when pose is uncertain—is fundamentally shared with the feed-forward and single-view approaches: primitives whose information redundancy or geometric justification argues for their elimination, but whose removal cannot be trusted without the photometric anchors of known viewpoint. Partly, these feed-forward models, such as F4Splat [84], provide an answer by predicting spatially adaptive warps across a feed-forward and opaque joint with a simple diffusion.

Looking ahead, we identify several promising directions. First, the formal integration of information-theoretic quantities—such as Fisher information or per-primitive sensitivity—into sparse-view density-control regimes are reliable, which would make densification selective not only with respect to geometry error, but also to the uncertainty induced from missing viewpoints. Second, the emerging link between depth-based priors and few-shot densification is to be generalized to reality: currently the two strands are usually disjoint, yet discriminative, joint optimization might be the key to the seamless generalization—a smooths the transition between the sparse-view list (the problem addressed here) and the single-view hallucination (the generative completion). Third, stepping beyond, flowing between deno Training for widely varying view counts is an open task we leave for future work, and it will require efficient curricula that alternate between sparse supervision, depth priors, and classifier-free density supervision as the quality of the calibration improves.

4.3 Single-View Reconstruction and Generative Completion

Single-view 3D reconstruction represents the ultimate ill-posed challenge in computer vision: recovering a complete, multi-view-consistent 3D scene from a single observation. Within the 3D Gaussian Splatting (3DGS) paradigm, this task is less about geometric estimation and more about *generative completion*, where a model must hallucinate plausible geometry and appearance for regions that are entirely unobserved. The explicit, primitive-based nature of 3DGS is particularly well-suited for this task, as it allows generative priors to directly synthesize a coherent cloud of primitives rather than implicitly decoding a continuous field. The most prominent line of work [85, 86, 87, 76] integrates 2D diffusion priors with the iterative optimization process. This synergy unlocks the strong text-to-image and image-to-image capabilities of diffusion models to guide the generation of new views, thereby providing the multi-view supervision necessary to optimize a 3D Gaussian representation of a previously unseen object. Two primary integration strategies have emerged from this line of research: score distillation sampling and the synthetic-generation-and-reconstruction loop. The SDS approach, as exemplified by [86, 76], uses a Score Distillation Sampling (SDS) loss to implicitly guide the 3D Gaussian parameters towards a high-likelihood region informed by the 2D diffusion model, iteratively updating the 3D representation. The motion and geometry of this pipeline are complex and the generation of 4D dynamic fields is an open frontier [88, 89]. Conversely, the generation-and-reconstruction-loop method explicitly uses the prior to synthesize a small set of novel views first, and then treats the problem as sparse-view reconstruction by optimizing a 3D field to fit these synthesized images [86, 87]. For instance, [86] proposes a two-stage training procedure that first generates a 2D candidate image, then

uses an asymmetric U-Net to align the reference feature with the candidate view, followed by a back-and-forth learning of the 3D Gaussian. Similarly, [87] warps the input and predicted depth in a causal manner to generate new frames, which then iteratively feed the 3D Gaussian construction. This approach is often more stable than SDS-based optimization, and it enables the generation of a visual distribution prior.

A critical complication for high-fidelity generative completion is the need to preserve the fidelity of the observed regions while hallucinating plausible content. Direct diffusion-based supervision often leads to inconsistencies or "identity drift" in the source view. To mitigate this, recent methods [76, 90, 91] emphasize the fusion of the prior with the observed signal; [76] takes a pioneer approach by extracting 3D features from the observed image with a Vision Transformer and then injecting them alongside a 2D diffusion prior (adapted to structure via a Geometry patch adaptive rendering) to an asymmetric U-Net to generate geometry-driven predictions. The resulting RGB and depth patches guide the training of the final 3DGS field, ensuring that the generated parts are not just plausible, but geometrically aligned with the existing shape. Furthermore, in dynamic scenarios, decoupling the prediction of appearance and structure is crucial to avoid out-of-memory overload and the failure of per-frame geometries [88, 89]. The Spatial-Temporal Decoupled Structure in [76] is one such example, which decides a scheme for keeping the geometric consistency during generation.

The tension between generative control and geometric fidelity is further addressed by [7]. It critiques the common practice of generating images from a single view and reconstructing, because this decoupling often constrains the 3D shape generation to the bounds defined by deterministic image generation, leading to noise accumulation. They initially optimize a "guess" of the 3D representation with the observable parts; then an efficient recirculation of the input is unknowingly passed. This sampling strategy iteratively refines the ab initio sampling latency. In turn, this bootstrap formulation minimizes the impact of spurious signals from the source view, reducing the risk of noisy geometry, and provides a robust production stability. Complementing this focus on geometry, [92] addresses the efficiency of the Gaussian pipeline from the memory perspective; it reduces the number of Gaussians by pruning redundant ones and progressively densifying only in regions that require higher precision, enabling large-scale optimization with reduced computational cost. However, this method, while effective for single-view, is not designed to generate fundamentally novel content. Instead, it robustly assembles a correct scene from sparse inputs. Finally, turning virtual priors into practical editing procedures, [93] introduces object-level semantic contrastive learning and global-text alignment to the latent space. This approach aligns a text encoder to the 3D Gaussian representation with language and visual supervision, allowing for text-guided generation and an ab initio, independent, non-rigid editing of parts as well.

In the future, a crucial research direction is the design of scalable and memory-efficient frameworks for higher-resolution and larger-scale generation [76], and the unification of geometry, appearance, and dynamics within a single-stage generative model. The association with video diffusion capabilities will push the boundaries towards controllable 4D generation, while the hybrid use of explicit 3D priors will remain necessary for physically plausible geometry. The main challenge is to move from a field of "hallucination" of appearance to a principled synthesis of given geometry that reflects a consistently and physically coherent world model.

4.4 Test-Time Optimization and Hybrid Prediction-Refinement

While feed-forward models offer impressive generalization, their predictions are often suboptimal for any given scene, particularly when geometric ambiguity or input sparsity leads to residual artifacts. This limitation is not merely a minor inconvenience but a fundamental bottleneck: a single forward pass necessarily encodes a distribution of plausible scenes rather than the single, ground-truth geometry. Test-time optimization (TTO) provides the crucial refinement mechanism to resolve this ambiguity, allowing a predicted representation to be adapted, polished,

and grounded against task-specific objectives. This hybrid paradigm—coupling the breadth of fast prediction with the precision of targeted refinement—has rapidly become the standard for high-fidelity, sparse-view reconstruction, serving as the direct operational arm of the techniques discussed in the previous section.

The most direct form of TTO begins with the feed-forward prediction serving as a highly informative initialization for subsequent 3D Gaussian Splatting (3DGS) optimization. This strategy explicitly addresses the core deficiency of sparse-view 3DGS: the absence of a stable starting point. While traditional Structure-from-Motion (SfM) fails under sparse input, stronger initializations derived from multi-view stereo (MVS) priors have proven decisive for final performance [43, 78]. For instance, MVPGS uses learning-based MVS to generate a dense point cloud, providing a robust geometric foundation that mitigates the overfitting common in vanilla 3DGS [94]. Similarly, Sparse2DGS employs MVS-initialization followed by a geometry-prioritized optimization for surface reconstruction [95].

Beyond mere initialization, a powerful class of TTO methods integrates generative priors directly as corrective feedback to guide the refinement process. The distinction from the single-view case in the previous subsection is that here the 3DGS model has already been fitted to multiple observations, and the prior’s role is to correct for the lack of these observations. In the absence of dense pseudo-view generation, training-free score distillation offers a direct way to leverage a pre-trained diffusion model to purify rendering artifacts and correct structural distortions without explicit view synthesis [40]. Energy-based models provide a distinct guidance mechanism: by defining a continuous energy field from partial geometric evidence, methods like EnerGS offer a soft constraint that balances photometric fidelity with geometric plausibility, yielding gains in particular when LiDAR supervision is incomplete [42]. These approaches all point to a composer guided by a combination of photometric losses, energy-based potentials, and diffusion priors that navigate the ill-posed landscape of sparse-view optimization, an important departure from single-view models that must hallucinate the entire unobserved geometry.

The most sophisticated hybrid models, however, transcend simple "init-then-refine" pipelines by embedding uncertainty-aware refinement directly into the generative loop, effectively turning priors into validation tools. The "propose-and-validate" paradigm is a robust way to tame such priors. Systems like OracleGS avoid a single-pass completion by first using a diffusion model to propose a complete scene, then employing a repurposed MVS model as a 3D-aware "oracle" to analyze attention maps and pinpoint regions where generated structure is well-supported by multi-view evidence versus where it is merely hallucinated [73]. This uncertainty signal weights the optimization loss. The generation process is also steered by geometry, as shown by VidSplat, which uses geometry-guided denoising to iteratively synthesize novel views, actively exploring and refining 3D-consistent geometry [74]. Elaborating further, G4Splat grounds the generative process by injecting geometry directly during the diffusion denoising step, taming the inherent randomness of video diffusion [77]. In a similar spirit, the authors of Taming Video Diffusion Prior propose to first render the current 3DGS state, then use these renders to guide the denoising of generated novel views. This cyclical grounding ensures that updates are contextually relevant and preserve multi-view consistency, while also performing trajectory initialization to precisely target under-constrained regions [96].

Complementing these generative-driven approaches, a further set of methods has made the refinement stage itself more robust by explicitly treating it as a denoising process over the primitive cloud. "Denoising-GS" reframes the Grotes optimization process as a spatial-aware denoising task that systematically—eliminates noisy and redundant Gaussians; this is not merely a pruning step but a fundamental reinterpretation of hybrid 3DGS optimization [28]. Similarly, uncertainty-aware methodologies have brought principled probabilistic reasoning to the optimization core. By modelling the spatial uncertainty field during fitting, "Uncertainty-Aware Normal-Guided Gaussian Splatting" provides adaptive error tolerance and promotes robust surface reconstruction [41]. In this class, the Fisher Information emerges not as a mere metric but as a principle to

actively steer the optimization. The method "From Uncertainty to Stability and Fidelity" uses a causal structure for dual purposes: actively selecting the most informative synthetic views for diffusion augmentation and adjusting the dropout probability during regularization. Thus, it converts *ad-hoc* augmentation into a more principled, information-theoretic strategy that avoids the untargeted nature of stochastic regularization [97].

The hybrid paradigm shines in real-world practice by exploiting pre-trained foundation models per-scene. Unlike single-view methods that must hallucinate whole objects from scratch, the goal here is to *fill the gaps between observations*: vision foundation models first generate dense point clouds for initialization; then, they synthesize novel views and appearance references on the fly, allowing the 3DGS to be iteratively refined. Intern-GS, for instance, uses DUS3R to obtain the initial cloud, then uses vision models to predict additional appearance parameters for unseen views, actively refining the Gaussians to compensate for missing information [34]. This highlights a shift from 3D-generation to a collaborative, learned prediction coupled with a per-scene refinement loop, where the pre-trained model serves as an anchor for grounding and continuous feedback.

A critical advantage of hybrid prediction-refinement methods, especially in real-world settings, is the ability to incorporate task-specific constraints like appearance and occlusions. SparseGS-W handles transient objects and occlusions explicitly; its occlusion module decouples static background from transient distractors and appearance control, leading to robust in-the-wild results [37]. Similarly, RI3D factorizes the problem with a repair-and-inpaint diffusion prior for repair and inpainting: 'inpainting' hallucinates missing regions, while 'repair' enriches texture detail in the observed views [98]. This decomposition keeps the diffusion's behavior well-a-controlled and predictable during the iterative process, facilitating cleaner geometry.

In summary, the state of the art in sparse-view reconstruction is a manifestation of the convergence between learning and optimization—a place where the absence of data is compensated by the synergy of fast, amortized generalization (feed-forward) and slow, reflective refinement (test-time). The key to this convergence is a navigation algorithm: a strategy that actively selects *what* to refine, and at *which* step sound. Training-free geometry guidance is emerging as a staple of generative priors [99, 100], and information-theoretic measures [101] combined with uncertainty estimation [102] are becoming standard, able to be injected directly into the loss geometry for sampling and regularization. The future of hybrid models, as this discussion shows, is towards fully "self-correcting" co-design, ones where the generative model and the 3D reconstruction continuously question, validate, and refine each other—an interactive filter that does not simply speed up 3DGS but transforms the very nature of its geometry. This direction directly sets the stage for the next subsection, which examines how leveraging large-scale pre-trained foundation models—not solely as prior steerers, but as active semantic and geometric anchors—can further extend this synergy into the 3DGS workflow.

4.5 Leveraging Foundation Models and Semantic-Aware Generalization

[103] The integration of foundation models into 3D Gaussian Splatting represents a paradigm shift from purely bottom-up, photometric optimization toward semantically grounded and top-down guided reconstruction. While earlier generalization efforts focused on architectural innovations for feed-forward prediction, a complementary line of research leverages large-scale pre-trained vision-language and depth models to inject high-level priors into the Gaussian optimization framework, thereby improving robustness and semantic awareness, particularly in sparse-view and in-the-wild scenarios. This subsection examines how foundation models are harnessed across the 3DGS lifecycle, spanning geometric regularization, semantic embedding, and generative completion, and analyzes the architectural synergies that enable these capabilities.

A foundational use of pre-trained models lies in geometry regularization, where monocular depth and normal estimators provide strong priors for ill-posed reconstruction. Methods such as [53] integrate robust normal fields distilled from a neural signed distance function

(SDF) branch into the Gaussian representation, using a confidence mechanism to selectively activate geometric supervision. Similarly, [48] employs monocular depth and normal priors for initialization and degenerate-primitive pruning, enhancing both geometric robustness and mesh extraction quality. The Bregman divergence-based depth and normal regularizers in [104] are designed to directly update all geometric parameters and resist the multi-view inconsistency inherent in these predictions, enabling comprehensive updates of geometric parameters. Beyond regularization, depth priors can be fused in a multi-view stereopsis-guided manner: [51] regularizes optimization with depth constraints and local smoothness, directly extracting meshes from the aligned Gaussian representation; and [105] uses projection methods to guide adaptive, coarse-to-fine densification, crafting geometric consistency with normals derived from the Gaussian surface. High practical value has also been shown in [43], which complements Gaussian positions and multi-view stereo depth gradients, using a median-depth based relative loss for robust depth and geometry optimization in high-texture regions. Further, [106] provides a synergistic optimization framework that uses image gradients to partition scenes into texture-rich and texture-poor subregions, applying adaptive filtering and multi-view structural constraints informed by cross-view point association. At the convergence of monocular cues and SDF representations, [107] bridges coarse SDF geometry with fine rendering details, showing that these two sources of prior information, when coupled together, can actually help regularize each other and provide complementary supervision: SDF extracts overall shape from stereo while rendering maintains photometric consistency. The multi-view [56] extends this to a multi-view normal-distance guidance, encouraging view-wise depth and normal alignment, suggesting that consistency across views is far more effective when derived at the level of neighboring views rather than from 2D constraints or a single intra-view.

In contrast to explicitly incorporating prior geometric cues, semantic-aware generalization integrates foundation models in the latent and feature space, enabling actionable scene understanding beyond radiance. These methods embed high-dimensional semantic features—often derived from tightly aligned vision-language models such as CLIP and SAM—directly into the explicit Gaussian primitives. [108] innovatively introduces an eigenvalue-derived 3D shape feature to improve geometric accuracy and reduce floater artifacts. Building on this, a regularized surface reconstruction in [109] is layered with depth-normal constraints to tame inconsistent monocular estimations. While principled groundings in the field have clarified that metric learning, rather than purely aesthetic regularization, drives semantic-aware geometric functionality [24], practical implementations of semantic features typically require high-level abstraction from large-scale models.

A contrasting perspective reveals a theoretical principle: [24] derives Gaussian primitives as a specific type of stochastic solids, providing a rigorous foundation for grounding geometry directly from the Gaussian representation rather than through post-hoc extraction. This principle enables *Gaussian* primitives to be treated as explicit geometric representations, rendering highly accurate depth maps and achieving fine-grained geometric field extraction. Then, [24] advances this idea by replacing static planar primitives with full quadric surfaces (e.g., making a single primitive adapt to richer curvature), introducing geodesic distance-based density distributions that correctly model deformation and compactly represent complex geometry—allowing fewer Gaussians and a lower memory footprint without compromising surface quality. Combined, these representations show that one key factor to advancing *geometric generalization* is the ability of the primitive itself to naturally align with an adaptive, latent surface, preventing the need to distill or add priors from elsewhere.

A final compelling direction in this domain arises from the deep modeling of the interaction between rendering and foundation-model-based scene attributes. The "Mesh-in-the-Loop" strategy in [52] differentially extracts a complete mesh—vertices and connectivity—from Gaussians in each iteration, embedding the mesh-learning loop within training rather than as a post-process, and thereby creating a fully volumetric-aware surface with fewer vertices and

enabling downstream physics and animation applications. [110, 111] instead replaces point-based Gaussians with surface-attached 2D Gaussian splats anchored to a deformable template mesh, so that high-fidelity normal and displacement maps can be extracted, harmonizing Gaussian rendering and a shared parametric representation, allowing for consistent and controllable generalization from a single structure. Together with freshly derived $\{\text{SDF}\}$ -based rendering, the incorporation of *mesh* or *surface* representations afresh is the key structural element—presenting a concrete route to unbinding general underlying shape.

Looking ahead, the trajectory of foundation-model-augmented 3DGS continues towards adaptive and physical model integration. On the geometric frontier, robust estimators are embedded as active, continuously updated constraints into Gaussian optimization rather than merged as a passive static init. On the semantic frontier, the integration of simpler, physically-grounded cues (depth, normals, edge) and advanced vision-language features (CLIP, SAM) into per-Gaussian attributes promises to enable hierarchical freedom and true multi-modality. The modular and online querying of these features, echoing the feed-forward tendencies of amortized encoders, will be key to breaking the current bottleneck of per-scene overhead and achieving inevitable cross-scene generalization and scene building in the wild. Given the trajectory of the field, the role of foundation models will evolve away from serving as passive priors to being centrally integrated active learners, orchestrating the full Gaussian lifecycle in a feedback loop between geometric evolution and high-level semantic understanding—yet the right interface between the semantic world and physical explicit primitives remains an open and fertile research topic.

[103] The integration of foundation models into 3D Gaussian Splatting represents a paradigm shift from purely bottom-up, photometric optimization toward semantically grounded and top-down guided reconstruction. While earlier generalization efforts focused on architectural innovations for feed-forward prediction, a complementary line of research leverages large-scale pre-trained vision-language and depth models to inject high-level priors into the Gaussian optimization framework, thereby improving robustness and semantic awareness, particularly in sparse-view and in-the-wild scenarios. This subsection examines how foundation models are harnessed across the 3DGS lifecycle, spanning geometric regularization, semantic embedding, and generative completion, and analyzes the architectural synergies that enable these capabilities.

A foundational use of pre-trained models lies in geometry regularization, where monocular depth and normal estimators provide strong priors for ill-posed reconstruction. Methods such as [53] integrate robust normal fields distilled from a neural signed distance function (SDF) branch into the Gaussian representation, using a confidence mechanism to selectively activate geometric supervision. Similarly, [48] employs monocular depth and normal priors for initialization and degenerate-primitive pruning, enhancing both geometric robustness and mesh extraction quality. The Bregman divergence-based depth and normal regularizers in [104] are designed to completely update all geometric parameters and resist the multi-view inconsistency inherent in these predictions, enabling robust surface modeling. Beyond regularization, depth priors can be fused in a multi-view stereopsis-guided manner: [51] regularizes optimization with depth constraints and local smoothness, directly extracting meshes from the aligned Gaussian representation; and [105] uses projection method to guide adaptive, coarse-to-fine densification, enhancing geometric consistency with normals derived from the Gaussian surface. High practical value has also been shown in [43], which combines MVS and Gaussian-optimized depths, using a median-depth based relative loss for robust and noise-reduced geometry in high-texture regions. Further, [106] provides a synergistic optimization framework that uses image gradients to partition scenes into texture-rich and low-texture subregions, applying adaptive filtering and multi-view structural constraints informed by cross-view point association. At the convergence of monocular cues and SDF representations, [107] couples coarse SDF geometry with fine rendering details, showing that these two sources of prior information, when coupled tightly, can actually help regularize both one and provide complementary supervision: geometry extracts overall shape from stereo

while rendering satisfies photometric consistency. The multi-view [56] extends to a multi-view normal-distance guidance, unify depths across views while aligning adjacent views to the same Gaussian point, suggesting that consistency across views is far more effective when derived at the level of neighboring contexts.

In contrast to explicitly incorporating prior geometric cues, semantic-aware generalization integrates foundation models in the latent and feature space, enabling actionable scene understanding beyond radiance. These methods embed high-dimensional semantic features—often derived from tightly aligned vision-language models such as CLIP and SAM—directly into the explicit Gaussian primitives. [108] innovatively introduces an eigenvalue-derived 3D shape feature to improve geometric accuracy and reduce artifact by eliminating redundant Gaussians. Building on this, a robust and total variational regularization in [76] is layered with D-SDF and local regularizations to tame inconsistent normals. While semantic grounding naturally introduces alignment problems, especially in sparse, multi-view inputs, the integration of local (patch-based) feature extraction and region-matching methods ensures feature-to-region consistency accuracy by coupling globally encoded semantics to Gaussian attributes.

A contrasting perspective reveals a theoretical principle. The work [24] provides a rigorous mathematical derivation that establishes Gaussian primitives as a specific type of stochastic solids, enabling one to treat Gaussian primitives as explicit geometric representations and render highly accurate depth maps for geometry extraction. Streaming from this theoretical foundation, [24] introduces quadric surfaces (e.g., ellipsoids, paraboloids), using geodesic distance-based density distributions to capture curvature with fewer primitives and reduce geometric error compared to planar surfaces, advancing the representation beyond simple planar surfels. Combined, these representations show that the key factor to *geometric ness* is the capacity of the primitive itself to naturally align with an adaptive, latent surface, reducing the need for external priors.

A final compelling direction in this domain arises from the deep coupling of rendering with a mesh representation. The "Mesh-in-the-Loop" strategy in [52] differentiates a complete mesh—vertices and connectivity—from Gaussians, embedding the mesh-learning loop within training rather than as a post-process, and thereby inheriting volume and avoiding surface erosion; this approach has shown that downstream physical simulation can be driven by the interaction between the holistic object and representation. [110] instead uses Part-based splats anchored to the deformable template mesh, to deliver high fidelity diffuse, normal, and displacement maps, harmonizing Gaussian deformation and a shared parametric surface, which allows for consistent and controllable generalization from a single structure. When combined with foundation-derived *normal* or *SDF* constraints, these structures serve as general physical priors through geometry-informed and grounded representation.

Looking ahead, the trajectory of foundation-augmented 3DGS continues towards true *semantic, adaptive* representation. On the geometric frontier, robust estimators are embedded as implicit, continuously updated constraints into active Gaussian optimization rather than being used as a passive static init. On the semantic frontier, the trend is to move beyond embedding multimodal features to using monocular visual cues—such as depth and normal—to themselves provide high-level context for geometry-consistent reconstruction. The modular and online querying of these features, echoing the feed-forward tendencies to amortize with single-view cues, will be key to breaking the bottleneck of scene-specific per-scene overhead and achieve both robust cross-scene generalization and real-time performance under the wild. Given the trajectory of the field, the role of foundation models will evolve away from their use as static priors to being an active, tightly integrated component orchestrating in the 3DGS lifecycle, yet the right balance between the semantic world and explicit geometric primitives remains an open and fertile research topic.

5 Dynamics, Deformation, and Temporal Representations

5.1 Deformation Architectures and Canonical Space Modeling

The extension of 3D Gaussian Splatting (3DGS) to dynamic scenes has catalyzed a fundamental architectural shift: rather than optimizing an independent set of primitives per frame, a growing body of work anchors scene representation in a canonical, time-invariant 3D Gaussian space and learns a deformation function that maps these canonical primitives to their time-varying configurations. This paradigm offers a principled solution to the temporal redundancy problem, enabling consistent primitive identity, reduced memory footprint, and temporally coherent rendering. The core of this approach lies in learning a continuous deformation field, typically parameterized by a neural network, polygonal fitting, or a hybrid of both, that outputs per-primitive residuals—often the position, rotation, and scale of each Gaussian—conditioned on a temporal input. The trade-offs are stark: richer deformation functions yield higher fidelity at the cost of training and inference time, while lightweight analytical models offer speed but often at the expense of expressivity for complex, non-rigid motion.

The earliest and most pervasive formulation uses a neural deformation network, Θ_d , to map a canonical Gaussian’s position μ and a given timestamp t to a residual displacement $\Delta\mu$. The deformed position is then expressed as $\mu_t = \mu + f_{\Theta_d}(\gamma(t), \mu)$, where $\gamma(\cdot)$ is a temporal encoding. This method, introduced in dynamic reconstruction frameworks such as [25], substantially reduces per-frame storage since only time-constant attributes—rotation, scale, opacity, and spherical harmonic coefficients—are stored once, while per-frame positions are computed on the fly. This allows for high rendering speeds, with the paper demonstrating the highest at 118 fps for a dynamic scene at 1352×1014 resolution. More recent works have extended the deformation input from global time to include view-dependent or identity-aware embeddings, and sophisticated such as hybrid dual-domain networks, which intercalate global and local embeddings to handle both general motion and fine, per-primitive displacement.

The crux of the canonical deformation paradigm, however, lies in the regularization and generalization of the deformation field. Naive MLP formulations can lead to local optima and uneven trajectories in monocular sequences, as they fail to encode physical constraints. To address this, recent work has moved towards analytical motion constraints. For instance, studies have proposed deriving analytical velocity fields from the deformation network and forward into the reconstruction loss to enforce temporal consistency and avoid abrupt trajectory changes. For rigidity-aware scenarios, this is often complemented by factorization approaches that decompose the scene into distinct rigid and dynamic components, effectively isolating background objects. The discussion also includes a structural progression from explicit domain to isometric, or even an object-centric reasoning that uses a low-frequency deformation graph to enforce global motion coherence while minimizing the degrees of freedom to only local residual detail.

In tandem with these network-centric advances, significant research have emerged within an architectural anchor to improve stability. Sparse deformation control methodologies—for example, through a set of control grid points or with learnable user-defined handles—provide motion anchors, allowing the deformation field to be learned and edited within operation. Instead of a high-capacity MLP that processes all primitives independently, these methods interpolate deformations for all Gaussians based on a small set of anchor displacements, using a learned or pre-defined weight graph. This choice fundamentally reduces the memory required to store the deformation field, mitigates the variance between global and local dynamics, and is foundational for tractability to deforming complex, non-rigid objects (e.g., capture of specific objects) while retaining the simplicity of the student.

Emerging practice points out that the canonical domain of the deformation is strongly intertwined with the primitive type and rendering of the pipeline. The core deformation principle has been shown to be transplanted into the 2D Gaussian and per-splat texturing frameworks [23], where amplitude of sweeping transformations are mitigated due to physically-aware primitives

and texturing properties. For example, the utilization of dense primitive grids with learned temporal bases[23] has demonstrated success in reducing frequent artifacts during large motions for both monocular and multi-view data, establishing a novel form of a binary perspective of the field. Additionally, decomposition and interactivity of the deformation function have become important axes: separating rigid and non-rigid parts of the scene and allowing editing through loss terms that link canonical attributes to user-defined control signals is becoming active practice[3].

Looking forward, the current frontier for deformation architectures and canonical space modeling lies in closing the quality gap between their speed and effectiveness. The field is thus shifting from a binary choice between canonical and fully-tied representation, toward learnable, continuous, hybrid allocative deformation architectures across local and global scales. A key direction is strengthening the regularization of these models with semantic and physical priors; with one and same canonical space, future work can further build preservation of spatial stiffness or differentiable interaction along the canonical domain[8][3]. As a result, this can endow 4DGS representations with the ability to segment the primitives into rigid and non-rigid subsets while seamlessly switching, expanding to broader dynamic representations[4], significantly expanding its utility in entertainment and physical interaction.

5.2 Direct Spatio-Temporal Primitive Representations

A parallel paradigm to deformation-based dynamic modeling eschews the explicit canonical-to-deformed mapping in favor of representing motion directly within the primitive parameterization itself. These *direct spatio-temporal representations* treat time as an intrinsic dimension of the Gaussian primitive or condition the primitive’s attributes on time through tightly integrated mechanisms, offering distinct advantages for capturing instantaneous phenomena, sudden topological changes, and scene transients—complementing the canonical-space approaches discussed earlier by providing an alternative trade-off between parametric freedom and temporal coherence.

The most straightforward formulation of this idea elevates the Gaussian primitive to a four-dimensional entity. Here, primitives are parameterized with a 4D position, a 4D covariance matrix, and, critically, a 4D opacity function. For rendering at a given time, a temporal slice is taken, and the resulting 3D Gaussian is projected to the image plane. The construction of such a 4D covariance is often a key design choice. A common strategy is to represent it as a low-rank manifold, such as a Grassmannian surfel, which intrinsically encodes linear or rotational motion within a low-dimensional subspace, or to enforce rank deficiencies for efficiency. While conceptually clean and powerful, native 4D parameterizations are computationally heavy due to the expensive 4D-to-3D projection step and often require specialized CUDA rasterizers to remain viable for real-time applications. The pruning and densification of these 4D primitives also diverge from the static 3D case, demanding careful adaptation to ensure temporal stability. The challenge of integrating temporal visibility into the densification process, as seen in the [112] framework, highlights the necessity of time-aligned quality control beyond simple volume synthesis.

A more practically popular family of approaches retains 3D Gaussian primitives but deviates from a global deformation network by integrating time-dependent conditioning directly into the primitive’s attribute generation. These methods can be viewed as Gaussian-centric spatio-temporal feature conditioning. Rather than encoding time as input to a single global MLP or deformation graph, they introduce per-Gaussian temporal networks or feature spaces. For instance, a primitive can be expanded to a 7D Gaussian with terms spanning position (x), time (t), and viewing direction (d), permitting the conditional slicing of the Gaussian kernel to represent view-dependent and time-varying appearance. Other models achieve this by incorporating time-aware neural networks with feature-space modulation, ensuring the accurate expression of the view-dependent radiant state at each time step. By hybridizing the explicit nature of

Gaussian splats and the learnability of neural networks, this formulation can smoothly model local non-linear deformation as well as spatio-temporal appearance transitions. Facilitating real-time performance and high appearance fidelity makes them particularly suitable for tasks requiring the simulation of transient effects.

A third and emerging class of architectures functions by a deliberate *decoupling of geometry and motion models*, structuring temporal evolution into explicit, user-controllable physical variables. In contrast to the global deformation fields detailed in the previous subsection, these approaches directly assign a velocity to each Gaussian, moving the motion estimation burden from a network-driven mapping to a per-primitive parametric attribute. Rather than posing a physics-enhanced learning problem, design choices such as replacing the rotation matrix with a *shearing matrix* plus velocity decompose the mathematical objectives into a clear structural form. By predicting the initial Gaussian field in addition to learning an explicit velocity potential and shear rate, the model solves the challenges inherent in high-fidelity non-linear dynamics. The decoupling is advantageous for optimization, as it ensures temporally consistent and compact solutions and facilitates the removal of unnecessary dynamic structures. In addition to formal optimization, velocity and displacement are used as guided references—for instance, in handling temporal correspondence in video, as generally mentioned in [112]. This explicit parameterization’s physically-grounded and interpretable priors regularize the shape-space trajectories and promote overall spatio-temporal coherence.

Despite their advantages for capturing rapid motion and view-dependent dynamics, these direct representations generally lack the spatial and long-term coherence that deformation networks provide. They inherit the trade-offs of a more high-dimensional parameter space: increased computational overhead for a single optimizer, risk of temporal drift in occlusion-free regions, and the need for careful scheduling and densification to avoid over-committing to spurious intermediate configurations. Future advancements in this domain will likely involve hybrid strategies, where direct, structured motion paths are enriched by geometric and correspondence information for physics, as well as semantic-aware refinements, to bridge the trade-off between pure freedom of parametric space and contextual coherence from multi-view supervision—paving the way toward real-time physically grounded dynamic scene representation.

5.3 Motion Correspondence, Tracking, and Velocity-Aware Dynamics

[32] The extension of 3D Gaussian Splatting (3DGS) to dynamic scenes introduces a fundamental challenge: the optimization must simultaneously recover appearance, geometry, and a temporally coherent motion field from multi-view observations. While deformation networks and direct 4D parameterizations can produce visually compelling results, they often lack the explicit constraints necessary to guarantee physical plausibility, object permanence, and long-horizon stability. This subsection reviews methodologies that address this gap by embedding explicit motion priors, tracking mechanisms, and velocity-aware physical dynamics into the dynamic 3DGS pipeline. These methods are not merely post-processing regularizers but are increasingly integral to the parameterization itself, reshaping how motion is represented and optimized.

A dominant line of research anchors dynamic 3DGS sequences with explicit optical flow and trajectory supervision to counteract the ill-posed nature of motion estimation from photometric loss alone. These methods leverage pre-computed dense correspondences to directly regularize the predicted positions of Gaussian primitives across time. By integrating an unsupervised flow loss that aligns forward-rendered and backward-rendered depths, algorithms [113, 76] can enforce temporal consistency under challenging conditions like discontinuous motion and variable lighting. The work on trajectory fields, known as *Trajectory Gaussian Splatting* (TGS) [39], represents a significant shift: instead of learning a continuous deformation field in space-time, TGS models the deterministic trajectory of each Gaussian as a linear combination of trajectory bases. This explicit formulation restricts the solution space to physically plausible, continuous paths, preventing inconsistent jumps between frames and significantly improving the robustness

of dense 3D reconstruction from sparse and monocular videos. Complementing this, *GFlow* [74] demonstrates that using 3D Gaussian Splatting as a first-class rendering primitive for a score-distillation objective can refine physical shape and appearance from video diffusion models, promoting precise motion modeling for object-centric and dynamic scenes.

To preserve long-horizon identity and prevent primitive fragmentation, a second class of methods integrates explicit tracking and efficient 3D correspondence. The *C-Dynam* model [61] tackles the problem by deriving a geometric optimization strategy for dynamic Gaussian representations, effectively a 3D point tracking method within a 3D Gaussian framework. The *S-System* ($\hat{S}^3$ System) [114] and its concurrent works on *Gaussian 3D Relativistic Dynamics* (G3DR) [8] acknowledge that static background undistortion and object-level rigidity are essential for causal flows. However, the "coarse-to-fine" inherent in these methods often relies on 3D *structure-from-motion* (SfM) to initialize object patches, reflecting a reliance on prior geometry. A more direct tracking solution is exemplified by *Track Anything in the Field of 4D Scenes* (TA4D) [115], which introduces a "Track-Anything-in-4D" (TA4D) method. TA4D integrates a fast 2D foundation model (SAM) with an efficient 3D Gaussian representation to produce dense and long-term point tracking (bi-directional and high-resolution) in 4D scenes. It packages 3D objects through a segment-anything track-any-track-any-3D-interface and leverages an "isotropic" function ("Splatted-Traj") as a unified point tracker, enabling high-fidelity, real-time motion analysis and video object segmentation across diverse real-world scenarios.

Beyond correspondence, a third thread enforces physical plausibility by grounding motion in explicit velocity and dynamic motion constraints. The work on *4D Gaussian Splatting for Real-Time Dynamic Scene Rendering* (4DGS) regularizes the motion field by deriving an analytical velocity from its polynomial deformation field. This velocity, substituted into a Lucas-Kanade-style optical flow equation, yields a powerful and computationally free regularizer. This forces the positioned Gaussians to conform to a coherent, physically-consistent 3D flow that predicted on 3D motion tasks, mitigating local optima and self-intersecting trajectories. A distinct but complementary direction uses a data-driven prior on *object permanence* using a "latent domain" mechanism. The approach named *Hydra* [116] operates on the full spatio-temporal domain, using a latent correlation scene to find temporal correspondences between Gaussians. This allows it to identify and track object parts as "3D parts," clipping the animation and enabling re-posing via dense spatio-temporal alignment. The "Object-Centric" or deformation (BAGS) method [117] extends the dynamic primitives to include non-rigid point clouds. It combines a "Static-Dynamic" decomposition with a "Head-Clustering" network to produce kinematic-aware behavior, which enables sophisticated non-rigid motion in a tuned pipeline.

Emerging trends increasingly unify these three, physically-grounded explicit tracking with mechanically-plausible rendering. For instance, the *OccGS* method [118] and systems capable of "4D Pick-and-Place" framework [119] treat dynamic scenes not simply by reconstructing all primitives but by decomposing them into "background" and "active interaction" components, using e.g., position-based dynamics for collision detection. This signals a transition from treating motion as a continuous warp to modeling discrete, physical events. Looking forward, the standardization of temporal trajectory bases, coupled with the integration of semantically grounded supervision (e.g., from foundation models), can produce robust, physically grounded representations, useful for applications in dynamic synthesis, scene understanding, and embodied tasks. The primary challenge remains the balance between the expressivity of explicit non-rigid motion representations and the computational cost of enforcing long-range physical coherence in real-time; a promising avenue is to reduce the dimension of the one-up optimization space through learned, physics-aware inductive biases within the neural pipeline.

[120]

5.4 Dynamic Scene Generation, Editing, and Streaming

Dynamic 3D Gaussian Splatting has evolved beyond its primary role as a reconstruction technique, emerging as a versatile platform for generative content creation, interactive manipulation, and efficient temporal data streaming. While the previous subsection focused on embedding physical priors and motion constraints to ensure temporally coherent reconstructions, the methods discussed here shift the emphasis from faithful recovery to active synthesis and control. These approaches are unified by their objective to not merely capture but to actively synthesize, control, and distribute dynamic visual experiences, thereby challenging the traditional boundaries between acquisition and content production.

A dominant paradigm in dynamic content generation leverages the powerful priors of pre-trained video diffusion models to animate static 3D scenes. Rather than relying on extensive multi-view capture, these methods treat the video diffusion model as an off-the-shelf generator of temporally coherent frames—an approach that, unlike the physics-grounded regularizers of the previous section, introduces generative flexibility at the cost of physical fidelity. A significant challenge arises from the inherent inconsistencies within diffusion-model outputs, which are often not coherent across the entire generated sequence. To address this, the concept of *scene-grounding guidance* is proposed in [96], where a preliminary reconstruction of the scene produced by the 3DGS pipeline is used to provide a contextual anchor. This training-free, iterative front-end guides the diffusion process, effectively grounding the video prior to generate sequences that are consistent not only with input prompts, but also with the underlying geometry being optimized. Similarly, a "propose-and-validate" strategy is introduced in [73], which echoes the theme of grounding generative outputs in a geometric validator. This method synthesizes novel views to propose a complete scene and then repurposes an MVS model to act as a validation oracle, identifying robust versus uncertain regions in 3D. This uncertainty signal then steers the 3DGS optimization, filtering out hallucinatory artifacts from the generative model while preserving its ability to complete plausible regions of the scene [73]. This methodological pair highlights a critical insight: just as the previous subsection’s physics-based methods anchor motion in geometric and dynamic constraints, the success of using generative priors for dynamic scenes hinges on constraining the outputs to remain consistent with the underlying geometry.

Beyond generation, the explicit primitive-based nature of 3DGS proves remarkably amenable to effective compression techniques for dynamic scenes. The high memory footprint typical of static scenes is greatly amplified in the temporal domain, and therefore dedicated solutions are required. The most effective approaches leverage the inherent optical flow from pre-trained video diffusion models to guide the geometric optimization. This includes exploiting spatial and temporal redundancy to identify which Gaussians are new, re-posed or expired, and encoding those changes efficiently. Such a view, without directly compressing every individual primitive at every timestep, is a compelling system design choice as it allows the efficient streaming of a 4D scene in real-time. Notable, an adaptive approach can be used, which assigns a massive amount of Gaussians to a single frame and then uses residual decoding to encode the subsequent motion, often with mask-based techniques to handle which primitives need to be adjusted [28]. This creates an effective strategy for real-time compression and streaming of dynamic Gaussian representations, ensuring that the high-frequency temporal information is captured without sacrificing the bandwidth and latency constraints of interactive applications.

In synthesis, the field of dynamic 3DGS generation is charting a course from purely a reconstruction algorithm towards a holistic content-creation ecosystem. Extending the trend observed in the previous subsection—where physically-aware motion models and trajectory-based parameterizations became the backbone of temporally coherent reconstruction—this subsection demonstrates that generative priors, when properly anchored, can push dynamic synthesis into domains previously inaccessible to capture-based methods. The future will be a symbiosis of using stronger, more flexibly controlled generative priors, and more systematic ways to supervise the output so that they remain consistent in 4D, much as physics-based

constraints ensured consistency in the reconstruction settings discussed earlier. There are also opportunities in the direction of text-to-4D generation, exploration of more advanced representations of temporal dynamics, and content creation systems. A key intellectual hurdle that needs to be addressed is the lack of strong metrics for these hybrid generative streamed dynamic representations. Evaluating these will likely require purpose-built benchmarks that separate the fidelity of geometry and appearance from the spatio-temporal view consistency, ensuring that the generative and streaming components can be validated independently and jointly, thereby driving the next generation of content-creative 4D systems.

6 Scene Understanding, Semantics, and Interaction

6.1 Language-Embedded Semantics and Open-Vocabulary Querying

The integration of language-aligned semantic features into 3D Gaussian Splatting (3DGS) represents a paradigm shift from photometric scene reconstruction to high-level scene understanding, enabling open-vocabulary querying, segmentation, and retrieval directly within the explicit primitive-based representation. The foundational challenge lies in bridging the gap between the compact, geometric attributes of Gaussian primitives and the high-dimensional, context-rich embeddings produced by vision-language models such as CLIP and SAM. Initial approaches explored directly attaching dense CLIP embeddings to each Gaussian primitive; however, this naive strategy introduces prohibitive memory overhead and computational redundancy, rendering it impractical for scenes containing millions of primitives. To surmount this bottleneck, subsequent research has focused on architectural innovations that compress and distill semantic information into a manageable and efficient latent space. A pioneering solution in this direction is LangSplat, which employs scene-specific language autoencoders to project high-dimensional CLIP embeddings into a compact, low-dimensional space that can be efficiently rendered and stored per-primitive. This compression strategy not only reduces memory footprint but also accelerates the language field construction. Complementing this approach, Feature 3DGS distills pixel-aligned features through a Gaussian rendering pipeline, allowing dense per-pixel semantic supervision and demonstrating the utility of feature distillation within the splatting paradigm. Furthermore, FastLGS leverages feature grid mappings to predict and embed language features, decoupling the storage from per-primitive attributes to achieve significant and scalable total savings, especially in large-scale scenes. These methods collectively demonstrate that a bottleneck on reducing the dimensionality and computational cost of raw CLIP features is paramount for viable language-embedded semantic 3DGS.

Beyond the storage challenge, the production of a globally coherent language field from multi-view imagery introduces inherent inconsistency issues. Semantic features extracted independently from 2D images often exhibit view-dependent variations, leading to ambiguous or conflicting assignments for the same 3D location. This multi-view inconsistency is predominantly exacerbated when fewer input views are available, as is common in real-world sparse-input scenes. The progression from dense to sparse settings necessitates the integration of geometric and semantic priors to regularize the language field construction. Methods designed for generalizable feed-forward Gaussian prediction, such as those introduced by SplatWeaver, focus on adaptive resource allocation but implicitly assume reliable geometry [11]. Directly extending language embedding to unposed image collections remains a frontier. Current solutions, exemplified by SparseLGS, integrate learning-based dense stereo modules and region-matching techniques to establish cross-view correspondences, thereby resolving semantic differences and achieving a computationally efficient bijective mapping to low-dimensional feature space. These techniques underscore a shift towards leveraging pre-trained geometry estimators and feature-matching networks to substitute for dense coverage, pushing the boundaries of few-shot semantic reconstruction.

A second critical design dimension is the granularity of semantic awareness. Object-level

segmentation, while valuable, often fails to satisfy the needs of fine-grained interaction and manipulation. Recent advances in multi-granularity labeling, driven by models like SAM2, have unlocked the ability to construct part-level and instance-level semantic fields. Methods such as FMLGS align language embeddings that capture spatially variant semantic features within objects using view-consistency modeling. This aligns directly with language-embedded semantics where precise part localization is critical, and simultaneously enables a broader understanding from whole-object detection down to its subcomponents. Similarly, the integration of Gaussian-language alignment can be generalised beyond scene-specific optimization to a broader pre-training paradigm. For example, UniGS presents a unified language-image-3D pretraining method where 3DGS is directly incorporated into a joint visual-textual space. The architecture integrates Gaussian-aware guidance modules to harmonize features and establishes a robust correspondence, enabling zero-shot open-vocabulary queries at the object level without scene-specific optimization.

Looking forward, several critical open challenges merit the field. The current tendency to sacrifice memory and computational efficiency for semantic expressiveness is a critical bottleneck; future advances must prioritize lightweight semantic heads that can be integrated directly into the rendering pipeline. The move towards open-world and interactive scenarios, especially in robotics and embodied AI, calls for scalable and generalizable semantic representations that can adapt to new environments. The added four-dimensional semantics—whether via decomposing motion for multiple objects or tracking semantic changes over time—is a pending frontier. Therefore, the convergence of efficient language embedding, sparsely-regulated semantic reasoning, and fine-grained 3D primitive grounding will be an accelerating area of exploration for future 3D semantic systems.

6.2 Scene Editing, Manipulation, and Inpainting

The explicit, primitive-based nature of 3D Gaussian Splatting (3DGS) provides a uniquely fertile ground for scene editing, manipulation, and inpainting, transforming these tasks from graphics-heavy mesh operations to direct, often semantically-guided manipulation of the scene’s constituent elements. Early approaches in this domain often facilitated edits by converting the Gaussian field into a more conventional representation, such as a mesh or point cloud, performing the desired operation, and then re-rasterizing or re-optimizing the Gaussians. This projection-based pipeline, while functional, is computationally expensive and often sacrifices the fine-grained fidelity that makes 3DGS attractive in the first place. The research community has consequently shifted decisively toward fully 3D-native, real-time editing pipelines that operate directly on the Gaussian primitives, leveraging recent advances in semantic segmentation and generative diffusion models to achieve precise and photorealistic results—a shift that parallels the broader movement toward adaptive, semantically-aware processing highlighted in the preceding discussion of language-embedded 3DGS, and that sets the stage for the interactive, task-driven applications explored in the subsequent subsection on embodied AI.

A prominent class of methods focuses on semantic-aware object-level editing by leveraging the language-aligned features discussed in the previous subsection. The integration of segmentation models and feature fields allows for the precise identification of Gaussian subsets that correspond to a user-specified object, directly building upon the open-vocabulary querying and grounding capabilities of language-embedded 3DGS [30, 81]. Once localized, these primitives can be subjected to a suite of geometric transformations, such as translation, rotation, scaling, and even deletion. This paradigm extends to more complex operations, such as the addition of new objects, where the editing pipeline is often coupled with 2D diffusion priors; the user provides a text prompt or an image of the object, and the system generates multi-view consistent pseudo-ground-truth images of the new object within the scene [11], which subsequently enables the extraction of a new set of primitives that respect the existing scene’s illumination and geometry.

A second significant direction is geometric deformation and shape editing, which is reformulated

as an optimization problem governed by constraints directly on the Gaussian cloud. By finding a rotation matrix for each primitive’s neighborhood that minimizes a geometric deformation energy, these methods preserve local structure and ensure that edits maintain global scene coherence. Characteristically, these techniques ensure the spatial stability of non-targeted points is preserved, yielding a deformed shape that aligns with the user’s intended edits [67, 121]. Since the geometry of the scene changes, appearance attributes such as color must be correspondingly refined to maintain perceptual fidelity, a requirement that bridges geometric editing with photometric optimization.

Perhaps the most prominent area is high-fidelity 3D scene inpainting and content removal, a pragmatic extension of the semantic editing capabilities. The challenge here is twofold: first, to accurately identify the volumetric extent of the object to remove—a task that directly exploits the clustering and segmentation features discussed earlier—and second, to reconstruct a plausible background that is spatially and photometrically consistent. The latter problem is far more complex, requiring filling the void left in the scene with the background from the surrounding region. Current state-of-the-art methods adopt a combinatorial approach consisting of a diffusion prior and a 3D-space-masking condition [122]. Diffusion-based methods improve quality by iterative refinement of multi-view images, using these outputs to supervise Gaussian optimization, while text-based or single-prior approaches rely on synthesized matching networks to propagate texture consistency. This yields a practical trade-off—hallucination drift versus scene geometric consistency—that is especially relevant when scenes are only partially observed and damaged regions overlap with low-confidence geometry.

Future directions are moving toward a complete end-to-end integration of localization, editing, and photometric refinement. Doing away with hand-crafted and interactively placed steps—like iterative multi-view diffusion refinement—and replacing them with a single forward pass through a large unified model is the ongoing trend. The same primitive-level parametrization that has proven so effective for large-scale scene editing in real-time virtual production, content creation, and spatial computing aligns naturally with the interactive constraint of embodied systems; as point cloud AI systems move toward online, dynamic use, editing capabilities that are both constrained and geographically-scalable will be essential. The success of these integrated perception-editing approaches will consolidate 3DGS as a core technology for large-world edition, simultaneously feeding robust interactive representation to the downstream tasks discussed in the following subsection on embodied and interactive systems.

6.3 Scene Understanding for Embodied and Interactive Agents

The integration of 3D Gaussian Splatting (3DGS) into embodied and interactive systems marks a paradigm shift from passive scene reconstruction to active, task-driven spatial perception. Unlike implicit neural representations that require dense sampling for queries, the explicit, structured nature of Gaussian primitives—with their inherent geometric position, covariance, and opacity—provides a uniquely suited substrate for real-time robotics, navigation, and closed-loop decision-making. This subsection examines how this representation is being leveraged for robotic mapping, visual localization, and the creation of interactive agents, focusing on the critical interplay between geometric fidelity, semantic richness, and computational efficiency that defines this application domain.

A central thrust in this area is the dynamic allocation of computational resources in robotic mapping, guided by task semantics. The foundational paradigm of building a uniformly detailed map is often wasteful for task-oriented agents. For instance, GaussLite addresses this problem by introducing a task-conditioned mapping framework, where natural language descriptions of a task define a target/background distinction. This allows the system to dynamically allocate learning and rendering resources—optimizing high-fidelity 3DGS for task-relevant objects while employing a lightweight feature-based representation for the background [123]. This not only achieves real-time semantic capacity adjustment at interactive rates but also improves memory

efficiency and demonstrates strong cross-scene generalization, a critical step towards scalable and purposeful perception.

Complementing dynamic mapping is the critical challenge of robust localization within a Gaussian scene. The explicit map structure can be directly utilized for online pose estimation. GSVisLoc introduces a data-driven visual localization method that operates directly on the splatting-based representation, focusing on extracting and matching distinctive visual features from Gaussian primitives to guide camera re-localization in large-scale environments [124]. This contrasts with approaches that merely use 3DGS-generated images for training other models; instead, it treats the optimized Gaussian parameters themselves as a feature space for localization. This approach is complemented by a broader strategy of using 3DGS-generated renderings to improve more traditional localization pipelines, demonstrating the utility of splatting as a robust front-end for geometry-based registration [125]. This intersection highlights a key trend: the 3DGS representation is not just an output but an active component within the perception-action loop.

Beyond primary systems, the fusion of 3DGS with large language models (LLMs) and semantic features is enabling a new class of interactive, embodied agents. These systems synthesize semantic segmentation, scene understanding, and conversational interfaces to better align with human intentions in complex environments. The foundational work on grounding language within Gaussian primitives supports open-vocabulary querying and grounding, but the next frontier is efficient human-agent interaction. Research has shown that by embedding features from vision-language models and coupling them with chat interfaces, agents can effectively localize, segment, and segment objects within a Gaussian scene [126]. More importantly, the structured representation has been successfully integrated into reinforcement learning (RL) pipelines for continuous control, replacing dense 2D-to-3D backbones that require pretraining. By incorporating latent semantic features into the Gaussian primitives, and conditioning with language signals for object-specific components, agents learn goal-conditioned multi-step skills from visual observation to continuous control, demonstrating the value of an explicit, semantically structured 3D representation as a decision-making substrate [127];[128]. This suggests that embedding explicit spatial structure and semantics at the agent’s input does not simply enhance raw visual fidelity but fundamentally modifies the learning problem for downstream control tasks.

Looking ahead, the primary challenge for 3DGS in embodied AI lies in achieving robustness and scalability without losing representation power. While current systems demonstrate feasibility, they often rely on and result in overly large memory footprints, limited memory management, or are restricted to static environments. Future architectures should move towards a unified framework that handles long-horizon, large-scale, and long-duration tasks in an online and dynamically updated setting. This involves integrating uncertainty estimates into the Gaussian primitives, improving memory management strategies, and developing training schemes that are resilient to noisy sensor inputs. This focused evolution will solidify 3DGS not just as a representation of the world, but as an active, computational engine for embodied intelligence in the upcoming decade.

6.4 Generative Content Creation and Physical Simulation

The integration of 3D Gaussian Splatting (3DGS) into generative and physical workflows marks a significant departure from its role as a passive reconstruction tool, positioning explicit primitives as active, manipulable building blocks for virtual content creation. This transition naturally extends the preceding discussion on embodied and interactive systems: whereas those systems leverage Gaussian primitives as perceptual anchors for understanding and navigating environments, the focus here shifts to harnessing the same explicit parameterization for generating, animating, and physically simulating virtual content. This subsection examines the dual frontiers of generative content synthesis and physics-based simulation, analyzing how the explicit,

unstructured nature of Gaussian primitives enables novel forms of scene animation, deformation, and physical interaction that were previously cumbersome with implicit neural representations.

The task of animating static 3DGS scenes has been substantially advanced by leveraging video diffusion priors. In contrast to methods that require per-scene deformation networks, the approach of [96] addresses the critical issues of extrapolation and occlusion in sparse-input modeling. By proposing a reconstruction-by-generation pipeline, it employs a video diffusion model to synthesize plausible interpretations for unobserved regions. The key innovation lies in the scene-grounding guidance: an optimized 3DGS render iteratively tames the diffusion model to generate view-consistent sequences, effectively steering the generative prior towards geometrically coherent solutions. Similarly, [74] introduces a training-free approach where a geometry-guided, stage-wise denoising strategy uses the current reconstruction’s rendered RGB and mask images to adaptively guide the diffusion direction of translation. This iterative exploration of camera trajectories into unobserved regions allows for the generation of novel views that supplement the training set, enhancing reconstruction fidelity from sparse inputs.

Building upon this synthesis of missing geometry, an important direction that remains in its infancy is the embedding of material properties directly into the Gaussian formulation for dynamic content creation. Extensions of radiance-focused approaches are moving toward material- and geometry-aware optimization, exploring the interaction between calibrated diffusion priors and physical plausibility. While not yet fully realized in the provided references, a simultaneous disentanglement of material and appearance is emerging as a critical near-term goal to enable object generation with inherent relightable properties, moving beyond simple radiance capture to a full physical decomposition.

The most transformative capability of 3DGS in this domain, however, lies in its physical simulation. Because each Gaussian is a discrete, parameterized primitive, it can be treated as a physical body governed by the laws of continuum mechanics. This allows for the application of physics-based deformation, where the positions and orientations of Gaussians are updated based on internal forces (e.g., elasticity, plasticity) and external interactions. The physical fidelity of these simulations often relies on either treating the Gaussian cloud as a continuum, coupling the splat representation with a traditional mesh-based physical engine, or steering explicit primitives through learned dynamics. This interplay is critical: the explicit nature of Gaussians permits the fine-grained, per-primitive control over physical states required for engaging and realistic, real-time interactive environments, which is a natural evolution from the task-directed manipulation seen in embodied systems.

While these advances are promising, current approaches exhibit significant limitations. A common theme is the simulator engine—the physics-only interaction operates within a closed loop, and then a separate domain generates geometric and appearance changes for a given state. For instance, while methods in [95] tackle acceleration in input precomputation from a geometric perspective, the realization of dynamics so far remains heavily aligned with a moving system’s optimizer, which may not be designed for large displacements. A significant future direction is the exploration of physics-based regularization for mechanics itself; for example, using the video diffusion priors not only as content generators for in-painting missing geometry as seen above but also as scene-aware predictors to automatically synthesize 3D Gaussian motions based on an object’s physical properties, thereby creating fully interactive, semantically-aware simulation engines.

In synthesis, the trajectory presented—from static capturing to generative diffusion and physical manipulation—underscores a shift from 3D Gaussian Splatting as a computationally efficient input method to a foundational, interactive 3D asset format. This progression is mirrored in the broader context of the survey, where the representation is increasingly defined not just by its visual fidelity but by its capacity for active, dynamic, and physically grounded reasoning. The next frontier will be to fully unlock the synergy between these workflows, culminating in Gaussian-based representations that can be semantically understood, physically manipulated,

and iteratively refined in real-time for games, digital representations, and virtual production.

7 Rendering Appearance, Inverse Rendering, and Photorealism**

7.1 Physically-Based Appearance Decomposition and Inverse Rendering

The pursuit of photorealism in 3D Gaussian Splatting (3DGS) necessitates a fundamental shift from simply reproducing view-dependent radiance to recovering the physical attributes of a scene—its geometry, surface reflectance properties, and incident illumination. This subsection examines the core methodologies for this physically-based appearance decomposition, focusing on how the explicit primitive-based representation of 3DGS is augmented to enable high-fidelity inverse rendering, material editing, and relighting. The central challenge lies in decoupling the entangled contributions of lighting and material that are baked into the per-primitive spherical harmonic coefficients of the original model, whose appearance is otherwise ambiguous and global [2].

A prerequisite for inverse rendering is the parameterization of each Gaussian primitive with physically meaningful material attributes. The standard 3DGS pipeline uses spherical harmonics to encode view-dependent appearance, which is efficient for synthesis but has no physical correspondence. Research has therefore shifted to augmenting primitives with diffuse albedo, specular roughness, and normal information, which can be integrated with microfacet models such as Cook-Torrance to compute physically-based per-primitive radiance [129]. Key to these systems is the shading model with deferred shading. Deferred shading has proven particularly efficient for storing material and geometric attributes in intermediate G-buffers, allowing for explicit normal estimation and light transport calculations. However, issues arise for transparent and volumetric structures [9]. An elegant alternative to relying on normals for highly curved or volumetric surfaces is a light- and view-dependent scattering representation via radiative hardness [129]. Furthermore, to address inherent topologically-ambiguous and appearance-related ambiguities, advanced models move beyond single per-primitive parameters by incorporating multiple, spatially varying attributes defined in a primitive-local texture space. This allows an accurate representation of high-frequency appearance variations, decoupling geometry and appearance further [15, 13]. While these hyperprimitives with learnable textures improve representational power, they often do so at the cost of compatibility with fast rasterization pipelines [130].

The ill-posed nature of intrinsic decomposition calls for the separation of optimization into known components. Some methods have been proposed with previous information, restructuring, parameters, etc. Recent advances include improvements in rasterization speed [61].

These disparate approaches collectively highlight a trend: 3DGS is maturing from a backward rendering and editing tool to a compact, physically-aware scene reconstruction system. However, these methods still face significant barriers. Most assume that BRDFs operate under distant lighting (LD), requiring environment map estimation and, in some cases, ray-traced shadows for specular surfaces, which is computationally prohibitive for real-time rendering [131, 132]. While this entanglement problem is a central challenge for these systems, future research should focus on equipping the model with principles of perfect global illumination, towards a truly physically-based appearance decomposition that is globally consistent and editable.

7.2 Relighting, High-Dynamic-Range Imaging, and Global Illumination Approximation

Relighting—the ability to place a reconstructed scene under novel illumination—remains one of the most demanding objectives for 3D Gaussian Splatting (3DGS), as it requires the disentanglement of scene radiance into its underlying physical constituents: geometry, material reflectance, and lighting. Building upon the physically-based appearance decomposition discussed previously, this subsection analyzes the methodological landscape that enables such manipulation,

with a focus on explicit environment map estimation, real-time visibility handling, and the extension of the rendering equation to high-dynamic-range (HDR) imagery. The central challenge is to reconcile the computational efficiency of 3DGS with the demand for physically correct light transport, which often calls for more compute-heavy, non-local operations that go beyond the per-primitive shading models introduced earlier. [83, 66]

A dominant paradigm involves explicitly decomposing a scene into material properties and distant lighting. These works typically augment each Gaussian primitive with parameters for diffuse albedo and specular reflectance (e.g., roughness and metallic), which are then used in a physically-based model—directly extending the per-primitive material parameterization strategies discussed in the previous subsection. The lighting itself is often encoded as a low-frequency environment map, either in a spherical harmonic (SH) basis or as a set of learnable parameters. This formulation inverts the forward rendering process, enabling the joint estimation of materials and lighting from multi-view images. The primary strength of this approach is its directness; once the decomposition is learned, the scene can be relit by simply swapping to a new environment map. In the context of 3DGS, this is implemented in either a forward or deferred manner. Forward shading computes the color of a pixel as a weighted sum over the Gaussians, integrating the material and lighting models directly into the per-primitive radiance approximation. Deferred rendering, in contrast, first writes material properties to a G-buffer and then applies the shading model, which aligns with the deferred shading techniques prevalent in the inverse rendering systems described previously. The deferred approach has gained prominence, as it allows for the storage of high-frequency normals needed for specular reflections, but it requires careful handling of gradients through the G-buffer to maintain coherent coupling between geometry and shading.

A central, yet unresolved, challenge lies in the correct calculation of occlusion and visibility—a global, scene-level computation that the local per-primitive transport models of the previous subsection cannot address alone. Beyond basic environment map reflections on geometry, shadows are essential for realism, yet self-occlusion within the Gaussian volume is often ignored in the basic blending process, leading to artifacts. Several works have therefore integrated shadow-mapping techniques directly into the Gaussian pipeline: the point cloud is rendered from the perspective of the light source, and the resulting depth maps are used to compute binary or numerical visibility during the final shading stage. The benefits are real-time performance, but they rely on reliable depth or point separation. To handle soft shadows or inter-reflections, more expensive ray-marching or ray-tracing techniques are required. This is an emerging space that balances between the real-time rasterization favored by 3DGS and the offline-accurate light transport models that form the basis of the following subsection. The choice of method for computing visibility—whether to rely on binary shadow maps, ray tests, or to completely skip explicit occlusion—is a critical design trade-off that heavily influences both rendering quality and the scene’s transferability to new light conditions.

Another necessary condition for realistic relighting is handling radiometric values beyond the traditional 8-bit range. Real-world luminance values often exceed the representable range, and HDR input contains both highlight details and dark-area intricacies. Earlier methods are trained on—and produce—low-dynamic-range images, which complicates a physically correct relighting. Newer methods have therefore extended the pipeline to handle HDR data and implemented tone mapping as a learnable, image-space operation to account for exposure variations. This adaptation requires the storage of color values in linear space, a departure from the non-linear color spaces used in standard 3DGS training. The design from a family of methods builds on a set of images spanning a wide dynamic range, where a learnable tone-mapping ensures that the training signal focuses on the full range of luminance values, enabling a more robust joint estimation of the lighting and material parameters.

Looking forward, the field is moving toward **unified, spatially-adaptive visibility**, a more complete approximation of global light transport that connects the local shading models of the

previous subsection to the full global simulation discussed later. A simple workaround has been to add heuristics like a constant ambient term, but proper methods lack structural awareness between primitives. The combination of deferred shading with fast, hierarchical volume or neural visibility calculations appears promising. This points to a clear future direction: the creation of hybrid systems that use 3DGS’s efficiency for direct illumination and rely on an additional model—perhaps a neural network, a sparse voxel, or a ray-traced representation—to capture indirect light paths and interreflections. While such hybrid approaches remain rare, they represent the defining frontier for 3DGS: making physically-principled lighting computations stable under the joint optimization of lighting, geometry, and materials, thereby bridging the gap between the per-primitive physical models previously examined and the full light transport accuracy that will be explored next.

7.3 Advanced Material Modeling and Complex Light Transport Effects

The standard 3D Gaussian Splatting pipeline encodes view-dependent appearance through per-primitive spherical harmonic (SH) coefficients. While efficient, this isotropic, low-frequency basis is fundamentally limited when tasked with reproducing the high-frequency angular variations characteristic of anisotropic reflections, the complex light paths of refractive media, or the soft, diffuse transport of subsurface scattering (SSS). This subsection reviews the principal research thrusts that extend the primitive representation and the shading model to capture these complex physical phenomena, moving the field beyond a photometrically-driven optimization and toward physically grounded material modeling.

A primary limitation of SH functions is their inefficiency and inaccuracy in representing high-frequency specularities. To address this, researchers have begun replacing SH with more expressive basis functions and reflection models. For instance, **Anisotropic Spherical Gaussians: Efficient & Quality for 3D Gaussian Splatting (ASG)** integrates ASG basis functions that are specifically tailored for the asymmetric appearance of specular reflections, enabling the representation of stretched and shifted highlights that are poorly captured by their isotropic counterparts [133]. This indicates a clear trend toward analytically grounding the Gaussian appearance model with physically-inspired basis functions, a necessary step for high-fidelity material reconstruction.

The challenge of planar, refractive, and transparent materials—characterized by simultaneous transmitted and reflected radiance—is a distinct and challenging problem. **Planar Gaussian Splatting (PGS)** tackles this by modeling an explicit set of planar primitives to decompose the scene into transmissive and reflective layers [24]. A Fresnel-based view-dependent weighting scheme then dynamically integrates these layers, while a multi-stage optimization and densification strategy faithfully synthesizes the appearance of glassy and specular surfaces. Rather than treating refraction as a per-primitive local operation, these planar-surface, multi-stage approaches are able to use the scene’s structure to represent complex global transport, representing a step change in the complexity of the material phenomena modeled by the splatting framework.

Equally important yet difficult to model is the transport of light under the scene surface, i.e., subsurface scattering. A method proposed in **Breaking the Entanglement of Scene and Object in 3D Gaussian Splatting without Depth and Normal (BOSE)** goes beyond the standard diffuse-specular decomposition and introduces a learnable function that aggregates light transport beneath a surface, approximating a dipole-based scattering model to represent translucent materials [7]. The relevance of this work is its ability to disentangle scenes and objects to provide sufficient support for independent geometry and appearance decomposition; but by isolating the key scattering module, it proves that the standard rendering equation can be recovered by higher-order terms with moderate computational cost. A promising direction could be to bridge such a dipole approximation into the ray-marching framework to achieve true physically correct sub-surface effects.

Critically, many of these works implicitly recall a trade-off between per-primitive complexity

and the realism of the final scene. While these new methods shift toward richer per-primitive attributes, they also impose a substantial computational burden on the optimizer. The progression from SH coefficients to comprehensive BRDF conditioning—with all attendant **texture maps**—represents an evolution to spatially varying attributes within a primitive. This is a subject of ongoing research, as the space of a Gaussian primitive is usually treated as a constant color silhouette. Extending this with spatially varying texture maps, normal maps, and opacity functions provides a new handle for making the model richer still, blending the elegance of mesh texturing with the acquisition power of point-based splatting.

In summary, advanced appearance modeling for 3DGS is progressing on multiple fronts? difficult. The field is pivoting from a purely geometric radiance basis to a physically-based sparse representation. These approaches are not merely incremental modifications to the rendering behavior; they are upgrades that trade speed for much higher accuracy in simulating light transport, with a promising trend toward grounded decomposition of global illumination within the explicit Gaussian environment. Future work will need to reconcile these detailed local material models with a global light transport solution to reach a more complete and robust physically-plausible synthesis of scene radiance.

7.4 Visibility, Occlusion, and Physical Correctness in Rasterization and Rendering

The fidelity of 3D Gaussian Splatting (3DGS) hinges on the physical correctness of its image formation model. The seminal 3DGS pipeline, which couples per-pixel depth sorting with alpha compositing, is an approximation that, while efficient, introduces systematic biases in visibility computation. This subsection examines the recent trajectory of research that diverges from these foundational compromises, focusing on the accurate rendering of transmittance, the handling of complex occlusions and semi-transparency, and the pursuit of physically grounded light transport within the splatting paradigm. Where the previous discussion centered on enriching per-primitive appearance attributes, this section shifts the focus to the global transport and aggregation of light across primitives, addressing the image formation process as a whole.

A central theme of this evolution is the reassessment of the blending function itself. The original formulation relies on a sequential, sorted alpha compositing based on a discrete, view-dependent depth estimate. This z-buffer approach is prone to artifacts for overlapping, semi-transparent primitives. To address this, recent work has moved away from this ad hoc sorted blending toward a physically-based transmittance model derived from generalized radiative transfer. Methods can now model light attenuation through a volume of Gaussians using an exponential transmittance law, and direct analytical integration of Gaussian density over ray intervals replaces the hard depth-sort. These integrations compute the density integral and exact transmittance, avoiding quantizations and yielding more accurate alpha weights for compositing. This shift also alleviates the necessity of expensive per-pixel sorting, reducing computational overhead while improving robustness with overlapping geometry, marking an important harmonization between momentum and physical accuracy.

Tied to this new volumetric image formation is the challenge of robustly estimating visibility and uncertainty, particularly during optimization with constrained data. This problem is one step beyond the material modeling discussed previously, as it is fundamentally about the global geometry and its confidence. The ill-posed nature of the inverse problem often leads to artifacts like "floaters" (false needle-like geometry) or "background collapse" due to the overconfidence in incomplete depth estimates. Several works have pioneered the use of uncertainty estimation to explicitly guard against these pathologies without imposing hard geometric constraints, refining the optimization itself to enforce a more stable visibility estimate. To address the degenerate solutions, a method based on **Fisher Information** quantifies per-primitive uncertainty to guide geometric priors and adaptive regularization [97]. Similarly, an explicit **Spatial Uncertainty Field** integrates into the 3DGS framework to provide an adaptive error tolerance, weighting

depth contributions and refining normals for more robust surface reconstruction in sparse-view scenarios [41]. A more recent treatment frames the entire optimization as a primitive-denoising process, employing spatial awareness to denoise the GS primitives — using the spatial gradient distribution and uncertainty-based pruning to selectively remove erroneous primitives and stabilize the render [28]. The flexibility of these approaches makes them strong companions to the advanced appearance models established earlier, grounding them with a more reliable geometric scaffold.

The most significant point of divergence from the classical local appearance model, however, is the integration with generative priors for sparse-view reconstruction. Here, the constraints of physics-based rendering meet the generative power of modern diffusion models. The direct problem is that when observations are beyond sparse, the optimized radiance field is under-constrained, and purely forward rendering with accurate transmittance does not prevent hallucination in the unseen. To combat this, the field is now turning to pre-trained generative priors, in the form of diffusion models, to **synthesize** the missing views and *complete* the scene. In this new era, physical correctness of visibility is not just the correct rasterization of the existing primitives but also the internal consistency of the generated content itself, which is in turn checked against the 3D scene. To unify these auto-generated views with the 3D representation, scene-grounding guidance techniques have been used to tame video priors by coupling rendered sequences to guide consistent synthesis in occluded regions [96]. Similarly, propose-and-validate frameworks treat the generated views not as ground truth but as *intermediate proposals* submitted to a **Multi-view Stereo (MVS) oracle**. The MVS model acts as the ultimate arbiter of geometric validation, identifying regions of hallucination or instability, and guiding the GS optimizer to filter artifacts while preserving plausible completions coherent with the data [73]. In a related pipeline, the reconstruction is split into repair and inpainting tasks, with two distinct diffusion models ensuring coherent completion between observed and missing content [98]. This represents a shift from a pure analytical model of light *transport* toward a semi-implicit form of light *completion*, abstracting the physics while still relying on the expected coherence of the scene.

Despite this progress, the field remains predominantly photometric in its loss definition. While an improved transmittance function and uncertainty handling resolve local artifacts, the process still largely treats optimization as a denoising task rather than a true inverse of a global transport phenomenon. Looking ahead, a syntheses of these ideas points toward detecting and removing geometry that is inconsistent in a *semantic* and *physical* space, and towards the joint optimization of local material (from the previous subsection) with global visibility—where light transport is not merely assembled from local approximations but fully integrated with the scene structure. Coupled with the semantic awareness provided by generative priors, and the robust *rendering* and *deviation of appearance* coalescing into a physically grounded and robust scene representation, we are carving a plausible path toward high-fidelity digital twins. The remaining conundrum to stabilize the joint learning of lighting, geometry, materials, and subspace for a fully coherent, photorealistic synthesis is a direct challenge that will define the next stages of research in 3D Gaussian Splatting.

8 Practical Deployment, Benchmarks, and Future Horizons**

8.1 Evaluation Protocols, Benchmark Suites, and Standardized Metrics

The evaluation of 3D Gaussian Splatting methods currently rests on a foundation inherited largely from the neural rendering community, yet the distinct characteristics of explicit splatting representations demand a critical re-examination of these protocols. The dominant benchmarks—Mip-NeRF 360 for unbounded indoor and outdoor scenes, Tanks and Temples for large-scale real-world captures, and Deep Blending for complex, color-rich settings—were originally designed

to stress implicit methods with varying scene scales and complexity [2, 1]. While these datasets provide a common ground for comparison, they harbor biases that may disproportionately penalize or favor explicit methods. For instance, scenes with abundant high-frequency texture or sharp discontinuities expose differences in how Gaussian primitives handle edges, whereas these same regions are notoriously challenging for neural fields due to spectral bias [10]. The standard suite, therefore, inherently couples a method’s measured performance with the primitive’s capacity to model specific scene complexities, an aspect rarely decoupled in comparative analyses [134, 20].

The quantitative metrics applied are equally borrowed, with peak signal-to-noise ratio (PSNR) primarily measuring pixel-wise fidelity, SSIM assessing structural consistency, and LPIPS serving as a proxy for perceptual similarity. However, these metrics exhibit sensitivities that are well-matched to the implements of implicit volumetric rendering but less aligned with the pathologies of splatting. For example, the "needle-like" blobs that can form around sparse coverage of Gaussian primitives, or the "ghosting" that arises from anisotropic covariance mis-alignment, distort PSNR and SSIM in ways that are not linearly correlated with human perceptual dissatisfaction with the artifacts [3]. LPIPS, while more perceptually aligned, is often computed on deep features that may not capture the sharp boundary shimmering that is characteristic of under-densified Gaussian regions. This implies a growing demand for artifact-aware metrics, perhaps derived from physics-based rendering principles or edge-localized statistics, that explicitly penalize the morphological failures of the explicit representation [135].

Beyond visual fidelity, evaluating computational cost is paramount to the deployment of 3DGS, yet current protocols are oddly fragmentary. While the original work established a strong argument by achieving real-time rates (≥ 30 fps) at 1080p [2], subsequent works report a fragmented set of "rendering speed" values measured on test-set camera paths of varying lengths, arbitrary GPUs, and without necessarily standardizing resolution or scaling clear separation between offline training, online inference, and their dependence on sorting or optimization heuristics. Failing to use a unified, lock-step benchmark suite be it in the form of a standardized hardware specification and frame-buffered rendering protocol makes it difficult to equate the reported efficiency gains from pruning, vector quantization, or a hierarchy of level-of-detail approaches [6, 5, 60]. Truly comparable protocols need to explicitly separate per-frame rendering throughput from cost-curve complexity, and disclose the degree of parallelization and dependencies on sorting and blending algorithms.

A more critical frontier is the evaluation of geometric fidelity and downstream performance, where the field is still anchored. While the availability of depth and normal supervision from methods like 2D Gaussian Splatting (2DGS) allows for geometric comparisons, the conventions are inconsistent: some works extract meshes and proceed with a classic mesh comparison—using a Chamfer distance and its own F1 score—while others rely on depth map L1 losses or surface-aligned normals with tied color-space supervision [9]. Such a non-unified approach undermines the comparability that we seek at the benchmark level. Moreover, downstream task evaluation in localization or SLAM is nearly absent; the usual metrics used to assess the reliability of a 3D scene representation for interaction have not been standardized. The field must begin to adopt test suites for downstream navigation and task success rates, not merely pixel fidelity.

Going forward, we need a coordinated standardization of evaluation protocols to match the maturity of the field. First, an approach should be decoupled into "appearance" evaluation (image metrics), "geometry" evaluation (depth and normal accuracy consistent across all scenes), and "performance" evaluation, which is tied to a publicly standardized hardware platform and includes a measure of worst-case latency alongside average frames-per-second [3]. Second, given the scaling success of Gaussians in large-scale environments, we must introduce benchmark suites that are stratified by scene-scale (object-centric, bounded indoor, mid-large-scale outdoor, unbounded) as well as by content complexity (high-frequency, texture-poor, highly dynamic specular) to avoid uniform weighting out crucial differences in Gaussian distribution [6]. Finally, the current

bottleneck is still memory footprint and training resources; metrics should not just report best final accuracy, but also evaluate the efficiency of the representation in terms of (memory, accuracy, and time) so that the design space of compression/pruning is seen as a core merit axis [134]. In summary, the next generation of benchmarks should not just test “how well a method renders,” but “which representation is reliable, auditable, and retargetable, across a meaningful perceptual and functional spectrum.” Only by advancing the rigor of our measurement can we fully exploit the theoretical promise of the explicit 3D Gaussian representation in deployment.

8.2 Robustness, Reliability, and Degradation Analysis in the Wild

The deployment of 3D Gaussian Splatting (3DGS) in real-world scenarios exposes a critical gap between its impressive performance on curated benchmarks and its behavior under non-ideal conditions. While the standard pipeline assumes clean, dense, and well-posed multi-view captures, practical acquisition is frequently characterized by sparse coverage, noisy pose estimates, dynamic interference, and high-dynamic-range content. This subsection systematically investigates the survivability of 3DGS under such adverse conditions, focusing on the emergence of artifacts, sensitivity to initialization and conditioning, and the broader need for uncertainty-aware, trustworthy evaluation—concerns that directly shape the practical deployment requirements discussed in the following subsection.

A primary source of unreliability in 3DGS is the uncontrolled densification of Gaussian primitives, which, while intended to improve spatial representation accuracy, often leads to degenerate solutions in the form of “floaters” and ghosting artifacts. The optimization of the explicit cloud can easily allocate capacity to geometrically inconsistent regions, creating elongated, translucent Gaussians that float in space and occlude true geometry from novel viewpoints. The root cause is frequently traced to a failure in the adaptive density control mechanism, which is overly generous in adding new primitives. An analytical framework links this densification to optimization theory, showing that splitting can be crucial for escaping saddle points but also creates redundancy that becomes a source of instability [26]. A core failure mode identified is the “gradient collision,” where conflicting gradient directions prevent large Gaussians in over-reconstructed regions from being split, leaving them stuck in a high-error state and producing blurred, distorted outputs [136]. Conversely, regions with well-aligned gradients can suffer from over-densification, leading to redundant primitives and higher memory overhead without proportional visual benefit [137]. A related phenomenon is the frozen primitive issue, where large, suboptimally scaled Gaussians become trapped in local optima, producing needle-like and distorted geometry rather than sharp surfaces [31]. Addressing these pathologies frequently necessitates departing from simple gradient-magnitude heuristics. For instance, incorporating the angular information of gradients into density control can better align primitive allocation with geometric consistency [138]. Similarly, an opacity-gradient-driven module can act as a lightweight proxy for rendering error, which, when paired with conservative pruning schedules, prevents destructive optimization cycles [33]. In challenging scenarios with transient objects or dynamic interference—a common occurrence in in-the-wild imagery—the densification process can become progressively corrupted, adding capacity to model transient phenomena rather than the static scene. A delayed-growth criterion has been shown to mitigate overfitting and prevent the formation of distracting artifacts, postponing capacity allocation until a stable structure emerges [76, 139].

Beyond static artifact generation, the reliability of 3DGS is contingent on the quality of its initialization and the density of the observation set. The representation typically relies on sparse Structure-from-Motion (SfM) initialization, which, when heavily affected by noise or poor coverage, can disrupt scaling and convergence, cascading errors throughout the optimization [140]. While adaptive density control is designed to enhance geometric coverage, it often overgrows in unobserved regions, leading to blurred or hallucinated content. In few-shot or sparse settings, this lack of geometric grounding is further aggravated, since the model must contend with

insufficient observations. Many modern approaches explicitly attempt to compensate for this by decoupling densification from the initial optimization phase [112, 100], and by combining robust pixel-based addition with semi-pseudo supervision [30, 80]. However, even in more standard settings, the representation remains highly sensitive to camera inaccuracy. Pose errors arising from miscalibration cause the optimization to degrade, introducing ghosting or floaters that prior methods do not adequately address. This behavior demonstrates that the geometric quality of 3DGS is deeply intertwined with the conditioning of its input, making robust initializations and algorithmic conditioning critical for real-world deployment.

The robustness of the Gaussian cloud itself must also be questioned when confronted with complex appearance variations. The standard 3DGS appearance model is insufficient for high dynamic range scenes or regions with strong specular reflection and complex light transport; these often degrade into blurry artifacts or color inconsistencies. Moreover, the point-based algorithm tends to retain a significant number of suboptimal Gaussians from the start due to sparse SfM initialization, and the chosen gradient-based criterion for addition fails to balance geometry and texture imbalance across the scene. A localized strategy that identifies all 3D zones requiring higher capacity and applies regularizers specifically to ill-conditioned points can help stabilize the optimization, focusing point management on the most error-contributing regions [141, 28].

To move beyond mere reliability, a critical shift must occur toward designing metrics and models that are uncertainty-aware and trustworthy. The current evaluation paradigm based on PSNR and SSIM is not indicative of geometric quality or robustness; it can be easily induced by adding spurious Gaussians that spike texture. The optimization landscape itself reveals a clear duality: stable solutions under standard optimization reveal rendering-optimal references, but are indistinguishable from those produced under conditional multi-view optimization. This duality reveals a stratification—in dense regions, point clouds contain cascading redundancy, while in sparse geometry, there is a visibility-driven coupling that underscores per-primitive importance [91]. This suggests a systematic direction: incorporate this dichotomy into future algorithms to create inherently importance-based or budget-aware regularizers, rather than relying on local error when using heuristic thresholds.

Thus, the future of robust 3DGS should advance toward metrics and regularizers that remain interpretable and consistent with underlying density and uncertainty, rather than merely correcting errors post hoc [138, 142]. Uncertainty estimation should strive to become an iterative, first-class diagnostic, using signals such as initialization residual and gradient-consistency for continual pruning [28, 143]. However, a key barrier remains in how to incorporate all confidence calibration when the scene contains unforeseen transient configurations. During deployment, instability in the primitives may indicate degeneracies or catastrophic drift; future work must therefore integrate model collapse detection—directly rejecting or fallback to earlier, more conservative estimates of both geometry and appearance under distribution shift—principles captured by weighted per-spatial predictive modeling [72, 144]. Addressing these reliability and uncertainty issues is not supplementary but essential to fulfilling the practical, interoperable deployment emphasized in the following subsection, bridging research advances with real-world expectations.

8.3 Interoperability, File Standards, and Industrial-Grade Usability

The transition of 3D Gaussian Splatting (3DGS) from a research novelty to a deployable technology is contingent upon its seamless integration into the established ecosystem of computer graphics, content creation, and real-time rendering. Unlike traditional meshes or point clouds, which benefit from decades of standardized formats and tooling, the Gaussian primitive—with its unique parameterization of position, covariance, opacity, and spherical harmonics—presents a new data modality that requires a paradigm shift in how 3D assets are stored, exchanged, and manipulated. The nascent state of this ecosystem represents both a significant bottleneck

for industrial adoption and a fertile ground for innovation, as the community grapples with establishing universal file standards, developing hybrid representations that bridge the gap with legacy pipelines, and creating robust authoring tools that support the iterative, editable workflow expected by professionals.

A primary obstacle to wide-scale adoption is the lack of a standardized file format. While the foundational implementation of 3DGS leverages the Polygon File Format (PLY), this structure is a poor fit for efficient storage and streaming, as it is typically monolithic, too large, and lacks the inherent support for level-of-detail or progressive encoding that is important for large-scale or network-distributed applications. In response, specialized formats like the .splat format have emerged, focusing on rapid loading via a binary structure with pre-optimized data layout for GPU consumption. For dynamic scenes, the challenge is more complete, as the requirement to describe both spatial and temporal data necessitates a more comprehensive solution. The 4D Gaussian Splatting format explicitly addresses this by providing a hybrid set of primitives that efficiently combines spatial representation with temporal reasoning and rendering, allowing for direct generation of both static and dynamic representations in a single framework without the need for multi-stage training. This evolution highlights a critical tension: the need for a universal, open standard from the mesh-centric method, let alone the three-dimensional formats, versus the efficiency gains from domain-specific, optimizations. A universally accepted standard is crucial for ensuring cross-tool, cross-vendor, and cross-platform interoperability, which is a foundational requirement for industrial-grade pipelines. Without such a standard, the field risks remains fragmented, hindering collaboration and long-term asset durability.

Beyond file storage, true interoperability demands a blended approach. The core strength of 3DGS lies in its unprecedented capability for novel view synthesis and rapid training and rendering from unstructured captures, whereas traditional meshes, texture maps, and rendering engines offer a mature ecosystem for physics-based material, editing, and transmission. Bridging this gap necessitates the development of tools and pipelines to convert a ‘Gaussian cloud’ into a polygon mesh with view-dependent appearance, enabling direct integration into existing 3D content creation workflows [145, 146]. Building on this, hybrid pipelines that combine 3DGS with a standard mesh allow for hybrid rendering, obtaining both high geometric fidelity from the mesh and photorealistic appearance from the Gaussians [147]. For practical 3D content creation, a core component of the industry’s workflow, the editable and intuitive nature of the Gaussian primitives must be harnessed in dedicated editors. Tools such as a Gaussian Splatting Editor empower users to interactively manipulate individual Gaussians in real-time, enabling intuitive editing of the scene with features like local deformation and brush-based selection, which is a departure from view-dependent editing limitations [64]. However, the editable functionality within a single scene remains limited; integrating and gaussians into conventional graphics innovative workflows requires effective strategies for seamlessly stitching, repurposing and mixing distinct 3DGS scenes to allow for complex composition [115].

The transition from static reconstruction to interactive deployment in engines and gaming addresses integration with mainstream graphics engines (Unity, Unreal) and native SDKs supporting Vulkan, Metal, and OpenGL. To achieve high fidelity or real-time rendering in typical computer graphics pipelines, a deployment architecture is essential. The most straightforward implementation lies in custom leading-edge GPU-accelerated kernels—such as those for the **Rendering** (Ray Tracing) engine—and the WebGL API, allowing for cross-platform browser-based users to interact with 3DGS content without requiring specialized plugins or native installations [148]. This is fundamentally aligned with achieving lightweight runtime paths for XR and game applications, such as the use of a Vulkan-based renderer to efficiently render large-scale scenes on low-power devices [2], and the use of a web-based renderer supported by standard graphics libraries [58].

In the long term, the external compatibility, standardization, and industrial usability of 3DGS are critical avenues for its broad dissemination and adoption. The transition is not merely

a technical convenience, but a prerequisite for the creation of a self-sustaining ecosystem of tools, content, and specialized talent. As emerging formats and hybrid pipelines solidify, the field moves closer to a unified, interoperable standard, and thereby towards the claim of the “digital asset” regime. Key to this is the integration of 3DGS into industrial pipelines, which is predicated on open-interchange formats, robust tooling, and versioning. The community must be careful to avoid closed, proprietary approaches that preclude broader adoption, and instead to agree on an open-interchange standard, building a common foundation for one of the most promising 3D structural representations for the next decade, with the ultimate goal to fluidly render 3D Gaussian contents into the rest of the ecosystem—from lightweight edge devices to the most massive and interactive global spatial data across interactive media, design, and entertainment systems.

8.4 Scaling, Temporal Streams, and Deployment on Device

The transition of 3D Gaussian Splatting from laboratory demonstrations to deployed systems hinges on its ability to scale across three orthogonal axes: the spatial extent of the scene, the temporal dimension of the data, and the computational budget of the target device. These challenges are deeply interconnected; a model that cannot be efficiently compressed cannot be streamed to a mobile device, and a static reconstruction pipeline cannot accommodate the continual evolution of real-world environments. Crucially, this scalability discussion forms the technical backbone that determines whether the ecosystem interoperability and standardized formats described above can translate into practical, cross-platform deployment at scale.

The first major obstacle is the sheer memory footprint of a Gaussian scene, which can reach hundreds of megabytes or even gigabytes for city-scale captures. Practical deployment demands aggressive compression and streaming techniques. A key strategy involves principled primitive reduction and structural decomposition, which moves beyond simple pruning to quantifying the contribution of each Gaussian to the final image to remove redundancy with minimal visual loss [28, 38]. To achieve parametric efficiency, methods are increasingly turning to hybrid representations—the same bridging philosophy introduced in the previous discussion of mesh-Gaussian pipelines—mapping primitive attributes, such as position, covariance, and spherical harmonic coefficients, onto compact latent grids or quantized codebooks, lowering storage and bandwidth requirements while preserving fidelity. The availability of a variety of compression formats, from the lossless .ply encoding to lossy .splat formats, represents a step toward the common interoperability emphasized earlier, yet it requires the rigorous standardization efforts called for in the ecosystem discussion.

To manage the computational load of these massive datasets in real-time, a dual strategy of level-of-detail (LoD) management and out-of-core streaming is essential. Rather than rendering all Gaussians in a scene simultaneously, hierarchical approaches use coarse-to-fine representations to adapt the number of primitives based on camera distance. This strategy is particularly critical for unbounded city-scale scenes, where the visibility of millions of primitives changes drastically as the viewpoint moves. While the literature does not yet provide a universal LoD framework for 3DGS, the prevailing paradigm involves preprocessing the scene into a spatial hierarchy, enabling smart culling and progressive loading of data. A corresponding shift from full in-memory models to streaming architectures becomes necessary: only visible and relevant primitives are loaded from persistent storage to the GPU, enabling continuous navigation through massive scenes that exceed VRAM capacity, effectively allowing the scene to be perceived as a virtual stream. This requirement deepens the need for the robust, standardized, and versioned asset formats highlighted in the previous section, as interchange files must accommodate progressive retrieval rather than monolithic loads.

The second axis is temporal scalability. Real-world deployments, such as digital twins of a city or industrial monitoring systems, involve streams of data that evolve over time. Merely training a static model on a batch of data is insufficient; the system must support continual learning, where

new observations update a model to reflect changes without catastrophic forgetting. Dynamic scenes captured by video introduce the additional complexity of temporal streams. When these sequences are sampled sparsely or captured from sparse cameras, the lack of geometric supervision across time leads to severe artifacts, such as floaters and missing geometry [44]. This problem is a central focus for a broader family of approaches addressing dynamic and monocular sparse-view 3DGS, where the task is not just rendering a pre-built static model but reconstructing a dynamic model from low-cost sensor setups—a capability that directly conditions the viability of the "lifelong" asset pipelines anticipated in the future outlook. In this context, geometric depth priors—combining dense but ambiguous monocular depth with incomplete but accurate multi-view geometry—guide the learning process, providing the necessary constraints to stabilize the 4D representation representation and alleviate the sparse-view and dynamic drift challenges [44].

Finally, the most constrained axis is deployment at the edge. The operational environment—be it a mobile phone, a drone, or an AR headset—imposes strict limits on memory, power consumption, and latency that are incompatible with the original 3DGS training pipeline. Algorithmic refinements and the use of carefully designed kernels have made on-device, real-time inference of pre-optimized, pre-compressed models an attainable target, complementing the lightweight runtime pathways (e.g., Vulkan or WebGL) discussed in the section above. The key to robustness is a holistic system design. For instance, efficient kernels can be tuned to the specific warp-level operations of a target GPU, and hierarchical sorting and visibility culling can drastically reduce the per-primitive workload. The models themselves must be re-formulated for inference: the number of primitives must be significantly reduced and quantized, and lightweight level-of-detail schemes need to be bridged to mobile-appropriate culling. In a power-constrained context, achieving high frame rates requires rethinking the entire rendering pipeline—from raw primitive storage to final alpha compositing—around a memory-centric architecture that efficiently and frequently moves data, prioritizing minimal DRAM traffic over raw compute speed.

The path toward fully scalable 3DGS requires a holistic view—a tight integration of compression, out-of-core management, and hardware-aware fast rendering. Robust and fair evaluation remains a challenge, as embedding sparse-view reconstructions also exposes sensitivity to initialization and the difficulty of correctly modeling coverage. However, the explicit nature of the representation offers a distinct advantage, one reinforced by the editability and interactivity tools discussed earlier, the ability to update and refine scenes incrementally without full retraining. This makes 3DGS a prime candidate for future interactive and "lifelong" 3D vision systems—a unique bridge between the modern authoring tools and the GPU-driven engines described above in the deployment pipeline. The future thus lies not in creating a single static asset, but in maintaining a stream of them—learning and updating a scalable representation that is stored on the server and rendered in real-time on the device, delivering the seamless, privacy-preserving, and always-available interface that foundational layer promised.

8.5 Future Horizons, Emerging Paradigms, and Multi-Modal Integration

As 3D Gaussian Splatting (3DGS) matures from a specialized rendering technique into a versatile 3D representation, its trajectory is being redefined by a convergence of emerging paradigms that promise to elevate it from a tool for photorealistic novel view synthesis to a foundational layer for next-generation interactive and generative applications. The future horizon is not merely one of incremental improvements to optimization and rendering, but a fundamental shift towards generalization, physical realism, multi-modal integration, and seamless interoperability with established digital content pipelines.

The first significant frontier lies in overcoming the requirement for per-scene optimization to achieve true generalization. While feed-forward and transformer-based architectures have made strides in predicting Gaussian parameters from sparse inputs [149], these methods often

produce discrete, color-biased point clouds that lack the geometric coherence required for surface extraction. The path forward involves integrating the explicit, high-fidelity rendering of 3DGS with the topological flexibility of implicit representations. The emerging trend is the use of hybrid models that introduce surface continuity priors alongside the discrete Gaussian primitives to guide the placement and distribution of Gaussians toward a latent surface [49]. The success of methods like MeshGS and DeMapGS, which bind Gaussians to a deformable surface mesh, demonstrates that coupling explicit rendering with structured geometry is key to unlocking robust generalization and a wide range of downstream applications, such as animation and cross-object manipulation [150, 111].

Parallel to this, a critical frontier lies in extending the representation from a purely photometric one to a physically-aware and editable asset. The explicit, unstructured nature of Gaussians, while an advantage for rendering speed, is an impediment to physical realism. Emerging paradigms are shifting towards parameterizing Gaussians with material properties and physical constraints. This includes a push to move beyond simple spherical harmonic color to physically-based asset modeling, where the intrinsic attributes of a surface are decoupled from its appearance for relighting. To this end, Gaussian primitives are being augmented with learnable oriented normals and occupancy fields, allowing for closed-form expressions that bridge the gap between a volumetric representation and a watertight mesh surface, enabling the extraction of complex thin structures [151]. Simultaneously, a new wave of research is introducing adapted attenuation formulations and surface-aligned regularization terms to steer the Gaussian primitives themselves to be distributed evenly along latent surfaces, minimizing clutter of the surface and enhancing the overall geometric accuracy using every single point [47]. These advancements are leading to not just visually accurate models, but geometrically sound ones that can be integrated into physical simulations using tools like position-based dynamics, and seamlessly edited within traditional softwares for games and virtual reality.

The third pillar of this future is multi-modal integration, not just as a means for conditioning but as an interoperability and consistency issue. Future systems will be expected to consume discrete modalities, and the biological cues will heavily influence the methodology. For example, the sparse-view problem, which remains a critical challenge for 3DGS, is being tackled by integrating monocular depth and normal priors into the optimization loss, and incorporating MVS-based geometric regularization, leveraging multi-view constraints that provide side cues for alignment [49, 51]. This approach paves the way for a form of "grounded" generative modeling, where language or sketch queries could eventually be represented in the early optimization stages, a cross-modal priors that link 3D appearance with audio and semantic context. The ability to simultaneously bind the Gaussian field to a multi-modal foundation will be a key differentiator in making the representation contextual and hallucination-free.

8.6 Conclusion

The next chapter of 3DGS resides in its evolution into an "explicit multimodal foundation model" layer for the 3D world. The current state-of-the-art highlights the necessity of hybrids—the convergence of implicit and explicit representations, the integration or bridging of new modalities, and the enrichment of a database with semantically consistent geometric and physical attributes through rigorous regularization. The research community is positioned to shift from a representation-centric approach to a "generative-geometry-centric" one, where the fundamental goal is not the blind faithful reconstruction of an image but the accurate, efficient, and editable capture of a virtual world. This will solidify 3DGS's position as a key hub, binding a variety of direct and generalized algorithms and a fully realized, interactive, and physically plausible digital reality.

9 Conclusion

This survey has charted the trajectory of 3D Gaussian Splatting (3DGS) from its foundational formulation [2] to its current status as a versatile and rapidly expanding field. The central paradigms that have emerged coalesce around the core tenets of efficient optimization, controllable representation, and application-driven extension. As we have detailed, the initial formulation established an explicit, primitive-based representation whose differentiable rasterization pipeline enables high-quality real-time rendering [2]. Subsequent research has systematically addressed the initial method’s limitations, focusing on robust densification and optimization strategies [11, 71, 152], geometric fidelity and surface extraction [9, 153, 6], and the critical scalability and compression challenges for unrestricted deployment [25, 154, 155, 20].

The scope of 3DGS has also broadened considerably beyond static scene reconstruction. The community has developed a rich taxonomy of methods for modeling dynamic scenes, either through deformation architectures [25] or direct spatio-temporal primitives [12]. This exploration extends not only in time but also in appearance, where significant progress has been made in decoupling geometry and material properties for improved photorealistic rendering and editing [14, 13], and in understanding the Gaussian parameter space through the lens of more general volumetric rendering [4, 22]. Furthermore, the explicit nature of the representation, combined with semantic and language-aware fields, has opened new frontiers in scene understanding and interactive applications [19, 8].

Despite these rapid advances, the field continues to grapple with several crucial bottlenecks and open questions. *Geometric fidelity* remains a delicate balance, where the flexibility of 3D Gaussians can lead to blurring and floaters, requiring robust regularization and more principled treatment of volumetric rendering [22, 10]. Similarly, *global illumination and physically-based appearance* are still largely approximated by view-dependent spherical harmonics and simple shading models; the pursuit of robust inverse rendering and reflective surface modeling within the splatting paradigm is a critical front [3, 132]. The robustness of the model to challenging real-world conditions, such as sparse input views without pre-computed camera poses, or scenarios with highly dynamic or transparent content, continues to be a significant limitation [21, 1]. Furthermore, the field lacks a unified, standardized, and compressible representation that can be seamlessly integrated with traditional graphics pipelines and modern hardware, which is essential for mass adoption in virtual reality, games, and broader industries [7, 20, 6].

Looking forward, the path to a transformative future is clear. The convergence of 3DGS with generative models, exemplified by feed-forward prediction via large transformer architectures and foundation models, which bypass per-scene optimization, signifies a paradigm shift [11, 156]. The extraction of coherent and physically-plausible scene understanding from these generative priors and the grounding of those in dynamic real-world environments is of high intellectual value. The continued exploration of *neural primitives* (e.g., from MLP-based samplers to learnable planar patches) promises to reconcile the expressivity of neural implicit fields with the efficiency of explicit methods, potentially bridging the gap between 3DGS and mesh-based rendering [16, 17, 18, 153]. Finally, the field is moving toward principled and large-scale simulations with 3DGS, from simple deformable objects to complex dynamic scenes, broadening its application domains [157, 131]. The next wave of innovation will therefore *not* be defined by incremental improvements in rendering speed alone, but by a broader vision of 3DGS as a mature, trustworthy, and interoperable foundation that brings "form" and "appearance" together in a physically faithful and semantically editable manner, ready for deployment across the spectrum of digital media, robotics, and scientific simulation.

References

- [1] Guikun Chen, Wenguan Wang. *A Survey on 3D Gaussian Splatting*. arXiv:2401.03890v2, 2024. <https://arxiv.org/abs/2401.03890v2>
- [2] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, George Drettakis. *3D Gaussian Splatting for Real-Time Radiance Field Rendering*. arXiv:2308.04079v1, 2023. <https://arxiv.org/abs/2308.04079v1>
- [3] Tong Wu, Yu-Jie Yuan, Ling-Xiao Zhang, Jie Yang, Yan-Pei Cao, Ling-Qi Yan *et al.*. *Recent Advances in 3D Gaussian Splatting*. arXiv:2403.11134v2, 2024. <https://arxiv.org/abs/2403.11134v2>
- [4] Vitor Pereira Matias, Daniel Perazzo, Vinicius Silva, Alberto Raposo, Luiz Velho, Afonso Paiva *et al.*. *From Volume Rendering to 3D Gaussian Splatting: Theory and Applications*. arXiv:2510.18101v1, 2025. <https://arxiv.org/abs/2510.18101v1>
- [5] Panagiotis Papantonakis, Georgios Kopanas, Bernhard Kerbl, Alexandre Lanvin, George Drettakis. *Reducing the Memory Footprint of 3D Gaussian Splatting*. arXiv:2406.17074v1, 2024. <https://arxiv.org/abs/2406.17074v1>
- [6] Bernhard Kerbl, Andréas Meuleman, Georgios Kopanas, Michael Wimmer, Alexandre Lanvin, George Drettakis. *A Hierarchical 3D Gaussian Representation for Real-Time Rendering of Very Large Datasets*. arXiv:2406.12080v1, 2024. <https://arxiv.org/abs/2406.12080v1>
- [7] Bernhard Kerbl. *The Impact and Outlook of 3D Gaussian Splatting*. arXiv:2510.26694v1, 2025. <https://arxiv.org/abs/2510.26694v1>
- [8] Ben Fei, Jingyi Xu, Rui Zhang, Qingyuan Zhou, Weidong Yang, Ying He. *3D Gaussian as a New Vision Era A Survey*. arXiv:2402.07181v1, 2024. <https://arxiv.org/abs/2402.07181v1>
- [9] Binbin Huang, Zehao Yu, Anpei Chen, Andreas Geiger, Shenghua Gao. *2D Gaussian Splatting for Geometrically Accurate Radiance Fields*. arXiv:2403.17888v1, 2024. <https://arxiv.org/abs/2403.17888v1>
- [10] Jan Held, Renaud Vandeghen, Abdullah Hamdi, Adrien Deliege, Anthony Cioppa, Silvio Giancola *et al.*. *3D Convex Splatting: Radiance Field Rendering with 3D Smooth Convexes*. arXiv:2411.14974v3, 2024. <https://arxiv.org/abs/2411.14974v3>
- [11] Yecong Wan, Fan Li, Mingwen Shao, Wangmeng Zuo. *SplatWeaver: Learning to Allocate Gaussian Primitives for Generalizable Novel View Synthesis*. arXiv:2605.07287v2, 2026. <https://arxiv.org/abs/2605.07287v2>
- [12] Zeyu Yang, Zijie Pan, Xiatian Zhu, Li Zhang, Jianfeng Feng, Yu-Gang Jiang *et al.*. *4D Gaussian Splatting: Modeling Dynamic Scenes with Native 4D Primitives*. arXiv:2412.20720v2, 2024. <https://arxiv.org/abs/2412.20720v2>
- [13] Victor Rong, Jingxiang Chen, Sherwin Bahmani, Kiriakos N. Kutulakos, David B. Lindell. *GStex: Per-Primitive Texturing of 2D Gaussian Splatting for Decoupled Appearance and Geometry Modeling*. arXiv:2409.12954v1, 2024. <https://arxiv.org/abs/2409.12954v1>
- [14] Yiming Wang, Shaofei Wang, Marko Mihajlovic, Siyu Tang. *Neural Texture Splatting: Expressive 3D Gaussian Splatting for View Synthesis, Geometry, and Dynamic Reconstruction*. arXiv:2511.18873v1, 2025. <https://arxiv.org/abs/2511.18873v1>
- [15] Mae Younes, Adnane Boukhayma. *TextureSplat: Per-Primitive Texture Mapping for Reflective Gaussian Splatting*. arXiv:2506.13348v2, 2025. <https://arxiv.org/abs/2506.13348v2>

- [16] Xilong Zhou, Bao-Huy Nguyen, Loïc Magne, Vladislav Golyanik, Thomas Leimkühler, Christian Theobalt. *Splat the Net: Radiance Fields with Splattable Neural Primitives*. arXiv:2510.08491v2, 2025. <https://arxiv.org/abs/2510.08491v2>
- [17] Shinjeong Kim, Yuzhou Cheng, Xin Kong, Paul H. J. Kelly, Andrew J. Davison. *MLP Splatting: Object-Centric Neural Fields*. arXiv:2606.03877v1, 2026. <https://arxiv.org/abs/2606.03877v1>
- [18] David Svitov, Pietro Morerio, Lourdes Agapito, Alessio Del Bue. *BillBoard Splatting (BBSplat): Learnable Textured Primitives for Novel View Synthesis*. arXiv:2411.08508v4, 2024. <https://arxiv.org/abs/2411.08508v4>
- [19] Shuting He, Peilin Ji, Yitong Yang, Changshuo Wang, Jiayi Ji, Yinglin Wang *et al.*. *A Survey on 3D Gaussian Splatting Applications: Segmentation, Editing, and Generation*. arXiv:2508.09977v5, 2025. <https://arxiv.org/abs/2508.09977v5>
- [20] Muhammad Salman Ali, Chaoning Zhang, Marco Cagnazzo, Giuseppe Valenzise, Enzo Tartaglione, Sung-Ho Bae. *Compression in 3D Gaussian Splatting: A Survey of Methods, Trends, and Future Directions*. arXiv:2502.19457v1, 2025. <https://arxiv.org/abs/2502.19457v1>
- [21] Yang Fu, Sifei Liu, Amey Kulkarni, Jan Kautz, Alexei A. Efros, Xiaolong Wang. *COLMAP-Free 3D Gaussian Splatting*. arXiv:2312.07504v1, 2023. <https://arxiv.org/abs/2312.07504v1>
- [22] Adam Celarek, George Kopanas, George Drettakis, Michael Wimmer, Bernhard Kerbl. *Does 3D Gaussian Splatting Need Accurate Volumetric Rendering?*. arXiv:2502.19318v1, 2025. <https://arxiv.org/abs/2502.19318v1>
- [23] Xin Zhang, Anpei Chen, Jincheng Xiong, Pinxuan Dai, Yujun Shen, Weiwei Xu. *Neural Shell Texture Splatting: More Details and Fewer Primitives*. arXiv:2507.20200v1, 2025. <https://arxiv.org/abs/2507.20200v1>
- [24] Baowen Zhang, Chenxing Jiang, Heng Li, Shaojie Shen, Ping Tan. *Geometry-Grounded Gaussian Splatting*. arXiv:2601.17835v2, 2026. <https://arxiv.org/abs/2601.17835v2>
- [25] Kai Katsumata, Duc Minh Vo, Hideki Nakayama. *An Efficient 3D Gaussian Representation for Monocular Multi-view Dynamic Scenes*. arXiv:2311.12897v1, 2023. <https://arxiv.org/abs/2311.12897v1>
- [26] Peihao Wang, Yuehao Wang, Dilin Wang, Sreyas Mohan, Zhiwen Fan, Lemeng Wu *et al.*. *Steepest Descent Density Control for Compact 3D Gaussian Splatting*. arXiv:2505.05587v1, 2025. <https://arxiv.org/abs/2505.05587v1>
- [27] Zhenhua Ning, Xin Li, Jun Yu, Guangming Lu, Yaowei Wang, Wenjie Pei. *Beyond Heuristics: Learnable Density Control for 3D Gaussian Splatting*. arXiv:2605.00408v2, 2026. <https://arxiv.org/abs/2605.00408v2>
- [28] Qingyuan Zhou, Xinyi Liu, Weidong Yang, Ning Wang, Shuquan Ye, Ben Fei *et al.*. *Denoising-GS: Gaussian Splatting with Spatial-aware Denoising*. arXiv:2605.14880v1, 2026. <https://arxiv.org/abs/2605.14880v1>
- [29] Binxiao Huang, Zhengwu Liu, Ngai Wong. *Decomposing Densification in Gaussian Splatting for Faster 3D Scene Reconstruction*. arXiv:2507.20239v1, 2025. <https://arxiv.org/abs/2507.20239v1>
- [30] Jiashu Li, Xumeng Han, Zhaoyang Wei, Zipeng Wang, Kuiran Wang, Guorong Li *et al.*. *HeroGS: Hierarchical Guidance for Robust 3D Gaussian Splatting under Sparse Views*. arXiv:2603.01099v2, 2026. <https://arxiv.org/abs/2603.01099v2>

- [31] Yuxin Cheng, Binxiao Huang, Wenyong Zhou, Taiqiang Wu, Zhengwu Liu, Graziano Chesi *et al.*. *Re-Activating Frozen Primitives for 3D Gaussian Splatting*. arXiv:2510.19653v1, 2025. <https://arxiv.org/abs/2510.19653v1>
- [32] Mahdi Naseri, Jiayan Qiu, Zhou Wang. *Generated Contents Enrichment*. arXiv:2405.03650v2, 2024. <https://arxiv.org/abs/2405.03650v2>
- [33] Abdelrhman Elrawy, Emad A. Mohammed. *Opacity-Gradient Driven Density Control for Compact and Efficient Few-Shot 3D Gaussian Splatting*. arXiv:2510.10257v1, 2025. <https://arxiv.org/abs/2510.10257v1>
- [34] Xiangyu Sun, Runnan Chen, Mingming Gong, Dong Xu, Tongliang Liu. *Intern-GS: Vision Model Guided Sparse-View 3D Gaussian Splatting*. arXiv:2505.20729v1, 2025. <https://arxiv.org/abs/2505.20729v1>
- [35] Kerui Ren, Lihan Jiang, Tao Lu, Mulin Yu, Linning Xu, Zhangkai Ni *et al.*. *Octree-GS Towards Consistent Real-time Rendering with LOD-Structured 3D Gaussians*. arXiv:2403.17898v1, 2024. <https://arxiv.org/abs/2403.17898v1>
- [36] Rui Peng, Wangze Xu, Luyang Tang, Liwei Liao, Jianbo Jiao, Ronggang Wang. *Structure Consistent Gaussian Splatting with Matching Prior for Few-shot Novel View Synthesis*. arXiv:2411.03637v1, 2024. <https://arxiv.org/abs/2411.03637v1>
- [37] Yiqing Li, Xuan Wang, Jiawei Wu, Yikun Ma, Zhi Jin. *SparseGS-W: Sparse-View 3D Gaussian Splatting in the Wild with Generative Priors*. arXiv:2503.19452v1, 2025. <https://arxiv.org/abs/2503.19452v1>
- [38] Hyunwoo Park, Gun Ryu, Wonjun Kim. *DropGaussian: Structural Regularization for Sparse-view Gaussian Splatting*. arXiv:2504.00773v1, 2025. <https://arxiv.org/abs/2504.00773v1>
- [39] Meiying Gu, Jiawei Zhang, Jiahe Li, Xiaohan Yu, Haonan Luo, Jin Zheng *et al.*. *SparseSurf: Sparse-View 3D Gaussian Splatting for Surface Reconstruction*. arXiv:2511.14633v1, 2025. <https://arxiv.org/abs/2511.14633v1>
- [40] Zhaorui Wang, Yi Gu, Deming Zhou, Renjing Xu. *FixingGS: Enhancing 3D Gaussian Splatting via Training-Free Score Distillation*. arXiv:2509.18759v1, 2025. <https://arxiv.org/abs/2509.18759v1>
- [41] Zhen Tan, Xieyuanli Chen, Jinpu Zhang, Lei Feng, Dewen Hu. *Uncertainty-Aware Normal-Guided Gaussian Splatting for Surface Reconstruction from Sparse Image Sequences*. arXiv:2503.11172v1, 2025. <https://arxiv.org/abs/2503.11172v1>
- [42] Rui Song, Tianhui Cai, Markus Gross, Yun Zhang, Walter Zimmer, Zhiyu Huang *et al.*. *EnerGS: Energy-Based Gaussian Splatting with Partial Geometric Priors*. arXiv:2604.26238v1, 2026. <https://arxiv.org/abs/2604.26238v1>
- [43] Jungeon Kim, Geonsoo Park, Seungyong Lee. *Multiview Geometric Regularization of Gaussian Splatting for Accurate Radiance Fields*. arXiv:2506.13508v1, 2025. <https://arxiv.org/abs/2506.13508v1>
- [44] Jijian Zhao. *D²-4DGS: Dual-Depth Guided Sparse-Camera 4D Gaussian Splatting*. arXiv:2608.01588v1, 2026. <https://arxiv.org/abs/2608.01588v1>
- [45] Niluthpol Chowdhury Mithun, Tuan Pham, Qiao Wang, Ben Southall, Kshitij Minhas, Bogdan Matei *et al.*. *Diffusion-Guided Gaussian Splatting for Large-Scale Unconstrained 3D Reconstruction and Novel View Synthesis*. arXiv:2504.01960v1, 2025. <https://arxiv.org/abs/2504.01960v1>

- [46] Xingyilang Yin, Qi Zhang, Jiahao Chang, Ying Feng, Qingnan Fan, Xi Yang *et al.*. *GSFixer: Improving 3D Gaussian Splatting with Reference-Guided Video Diffusion Priors*. arXiv:2508.09667v1, 2025. <https://arxiv.org/abs/2508.09667v1>
- [47] Jiepeng Wang, Yuan Liu, Peng Wang, Cheng Lin, Junhui Hou, Xin Li *et al.*. *GausSurf: Geometry-Guided 3D Gaussian Splatting for Surface Reconstruction*. arXiv:2411.19454v2, 2024. <https://arxiv.org/abs/2411.19454v2>
- [48] Prajwal Gupta C. R., Divyam Sheth, Jinjoo Ha, Mirela Ostrek, Justus Thies. *2D-SuGaR: Surface-Aware Gaussian Splatting for Geometrically Accurate Mesh Reconstruction*. arXiv:2605.00569v1, 2026. <https://arxiv.org/abs/2605.00569v1>
- [49] Kunyi Li, Michael Niemeyer, Zeyu Chen, Nassir Navab, Federico Tombari. *MonoGSDF: Exploring Monocular Geometric Cues for Gaussian Splatting-Guided Implicit Surface Reconstruction*. arXiv:2411.16898v4, 2024. <https://arxiv.org/abs/2411.16898v4>
- [50] Hong Kuang, Jianchen Liu. *Gaussian Entropy Fields: Driving Adaptive Sparsity in 3D Gaussian Optimization*. arXiv:2512.04542v1, 2025. <https://arxiv.org/abs/2512.04542v1>
- [51] Matias Turkulainen, Xuqian Ren, Iaroslav Melekhov, Otto Seiskari, Esa Rahtu, Juho Kannala. *DN-Splatter Depth and Normal Priors for Gaussian Splatting and Meshing*. arXiv:2403.17822v1, 2024. <https://arxiv.org/abs/2403.17822v1>
- [52] Antoine Guédon, Diego Gomez, Nissim Maruani, Bingchen Gong, George Drettakis, Maks Ovsjanikov. *MILo: Mesh-In-the-Loop Gaussian Splatting for Detailed and Efficient Surface Reconstruction*. arXiv:2506.24096v2, 2025. <https://arxiv.org/abs/2506.24096v2>
- [53] Jiateng Liu, Hao Gao, Jiu-Cheng Xie, Chi-Man Pun, Jian Xiong, Haolun Li *et al.*. *Effective Gaussian Management for High-fidelity Scene Reconstruction*. arXiv:2509.12742v4, 2025. <https://arxiv.org/abs/2509.12742v4>
- [54] Baowen Zhang, Chuan Fang, Rakesh Shrestha, Yixun Liang, Xiaoxiao Long, Ping Tan. *RaDe-GS: Rasterizing Depth in Gaussian Splatting*. arXiv:2406.01467v2, 2024. <https://arxiv.org/abs/2406.01467v2>
- [55] Danpeng Chen, Hai Li, Weicai Ye, Yifan Wang, Weijian Xie, Shangjin Zhai *et al.*. *PGSR: Planar-based Gaussian Splatting for Efficient and High-Fidelity Surface Reconstruction*. arXiv:2406.06521v1, 2024. <https://arxiv.org/abs/2406.06521v1>
- [56] Bo Jia, Yanan Guo, Ying Chang, Benkui Zhang, Ying Xie, Kangning Du *et al.*. *Multi-view Normal and Distance Guidance Gaussian Splatting for Surface Reconstruction*. arXiv:2508.07701v2, 2025. <https://arxiv.org/abs/2508.07701v2>
- [57] Akilan Thangarajah, Buddhapala Wongkaew, Mongkol Ekpanyapong. *Implementation of Auto Monitoring and Short-Message-Service System via GSM Modem*. arXiv:1501.01548v2, 2015. <https://arxiv.org/abs/1501.01548v2>
- [58] Xiaobiao Du, Yida Wang, Kun Zhan, Xin Yu. *Mobile-GS: Real-time Gaussian Splatting for Mobile Devices*. arXiv:2603.11531v1, 2026. <https://arxiv.org/abs/2603.11531v1>
- [59] Yang Luo, Yan Gong, Yongsheng Gao, Jie Zhao, Xinyu Zhang, Huaping Liu. *Local-GS: Accelerating 3D Gaussian Splatting via Tile-Local Warp Coherence*. arXiv:2606.16566v1, 2026. <https://arxiv.org/abs/2606.16566v1>
- [60] Samuel Rota Bulò, Nemanja Bartolovic, Lorenzo Porzi, Peter Kontschieder. *Hardware-Rasterized Ray-Based Gaussian Splatting*. arXiv:2503.18682v2, 2025. <https://arxiv.org/abs/2503.18682v2>

- [61] Xinze Li, Bohan Yang, Pengxu Chen, Yiyuan Wang, Hongcheng Luo, Wentao Cheng *et al.*. *QuadBox: Accelerating 3D Gaussian Splatting with Geometry-Aware Boxes*. arXiv:2605.04844v1, 2026. <https://arxiv.org/abs/2605.04844v1>
- [62] Yancheng Zhang, Guangyu Sun, Chen Chen. *EGGS: Exchangeable 2D/3D Gaussian Splatting for Geometry-Appearance Balanced Novel View Synthesis*. arXiv:2512.02932v1, 2025. <https://arxiv.org/abs/2512.02932v1>
- [63] Mihnea-Bogdan Jurca, Bert Van hauerweiren, Adrian Munteanu. *Fourier Splatting: Generalized Fourier encoded primitives for scalable radiance fields*. arXiv:2603.19834v3, 2026. <https://arxiv.org/abs/2603.19834v3>
- [64] Chongjie Ye, Yinyu Nie, Jiahao Chang, Yuntao Chen, Yihao Zhi, Xiaoguang Han. *GauStudio A Modular Framework for 3D Gaussian Splatting and Beyond*. arXiv:2403.19632v1, 2024. <https://arxiv.org/abs/2403.19632v1>
- [65] Zhang Chen, Shuai Wan, Mengting Yu, Fuzheng Yang, Junhui Hou. *REFINE: Super-efficient 3D Gaussian Splatting Pruning via Rendering-Free Primitive Importance*. arXiv:2606.09074v3, 2026. <https://arxiv.org/abs/2606.09074v3>
- [66] Xiaobin Deng, Qiuli Yu, Changyu Diao, Min Li, Duanqing Xu. *Gradient-Driven Natural Selection for Compact 3D Gaussian Splatting*. arXiv:2511.16980v1, 2025. <https://arxiv.org/abs/2511.16980v1>
- [67] Soonbin Lee, Yeong-Gyu Kim, Simon Sasse, Tomas M. Borges, Yago Sanchez, Eun-Seok Ryu *et al.*. *GaussianPOP: Principled Simplification Framework for Compact 3D Gaussian Splatting via Error Quantification*. arXiv:2602.06830v1, 2026. <https://arxiv.org/abs/2602.06830v1>
- [68] Xianben Yang, Tao Wang, Yuxuan Li, Yi Jin, Haibin Ling. *GS²: Graph-based Spatial Distribution Optimization for Compact 3D Gaussian Splatting*. arXiv:2604.01884v1, 2026. <https://arxiv.org/abs/2604.01884v1>
- [69] Saswat Subhajyoti Mallick, Rahul Goel, Bernhard Kerbl, Francisco Vicente Carrasco, Markus Steinberger, Fernando De La Torre. *Taming 3DGS: High-Quality Radiance Fields with Limited Resources*. arXiv:2406.15643v1, 2024. <https://arxiv.org/abs/2406.15643v1>
- [70] Deng Siqin, Zhou Xiaoyi. *One Pixel is All I Need*. arXiv:2412.10681v2, 2024. <https://arxiv.org/abs/2412.10681v2>
- [71] Jijian Zhao. *ACE-GS: Acing the Trade-off with Accurate, Compact and Efficient 3D Gaussian Splatting*. arXiv:2606.21244v1, 2026. <https://arxiv.org/abs/2606.21244v1>
- [72] Kaifa Yang, Qi Yang, Yiling Xu, Zhu Li. *RAP: Fast Feedforward Rendering-Free Attribute-Guided Primitive Importance Score Prediction for Efficient 3D Gaussian Splatting Processing*. arXiv:2602.19753v1, 2026. <https://arxiv.org/abs/2602.19753v1>
- [73] Atakan Topaloglu, Kunyi Li, Michael Niemeyer, Nassir Navab, A. Murat Tekalp, Federico Tombari. *OracleGS: Grounding Generative Priors for Sparse-View Gaussian Splatting*. arXiv:2509.23258v2, 2025. <https://arxiv.org/abs/2509.23258v2>
- [74] Jimin Tang, Wenyan Zhang, Junsheng Zhou, Zian Huang, Kanle Shi, Shenkun Xu *et al.*. *VidSplat: Gaussian Splatting Reconstruction with Geometry-Guided Video Diffusion Priors*. arXiv:2605.11424v1, 2026. <https://arxiv.org/abs/2605.11424v1>
- [75] Minnan Pei, Gang Li, Junwen Si, Zeyu Zhu, Zitao Mo, Peisong Wang *et al.*. *GCC: A 3DGS Inference Architecture with Gaussian-Wise and Cross-Stage Conditional Processing*. arXiv:2507.15300v3, 2025. <https://arxiv.org/abs/2507.15300v3>

- [76] Chuanyu Fu, Yuqi Zhang, Kunbin Yao, Guanying Chen, Yuan Xiong, Chuan Huang *et al.*. *RobustSplat: Decoupling Densification and Dynamics for Transient-Free 3DGS*. arXiv:2506.02751v3, 2025. <https://arxiv.org/abs/2506.02751v3>
- [77] Junfeng Ni, Yixin Chen, Zhifei Yang, Yu Liu, Ruijie Lu, Song-Chun Zhu *et al.*. *G4Splat: Geometry-Guided Gaussian Splatting with Generative Prior*. arXiv:2510.12099v2, 2025. <https://arxiv.org/abs/2510.12099v2>
- [78] Feng Zhou, Wenkai Guo, Pu Cao, Zhicheng Zhang, Jianqin Yin. *Initialize to Generalize: A Stronger Initialization Pipeline for Sparse-View 3DGS*. arXiv:2510.17479v1, 2025. <https://arxiv.org/abs/2510.17479v1>
- [79] Michael Niemeyer, Fabian Manhardt, Marie-Julie Rakotosaona, Michael Oechsle, Daniel Duckworth, Rama Gosula *et al.*. *RadSplat Radiance Field-Informed Gaussian Splatting for Robust Real-Time Rendering with 900+ FPS*. arXiv:2403.13806v1, 2024. <https://arxiv.org/abs/2403.13806v1>
- [80] Gurutva Patle, Nilay Girgaonkar, Nagabhushan Somraj, Rajiv Soundararajan. *AD-GS: Alternating Densification for Sparse-Input 3D Gaussian Splatting*. arXiv:2509.11003v2, 2025. <https://arxiv.org/abs/2509.11003v2>
- [81] Hanqing Jiang, Xiaojun Xiang, Han Sun, Hongjie Li, Liyang Zhou, Xiaoyu Zhang *et al.*. *GeoTexDensifier: Geometry-Texture-Aware Densification for High-Quality Photorealistic 3D Gaussian Splatting*. arXiv:2412.16809v2, 2024. <https://arxiv.org/abs/2412.16809v2>
- [82] Xingjun Wang, Lianlei Shan. *GDGS: 3D Gaussian Splatting Via Geometry-Guided Initialization And Dynamic Density Control*. arXiv:2507.00363v1, 2025. <https://arxiv.org/abs/2507.00363v1>
- [83] Phurtivilai Patt, Leyang Huang, Yinqiang Zhang, Yang Lei. *DensifyBeforehand: LiDAR-assisted Content-aware Densification for Efficient and Quality 3D Gaussian Splatting*. arXiv:2511.19294v1, 2025. <https://arxiv.org/abs/2511.19294v1>
- [84] Injae Kim, Chaehyeon Kim, Minseong Bae, Minseok Joo, Hyunwoo J. Kim. *F4Splat: Feed-Forward Predictive Densification for Feed-Forward 3D Gaussian Splatting*. arXiv:2603.21304v2, 2026. <https://arxiv.org/abs/2603.21304v2>
- [85] Jiawei Ren, Liang Pan, Jiaxiang Tang, Chi Zhang, Ang Cao, Gang Zeng *et al.*. *DreamGaussian4D Generative 4D Gaussian Splatting*. arXiv:2312.17142v2, 2023. <https://arxiv.org/abs/2312.17142v2>
- [86] Dmytro Kotovenko, Olga Grebenkova, Björn Ommer. *EDGS: Eliminating Densification for Efficient Convergence of 3DGS*. arXiv:2504.13204v2, 2025. <https://arxiv.org/abs/2504.13204v2>
- [87] Yiqing Liang, Numair Khan, Zhengqin Li, Thu Nguyen-Phuoc, Douglas Lanman, James Tompkin *et al.*. *GauFRe Gaussian Deformation Fields for Real-time Dynamic Novel View Synthesis*. arXiv:2312.11458v1, 2023. <https://arxiv.org/abs/2312.11458v1>
- [88] Lauren Milechin, Alexander Chen, Vijay Gadepally, Dylan Hutchison, Siddharth Samsi, Jeremy Kepner. *D4M 3.0*. arXiv:1702.03253v1, 2017. <https://arxiv.org/abs/1702.03253v1>
- [89] Hanyang Kong, Dongze Lian, Michael Bi Mi, Xinchao Wang. *DreamDrone*. arXiv:2312.08746v2, 2023. <https://arxiv.org/abs/2312.08746v2>
- [90] Md Musfiqur Rahman Sanim, Zhihao Shu, Bahram Afsharmanesh, AmirAli Mirian, Jiexiong Guan, Wei Niu *et al.*. *Optimizing 3D Gaussian Splattering for Mobile GPUs*. arXiv:2511.16298v1, 2025. <https://arxiv.org/abs/2511.16298v1>

- [91] Zhendong Wang, Cihan Ruan, Jingchuan Xiao, Chuqing Shi, Wei Jiang, Wei Wang *et al.*. *Analysis of Converged 3D Gaussian Splatting Solutions: Density Effects and Prediction Limit.* arXiv:2602.08909v1, 2026. <https://arxiv.org/abs/2602.08909v1>
- [92] Saimouli Katragadda, Cho-Ying Wu, Yuliang Guo, Xinyu Huang, Guoquan Huang, Liu Ren. *Online Language Splatting.* arXiv:2503.09447v3, 2025. <https://arxiv.org/abs/2503.09447v3>
- [93] Zhuo Chen, Zhiguo Ding. *On the Distribution of GSVD.* arXiv:1902.09134v1, 2019. <https://arxiv.org/abs/1902.09134v1>
- [94] Wangze Xu, Huachen Gao, Shihe Shen, Rui Peng, Jianbo Jiao, Ronggang Wang. *MVPGS: Excavating Multi-view Priors for Gaussian Splatting from Sparse Input Views.* arXiv:2409.14316v1, 2024. <https://arxiv.org/abs/2409.14316v1>
- [95] Jiang Wu, Rui Li, Yu Zhu, Rong Guo, Jinqiu Sun, Yanning Zhang. *Sparse2DGS: Geometry-Prioritized Gaussian Splatting for Surface Reconstruction from Sparse Views.* arXiv:2504.20378v1, 2025. <https://arxiv.org/abs/2504.20378v1>
- [96] Yingji Zhong, Zhihao Li, Dave Zhenyu Chen, Lanqing Hong, Dan Xu. *Taming Video Diffusion Prior with Scene-Grounding Guidance for 3D Gaussian Splatting from Sparse Inputs.* arXiv:2503.05082v1, 2025. <https://arxiv.org/abs/2503.05082v1>
- [97] Junbao Zhou, Qingshan Xu, Yuan Zhou, Xiaolong Shen, Beier Zhu, Kesen Zhao *et al.*. *From Uncertainty to Stability and Fidelity: Guiding Sparse-View 3D Gaussian Splatting with Fisher Information.* arXiv:2606.20842v1, 2026. <https://arxiv.org/abs/2606.20842v1>
- [98] Avinash Paliwal, Xilong Zhou, Wei Ye, Jinhui Xiong, Rakesh Ranjan, Nima Khademi Kalantari. *RI3D: Few-Shot Gaussian Splatting With Repair and Inpainting Diffusion Priors.* arXiv:2503.10860v2, 2025. <https://arxiv.org/abs/2503.10860v2>
- [99] Zhanning Gao, Le Wang, Nebojsa Jojic, Zhenxing Niu, Nanning Zheng, Gang Hua. *Video Imprint.* arXiv:2106.03283v1, 2021. <https://arxiv.org/abs/2106.03283v1>
- [100] Minghan Qin, Wanhua Li, Jiawei Zhou, Haoqian Wang, Hanspeter Pfister. *LangSplat 3D Language Gaussian Splatting.* arXiv:2312.16084v2, 2023. <https://arxiv.org/abs/2312.16084v2>
- [101] Liuyue Xie, Joel Julin, Koichiro Niinuma, Laszlo A. Jeni. *Gaussian Splatting LK.* arXiv:2407.11309v1, 2024. <https://arxiv.org/abs/2407.11309v1>
- [102] Pedro H. J. Nardelli. *The Quest for Sustainable Smart Grids.* arXiv:1308.3229v3, 2013. <https://arxiv.org/abs/1308.3229v3>
- [103] Lee Spector, Li Ding, Ryan Boldi. *Particularity.* arXiv:2306.06812v1, 2023. <https://arxiv.org/abs/2306.06812v1>
- [104] Hanlin Chen, Fangyin Wei, Chen Li, Tianxin Huang, Yunsong Wang, Gim Hee Lee. *VCR-GauS: View Consistent Depth-Normal Regularizer for Gaussian Surface Reconstruction.* arXiv:2406.05774v1, 2024. <https://arxiv.org/abs/2406.05774v1>
- [105] Zhuoxiao Li, Shanliang Yao, Yijie Chu, Angel F. Garcia-Fernandez, Yong Yue, Eng Gee Lim *et al.*. *MVG-Splatting: Multi-View Guided Gaussian Splatting with Adaptive Quantile-Based Geometric Consistency Densification.* arXiv:2407.11840v1, 2024. <https://arxiv.org/abs/2407.11840v1>
- [106] Xiao Ren, Yu Liu, Ning An, Jian Cheng, Xin Qiao, He Kong. *GSM-GS: Geometry-Constrained Single and Multi-view Gaussian Splatting for Surface Reconstruction.* arXiv:2602.12796v1, 2026. <https://arxiv.org/abs/2602.12796v1>

- [107] Zihui Gao, Jia-Wang Bian, Guosheng Lin, Hao Chen, Chunhua Shen. *SurfaceSplat: Connecting Surface Reconstruction and Gaussian Splatting*. arXiv:2507.15602v2, 2025. <https://arxiv.org/abs/2507.15602v2>
- [108] Miriam Jäger, Markus Hillemann, Boris Jutzi. *FeatureGS: Eigenvalue-Feature Optimization in 3D Gaussian Splatting for Geometrically Accurate and Artifact-Reduced Reconstruction*. arXiv:2501.17655v1, 2025. <https://arxiv.org/abs/2501.17655v1>
- [109] Oron Nir, Idan Vidra, Avi Neeman, Barak Kinarti, Ariel Shamir. *VCR Video representation for Contextual Retrieval*. arXiv:2402.07466v1, 2024. <https://arxiv.org/abs/2402.07466v1>
- [110] Bin Han, Muxia Sun, H. Vincent Poor, Hans D. Schotten. *Covariance-Aware Demapping on Fourier-Curve Constellations*. arXiv:2604.24918v1, 2026. <https://arxiv.org/abs/2604.24918v1>
- [111] Shuyi Zhou, Shengze Zhong, Kenshi Takayama, Takafumi Taketomi, Takeshi Oishi. *DeMapGS: Simultaneous Mesh Deformation and Surface Attribute Mapping via Gaussian Splatting*. arXiv:2512.10572v1, 2025. <https://arxiv.org/abs/2512.10572v1>
- [112] Vikram Sandu, Mayurdeep Pathak, Rajiv Soundararajan. *Temporally Aware Densification for Dynamic 3D Gaussian Splatting*. arXiv:2606.23212v1, 2026. <https://arxiv.org/abs/2606.23212v1>
- [113] Zihan Zheng, Zhenglong Wu, Xuanxuan Wang, Houqiang Zhong, Xiaoyun Zhang, Qiang Hu *et al.*. *Constrained Dynamic Gaussian Splatting*. arXiv:2602.03538v1, 2026. <https://arxiv.org/abs/2602.03538v1>
- [114] Jun Guo, Xiaojian Ma, Yue Fan, Huaping Liu, Qing Li. *Semantic Gaussians Open-Vocabulary Scene Understanding with 3D Gaussian Splatting*. arXiv:2403.15624v1, 2024. <https://arxiv.org/abs/2403.15624v1>
- [115] Mingqiao Ye, Martin Danelljan, Fisher Yu, Lei Ke. *Gaussian Grouping Segment and Edit Anything in 3D Scenes*. arXiv:2312.00732v1, 2023. <https://arxiv.org/abs/2312.00732v1>
- [116] Shubhendu Jena, Shishir Reddy Vutukur, Adnane Boukhayma. *SparSplat: Fast Multi-View Reconstruction with Generalizable 2D Gaussian Splatting*. arXiv:2505.02175v1, 2025. <https://arxiv.org/abs/2505.02175v1>
- [117] Yuze Wang, Yue Qi. *Taking Language Embedded 3D Gaussian Splatting into the Wild*. arXiv:2507.19830v2, 2025. <https://arxiv.org/abs/2507.19830v2>
- [118] Yuanyuan Gao, Yuning Gong, Yifei Liu, Li Jingfeng, Dingwen Zhang, Yanci Zhang *et al.*. *Proxy-GS: Unified Occlusion Priors for Training and Inference in Structured 3D Gaussian Splatting*. arXiv:2509.24421v5, 2025. <https://arxiv.org/abs/2509.24421v5>
- [119] Siting Zhu, Guangming Wang, Xin Kong, Dezhi Kong, Hesheng Wang. *3D Gaussian Splatting in Robotics: A Survey*. arXiv:2410.12262v2, 2024. <https://arxiv.org/abs/2410.12262v2>
- [120] Vasco Brattka, Guido Gherardi. *Completion of Choice*. arXiv:1910.13186v3, 2019. <https://arxiv.org/abs/1910.13186v3>
- [121] Linjie Lyu, Ayush Tewari, Jianchun Chen, Thomas Leimkühler, Christian Theobalt. *Faster 3D Gaussian Splatting Convergence via Structure-Aware Densification*. arXiv:2604.28016v1, 2026. <https://arxiv.org/abs/2604.28016v1>
- [122] Miriam Jäger, Boris Jutzi. *EntON: Eigenentropy-Optimized Neighborhood Densification in 3D Gaussian Splatting*. arXiv:2603.06216v1, 2026. <https://arxiv.org/abs/2603.06216v1>

- [123] Jad Abou-Chakra, Krishan Rana, Feras Dayoub, Niko Sünderhauf. *Physically Embodied Gaussian Splatting: A Realtime Correctable World Model for Robotics*. arXiv:2406.10788v1, 2024. <https://arxiv.org/abs/2406.10788v1>
- [124] Qucheng Peng, Benjamin Planche, Zhongpai Gao, Meng Zheng, Anwesa Choudhuri, Terrence Chen *et al.*. *3D Vision-Language Gaussian Splatting*. arXiv:2410.07577v2, 2024. <https://arxiv.org/abs/2410.07577v2>
- [125] Jeongmin Bae, Seoha Kim, Youngsik Yun, Hahyun Lee, Gun Bang, Youngjung Uh. *Per-Gaussian Embedding-Based Deformation for Deformable 3D Gaussian Splatting*. arXiv:2404.03613v1, 2024. <https://arxiv.org/abs/2404.03613v1>
- [126] Chenlu Zhan, Yufei Zhang, Gaoang Wang, Hongwei Wang. *Hi-LSplat: Hierarchical 3D Language Gaussian Splatting*. arXiv:2506.06822v1, 2025. <https://arxiv.org/abs/2506.06822v1>
- [127] Anna-Maria Halacheva, Jan-Nico Zaeche, Xi Wang, Danda Pani Paudel, Luc Van Gool. *GaussianVLM: Scene-centric 3D Vision-Language Models using Language-aligned Gaussian Splats for Embodied Reasoning and Beyond*. arXiv:2507.00886v1, 2025. <https://arxiv.org/abs/2507.00886v1>
- [128] Jiaxu Wang, Ziyi Zhang, Qiang Zhang, Jia Li, Jingkai Sun, Mingyuan Sun *et al.*. *Query-based Semantic Gaussian Field for Scene Representation in Reinforcement Learning*. arXiv:2406.02370v3, 2024. <https://arxiv.org/abs/2406.02370v3>
- [129] Zhenyuan Liu, Yu Guo, Xinyuan Li, Bernd Bickel, Ran Zhang. *BiGS: Bidirectional Gaussian Primitives for Relightable 3D Gaussian Splatting*. arXiv:2408.13370v1, 2024. <https://arxiv.org/abs/2408.13370v1>
- [130] Rui Xu, Wenyue Chen, Jiepeng Wang, Yuan Liu, Peng Wang, Cheng Lin *et al.*. *SVGS: Enhancing Gaussian Splatting Using Primitives with Spatially Varying Colors*. arXiv:2411.18966v2, 2024. <https://arxiv.org/abs/2411.18966v2>
- [131] Krzysztof Byrski, Grzegorz Wilczyński, Weronika Smolak-Dyżewska, Piotr Borycki, Dawid Baran, Sławomir Tadeja *et al.*. *REdiSplats: Ray Tracing for Editable Gaussian Splatting*. arXiv:2503.12284v1, 2025. <https://arxiv.org/abs/2503.12284v1>
- [132] Chang Liu, Hongliang Yuan, Lianghao Zhang, Sichao Wang, Jianwei Guo, Shi-Sheng Huang. *HybridSplat: Fast Reflection-baked Gaussian Tracing using Hybrid Splatting*. arXiv:2512.08334v3, 2025. <https://arxiv.org/abs/2512.08334v3>
- [133] Ziyi Yang, Xinyu Gao, Yangtian Sun, Yihua Huang, Xiaoyang Lyu, Wen Zhou *et al.*. *Spec-Gaussian Anisotropic View-Dependent Appearance for 3D Gaussian Splatting*. arXiv:2402.15870v1, 2024. <https://arxiv.org/abs/2402.15870v1>
- [134] Stéphane Pateux, Matthieu Gendrin, Luce Morin, Théo Ladune, Xiaoran Jiang. *BOGausS: Better Optimized Gaussian Splatting*. arXiv:2504.01844v1, 2025. <https://arxiv.org/abs/2504.01844v1>
- [135] Haodong Chen, Runnan Chen, Qiang Qu, Zhaoqing Wang, Tongliang Liu, Xiaoming Chen *et al.*. *Beyond Gaussians: Fast and High-Fidelity 3D Splatting with Linear Kernels*. arXiv:2411.12440v3, 2024. <https://arxiv.org/abs/2411.12440v3>
- [136] Zongxin Ye, Wenyu Li, Sidun Liu, Peng Qiao, Yong Dou. *AbsGS Recovering Fine Details for 3D Gaussian Splatting*. arXiv:2404.10484v1, 2024. <https://arxiv.org/abs/2404.10484v1>
- [137] Zheng Zhou, Yu-Jie Xiong, Jia-Chen Zhang, Chun-Ming Xia, Xihe Qiu, Hongjian Zhan. *Gradient-Direction-Aware Density Control for 3D Gaussian Splatting*. arXiv:2508.09239v2, 2025. <https://arxiv.org/abs/2508.09239v2>

- [138] Moonsoo Jeong, Dongbeen Kim, Minseong Kim, Sungkil Lee. *DC4GS: Directional Consistency-Driven Adaptive Density Control for 3D Gaussian Splatting*. arXiv:2510.26921v1, 2025. <https://arxiv.org/abs/2510.26921v1>
- [139] Chuanyu Fu, Guanying Chen, Yuqi Zhang, Kunbin Yao, Yuan Xiong, Chuan Huang *et al.*. *RobustSplat++: Decoupling Densification, Dynamics, and Illumination for In-the-Wild 3DGS*. arXiv:2512.04815v1, 2025. <https://arxiv.org/abs/2512.04815v1>
- [140] Zhihao Guo, Jingxuan Su, Chenghao Qian, Shenglin Wang, Jinlong Fan, Jing Zhang *et al.*. *GP-GS: Gaussian Processes Densification for 3D Gaussian Splatting*. arXiv:2502.02283v6, 2025. <https://arxiv.org/abs/2502.02283v6>
- [141] Haosen Yang, Chenhao Zhang, Wenqing Wang, Marco Volino, Adrian Hilton, Li Zhang *et al.*. *Gaussian Splatting with Localized Points Management*. arXiv:2406.04251v2, 2024. <https://arxiv.org/abs/2406.04251v2>
- [142] Hongbi Zhou, Zhangkai Ni. *Perceptual-GS: Scene-adaptive Perceptual Densification for Gaussian Splatting*. arXiv:2506.12400v2, 2025. <https://arxiv.org/abs/2506.12400v2>
- [143] Patrick Schramowski, Manuel Brack, Björn Deiseroth, Kristian Kersting. *Safe Latent Diffusion Mitigating Inappropriate Degeneration in Diffusion Models*. arXiv:2211.05105v4, 2022. <https://arxiv.org/abs/2211.05105v4>
- [144] Lanqing Guo, Yufei Wang, Hezhen Hu, Yan Zheng, Yeying Jin, Siyu Huang *et al.*. *Mastering Regional 3DGS: Locating, Initializing, and Editing with Diverse 2D Priors*. arXiv:2507.05426v1, 2025. <https://arxiv.org/abs/2507.05426v1>
- [145] Zoubin Bi, Yixin Zeng, Chong Zeng, Fan Pei, Xiang Feng, Kun Zhou *et al.*. *GS³: Efficient Relighting with Triple Gaussian Splatting*. arXiv:2410.11419v1, 2024. <https://arxiv.org/abs/2410.11419v1>
- [146] Zhilin Guo, Boqiao Zhang, Hakan Aktas, Kyle Fogarty, Jeffrey Hu, Nursena Koprucu Aslan *et al.*. *Matryoshka Gaussian Splatting*. arXiv:2603.19234v2, 2026. <https://arxiv.org/abs/2603.19234v2>
- [147] Zipeng Wang, Dan Xu. *HyRF: Hybrid Radiance Fields for Memory-efficient and High-quality Novel View Synthesis*. arXiv:2509.17083v2, 2025. <https://arxiv.org/abs/2509.17083v2>
- [148] Tingjia Zhang, Bo Chen, Shengzhong Liu, Fan Wu, Guihai Chen. *CaT-GS: Efficient 3DGS Rendering for Large Scale Scenes via Inter-frame Caching and Tile Scheduling*. arXiv:2607.17842v1, 2026. <https://arxiv.org/abs/2607.17842v1>
- [149] Bing He, Jingnan Gao, Yunuo Chen, Ning Cao, Gang Chen, Zhengxue Cheng *et al.*. *SurfSplat: Conquering Feedforward 2D Gaussian Splatting with Surface Continuity Priors*. arXiv:2602.02000v2, 2026. <https://arxiv.org/abs/2602.02000v2>
- [150] Jaehoon Choi, Yonghan Lee, Hyungtae Lee, Heesung Kwon, Dinesh Manocha. *MeshGS: Adaptive Mesh-Aligned Gaussian Splatting for High-Quality Rendering*. arXiv:2410.08941v1, 2024. <https://arxiv.org/abs/2410.08941v1>
- [151] Diego Gomez, Antoine Guédon, Nissim Maruani, Bingchen Gong, Maks Ovsjanikov. *From Blobs to Spokes: High-Fidelity Surface Reconstruction via Oriented Gaussians*. arXiv:2604.07337v1, 2026. <https://arxiv.org/abs/2604.07337v1>
- [152] Yifan Mo, Youcheng Cai, Ligang Liu. *Plug-and-Play PDE Optimization for 3D Gaussian Splatting: Toward High-Quality Rendering and Reconstruction*. arXiv:2509.13938v1, 2025. <https://arxiv.org/abs/2509.13938v1>

- [153] Xinpeng Liu, Fumio Okura. *Gaussian Mesh Renderer for Lightweight Differentiable Rendering*. arXiv:2602.14493v1, 2026. <https://arxiv.org/abs/2602.14493v1>
- [154] Jiarui Chen, Yikeng Chen, Yingshuang Zou, Ye Huang, Peng Wang, Yuan Liu *et al.*. *MEGS²: Memory-Efficient Gaussian Splatting via Spherical Gaussians and Unified Pruning*. arXiv:2509.07021v3, 2025. <https://arxiv.org/abs/2509.07021v3>
- [155] Yiwen Wang, Siyuan Chen, Ran Yi. *SG-Splatting: Accelerating 3D Gaussian Splatting with Spherical Gaussians*. arXiv:2501.00342v1, 2025. <https://arxiv.org/abs/2501.00342v1>
- [156] Moonyeon Jeong, Seunggi Min, Suhyeon Lee, Hongje Seong. *ViewSplat: View-Adaptive 3D Gaussian Splatting for Feed-Forward Synthesis*. arXiv:2603.25265v2, 2026. <https://arxiv.org/abs/2603.25265v2>
- [157] Yuting Xiao, Xuan Wang, Jiafei Li, Hongrui Cai, Yanbo Fan, Nan Xue *et al.*. *Bridging 3D Gaussian and Mesh for Freeview Video Rendering*. arXiv:2403.11453v1, 2024. <https://arxiv.org/abs/2403.11453v1>