

Adversarial Attacks on Multimodal Large Language Models: A Comprehensive Survey

Generated by SurveyForge

Contents

1	Introduction	2
2	Background and Fundamentals of Multimodal Large Language Models	2
2.1	Introduction to Multimodal Large Language Models	2
2.2	Architectural Overview of Multimodal Large Language Models	3
2.3	Training Objectives and Methods for Multimodal Large Language Models	4
2.4	Applications and Evaluation Metrics for Multimodal Large Language Models	5
2.5	Challenges and Future Directions for Multimodal Large Language Models	6
3	Types and Mechanisms of Adversarial Attacks	7
3.1	Classification of Adversarial Attacks	7
3.2	Methodologies for Generating Adversarial Examples	7
3.3	Evaluation Metrics for Adversarial Attacks	8
3.4	Defense Mechanisms Against Adversarial Attacks	9
3.5	Emerging Trends and Future Directions	10
4	Defense Mechanisms and Robustness Evaluation	11
4.1	Introduction to Defense Mechanisms	11
4.2	Adversarial Training for Defense	12
4.3	Input Preprocessing and Detection Methods	13
4.4	Evaluation Metrics for Defense Mechanisms	14
4.5	Emerging Trends and Future Directions	14
5	Applications and Case Studies of Adversarial Attacks	15
5.1	Introduction to Applications and Case Studies of Adversarial Attacks	15
5.2	Adversarial Attacks on Multimodal Models in Real-World Scenarios	16
5.3	Case Studies of Adversarial Attacks on Multimodal Large Language Models	17
5.4	Lessons Learned and Future Directions	17
5.5	Evaluation Metrics and Benchmarking for Adversarial Attacks	18
5.6	Defense Mechanisms and Countermeasures	19
6	Real-World Implications and Future Directions	20
6.1	Introduction to Real-World Implications	20
6.2	Emerging Trends and Research Opportunities	21
6.3	Open Challenges and Potential Gaps	22
6.4	Future Research Directions	23
6.5	Mitigation Strategies and Best Practices	23
6.6	Conclusion and Recommendations	24
7	Conclusion	25

1 Introduction

The emergence of multimodal large language models has revolutionized the field of artificial intelligence, enabling machines to process and generate content across multiple modalities, including text, images, audio, and video. However, this rapid advancement has also introduced new challenges, particularly in terms of security and robustness. Adversarial attacks, which involve manipulating input data to cause models to produce erroneous or undesirable outputs, have become a significant concern [1]. These attacks can have severe consequences, including compromising model reliability, undermining user trust, and even posing risks to safety and security [2].

Recent studies have highlighted the vulnerability of multimodal large language models to adversarial attacks, including jailbreaking attacks, which can bypass safety mechanisms and induce models to generate harmful or inappropriate content [3]. For instance, [4] demonstrated that vision-language models can be compromised using targeted adversarial examples, compromising their ability to provide accurate and reliable outputs. Similarly, [5] emphasized the importance of addressing backdoor threats in large language models, which can be exploited by adversaries to compromise model security.

The importance of understanding and mitigating adversarial attacks on multimodal large language models cannot be overstated. As these models become increasingly ubiquitous in various applications, including natural language processing, computer vision, and human-computer interaction, their vulnerability to attacks can have far-reaching consequences [6]. Furthermore, the complexity and multimodality of these models make them more challenging to defend against attacks, emphasizing the need for comprehensive and robust defense mechanisms [7].

Despite the growing concern about adversarial attacks, there is a lack of comprehensive surveys and analyses that provide a thorough understanding of the current state of research in this area. Most existing studies focus on specific aspects of adversarial attacks or defense mechanisms, without providing a broader perspective on the challenges and opportunities in this field [8]. To address this gap, this survey aims to provide a comprehensive overview of adversarial attacks on multimodal large language models, including their types, mechanisms, and evaluation metrics [9]. By synthesizing existing research and identifying emerging trends and challenges, this survey seeks to contribute to the development of more robust and secure multimodal large language models [10].

The remainder of this survey will delve into the specifics of adversarial attacks on multimodal large language models, including their classification, methodologies, and evaluation metrics [11]. We will also discuss the current state of defense mechanisms, including adversarial training, input preprocessing, and detection methods [12]. Furthermore, we will examine the applications and case studies of adversarial attacks on multimodal large language models, highlighting their impact on real-world systems and the need for more comprehensive testing and evaluation protocols [13]. By providing a thorough analysis of the current state of research in this area, this survey aims to facilitate the development of more robust and secure multimodal large language models, ultimately contributing to the advancement of artificial intelligence and its applications [14].

2 Background and Fundamentals of Multimodal Large Language Models

2.1 Introduction to Multimodal Large Language Models

The advent of multimodal large language models (MLLMs) has revolutionized the field of artificial intelligence, enabling machines to process and generate content across multiple modalities, such as text, images, audio, and video. As discussed in [15], MLLMs have demonstrated impressive

capabilities in various applications, including image captioning, visual question answering, and text-image generation. The architecture of MLLMs typically consists of a vision encoder, a language encoder, and a fusion module that integrates the outputs of the two encoders to produce a unified representation [1]. This unified representation enables MLLMs to capture complex relationships between different modalities and generate coherent and contextually relevant outputs.

Recent studies, such as [16] and [7], have explored the capabilities of MLLMs in various domains, including natural language processing, computer vision, and human-computer interaction. These studies have demonstrated the potential of MLLMs to improve the performance of various tasks, such as language translation, image classification, and dialogue generation. However, as noted in [17], MLLMs also pose significant security and privacy challenges, including the risk of adversarial attacks, data poisoning, and model inversion attacks.

The training of MLLMs typically involves a combination of masked language modeling, next-token prediction, and multimodal fusion techniques [18]. These techniques enable MLLMs to learn effective representations of different modalities and integrate them into a unified framework. However, as discussed in [11], the evaluation of MLLMs is a challenging task, requiring the development of new benchmarks and evaluation metrics that can capture the complexities of multimodal interactions.

Despite the challenges, MLLMs have the potential to revolutionize various applications, including education, healthcare, and entertainment [6]. As noted in [14], MLLMs can be used to improve the security and privacy of various systems, including chatbots, virtual assistants, and language translation systems. However, to realize the full potential of MLLMs, it is essential to address the security and privacy challenges associated with these models, including the development of robust defense mechanisms against adversarial attacks and data poisoning [19].

In conclusion, multimodal large language models have the potential to transform various applications, including natural language processing, computer vision, and human-computer interaction. However, to realize the full potential of these models, it is essential to address the security and privacy challenges associated with them, including the development of robust defense mechanisms against adversarial attacks and data poisoning [8]. Therefore, it is essential to stay up-to-date with the latest developments in this field and to contribute to the ongoing research efforts aimed at improving the security, privacy, and performance of MLLMs [20].

2.2 Architectural Overview of Multimodal Large Language Models

The architectural design of multimodal large language models (MLLMs) plays a crucial role in enabling these models to effectively process and integrate multiple forms of input, such as text, images, and audio. As discussed in the previous section, MLLMs have demonstrated impressive capabilities in various applications, including image captioning, visual question answering, and text-image generation, highlighting the importance of their architectural design. At the core of MLLMs are two primary components: vision encoders and language encoders, which work together to capture complex relationships between different modalities and generate coherent and contextually relevant outputs. The vision encoder is responsible for extracting relevant features from visual inputs, such as images or videos, while the language encoder processes textual inputs and generates human-like language outputs [21].

Recent studies have explored various architectures for integrating vision and language encoders, including the use of attention mechanisms and fusion techniques [22]. For example, the Macaw-LLM model uses a modality module to encode multi-modal data and a cognitive module to harness pre-trained large language models [22]. The choice of vision encoder is also critical in MLLMs, with popular options including convolutional neural networks (CNNs) and vision transformers (ViTs) [23]. In addition to the vision encoder, the language encoder is also a vital component of MLLMs, with many models using pre-trained large language models, such as BERT or RoBERTa, as their language encoder [16].

The integration of vision and language encoders is typically achieved through attention mechanisms or fusion techniques [24]. These techniques enable MLLMs to capture complex relationships between different modalities and generate coherent and contextually relevant outputs. By understanding the strengths and limitations of different architectures and components, researchers can develop more robust and efficient MLLMs that can perform a wide range of tasks with high accuracy [11]. This understanding is essential for improving the performance of MLLMs and addressing the security and privacy challenges associated with these models, which will be discussed in the following section. Future research directions may include the exploration of new architectures and techniques, such as the use of graph neural networks or multimodal attention mechanisms [15], which can further improve the performance and robustness of MLLMs.

2.3 Training Objectives and Methods for Multimodal Large Language Models

The training objectives and methods for multimodal large language models are crucial in enabling these models to effectively process and generate multimodal content, including text, images, and other forms of data. One of the primary training objectives for these models is masked language modeling, which involves masking certain tokens in the input sequence and predicting the masked tokens based on the context [21]. This approach has been extended to multimodal settings, where the model is trained to predict masked visual tokens in addition to textual tokens [15]. Another key training objective is next-token prediction, which involves predicting the next token in a sequence given the context [16].

In addition to these objectives, multimodal fusion techniques play a vital role in training multimodal large language models. These techniques involve combining the representations of different modalities, such as text and images, to create a unified representation that can be used for downstream tasks [22]. One popular approach for multimodal fusion is early fusion, which involves combining the representations of different modalities at an early stage of the model [25]. Another approach is late fusion, which involves combining the representations of different modalities at a later stage of the model [26].

Recent studies have also explored the use of intermediate fusion techniques, which involve combining the representations of different modalities at intermediate stages of the model [27]. These techniques have been shown to be effective in capturing the complex interactions between different modalities and improving the overall performance of the model [28]. Furthermore, the use of attention mechanisms has been found to be effective in multimodal fusion, as it allows the model to focus on the most relevant features of each modality when combining them [24].

The training methods for multimodal large language models also involve the use of large-scale datasets that contain a diverse range of multimodal content [29]. These datasets are typically constructed by combining existing datasets from different modalities and annotating them with labels that indicate the relationships between the different modalities [20]. The models are then trained on these datasets using a combination of supervised and self-supervised learning techniques [7].

In terms of evaluation metrics, multimodal large language models are typically evaluated using a combination of metrics that assess their performance on different modalities and tasks [30]. These metrics may include measures of accuracy, precision, recall, and F1-score, as well as metrics that assess the model’s ability to generate coherent and contextually relevant multimodal content [31]. Recent studies have also highlighted the importance of evaluating multimodal large language models using metrics that assess their robustness and generalizability to different datasets and tasks [32].

Overall, the training objectives and methods for multimodal large language models are critical in enabling these models to effectively process and generate multimodal content. By combining different modalities and using a range of training objectives and methods, these models can capture the complex interactions between different modalities and improve their overall performance [33]. As the field continues to evolve, it is likely that we will see the development

of new training objectives and methods that can further improve the performance of multimodal large language models [34].

2.4 Applications and Evaluation Metrics for Multimodal Large Language Models

Multimodal large language models have been successfully applied to a wide range of applications, including image captioning, visual question answering, text-image generation, and multimodal dialogue systems. As discussed in the previous section, the training objectives and methods for these models play a vital role in enabling them to effectively process and generate multimodal content. In image captioning, models are trained to generate descriptive captions for given images [21]. For instance, [35] proposes a method for vision-and-language pre-training, which achieves state-of-the-art performance on image captioning tasks. The success of these models in image captioning can be attributed to their ability to effectively integrate visual and textual information, which is a key aspect of multimodal large language models.

In addition to image captioning, multimodal large language models have also been applied to visual question answering, where models are required to answer questions based on the content of an image [15]. [22] presents a multi-modal language model that integrates image, audio, video, and text inputs, demonstrating its effectiveness in visual question answering tasks. Text-image generation is another application where models generate images based on textual descriptions [36]. [37] introduces a multi-modal model that combines the strengths of both discrete and continuous representations, achieving impressive results in text-to-image generation.

The evaluation of multimodal large language models is crucial to assess their performance and identify areas for improvement. Common evaluation metrics include BLEU and ROUGE scores for image captioning, accuracy and F1 scores for visual question answering, and Inception scores and Fréchet Inception distances for text-image generation [11]. [38] proposes a benchmark for evaluating multimodal large language models on multiple images, highlighting the need for more comprehensive evaluation protocols. Furthermore, multimodal dialogue systems involve generating human-like responses to user input, which can be in the form of text, images, or speech [39]. [40] proposes an adapter-based approach for efficient multimodal learning, demonstrating its potential in multimodal dialogue systems.

Recent studies have also focused on improving the efficiency and effectiveness of multimodal large language models. [41] proposes a data-efficient and compute-efficient multimodal large language model, achieving significant reductions in training time and computational costs. [42] introduces a mixture-of-depth adaptation approach for multimodal large language models, demonstrating its potential in improving model efficiency. Despite the significant progress made in multimodal large language models, there are still several challenges and limitations that need to be addressed, such as text dominance and cross-modal hallucinations [43]. [44] proposes a modality-adaptive decoding approach to mitigate cross-modal hallucinations in multimodal large language models.

As the field continues to evolve, it is essential to develop more comprehensive and robust multimodal large language models that can effectively integrate multiple modalities and generate human-like responses [45]. Future research directions include developing more efficient and effective evaluation metrics, improving model efficiency, and addressing the challenges and limitations of multimodal large language models [8]. The development of such models will have significant implications for a wide range of applications, including image captioning, visual question answering, text-image generation, and multimodal dialogue systems. Ultimately, the success of multimodal large language models will depend on the ability of the research community to collaborate, innovate, and address the complex challenges associated with these powerful and versatile AI systems.

2.5 Challenges and Future Directions for Multimodal Large Language Models

The development of multimodal large language models has achieved significant milestones in recent years, with models such as [46] and [47] demonstrating impressive capabilities in integrating visual and textual information. However, despite these advancements, several challenges and limitations persist, hindering the widespread adoption and reliability of these models in real-world applications. One of the primary concerns is the potential for biases in multimodal large language models, as highlighted in [48], which can lead to inconsistent performance across different datasets and tasks.

Another significant challenge is the need for more efficient and scalable training methods, as the current approaches often require substantial computational resources and large amounts of annotated data. [49] proposes a data-centric approach to address this issue, focusing on developing smaller but more effective models that can achieve comparable performance with less computational overhead. Furthermore, [50] emphasizes the importance of scaling models, data, and testing strategies to improve performance, highlighting the potential for significant advancements through careful model design and optimization.

The evaluation of multimodal large language models is another crucial aspect that requires attention, as existing benchmarks may not adequately capture the complexities of real-world scenarios. [11] provides a comprehensive overview of current benchmarks, emphasizing the need for more diverse and challenging evaluation protocols to assess the true capabilities of these models. In line with this, [13] introduces a novel benchmark focused on social media platforms, highlighting the importance of evaluating models in contexts that reflect real-world applications and user-generated content.

In addition to these challenges, ensuring the explainability and transparency of multimodal large language models is essential for building trust and understanding their decision-making processes. [10] offers a detailed analysis of current approaches to explainability, emphasizing the need for more sophisticated methods that can provide insights into the complex interactions between visual and textual components of these models.

Future directions for multimodal large language models include exploring new architectures and training methodologies that can better capture the nuances of human perception and cognition. [22] proposes a novel framework for integrating multiple modalities, highlighting the potential for more comprehensive and generalizable models. Moreover, [51] demonstrates the feasibility of developing models that can not only understand but also generate multimodal content, opening new avenues for applications in creative industries and human-computer interaction.

The development of more robust and efficient multimodal large language models will also depend on advances in areas such as computer vision, natural language processing, and human-computer interaction. Collaborations between researchers from these fields, as well as industry practitioners, will be crucial for addressing the complex challenges associated with multimodal AI. As noted in [8], the rapid evolution of multimodal large language models necessitates continuous review and synthesis of existing research to identify key trends, challenges, and future directions.

In conclusion, while multimodal large language models have made significant progress, there are several challenges and limitations that need to be addressed to ensure their reliable and efficient deployment in real-world applications. By acknowledging these challenges and pursuing innovative solutions, researchers can unlock the full potential of multimodal AI, leading to breakthroughs in areas such as human-computer interaction, healthcare, education, and entertainment. As [29] emphasizes, the development of comprehensive benchmarks and evaluation protocols will be essential for assessing the capabilities of these models and guiding future research. Ultimately, the success of multimodal large language models will depend on the ability of the research community to collaborate, innovate, and address the complex challenges associated with these powerful and versatile AI systems.

3 Types and Mechanisms of Adversarial Attacks

3.1 Classification of Adversarial Attacks

The classification of adversarial attacks on multimodal large language models is a crucial aspect of understanding and mitigating these threats. Adversarial attacks can be broadly categorized into two primary types: targeted and untargeted attacks [52]. Targeted attacks are designed to cause the model to produce a specific, incorrect output, and can be further divided into subcategories based on their goals, such as misclassification or data extraction [53]. Untargeted attacks, on the other hand, aim to cause the model to produce any incorrect output, without a specific goal in mind, and can be used to evaluate the model’s overall robustness [54].

Model evasion attacks are a type of targeted attack that involves manipulating the input to evade the model’s detection or classification capabilities [2]. These attacks can be used to bypass security mechanisms, such as content filters or spam detectors, and can have significant consequences in applications where security is critical [17]. Model extraction attacks, another type of targeted attack, involve extracting sensitive information from the model, such as its parameters or training data [5]. Data poisoning attacks, a type of untargeted attack, involve manipulating the training data to compromise the model’s performance or robustness [29].

The goals of adversarial attacks can vary widely, depending on the motivations of the attacker. Some attacks may be designed to compromise the model’s accuracy or robustness, while others may aim to extract sensitive information or bypass security mechanisms [14]. Understanding the goals and mechanisms of adversarial attacks is essential for developing effective defense mechanisms and evaluating the robustness of models against various types of attacks [8].

Recent studies have shown that multimodal large language models are vulnerable to a wide range of adversarial attacks, including those that exploit the model’s visual or textual inputs [55]. For example, [1] demonstrated that multimodal models can be compromised using adversarial images, while [3] showed that visual adversarial examples can be used to jailbreak aligned large language models. These findings highlight the need for robust defense mechanisms that can protect against a wide range of adversarial attacks [56].

In addition to understanding the types and goals of adversarial attacks, it is also essential to evaluate the effectiveness of defense mechanisms against these attacks [20]. This can involve assessing the model’s robustness against various types of attacks, as well as evaluating the performance of defense mechanisms, such as adversarial training or input preprocessing [12]. Recent studies have shown that defense mechanisms can be effective in mitigating the impact of adversarial attacks, but that there is still a need for more robust and effective defense strategies [7].

In conclusion, the classification of adversarial attacks on multimodal large language models is a critical aspect of understanding and mitigating these threats. By understanding the types and goals of adversarial attacks, as well as the mechanisms and effectiveness of defense mechanisms, researchers and practitioners can develop more robust and secure models that can protect against a wide range of adversarial attacks [57]. Future research should focus on developing more effective defense mechanisms, as well as evaluating the robustness of models against various types of attacks [9]. This will require a comprehensive understanding of the types and goals of adversarial attacks, as well as the mechanisms and effectiveness of defense mechanisms [11].

3.2 Methodologies for Generating Adversarial Examples

The generation of adversarial examples is a crucial step in understanding and evaluating the robustness of multimodal large language models, as it enables researchers to identify potential vulnerabilities and develop effective defense mechanisms. Building on the classification of adversarial attacks discussed earlier, the generation of adversarial examples can be used to assess the model’s robustness against various types of attacks, including targeted and untargeted

attacks. In the context of multimodal large language models, adversarial examples can be generated using a variety of methodologies, including gradient-based methods and gradient-free methods.

Gradient-based methods, such as the fast gradient sign method [19], are widely used for generating adversarial examples. These methods leverage the model’s gradients to optimize the adversarial example, often resulting in highly effective attacks that can be computationally efficient. For instance, the fast gradient sign method can be used to generate adversarial examples for image classification tasks by maximizing the loss function with respect to the input image. However, gradient-based methods can be limited by their reliance on model gradients, which may not always be accessible or reliable.

In contrast, gradient-free methods, such as genetic algorithms and evolutionary strategies [7], do not rely on gradient information and can be used when the model’s gradients are not accessible. These methods can be more computationally intensive than gradient-based methods but offer an alternative approach for generating adversarial examples. For example, genetic algorithms can be used to search for adversarial examples by iteratively evolving a population of candidate inputs. While gradient-free methods can be effective, they may require significant computational resources and may not always converge to optimal solutions.

Another approach for generating adversarial examples is to use hybrid methods that combine gradient-based and gradient-free methods. These methods can leverage the strengths of each approach to improve the success rate of adversarial examples. The choice of methodology depends on the specific application, the type of attack being considered, and the availability of model gradients.

The evaluation of adversarial examples is also an essential aspect of studying adversarial attacks on multimodal large language models. Evaluation metrics, such as attack success rate and model accuracy under attack, can be used to assess the effectiveness of adversarial examples [20]. These metrics can provide insights into the robustness of a model and the potential vulnerabilities that can be exploited by attackers. As discussed in the following section, the evaluation of adversarial attacks is crucial for understanding the model’s robustness and developing effective defense mechanisms.

In conclusion, the generation of adversarial examples is a critical component in the study of adversarial attacks on multimodal large language models. Various methodologies, including gradient-based and gradient-free methods, can be used to generate adversarial examples, each with their strengths and limitations. The evaluation of adversarial examples is also crucial, and different evaluation metrics can provide different insights into the robustness of a model. As the field of multimodal large language models continues to evolve, it is essential to continue to develop and evaluate new methodologies for generating adversarial examples and to improve our understanding of the vulnerabilities of these models. Future research directions may include the development of more robust models that can withstand adversarial attacks, as well as the exploration of new applications and domains where multimodal large language models can be used [11].

3.3 Evaluation Metrics for Adversarial Attacks

The evaluation of adversarial attacks on multimodal large language models is a crucial aspect of understanding their vulnerability and robustness. Various metrics have been proposed to assess the effectiveness of these attacks, including attack success rate, model accuracy under attack, and robustness metrics [7]. The attack success rate measures the percentage of successful attacks, providing a direct indication of the attack’s effectiveness and the model’s vulnerability. Model accuracy under attack, on the other hand, evaluates the model’s performance under adversarial conditions, helping to understand its robustness and ability to maintain accuracy in the face of attacks.

Robustness metrics, such as the robustness score or certification bounds, offer a more comprehensive view of the model’s ability to withstand adversarial perturbations, beyond just accuracy or attack success rate [11]. These metrics are essential for assessing the model’s reliability and security in various applications. Furthermore, [28] highlights the importance of evaluating the model’s performance under different attack scenarios, including targeted and untargeted attacks.

In addition to these metrics, [20] proposes the use of adversarial training to improve the model’s robustness. Adversarial training involves training the model on adversarial examples to enhance its ability to withstand attacks.

The choice of evaluation metric depends on the specific application and the type of attack being considered. For instance, [58] suggests that the attack success rate is a suitable metric for evaluating the effectiveness of attacks on text-rich image understanding tasks. On the other hand, [20] recommends using a combination of metrics, including model accuracy under attack and robustness metrics, to provide a comprehensive evaluation of the model’s robustness.

Recent studies have also explored the use of explainability techniques to understand the model’s behavior under adversarial attacks. For example, [59] proposes the use of saliency maps to identify the regions of the input that are most relevant to the model’s predictions.

In conclusion, the evaluation of adversarial attacks on multimodal large language models requires a comprehensive approach that considers multiple metrics and techniques. By selecting suitable evaluation metrics and using techniques such as adversarial training and explainability, researchers can gain a deeper understanding of the model’s robustness and vulnerability to adversarial attacks. Future research should focus on developing more effective evaluation metrics and techniques, as well as improving the robustness of multimodal large language models to adversarial attacks [8]. This will be crucial for ensuring the reliability and security of these models in various applications, including natural language processing, computer vision, and human-computer interaction.

3.4 Defense Mechanisms Against Adversarial Attacks

The development of effective defense mechanisms against adversarial attacks on multimodal large language models is a critical area of research, given the potential consequences of such attacks on model reliability and security. As [60] notes, multimodal large language models are susceptible to malicious attacks through visual inputs, highlighting the need for robust defense strategies. Building on the evaluation metrics and techniques discussed earlier, such as attack success rate, model accuracy under attack, and robustness metrics, defense mechanisms can be designed to enhance model robustness and reliability. One approach to mitigating these attacks is adversarial training, which involves training the model on adversarial examples to improve its robustness. This method has been shown to enhance model performance under attack, but it can be computationally expensive and may not provide comprehensive protection against all types of attacks.

In addition to adversarial training, other defense mechanisms, such as input preprocessing and detection methods, can be employed to identify and mitigate adversarial examples. Input preprocessing techniques, including data normalization, feature extraction, and dimensionality reduction, can reduce the impact of adversarial perturbations, but their effectiveness depends on the specific attack and model architecture. Detection methods, including statistical methods and machine learning-based approaches, can also be used to identify potential adversarial inputs, but may introduce additional computational overhead and require careful tuning of hyperparameters. Recent studies have explored the use of mixture-of-experts (MoE) architectures to improve the robustness of multimodal large language models [61]. By routing task-specific tokens through specialized experts, MoE models can reduce the risk of adversarial attacks while maintaining model performance.

The development of effective defense mechanisms against adversarial attacks on multimodal large language models is an ongoing challenge, and the evaluation of these mechanisms is crucial to ensuring their effectiveness. As [29] highlights, standardized benchmarks are needed to facilitate comparison and improvement of different approaches. Furthermore, the trade-off between model robustness and performance on benign inputs must be carefully considered, as overly defensive models may compromise their ability to perform well on legitimate tasks. Emerging trends in defense mechanisms against adversarial attacks include the use of explainable AI techniques to provide insights into model vulnerabilities [10] and the development of more efficient and scalable defense methods [8]. Additionally, research has focused on adapting defense mechanisms to specific application domains, such as autonomous driving [39], where the consequences of adversarial attacks can be particularly severe.

In conclusion, the development of effective defense mechanisms against adversarial attacks on multimodal large language models is a complex and multifaceted challenge. By synthesizing insights from various approaches, including adversarial training, input preprocessing, and detection methods, researchers can create more robust and reliable models. As the field continues to evolve, it is essential to prioritize the development of standardized evaluation benchmarks, efficient defense methods, and domain-specific adaptations to ensure the security and reliability of multimodal large language models in real-world applications, ultimately paving the way for the development of robust, secure, and reliable models that can effectively integrate and process multiple sources of information, as emphasized by [7]. This will enable the development of more robust and secure models, which is crucial for advancing the field and ensuring the reliability and security of multimodal AI systems, as highlighted in the subsequent section on emerging trends and future directions.

3.5 Emerging Trends and Future Directions

The field of adversarial attacks on multimodal large language models is rapidly evolving, with emerging trends and future directions focusing on the development of more robust and secure models, as well as the exploration of new applications and domains. As highlighted in [46], the integration of visual and textual information can enhance the performance of large language models, but also introduces new vulnerabilities to adversarial attacks. To address these challenges, researchers are exploring novel architectures and training methodologies, such as the use of multimodal fusion techniques and attention mechanisms, as discussed in [38].

One of the key emerging trends in this field is the development of more robust and secure multimodal large language models. As noted in [48], current models are not robust to image and text perturbations, and can be easily fooled by adversarial examples. To mitigate this, researchers are proposing new architectures and training methodologies, such as the use of adversarial training and input preprocessing techniques, as discussed in [62]. Additionally, the use of multimodal large language models for text-to-image generation and image-to-text generation is becoming increasingly popular, as highlighted in [36] and [63].

Another important direction in this field is the exploration of new applications and domains for multimodal large language models. As discussed in [58], these models have the potential to be applied to a wide range of tasks, including text-rich image understanding, visual question answering, and image captioning. Furthermore, the use of multimodal large language models for multimodal dialogue systems and human-computer interaction is becoming increasingly popular, as highlighted in [22] and [13].

The development of benchmarks and evaluation metrics for multimodal large language models is also an active area of research. As noted in [11], the lack of standardized benchmarks and evaluation metrics makes it challenging to compare the performance of different models and identify areas for improvement. To address this, researchers are proposing new benchmarks and evaluation metrics, such as the use of multimodal datasets and metrics that evaluate the performance of models on specific tasks, as discussed in [38] and [64].

In conclusion, the field of adversarial attacks on multimodal large language models is rapidly evolving, with emerging trends and future directions focusing on the development of more robust and secure models, as well as the exploration of new applications and domains. As highlighted in [8], the development of novel architectures, training methodologies, and evaluation metrics is crucial for advancing the field and ensuring the reliability and security of multimodal AI systems. By drawing on the insights and findings presented in [15], [7], and [20], researchers can develop more effective defense mechanisms and improve the overall performance of multimodal large language models. Ultimately, the future of multimodal AI systems depends on the ability to develop robust, secure, and reliable models that can effectively integrate and process multiple sources of information, as discussed in [51] and [65].

4 Defense Mechanisms and Robustness Evaluation

4.1 Introduction to Defense Mechanisms

The development of defense mechanisms against adversarial attacks on multimodal large language models is crucial for ensuring their reliability and security in various applications. As highlighted in [52], the vulnerability of these models to adversarial attacks can have significant consequences, including compromised safety and security. To address this issue, researchers have proposed various defense strategies, including adversarial training, input preprocessing, and detection methods. Adversarial training, as discussed in [12], involves training models on adversarial examples to improve their robustness. This approach has been shown to be effective in enhancing model robustness, but it can be computationally expensive and may not provide complete protection against all types of attacks.

Input preprocessing techniques, such as data normalization and feature extraction, can also be used to mitigate the impact of adversarial attacks. These methods, as described in [55], can help to reduce the sensitivity of models to adversarial perturbations. However, they may not be effective against all types of attacks and can potentially compromise model performance on benign inputs. Detection methods, including statistical methods and machine learning-based approaches, can be used to identify and flag potential adversarial inputs. These methods, as discussed in [66], can provide an additional layer of protection against adversarial attacks.

The effectiveness of defense mechanisms can be evaluated using various metrics, including attack success rate, model accuracy under attack, and robustness metrics. As highlighted in [11], the choice of evaluation metrics is critical in assessing the robustness of multimodal large language models. Researchers have also proposed various benchmarks and datasets to evaluate the robustness of these models, including [13] and [67]. These benchmarks and datasets provide a comprehensive framework for evaluating the robustness of multimodal large language models and identifying potential vulnerabilities.

Emerging trends in defense mechanisms include the use of multimodal fusion techniques and attention mechanisms to enhance model robustness. As discussed in [15], these techniques can help to improve model performance on multimodal tasks while reducing their vulnerability to adversarial attacks. However, they can also introduce new challenges, such as the need for careful tuning of hyperparameters and the potential for increased computational costs. Another emerging trend is the use of transfer learning and meta-learning to develop more robust and adaptable defense mechanisms. As highlighted in [7], these approaches can help to improve model performance on unseen tasks and datasets while reducing their vulnerability to adversarial attacks.

In conclusion, the development of effective defense mechanisms against adversarial attacks on multimodal large language models is a critical area of research. By understanding the strengths and limitations of different defense strategies and evaluating their effectiveness using robust benchmarks and datasets, researchers can develop more reliable and secure models for various

applications. As noted in [68], the future of multimodal large language models depends on their ability to withstand adversarial attacks and maintain their performance in the face of evolving threats. Therefore, continued research and development in this area are essential for ensuring the reliability and security of these models in real-world applications, as emphasized in [14] and [29].

4.2 Adversarial Training for Defense

Adversarial training has emerged as a prominent defense mechanism against adversarial attacks on multimodal large language models (MLLMs) [7]. As a crucial component of a comprehensive defense strategy, adversarial training involves training the model on adversarial examples to improve its robustness and resilience against targeted and untargeted attacks. This approach is particularly significant in the context of MLLMs, which are vulnerable to adversarial attacks due to their complexity and multimodal nature.

One of the primary advantages of adversarial training is its ability to improve model robustness without requiring significant modifications to the underlying architecture [16]. This makes it an attractive solution for defending against adversarial attacks, as it can be applied to existing models without incurring substantial computational costs or requiring extensive retraining [49]. Furthermore, adversarial training can be used in conjunction with other defense mechanisms, such as input preprocessing and detection methods, to provide a multi-layered defense against adversarial attacks. By combining these approaches, researchers can develop more robust and resilient MLLMs that can withstand a wide range of adversarial attacks.

However, adversarial training also has its limitations. One of the primary challenges associated with this approach is the increased computational cost of generating and incorporating adversarial examples into the training process [69]. This can be particularly problematic for large-scale MLLMs, which already require significant computational resources for training and deployment [70]. Additionally, adversarial training may not provide complete protection against all types of adversarial attacks, and the model may still be vulnerable to attacks that are tailored to exploit specific weaknesses in the defense mechanism [2]. Therefore, it is essential to carefully evaluate the effectiveness of adversarial training and identify potential areas for improvement.

Recent studies have explored various methodologies for generating adversarial examples and incorporating them into the training process. For example, [71] proposes a method for generating universal embeddings that can be used to improve the robustness of MLLMs against adversarial attacks. Similarly, [13] introduces a benchmark for evaluating the performance of MLLMs on social media platforms, which can be used to assess the effectiveness of adversarial training and other defense mechanisms. These advancements have contributed significantly to the development of more robust and resilient MLLMs, and they highlight the importance of continued research in this area.

Despite these advancements, there are still several open challenges and future directions in the area of adversarial training for defense. One of the primary challenges is the development of more efficient and effective methods for generating adversarial examples, which can be used to improve the robustness of MLLMs without incurring significant computational costs [67]. Additionally, there is a need for more comprehensive evaluation metrics and benchmarks, which can be used to assess the effectiveness of adversarial training and other defense mechanisms in a more systematic and thorough manner [11]. By addressing these challenges, researchers can continue to advance the state-of-the-art in adversarial training and defense, ultimately leading to the development of more secure and reliable MLLMs.

In conclusion, adversarial training has emerged as a promising defense mechanism against adversarial attacks on MLLMs. While it has its limitations and challenges, it offers a flexible and effective solution for improving model robustness and resilience. By exploring new methodologies for generating adversarial examples and incorporating them into the training process, researchers can continue to advance the state-of-the-art in adversarial training and defense. As the field

continues to evolve, it is likely that we will see the development of more sophisticated and effective defense mechanisms, which can be used to protect MLLMs against an increasingly wide range of adversarial attacks [10]. The development of these mechanisms will be crucial in ensuring the reliability and security of MLLMs in various applications, and it will require continued research and innovation in the area of adversarial training and defense.

4.3 Input Preprocessing and Detection Methods

Input preprocessing and detection methods have emerged as crucial complementary approaches to adversarial training in enhancing the robustness of multimodal large language models against adversarial attacks [22]. These methods focus on preprocessing the input data to reduce the impact of adversarial perturbations and detecting potential adversarial examples before they are fed into the model [26]. By doing so, they provide an additional layer of defense that can be used in conjunction with adversarial training to improve the overall security of multimodal large language models.

One key aspect of input preprocessing is data normalization, which involves scaling the input data to a common range to prevent features with large ranges from dominating the model [7]. This can help reduce the effectiveness of adversarial attacks that rely on manipulating the input data to cause the model to misbehave [30]. Another approach is feature extraction, which involves selecting a subset of the most relevant features from the input data to reduce the dimensionality of the input space [33]. This can make it more difficult for attackers to craft effective adversarial examples.

Detection methods, on the other hand, focus on identifying potential adversarial examples before they are fed into the model [72]. One common approach is to use statistical methods to detect anomalies in the input data that may indicate an adversarial attack [73]. Another approach is to use machine learning-based methods, such as training a separate model to detect adversarial examples [74]. These methods can be used in conjunction with input preprocessing techniques to provide a robust defense against adversarial attacks.

The effectiveness of input preprocessing and detection methods can be evaluated using a variety of metrics, including attack success rate, model accuracy under attack, and robustness metrics [11]. These metrics provide a comprehensive picture of the model’s performance under different attack scenarios and can be used to compare the effectiveness of different defense mechanisms [8]. By combining input preprocessing and detection methods with adversarial training, it is possible to develop robust multimodal large language models that can withstand a wide range of adversarial attacks [75].

In terms of future directions, there are several areas that require further research and development [76]. One key area is the development of more effective input preprocessing techniques that can reduce the impact of adversarial perturbations without compromising model performance [34]. Another area is the development of more robust detection methods that can accurately identify potential adversarial examples in real-time [77]. By addressing these challenges, it is possible to develop robust multimodal large language models that can be used in a wide range of applications, from natural language processing to computer vision [58].

Overall, input preprocessing and detection methods provide a crucial layer of defense against adversarial attacks on multimodal large language models [78]. By combining these methods with adversarial training, it is possible to develop robust models that can withstand a wide range of attacks and provide accurate and reliable performance in a variety of applications [79]. As the field continues to evolve, it is likely that we will see the development of even more effective defense mechanisms that can provide robust protection against adversarial attacks [32].

4.4 Evaluation Metrics for Defense Mechanisms

Evaluating the effectiveness of defense mechanisms against adversarial attacks on multimodal large language models is a critical aspect of ensuring the reliability and security of these models in various applications. As [7] highlights, the development of robust defense mechanisms is essential to mitigate the risks associated with adversarial attacks. To assess the effectiveness of these defense mechanisms, several evaluation metrics have been proposed, including attack success rate, model accuracy under attack, and robustness metrics such as the robust accuracy and the area under the receiver operating characteristic curve (AUROC) [11]. These metrics provide a comprehensive picture of the model’s performance under different attack scenarios and can be used to compare the effectiveness of different defense mechanisms.

One of the primary challenges in evaluating defense mechanisms is the lack of standardized evaluation protocols. Different studies often employ varying evaluation metrics and methodologies, making it difficult to compare the effectiveness of different defense mechanisms [20]. To address this issue, [38] proposes a comprehensive benchmark for evaluating the performance of multimodal large language models on multiple images. This benchmark provides a standardized evaluation protocol for assessing the effectiveness of defense mechanisms against adversarial attacks on multimodal large language models, which is crucial for developing robust and reliable models.

In addition to standardized evaluation protocols, it is also essential to consider both model performance on benign inputs and robustness against adversarial attacks when evaluating defense mechanisms [60]. A defense mechanism that improves robustness against adversarial attacks but compromises model performance on benign inputs may not be desirable in practice. Therefore, evaluation metrics that balance both aspects, such as the robust accuracy, are essential for assessing the effectiveness of defense mechanisms [80]. Furthermore, recent studies have also highlighted the importance of considering the modality imbalance issue in evaluating defense mechanisms [43].

To address the modality imbalance issue, [81] proposes a directed-token approach that captures visual and textual knowledge and enables the model to reconstruct the correct order of visual inputs. Moreover, [77] emphasizes the importance of visual representation alignment in evaluating defense mechanisms. The authors propose a regularization strategy that aligns the internal visual representations of multimodal large language models with those of pre-trained vision foundation models. This approach enables the model to retain critical visual details and complement additional visual knowledge from vision foundation models, thereby enhancing its ability to reason over complex visual inputs.

Overall, evaluating the effectiveness of defense mechanisms against adversarial attacks on multimodal large language models requires a comprehensive approach that considers multiple evaluation metrics and methodologies. By employing standardized evaluation protocols and considering both model performance on benign inputs and robustness against adversarial attacks, researchers can develop more effective defense mechanisms that ensure the reliability and security of multimodal large language models in various applications [30]. As [42] highlights, the development of robust defense mechanisms is an ongoing challenge that requires continued research and innovation in the field, ultimately leading to the development of more robust and reliable multimodal large language models that can withstand a wide range of adversarial attacks.

4.5 Emerging Trends and Future Directions

The development of defense mechanisms against adversarial attacks on multimodal large language models is an active area of research, with several emerging trends and future directions. One of the key trends is the integration of multimodal fusion techniques and attention mechanisms to enhance model robustness and adaptability to different types of adversarial attacks [38]. Another trend is the use of transfer learning and meta-learning to develop more generalizable

and adaptable defense mechanisms [22].

Recent studies have also explored the use of multimodal large language models for generating images with text prompts, such as [36]. This line of research has the potential to improve the robustness of models against adversarial attacks by generating more diverse and realistic images. Additionally, the development of benchmarks and evaluation metrics for multimodal large language models, such as [11], is crucial for assessing the effectiveness of defense mechanisms and identifying areas for improvement.

The use of multimodal large language models for fact-checking and misinformation detection is another emerging trend [82]. This application has significant implications for improving the reliability and trustworthiness of information on social media platforms and other online sources. Furthermore, the development of models that can effectively evaluate and mitigate the risks of hallucinations in multimodal large language models is an important area of research [83].

In terms of future directions, one potential area of research is the development of more robust and efficient defense mechanisms against adversarial attacks on multimodal large language models. This could involve exploring new architectures and training methods, such as [41], or leveraging techniques from other areas of machine learning, such as [84]. Another potential direction is the application of multimodal large language models to new domains and tasks, such as [85], which could lead to new insights and innovations in areas such as computer vision, natural language processing, and human-computer interaction.

Overall, the development of defense mechanisms against adversarial attacks on multimodal large language models is a complex and multifaceted challenge that requires a comprehensive and interdisciplinary approach. By exploring emerging trends and future directions, such as those outlined above, researchers and practitioners can work towards creating more robust, reliable, and trustworthy multimodal AI systems that can effectively mitigate the risks of adversarial attacks and improve the overall performance and safety of these models [8]. As the field continues to evolve, it is essential to prioritize the development of more effective defense mechanisms and to explore new applications and domains for multimodal large language models [64].

5 Applications and Case Studies of Adversarial Attacks

5.1 Introduction to Applications and Case Studies of Adversarial Attacks

The concept of adversarial attacks on multimodal large language models has garnered significant attention in recent years, with a growing body of research exploring the vulnerabilities of these models to various types of attacks [1]. Adversarial attacks refer to the process of crafting inputs that are designed to mislead or deceive a machine learning model, causing it to produce incorrect or undesirable outputs [5]. In the context of multimodal large language models, adversarial attacks can have significant consequences, including compromising the security and reliability of these models in real-world applications [6].

One of the primary applications of adversarial attacks on multimodal large language models is in the evaluation of their robustness and security [11]. By designing and implementing adversarial attacks, researchers can identify vulnerabilities in these models and develop more effective defense mechanisms [2]. For instance, [4] proposes a novel attack approach that leverages the visual and textual modalities to disrupt the model’s performance. Similarly, [86] introduces a flow-based multi-modal adversarial attack that can effectively induce video-based large language models to generate incorrect answers.

Case studies have also been conducted to demonstrate the effectiveness of adversarial attacks on multimodal large language models [87]. For example, [88] proposes a jailbreak attack that utilizes collaborative visual perturbation and textual steering to manipulate the model’s output. Additionally, [89] introduces a meticulous adversarial attack that can effectively disrupt the performance of vision-language pre-trained models.

The importance of evaluating the adversarial robustness of multimodal large language models cannot be overstated [90]. As these models are increasingly being deployed in real-world applications, including healthcare, education, and finance, it is crucial to ensure their security and reliability [14]. By developing and implementing effective defense mechanisms, researchers can mitigate the risks associated with adversarial attacks and ensure the safe and reliable operation of these models [56].

In conclusion, the study of adversarial attacks on multimodal large language models is a rapidly evolving field, with significant implications for the security and reliability of these models in real-world applications [8]. As researchers continue to develop and implement more effective defense mechanisms, it is essential to evaluate the adversarial robustness of these models and identify potential vulnerabilities [20]. By doing so, we can ensure the safe and reliable operation of multimodal large language models and mitigate the risks associated with adversarial attacks [91]. Future research directions include the development of more effective defense mechanisms, the evaluation of adversarial robustness in real-world applications, and the exploration of new attack approaches and methodologies [92].

5.2 Adversarial Attacks on Multimodal Models in Real-World Scenarios

The application of adversarial attacks on multimodal large language models in real-world scenarios has significant implications for various industries, including autonomous vehicles, healthcare, and education. As discussed in the previous section, the study of adversarial attacks on multimodal large language models is a rapidly evolving field, with significant implications for the security and reliability of these models in real-world applications [8]. In autonomous vehicles, for instance, adversarial attacks can compromise the perception and decision-making capabilities of multimodal models, leading to potential safety risks [39]. Similarly, in healthcare, adversarial attacks on medical imaging models can result in misdiagnosis or inappropriate treatment recommendations [87].

The potential impact of adversarial attacks on these systems is substantial, and it is essential to develop effective defense mechanisms to mitigate these risks. Recent studies have explored the use of adversarial training and input preprocessing techniques to improve the robustness of multimodal models. However, these approaches have limitations, and there is a need for more comprehensive and effective defense strategies. To address this challenge, researchers have proposed various approaches, including the development of multimodal models that can adapt to different scenarios and environments. For example, the use of mixture-of-experts (MoE) architectures has been proposed as a way to improve the robustness of multimodal models [27].

In addition to developing more robust models, the evaluation of multimodal models' robustness to adversarial attacks is also crucial. Recent benchmarks, such as MIBench [38], have been proposed to assess the performance of multimodal models in various scenarios. These benchmarks provide a comprehensive evaluation of multimodal models' abilities and can help identify areas for improvement. Furthermore, the development of more advanced attack detection methods, such as those using machine learning-based approaches, can help identify and mitigate adversarial attacks [67].

As the use of multimodal large language models in real-world scenarios becomes more widespread, it is essential to consider the ethical implications of these models and develop guidelines for their responsible development and deployment [10]. The potential consequences of adversarial attacks on these models can have significant consequences for human safety and well-being, particularly in applications such as autonomous vehicles or healthcare [93]. Therefore, it is crucial to prioritize the development of effective defense mechanisms and evaluation metrics to ensure the reliable and secure deployment of multimodal models in real-world applications [6].

In terms of future directions, there is a need for more research on developing effective defense mechanisms against adversarial attacks on multimodal models. This can involve exploring new architectures, such as those using attention mechanisms or graph-based approaches [30].

Additionally, the development of more advanced evaluation metrics and benchmarks can help assess the robustness of multimodal models and identify areas for improvement [11]. By exploring new architectures, developing more advanced attack detection methods, and evaluating multimodal models' robustness, we can improve the security and reliability of these models and enable their widespread adoption in various industries [18]. Ultimately, the application of adversarial attacks on multimodal large language models in real-world scenarios highlights the need for ongoing research and development in this area, to ensure the reliable and secure deployment of these models in various industries [94].

5.3 Case Studies of Adversarial Attacks on Multimodal Large Language Models

The realm of multimodal large language models has witnessed significant advancements, with these models demonstrating impressive capabilities in understanding and generating content across multiple modalities. However, the security and robustness of these models against adversarial attacks remain a pressing concern. This subsection delves into case studies of successful adversarial attacks on multimodal large language models, focusing on attacks targeting image captioning, visual question answering, and multimodal sentiment analysis.

Recent studies have highlighted the vulnerability of multimodal large language models to adversarial attacks. For instance, an attacker can cause the model to generate incorrect captions by manipulating the input image. Similarly, visual question answering models can be compromised using adversarial questions. These attacks not only undermine the reliability of multimodal large language models but also raise concerns about their potential applications in real-world scenarios.

Another significant concern is the attack on multimodal sentiment analysis models. An attacker can craft adversarial text and images to manipulate the model's sentiment analysis output. This vulnerability has far-reaching implications, particularly in applications where sentiment analysis is critical, such as customer service chatbots or social media monitoring tools.

The [22] study further emphasized the importance of evaluating the robustness of multimodal large language models against adversarial attacks, highlighting the need for more comprehensive testing and evaluation protocols.

In addition, the [8] study provided a comprehensive overview of the current state of multimodal large language models, highlighting the need for more research on the security and robustness of these models against adversarial attacks.

In conclusion, the case studies presented in this subsection highlight the significance of evaluating the robustness of multimodal large language models against adversarial attacks. As these models become increasingly ubiquitous in real-world applications, ensuring their security and reliability is crucial. Future research directions should focus on developing more comprehensive evaluation protocols, improving the transparency and interpretability of multimodal large language models, and designing more robust and secure architectures that can withstand adversarial attacks. By addressing these challenges, we can unlock the full potential of multimodal large language models and ensure their safe and reliable deployment in a wide range of applications.

5.4 Lessons Learned and Future Directions

The study of adversarial attacks on multimodal large language models has yielded valuable insights into the vulnerabilities and limitations of these models, building on the concerns raised in the previous discussion on the application of adversarial attacks in real-world scenarios. Through various case studies and applications, researchers have demonstrated the potential consequences of successful adversarial attacks, including compromised model performance, security breaches, and misuse of sensitive information [7]. These findings highlight the need for robust defense

mechanisms and countermeasures to mitigate the risks associated with adversarial attacks, which will be further discussed in the subsequent section on developing effective defense mechanisms and countermeasures.

One key lesson learned from the analysis of adversarial attacks is the importance of evaluating model robustness and security [11]. Current evaluation metrics and benchmarks often focus on model performance under normal operating conditions, neglecting the potential vulnerabilities that can be exploited by adversaries [38]. To address this limitation, researchers have proposed new evaluation metrics and benchmarks that specifically assess model robustness and security, such as the Multimodal Needle-in-a-haystack (MMNeedle) benchmark [95]. This emphasis on robustness evaluation is crucial for identifying the strengths and weaknesses of multimodal large language models and for developing effective defense strategies.

Another important finding is the need for more effective defense mechanisms against adversarial attacks. Current approaches, such as adversarial training and input preprocessing, have shown promise in improving model robustness, but they often come with significant computational costs and may not be effective against all types of attacks [60]. To address this challenge, researchers have proposed novel defense mechanisms, such as the use of mixture-of-experts architectures [27] and token-level supervision alignment methods [96]. These innovative approaches can help improve the robustness of multimodal large language models and mitigate the risks associated with adversarial attacks.

The analysis of adversarial attacks has also highlighted the importance of considering the intersection of multiple modalities, such as vision, language, and audio [85]. Multimodal models are particularly vulnerable to attacks that exploit the interactions between different modalities, and defense mechanisms must be designed to account for these complexities [81]. Furthermore, the study of adversarial attacks has emphasized the need for more transparent and explainable models, which can provide insights into the decision-making processes and vulnerabilities of multimodal models [10]. This increased transparency can facilitate the development of more effective defense mechanisms and improve the overall robustness of multimodal large language models.

In terms of future directions, researchers have identified several areas of focus for improving the security and robustness of multimodal large language models. These include the development of more effective defense mechanisms, such as those based on game-theoretic approaches [97] and the use of multimodal fusion techniques [98]. Additionally, there is a need for more comprehensive evaluation metrics and benchmarks that can assess model robustness and security in a more holistic manner [20]. The development of these metrics and benchmarks will be crucial for evaluating the effectiveness of defense mechanisms and identifying areas for improvement, as will be discussed in the following section on the evaluation of adversarial attacks on multimodal large language models.

Overall, the analysis of adversarial attacks on multimodal large language models has provided valuable insights into the challenges and opportunities associated with these models. By leveraging these findings, researchers can develop more robust and secure models that can effectively mitigate the risks associated with adversarial attacks and provide reliable performance in a wide range of applications [37]. As the field continues to evolve, it is essential to prioritize the development of more effective defense mechanisms, comprehensive evaluation metrics, and transparent models that can provide insights into the decision-making processes and vulnerabilities of multimodal models [71].

5.5 Evaluation Metrics and Benchmarking for Adversarial Attacks

The evaluation of adversarial attacks on multimodal large language models is a crucial aspect of ensuring the reliability and security of these models in various applications. To assess the effectiveness of these attacks, several evaluation metrics and benchmarking methods have been proposed [11]. These metrics can be broadly categorized into two types: those that measure the

success rate of the attack and those that evaluate the robustness of the model against the attack [99]. The attack success rate is typically measured using metrics such as the attack success rate, which calculates the percentage of successful attacks [38]. On the other hand, the robustness of the model is often evaluated using metrics such as the model’s accuracy under attack or the robustness score [100].

In addition to these metrics, several benchmarking methods have been proposed to evaluate the effectiveness of adversarial attacks on multimodal large language models [38]. These benchmarks typically consist of a set of tasks and datasets that are designed to test the model’s robustness against various types of attacks [13]. For example, the MMBench benchmark evaluates the model’s ability to withstand attacks on multiple images, while the MM-Soc benchmark assesses the model’s robustness in social media platforms [67].

Recent studies have also highlighted the importance of standardized evaluation protocols for comparing and improving defense mechanisms against adversarial attacks [8]. Standardized protocols can help ensure that the evaluation of adversarial attacks is fair and consistent, and can facilitate the comparison of different defense mechanisms [20]. Furthermore, the use of standardized protocols can also help identify the most effective defense mechanisms and provide insights into the strengths and limitations of different approaches [24].

The development of effective evaluation metrics and benchmarking methods is an ongoing area of research, with new metrics and benchmarks being proposed regularly [101]. For example, the LENS benchmark evaluates the model’s ability to reason about complex and real-world scenarios, while the MMR benchmark assesses the model’s ability to understand text-rich images [58]. These new metrics and benchmarks can provide valuable insights into the strengths and limitations of different models and defense mechanisms, and can help drive the development of more effective and robust multimodal large language models [102].

In conclusion, the evaluation of adversarial attacks on multimodal large language models is a critical aspect of ensuring the reliability and security of these models. The use of standardized evaluation metrics and benchmarking methods can help facilitate the comparison of different defense mechanisms and provide insights into the strengths and limitations of different approaches [103]. As the field continues to evolve, it is likely that new metrics and benchmarks will be proposed, and it is essential to stay up-to-date with the latest developments in this area [104]. By doing so, researchers and practitioners can work together to develop more effective and robust multimodal large language models that can withstand the threats posed by adversarial attacks [105].

5.6 Defense Mechanisms and Countermeasures

The development of effective defense mechanisms and countermeasures is crucial for ensuring the security and robustness of multimodal large language models against adversarial attacks. As highlighted in the previous discussion on evaluation metrics and benchmarking methods, assessing the effectiveness of adversarial attacks is essential, but it is equally important to develop strategies that can mitigate the impact of these attacks [11]. One approach to achieving this is through adversarial training, which involves training the model on adversarial examples to improve its robustness [23]. Adversarial training has been shown to be effective in improving the robustness of multimodal large language models, but it can be computationally expensive and may not always be feasible.

In addition to adversarial training, input preprocessing techniques can help to detect and mitigate adversarial attacks [15]. Input preprocessing techniques, such as data normalization and feature extraction, can help to reduce the impact of adversarial perturbations on the model’s performance. Furthermore, detection methods, such as anomaly detection and adversarial example detection, can be used to identify potential adversarial attacks [20]. These methods can be used to alert the system to potential attacks, allowing for timely countermeasures to be taken. Recent studies, such as [7], have also explored the use of multimodal fusion techniques to improve

the robustness of multimodal large language models. These techniques involve combining the outputs of multiple models to produce a more robust and accurate output.

Other approaches, such as [30], have investigated the use of ensemble methods to improve the robustness of multimodal large language models. Moreover, understanding the visual perception capabilities of large language models, as highlighted in [106], can be used to improve the robustness of multimodal large language models. The study of [80] also provides insights into the cross-modal alignment of large vision-language models, which can be used to develop more effective defense mechanisms. Additionally, addressing unimodal biases in multimodal large language models, as emphasized in [107], is crucial for developing effective defense mechanisms. The study of [44] also proposes a modality-adaptive decoding approach to mitigate cross-modal hallucinations in multimodal large language models.

Despite these advances, there are still significant challenges to be addressed in developing effective defense mechanisms and countermeasures for multimodal large language models. As [8] notes, the field of multimodal large language models is rapidly evolving, and new challenges and opportunities are emerging continuously. Therefore, it is essential to continue researching and developing new defense mechanisms and countermeasures to ensure the security and robustness of these models. This will ultimately lead to the development of more secure and reliable multimodal large language models that can be deployed in a wide range of applications, which is critical for ensuring the reliable and secure deployment of these models. The development of effective defense mechanisms and countermeasures will be crucial for addressing the emerging challenges and opportunities in this field, as highlighted by studies such as [108] and [109].

6 Real-World Implications and Future Directions

6.1 Introduction to Real-World Implications

The real-world implications of adversarial attacks on multimodal large language models are far-reaching and multifaceted, with significant consequences for privacy, security, and trust in AI systems. As [1] highlights, the vulnerability of these models to adversarial attacks poses a substantial threat to their reliability and safety. The integration of multiple modalities, such as vision and language, expands the attack surface, making it easier for adversaries to manipulate the model’s behavior. This is particularly concerning in applications where multimodal large language models are used to make critical decisions, such as in healthcare, finance, or education.

One of the primary concerns is the potential for adversarial attacks to compromise the privacy and security of sensitive information. For instance, [3] demonstrates how visual adversarial examples can be used to jailbreak aligned large language models, allowing attackers to access sensitive information or manipulate the model’s behavior. Similarly, [2] shows how multimodal large language models can be vulnerable to jailbreaking attacks, which can lead to the exposure of sensitive information or the manipulation of the model’s output.

The implications of these attacks are not limited to privacy and security; they also have significant consequences for trust in AI systems. As [8] notes, the reliability and safety of multimodal large language models are crucial for maintaining trust in AI systems. If these models are vulnerable to adversarial attacks, it can erode trust in their ability to make accurate and reliable decisions. This, in turn, can have significant consequences for the adoption and deployment of AI systems in critical applications.

Furthermore, the complexity of multimodal large language models makes it challenging to develop effective defense mechanisms against adversarial attacks. As [56] highlights, the development of robust defense mechanisms requires a deep understanding of the model’s architecture, the attack mechanisms, and the potential vulnerabilities.

In addition to the technical challenges, there are also significant societal implications of adversarial attacks on multimodal large language models. As [53] notes, the use of large language

models for illicit purposes, such as generating malicious content or spreading disinformation, can have significant consequences for individuals and society.

In conclusion, the real-world implications of adversarial attacks on multimodal large language models are significant and far-reaching. The vulnerability of these models to adversarial attacks poses a substantial threat to their reliability, safety, and trustworthiness. The development of effective defense mechanisms against these attacks is crucial for maintaining trust in AI systems and preventing the misuse of these models. As [11] highlights, the evaluation of multimodal large language models against adversarial attacks is essential for understanding their limitations and developing effective defense mechanisms.

6.2 Emerging Trends and Research Opportunities

The realm of adversarial attacks on multimodal large language models is rapidly evolving, with emerging trends and research opportunities presenting both challenges and potential breakthroughs. Recent studies, such as [15], have underscored the importance of understanding and mitigating these attacks to ensure the reliability and security of multimodal AI systems. Building on the previous discussion of the real-world implications of adversarial attacks, it is clear that the development of more robust and secure models is a key area of focus. This involves not only enhancing model architecture but also improving training methodologies to withstand sophisticated attacks, as demonstrated in [19].

One promising approach to achieving this goal is the integration of multimodal fusion techniques and attention mechanisms to enhance model robustness and adaptability to different types of adversarial attacks, as explored in [24]. Additionally, leveraging large language models as a foundation for multimodal models, such as in [26], offers a pathway to improving multimodal robustness by capitalizing on the strengths of language models. The incorporation of visual experts, as proposed in [110], can also contribute to more accurate and robust multimodal understanding, addressing some of the challenges highlighted in the preceding section.

Efficient and scalable training and deployment of multimodal large language models are also critical research directions. Techniques such as those presented in [49] and [41] aim to reduce the computational costs associated with multimodal model training, making these models more accessible for widespread adoption. Furthermore, the exploration of state space models, as in [111], offers an innovative approach to improving the efficiency and scalability of multimodal models, which is essential for addressing the open challenges and gaps in the field, as will be discussed in the following section.

The evaluation and benchmarking of multimodal large language models against adversarial attacks are also areas of active research, providing valuable insights into model robustness and identifying areas for improvement. Benchmarks like [38] and [13] are crucial for assessing model performance under adversarial conditions. Moreover, studies on the explainability and interpretability of multimodal models, such as [10], are essential for understanding how these models make decisions and for developing more trustworthy AI systems, which is a pressing concern in the context of real-world applications.

In real-world scenarios, the vulnerability of multimodal large language models to adversarial attacks poses significant risks, particularly in critical domains such as autonomous vehicles and healthcare, where model reliability is paramount. Research into [18] and [87] highlights the need for secure and robust models that can withstand adversarial manipulations. The development of models like [67] and [112] demonstrates the potential for creating more resilient and adaptable multimodal AI systems, which is aligned with the future research directions outlined in the subsequent section.

Future research directions include the exploration of novel architectures and training methodologies that can enhance the security and robustness of multimodal large language models. The integration of techniques from other fields, such as [94], can provide new insights into improving multimodal model performance under adversarial conditions. Additionally, the

development of more comprehensive and realistic benchmarks, as well as the investigation of the ethical and societal implications of adversarial attacks on multimodal AI, are essential for advancing the field towards more robust, secure, and trustworthy multimodal AI systems.

In conclusion, the emerging trends and research opportunities in adversarial attacks on multimodal large language models present a complex and dynamic landscape. Through continued innovation and collaboration, the development of more robust, secure, and trustworthy multimodal AI systems is not only feasible but also imperative for the widespread adoption and beneficial impact of these technologies in various sectors. As research in this area progresses, it is crucial to address the challenges posed by adversarial attacks, ensuring that multimodal large language models can reliably and securely process and generate information across different modalities, ultimately contributing to the advancement of artificial intelligence towards more human-like capabilities. Studies like [113] and [114] underscore the potential for multimodal models to achieve state-of-the-art performance across a variety of tasks, further emphasizing the need for robustness against adversarial attacks to safeguard their reliability and trustworthiness.

6.3 Open Challenges and Potential Gaps

The realm of adversarial attacks on multimodal large language models presents a complex and dynamic landscape, replete with challenges and gaps that necessitate urgent attention from the research community. As [15] underscores, the integration of multiple modalities, such as text, images, and audio, into a single framework, while offering unparalleled capabilities, also introduces novel vulnerabilities. One of the primary open challenges is the development of effective defense mechanisms that can comprehensively mitigate adversarial attacks without compromising the model’s performance on benign inputs.

Another significant gap in current research pertains to the evaluation metrics used to assess the effectiveness of adversarial attacks and the robustness of models. While metrics such as attack success rate and model accuracy under attack are commonly employed, as discussed in [11] and [20], there is a need for more comprehensive and standardized evaluation protocols that can accurately capture the nuances of multimodal interactions.

The emergence of new applications and domains for multimodal large language models also poses challenges, as these models are increasingly being applied in high-stakes areas such as healthcare, education, and autonomous vehicles. [7] and [85] highlight the potential of these models in various tasks, but also underscore the need for careful consideration of their limitations and potential biases.

Furthermore, the issue of explainability and interpretability in multimodal large language models is a pressing concern. As models become more complex and integrate multiple modalities, understanding how they arrive at their decisions and predictions becomes increasingly challenging. [10] discusses the importance of developing methods and tools that can provide insights into the decision-making processes of these models, which is essential for building trust and ensuring accountability in their applications.

In conclusion, the field of adversarial attacks on multimodal large language models is characterized by a plethora of open challenges and potential gaps, ranging from the development of effective defense mechanisms and better evaluation metrics to the need for more explainable and interpretable models. Addressing these challenges will require sustained research efforts, innovative solutions, and a commitment to advancing the field in a responsible and ethical manner. As [8] suggests, the future of multimodal large language models depends on our ability to navigate these complexities and create models that are not only powerful but also reliable, fair, and transparent. By pursuing this agenda, we can unlock the full potential of multimodal large language models and harness their capabilities to drive progress in a wide range of applications and domains.

6.4 Future Research Directions

The development of robust and secure multimodal large language models (MLLMs) is a crucial aspect of ensuring their reliability and security in various applications, as highlighted in the previous discussion on the challenges and gaps in the field of adversarial attacks on MLLMs. One potential direction for achieving this goal is the use of adversarial training, which involves training the model on adversarial examples to improve its robustness [35]. This approach has shown promise in enhancing the model’s ability to withstand adversarial attacks, and its effectiveness can be further evaluated using metrics and benchmarks such as those discussed in [20].

Recent studies have also emphasized the importance of evaluating the effectiveness of defense mechanisms against adversarial attacks on MLLMs, and have introduced new benchmarks and evaluation methods to assess model robustness and performance under attack, such as [38]. These efforts aim to identify potential vulnerabilities in MLLMs and provide a comprehensive understanding of their strengths and weaknesses. Furthermore, researchers are exploring new architectures and training methods for MLLMs, including the use of mixture-of-experts (MoE) models [27] and sparse attention mechanisms [42], which can potentially improve the model’s robustness and security.

In addition to these approaches, the integration of visual and language modalities in MLLMs can be improved through the use of techniques such as visual embedding alignment [80] and token-level visual-textual integration [96]. The development of more robust and secure MLLMs also relies on the creation of large-scale datasets and benchmarks that can effectively evaluate the performance of these models, as highlighted in [7] and [11]. These datasets and benchmarks play a crucial role in identifying areas for improvement and guiding the development of more effective and reliable MLLMs.

Ultimately, the future of MLLMs depends on the ability to address the challenges posed by adversarial attacks and to develop more robust and secure models, which can be achieved through the use of techniques such as adversarial training, multimodal fusion, and sparse attention mechanisms, as well as the creation of large-scale datasets and benchmarks. By exploring these directions, researchers can develop more effective and reliable MLLMs that can be used in a variety of real-world applications, from image captioning and visual question answering to text generation and dialogue systems [37]. As noted in [85], the development of MLLMs that can handle multiple modalities, including images, videos, and audio, is crucial for creating more generalizable and applicable models, and will be essential for mitigating adversarial attacks, as discussed in the following section.

6.5 Mitigation Strategies and Best Practices

Mitigating adversarial attacks on multimodal large language models is crucial for ensuring the reliability and security of these models in various applications. One effective mitigation strategy is adversarial training, which involves training the model on adversarial examples to improve its robustness [46]. Adversarial training can be performed using various methods, including fast gradient sign method and projected gradient descent [36]. Another approach is input validation, which involves validating the input data to detect and prevent adversarial attacks [38]. Input validation can be performed using various techniques, such as data normalization and feature extraction [115].

In addition to adversarial training and input validation, other mitigation strategies include detection methods, such as anomaly detection and adversarial example detection [48]. These methods can be used to detect and prevent adversarial attacks, and can be integrated into the model’s training pipeline [15]. Furthermore, techniques such as transfer learning and meta-learning can be used to improve the model’s robustness and adaptability to different types of adversarial attacks [7].

Best practices for mitigating adversarial attacks on multimodal large language models include

using a combination of mitigation strategies, such as adversarial training, input validation, and detection methods [11]. It is also important to regularly update and refine the model’s training data and parameters to ensure that it remains robust to evolving adversarial attacks [24]. Moreover, using techniques such as ensemble methods and uncertainty estimation can help to improve the model’s robustness and reliability [10].

Recent studies have shown that multimodal large language models can be vulnerable to adversarial attacks, and that mitigation strategies such as adversarial training and input validation can be effective in improving the model’s robustness [83]. However, more research is needed to develop more effective and efficient mitigation strategies, and to improve the model’s ability to detect and prevent adversarial attacks [13]. Additionally, the development of more robust and secure multimodal large language models requires a deeper understanding of the underlying mechanisms of adversarial attacks, and the development of more effective evaluation metrics and benchmarks [67].

In conclusion, mitigating adversarial attacks on multimodal large language models requires a comprehensive approach that includes a combination of mitigation strategies, such as adversarial training, input validation, and detection methods. By using these strategies and following best practices, such as regularly updating and refining the model’s training data and parameters, it is possible to improve the model’s robustness and reliability, and to ensure the security and integrity of multimodal AI systems [8]. Future research should focus on developing more effective and efficient mitigation strategies, and on improving the model’s ability to detect and prevent adversarial attacks [20].

6.6 Conclusion and Recommendations

The discussion on mitigating adversarial attacks on multimodal large language models has highlighted the importance of developing robust defense mechanisms and comprehensive evaluation metrics. As the previous section has emphasized, the development of effective mitigation strategies, such as adversarial training, input validation, and detection methods, is crucial for ensuring the reliability and security of multimodal large language models [11]. Furthermore, the need for systematic reviews and meta-reviews, as demonstrated by [8], underscores the importance of identifying emerging trends and challenges in the field.

In light of these findings, several recommendations can be made for future research and practice. Firstly, there is a need for more comprehensive and diverse benchmarks that can effectively evaluate the capabilities and limitations of multimodal large language models. As suggested by [13], benchmarks should be designed to assess the performance of models in real-world scenarios, including social media platforms. This will enable researchers to better understand the vulnerabilities of multimodal large language models to adversarial attacks, as demonstrated by [86], and develop more robust and secure models.

Moreover, the development of more robust defense mechanisms is crucial for mitigating the effects of adversarial attacks. [56] demonstrates the potential of using robust image encoders to improve the security of multimodal large language models. Additionally, [94] highlights the importance of developing more effective integration strategies for combining different modalities. The use of modality-adaptive decoding, as proposed by [44], can help to reduce cross-modal hallucinations and improve the overall performance of multimodal large language models.

To further enhance the robustness and security of multimodal large language models, it is essential to address unimodal biases, as highlighted by [107]. The use of causal frameworks, as proposed by this study, can help to identify and mitigate these biases, leading to more fair and robust models. Furthermore, [41] demonstrates the potential of using data-efficient and compute-efficient models to reduce the computational costs associated with training and deploying multimodal large language models. By drawing on the insights and recommendations provided by [20], [116], and [30], researchers and practitioners can work towards creating more effective and reliable multimodal large language models that can withstand the challenges of

adversarial attacks. Ultimately, the development of robust and secure multimodal large language models will require a comprehensive approach that addresses the challenges of benchmarking, integration strategies, unimodal biases, and adversarial attacks.

7 Conclusion

The conclusion of this comprehensive survey on adversarial attacks on multimodal large language models underscores the critical importance of understanding and mitigating these vulnerabilities. As highlighted in [52], the susceptibility of these models to adversarial perturbations poses significant risks to their reliability and security. The [117] study further emphasizes the need for robust defense mechanisms, given the potential for data poisoning and model evasion attacks.

Recent advancements in multimodal large language models, such as those discussed in [15], have achieved remarkable performance in various applications. However, these models' vulnerability to adversarial attacks, as demonstrated in [118], necessitates a thorough examination of their security implications. The [119] paper provides valuable insights into the language and position-specific vulnerabilities of these models, highlighting the need for tailored defense strategies.

The development of effective defense mechanisms against adversarial attacks is an active area of research, with studies like [12] and [53] proposing various approaches to mitigate these threats. The [5] paper provides a comprehensive overview of backdoor threats and defense strategies, emphasizing the importance of evaluating the effectiveness of these defenses.

The [66] study introduces a benchmark for evaluating the robustness of multimodal large language models against face spoofing and forgery detection attacks. This work highlights the need for standardized evaluation protocols to assess the security of these models. The [10] paper provides a thorough review of explainability and interpretability methods for multimodal large language models, emphasizing the importance of understanding model decisions and behaviors.

The [6] study explores the applications of large language models in cybersecurity, highlighting their potential to enhance security protocols and mitigate threats. However, the [14] paper emphasizes the unique security and privacy threats associated with large language models, necessitating careful consideration of these risks in their development and deployment.

In conclusion, the study of adversarial attacks on multimodal large language models is a rapidly evolving field, with significant implications for the security and reliability of AI systems. As highlighted in [8], a comprehensive understanding of the current state of research in this area is essential for advancing the field. The [29] paper provides a thorough review of multimodal benchmarks, emphasizing the need for standardized evaluation protocols to assess the performance and security of multimodal large language models.

Future research directions in this area include the development of more robust and effective defense mechanisms, as well as the exploration of new applications and domains for multimodal large language models. The [7] study highlights the potential of model composition for creating versatile multimodal models, while the [18] paper emphasizes the importance of evaluating the fact-checking capabilities of these models.

Ultimately, the security and reliability of multimodal large language models depend on the development of effective defense mechanisms and the careful consideration of potential vulnerabilities. As the field continues to evolve, it is essential to prioritize research in this area, ensuring that these models can be deployed safely and reliably in a wide range of applications. The [2] study and the [120] paper provide valuable insights into the vulnerabilities of these models, highlighting the need for ongoing research and development in this critical area.

References

- [1] *On the Adversarial Robustness of Multi-Modal Foundation Models*. DOI:10.1109/iccvw60793.2023.00395, 2023. <https://doi.org/10.1109/iccvw60793.2023.00395>
- [2] *Jailbreaking Attack against Multimodal Large Language Model*. arXiv:2402.02309, 2024. <https://arxiv.org/abs/2402.02309>
- [3] *Visual Adversarial Examples Jailbreak Aligned Large Language Models*. DOI:10.1609/aaai.v38i19.30150, 2024. <https://doi.org/10.1609/aaai.v38i19.30150>
- [4] *On Evaluating Adversarial Robustness of Large Vision-Language Models*. arXiv:2305.16934, 2023. <https://arxiv.org/abs/2305.16934>
- [5] *A Survey on Backdoor Threats in Large Language Models (LLMs): Attacks, Defenses, and Evaluation Methods*. DOI:10.53941/tai.2025.100003, 2025. <https://doi.org/10.53941/tai.2025.100003>
- [6] *Large Language Models in Cybersecurity: State-of-the-Art*. arXiv:2402.00891, 2024. <https://arxiv.org/abs/2402.00891>
- [7] *MM-LLMs: Recent Advances in MultiModal Large Language Models*. arXiv:2401.13601, 2024. <https://arxiv.org/abs/2401.13601>
- [8] *Surveying the MLLM Landscape: A Meta-Review of Current Surveys*. arXiv:2409.18991, 2024. <https://arxiv.org/abs/2409.18991>
- [9] *Survey of Adversarial Robustness in Multimodal Large Language Models*. arXiv:2503.13962, 2025. <https://arxiv.org/abs/2503.13962>
- [10] *Explainable and Interpretable Multimodal Large Language Models: A Comprehensive Survey*. arXiv:2412.02104, 2024. <https://arxiv.org/abs/2412.02104>
- [11] *A Survey on Benchmarks of Multimodal Large Language Models*. arXiv:2408.08632, 2024. <https://arxiv.org/abs/2408.08632>
- [12] *Baseline Defenses for Adversarial Attacks Against Aligned Language Models*. arXiv:2309.00614, 2023. <https://arxiv.org/abs/2309.00614>
- [13] *MM-Soc: Benchmarking Multimodal Large Language Models in Social Media Platforms*. arXiv:2402.14154, 2024. <https://arxiv.org/abs/2402.14154>
- [14] *Unique Security and Privacy Threats of Large Language Models: A Comprehensive Survey*. DOI:10.1145/3764113, 2025. <https://doi.org/10.1145/3764113>
- [15] *Multimodal Large Language Models: A Survey*. DOI:10.1109/bigdata59044.2023.10386743, 2023. <https://doi.org/10.1109/bigdata59044.2023.10386743>
- [16] *Large Language Models Meet NL2Code: A Survey*. DOI:10.18653/v1/2023.acl-long.411, 2023. <https://doi.org/10.18653/v1/2023.acl-long.411>
- [17] *Security and Privacy Challenges of Large Language Models: A Survey*. DOI:10.1145/3712001, 2025. <https://doi.org/10.1145/3712001>
- [18] *Multimodal Large Language Models to Support Real-World Fact-Checking*. arXiv:2403.03627, 2024. <https://arxiv.org/abs/2403.03627>
- [19] *A Survey of Attacks on Large Vision-Language Models: Resources, Advances, and Future Trends*. DOI:10.1109/tnnls.2025.3592935, 2025. <https://doi.org/10.1109/tnnls.2025.3592935>

- [20] *A Survey on Evaluation of Multimodal Large Language Models*. arXiv:2408.15769, 2024. <https://arxiv.org/abs/2408.15769>
- [21] *Neural Language Modeling with Visual Features*. arXiv:1903.02930, 2019. <https://arxiv.org/abs/1903.02930>
- [22] *Macaw-LLM: Multi-Modal Language Modeling with Image, Audio, Video, and Text Integration*. arXiv:2306.09093, 2023. <https://arxiv.org/abs/2306.09093>
- [23] *Large-scale Multi-modal Pre-trained Models: A Comprehensive Survey*. DOI:10.1007/s11633-022-1410-8, 2023. <https://doi.org/10.1007/s11633-022-1410-8>
- [24] *Model Composition for Multimodal Large Language Models*. arXiv:2402.12750, 2024. <https://arxiv.org/abs/2402.12750>
- [25] *From CLIP to DINO: Visual Encoders Shout in Multi-modal Large Language Models*. arXiv:2310.08825, 2023. <https://arxiv.org/abs/2310.08825>
- [26] *ChatBridge: Bridging Modalities with Large Language Model as a Language Catalyst*. arXiv:2305.16103, 2023. <https://arxiv.org/abs/2305.16103>
- [27] *MoME: Mixture of Multimodal Experts for Generalist Multimodal Large Language Models*. arXiv:2407.12709, 2024. <https://arxiv.org/abs/2407.12709>
- [28] *Visual Prompting in Multimodal Large Language Models: A Survey*. arXiv:2409.15310, 2024. <https://arxiv.org/abs/2409.15310>
- [29] *A Survey on Multimodal Benchmarks: In the Era of Large AI Models*. arXiv:2409.18142, 2024. <https://arxiv.org/abs/2409.18142>
- [30] *Eagle: Exploring The Design Space for Multimodal LLMs with Mixture of Encoders*. arXiv:2408.15998, 2024. <https://arxiv.org/abs/2408.15998>
- [31] *How Multimodal LLMs Solve Image Tasks: A Lens on Visual Grounding, Task Reasoning, and Answer Decoding*. arXiv:2508.20279, 2025. <https://arxiv.org/abs/2508.20279>
- [32] *Seeing is Understanding: Unlocking Causal Attention into Modality-Mutual Attention for Multimodal LLMs*. arXiv:2503.02597, 2025. <https://arxiv.org/abs/2503.02597>
- [33] *LaVi: Efficient Large Vision-Language Models via Internal Feature Modulation*. arXiv:2506.16691, 2025. <https://arxiv.org/abs/2506.16691>
- [34] *MHA2MLA-VLM: Enabling DeepSeek’s Economical Multi-Head Latent Attention across Vision-Language Models*. arXiv:2601.11464, 2026. <https://arxiv.org/abs/2601.11464>
- [35] *MLIM: Vision-and-Language Model Pre-training with Masked Language and Image Modeling*. arXiv:2109.12178, 2021. <https://arxiv.org/abs/2109.12178>
- [36] *Generating Images with Multimodal Language Models*. arXiv:2305.17216, 2023. <https://arxiv.org/abs/2305.17216>
- [37] *Transfusion: Predict the Next Token and Diffuse Images with One Multi-Modal Model*. arXiv:2408.11039, 2024. <https://arxiv.org/abs/2408.11039>
- [38] *MIBench: Evaluating Multimodal Large Language Models over Multiple Images*. arXiv:2407.15272, 2024. <https://arxiv.org/abs/2407.15272>
- [39] *A Survey on Multimodal Large Language Models for Autonomous Driving*. DOI:10.1109/wacvw60836.2024.00106, 2024. <https://doi.org/10.1109/wacvw60836.2024.00106>

- [40] *Multiway-Adapter: Adapting Multimodal Large Language Models for Scalable Image-Text Retrieval*. DOI:10.1109/icassp48485.2024.10446792, 2024. <https://doi.org/10.1109/icassp48485.2024.10446792>
- [41] *EE-MLLM: A Data-Efficient and Compute-Efficient Multimodal Large Language Model*. arXiv:2408.11795, 2024. <https://arxiv.org/abs/2408.11795>
- [42] *γ -MoD: Exploring Mixture-of-Depth Adaptation for Multimodal Large Language Models*. arXiv:2410.13859, 2024. <https://arxiv.org/abs/2410.13859>
- [43] *When Language Overrides: Revealing Text Dominance in Multimodal Large Language Models*. arXiv:2508.10552, 2025. <https://arxiv.org/abs/2508.10552>
- [44] *MAD: Modality-Adaptive Decoding for Mitigating Cross-Modal Hallucinations in Multimodal Large Language Models*. arXiv:2601.21181, 2026. <https://arxiv.org/abs/2601.21181>
- [45] *A survey on multimodal large language models*. DOI:10.1093/nsr/nwae403, 2024. <https://doi.org/10.1093/nsr/nwae403>
- [46] *Unifying Visual-Semantic Embeddings with Multimodal Neural Language Models*. arXiv:1411.2539, 2014. <https://arxiv.org/abs/1411.2539>
- [47] *Training and Evaluating Multimodal Word Embeddings with Large-scale Web Annotated Images*. arXiv:1611.08321, 2016. <https://arxiv.org/abs/1611.08321>
- [48] *Are Multimodal Models Robust to Image and Text Perturbations?*. arXiv:2212.08044, 2022. <https://arxiv.org/abs/2212.08044>
- [49] *Efficient Multimodal Learning from Data-centric Perspective*. arXiv:2402.11530, 2024. <https://arxiv.org/abs/2402.11530>
- [50] *Expanding Performance Boundaries of Open-Source Multimodal Models with Model, Data, and Test-Time Scaling*. arXiv:2412.05271, 2024. <https://arxiv.org/abs/2412.05271>
- [51] *DreamLLM: Synergistic Multimodal Comprehension and Creation*. arXiv:2309.11499, 2023. <https://arxiv.org/abs/2309.11499>
- [52] *On Evaluation of Adversarial Perturbations for Sequence-to-Sequence Models*. DOI:10.18653/v1/n19-1314, 2019. <https://doi.org/10.18653/v1/n19-1314>
- [53] *Use of LLMs for Illicit Purposes: Threats, Prevention Measures, and Vulnerabilities*. arXiv:2308.12833, 2023. <https://arxiv.org/abs/2308.12833>
- [54] *Exploring Vulnerabilities and Protections in Large Language Models: A Survey*. arXiv:2406.00240, 2024. <https://arxiv.org/abs/2406.00240>
- [55] *Probing the Robustness of Vision-Language Pretrained Models: A Multimodal Adversarial Attack Approach*. arXiv:2408.13461, 2024. <https://arxiv.org/abs/2408.13461>
- [56] *Robust-LLaVA: On the Effectiveness of Large-Scale Robust Image Encoders for Multi-modal Large Language Models*. arXiv:2502.01576, 2025. <https://arxiv.org/abs/2502.01576>
- [57] *Jailbreak Attacks and Defenses against Multimodal Generative Models: A Survey*. arXiv:2411.09259, 2024. <https://arxiv.org/abs/2411.09259>
- [58] *Multimodal Large Language Models for Text-rich Image Understanding: A Comprehensive Review*. arXiv:2502.16586, 2025. <https://arxiv.org/abs/2502.16586>
- [59] *Explaining Multi-modal Large Language Models by Analyzing their Vision Perception*. arXiv:2405.14612, 2024. <https://arxiv.org/abs/2405.14612>

- [60] *MLLM-Protector: Ensuring MLLM’s Safety without Hurting Performance.* arXiv:2401.02906, 2024. <https://arxiv.org/abs/2401.02906>
- [61] *A Brief Introduction to Generative Models.* arXiv:2103.00265, 2021. <https://arxiv.org/abs/2103.00265>
- [62] *Finetuned Multimodal Language Models Are High-Quality Image-Text Data Filters.* arXiv:2403.02677, 2024. <https://arxiv.org/abs/2403.02677>
- [63] *VisionLLM v2: An End-to-End Generalist Multimodal Large Language Model for Hundreds of Vision-Language Tasks.* arXiv:2406.08394, 2024. <https://arxiv.org/abs/2406.08394>
- [64] *MM-InstructEval: Zero-Shot Evaluation of (Multimodal) Large Language Models on Multimodal Reasoning Tasks.* arXiv:2405.07229, 2024. <https://arxiv.org/abs/2405.07229>
- [65] *AlignGPT: Multi-modal Large Language Models with Adaptive Alignment Capability.* arXiv:2405.14129, 2024. <https://arxiv.org/abs/2405.14129>
- [66] *SHIELD: an evaluation benchmark for face spoofing and forgery detection with multimodal large language models.* DOI:10.1007/s44267-025-00079-w, 2025. <https://doi.org/10.1007/s44267-025-00079-w>
- [67] *MME-RealWorld: Could Your Multimodal LLM Challenge High-Resolution Real-World Scenarios that are Difficult for Humans?.* arXiv:2408.13257, 2024. <https://arxiv.org/abs/2408.13257>
- [68] *Jailbreaking LLMs & VLMs: Mechanisms, Evaluation, and Unified Defense.* arXiv:2601.03594, 2026. <https://arxiv.org/abs/2601.03594>
- [69] *Uni-MoE: Scaling Unified Multimodal LLMs With Mixture of Experts.* DOI:10.1109/tpami.2025.3532688, 2025. <https://doi.org/10.1109/tpami.2025.3532688>
- [70] *SPHINX-X: Scaling Data and Parameters for a Family of Multi-modal Large Language Models.* arXiv:2402.05935, 2024. <https://arxiv.org/abs/2402.05935>
- [71] *E5-V: Universal Embeddings with Multimodal Large Language Models.* arXiv:2407.12580, 2024. <https://arxiv.org/abs/2407.12580>
- [72] *ControlMLLM: Training-Free Visual Prompt Learning for Multimodal Large Language Models.* arXiv:2407.21534, 2024. <https://arxiv.org/abs/2407.21534>
- [73] *MMEvol: Empowering Multimodal Large Language Models with Evol-Instruct.* arXiv:2409.05840, 2024. <https://arxiv.org/abs/2409.05840>
- [74] *Libra: Building Decoupled Vision System on Large Language Models.* arXiv:2405.10140, 2024. <https://arxiv.org/abs/2405.10140>
- [75] *X-VILA: Cross-Modality Alignment for Large Language Model.* arXiv:2405.19335, 2024. <https://arxiv.org/abs/2405.19335>
- [76] *From One-to-One to Many-to-Many: Dynamic Cross-Layer Injection for Deep Vision-Language Fusion.* arXiv:2601.10710, 2026. <https://arxiv.org/abs/2601.10710>
- [77] *Visual Representation Alignment for Multimodal Large Language Models.* arXiv:2509.07979, 2025. <https://arxiv.org/abs/2509.07979>
- [78] *Unleashing the Intrinsic Visual Representation Capability of Multimodal Large Language Models.* arXiv:2512.06281, 2025. <https://arxiv.org/abs/2512.06281>
- [79] *How Visual Representations Map to Language Feature Space in Multimodal LLMs.* arXiv:2506.11976, 2025. <https://arxiv.org/abs/2506.11976>

- [80] *Deciphering Cross-Modal Alignment in Large Vision-Language Models with Modality Integration Rate.* arXiv:2410.07167, 2024. <https://arxiv.org/abs/2410.07167>
- [81] *Directed-Tokens: A Robust Multi-Modality Alignment Approach to Large Language-Vision Models.* arXiv:2508.14264, 2025. <https://arxiv.org/abs/2508.14264>
- [82] *Multimodal Fact-Checking with Vision Language Models: A Probing Classifier based Solution with Embedding Strategies.* arXiv:2412.05155, 2024. <https://arxiv.org/abs/2412.05155>
- [83] *AVHBench: A Cross-Modal Hallucination Benchmark for Audio-Visual Large Language Models.* arXiv:2410.18325, 2024. <https://arxiv.org/abs/2410.18325>
- [84] *Optimizing Multimodal Language Models through Attention-based Interpretability.* arXiv:2511.23375, 2025. <https://arxiv.org/abs/2511.23375>
- [85] *From Image to Video, what do we need in multimodal LLMs?.* arXiv:2404.11865, 2024. <https://arxiv.org/abs/2404.11865>
- [86] *FMM-Attack: A Flow-based Multi-modal Adversarial Attack on Video-based LLMs.* arXiv:2403.13507, 2024. <https://arxiv.org/abs/2403.13507>
- [87] *Medical MLLM Is Vulnerable: Cross-Modality Jailbreak and Mismatched Attacks on Medical Multimodal Large Language Models.* DOI:10.1609/aaai.v39i4.32396, 2025. <https://doi.org/10.1609/aaai.v39i4.32396>
- [88] *JPS: Jailbreak Multimodal Large Language Models with Collaborative Visual Perturbation and Textual Steering.* DOI:10.1145/3746027.3754561, 2025. <https://doi.org/10.1145/3746027.3754561>
- [89] *MAA: Meticulous Adversarial Attack against Vision-Language Pre-trained Models.* arXiv:2502.08079, 2025. <https://arxiv.org/abs/2502.08079>
- [90] *Assessing Adversarial Robustness of Large Language Models: An Empirical Study.* arXiv:2405.02764, 2024. <https://arxiv.org/abs/2405.02764>
- [91] *Adversarial Confusion Attack: Disrupting Multimodal Large Language Models.* arXiv:2511.20494, 2025. <https://arxiv.org/abs/2511.20494>
- [92] *Hierarchical Refinement of Universal Multimodal Attacks on Vision-Language Models.* arXiv:2601.10313, 2026. <https://arxiv.org/abs/2601.10313>
- [93] *On opportunities and challenges of large multimodal foundation models in education.* DOI:10.1038/s41539-025-00301-w, 2025. <https://doi.org/10.1038/s41539-025-00301-w>
- [94] *Towards LLM-Centric Multimodal Fusion: A Survey on Integration Strategies and Techniques.* arXiv:2506.04788, 2025. <https://arxiv.org/abs/2506.04788>
- [95] *Multimodal Needle in a Haystack: Benchmarking Long-Context Capability of Multimodal Large Language Models.* arXiv:2406.11230, 2024. <https://arxiv.org/abs/2406.11230>
- [96] *SEA: Supervised Embedding Alignment for Token-Level Visual-Textual Integration in MLLMs.* arXiv:2408.11813, 2024. <https://arxiv.org/abs/2408.11813>
- [97] *mDPO: Conditional Preference Optimization for Multimodal Large Language Models.* arXiv:2406.11839, 2024. <https://arxiv.org/abs/2406.11839>
- [98] *LLMBind: A Unified Modality-Task Integration Framework.* arXiv:2402.14891, 2024. <https://arxiv.org/abs/2402.14891>
- [99] *Are We on the Right Way for Evaluating Large Vision-Language Models?.* arXiv:2403.20330, 2024. <https://arxiv.org/abs/2403.20330>

- [100] *Unveiling Uncertainty: A Deep Dive into Calibration and Performance of Multimodal Large Language Models*. arXiv:2412.14660, 2024. <https://arxiv.org/abs/2412.14660>
- [101] *LENS: Multi-level Evaluation of Multimodal Reasoning with Large Language Models*. arXiv:2505.15616, 2025. <https://arxiv.org/abs/2505.15616>
- [102] *CODIS: Benchmarking Context-Dependent Visual Comprehension for Multimodal Large Language Models*. arXiv:2402.13607, 2024. <https://arxiv.org/abs/2402.13607>
- [103] *Benchmarking Large Multimodal Models against Common Corruptions*. arXiv:2401.11943, 2024. <https://arxiv.org/abs/2401.11943>
- [104] *LLaVA-RE: Binary Image-Text Relevancy Evaluation with Multimodal Large Language Model*. arXiv:2508.05602, 2025. <https://arxiv.org/abs/2508.05602>
- [105] *Connector-S: A Survey of Connectors in Multi-modal Large Language Models*. arXiv:2502.11453, 2025. <https://arxiv.org/abs/2502.11453>
- [106] *Visual Perception by Large Language Model’s Weights*. arXiv:2405.20339, 2024. <https://arxiv.org/abs/2405.20339>
- [107] *Quantifying and Mitigating Unimodal Biases in Multimodal Large Language Models: A Causal Perspective*. arXiv:2403.18346, 2024. <https://arxiv.org/abs/2403.18346>
- [108] *Safety of Multimodal Large Language Models on Images and Text*. arXiv:2402.00357, 2024. <https://arxiv.org/abs/2402.00357>
- [109] *Improving Multimodal Large Language Models Using Continual Learning*. arXiv:2410.19925, 2024. <https://arxiv.org/abs/2410.19925>
- [110] *Incorporating Visual Experts to Resolve the Information Loss in Multimodal Large Language Models*. arXiv:2401.03105, 2024. <https://arxiv.org/abs/2401.03105>
- [111] *VL-Mamba: Exploring State Space Models for Multimodal Learning*. arXiv:2403.13600, 2024. <https://arxiv.org/abs/2403.13600>
- [112] *LaVy: Vietnamese Multimodal Large Language Model*. arXiv:2404.07922, 2024. <https://arxiv.org/abs/2404.07922>
- [113] *MammothModa: Multi-Modal Large Language Model*. arXiv:2406.18193, 2024. <https://arxiv.org/abs/2406.18193>
- [114] *SLAM-LLM: A Modular, Open-Source Multimodal Large Language Model Framework and Best Practice for Speech, Language, Audio and Music Processing*. DOI:10.1109/jstsp.2026.3653157, 2026. <https://doi.org/10.1109/jstsp.2026.3653157>
- [115] *Planting a SEED of Vision in Large Language Model*. arXiv:2307.08041, 2023. <https://arxiv.org/abs/2307.08041>
- [116] *MMBench: Is Your Multi-modal Model an All-around Player?*. arXiv:2307.06281, 2023. <https://arxiv.org/abs/2307.06281>
- [117] *Be Careful about Poisoned Word Embeddings: Exploring the Vulnerability of the Embedding Layers in NLP Models*. DOI:10.18653/v1/2021.naacl-main.165, 2021. <https://doi.org/10.18653/v1/2021.naacl-main.165>
- [118] *Rethinking Stealthiness of Backdoor Attack against NLP Models*. DOI:10.18653/v1/2021.acl-long.431, 2021. <https://doi.org/10.18653/v1/2021.acl-long.431>
- [119] *An Empirical Study on Adversarial Attack on NMT: Languages and Positions Matter*. DOI:10.18653/v1/2021.acl-short.58, 2021. <https://doi.org/10.18653/v1/2021.acl-short.58>

- [120] *Adversarial Robustness for Visual Grounding of Multimodal Large Language Models*. arXiv:2405.09981, 2024. <https://arxiv.org/abs/2405.09981>